

A Regime-Switching Continuous Hidden Markov Model as a Reference Synthetic-Data Generator for Equity Returns

Abdulrahman Alswaidan* Cade Jin† Jeffrey D. Varner‡

April 26, 2026

Abstract

Synthetic equity-return generators support stress testing, regulatory backtesting, and scenario design, but existing approaches struggle to reproduce the three canonical stylized facts of daily returns simultaneously: heavy-tailed marginals, negligible linear autocorrelation, and slow decay of the absolute-return autocorrelation function (ACF). We show that a continuous hidden Markov model (CHMM) at moderate state resolution reproduces all three within a single EM scaffold. The mechanism is spectral: the absolute-return autocovariance admits the closed-form mixture-of-eigenvalues identity $\rho_{|G|}(\tau) = \sum_{k \geq 2} w_k \lambda_k^\tau$ over the non-unit eigenvalues of the transition matrix $\mathbf{T}$, recasting the low- K limitation of regime-switching HMMs as a K -rank statement on $\mathbf{T}$. The same scaffold accommodates Gaussian (CHMM-N), Student-t (CHMM-t), and Laplace (CHMM-L) emissions through a single M-step swap. On ten years of daily SPDR S&P 500 ETF (SPY) data with a 2.5-year held-out window, CHMM at $K = 18$ attains 93.8% in-sample and 84.2% out-of-sample Kolmogorov–Smirnov pass rates and an absolute-return ACF mean absolute error within 0.003 of GARCH(1,1); VaR and Expected-Shortfall envelopes bracket observed values at 1% and 5% on both windows. CHMM occupies the joint-fit corner of the metric panel where i.i.d. baselines win on marginal fidelity and GARCH wins on volatility clustering. Multi-asset synthesis composes through a rank-based Student-t copula on the CHMM marginals: profile MLE selects $\nu^* = 6$, reproducing the off-diagonal correlation matrix on six SPY-member assets with MAE 0.027.

Keywords: Continuous Hidden Markov Model, Baum-Welch, Student-t Emissions, Stylized Facts, Volatility Clustering, Spectral ACF, Value-at-Risk, Synthetic-Data Generator.

1 Introduction

Synthetic generators of equity-return time series support stress testing, regulatory backtesting, scenario design, and the training of downstream machine-learning models on data that cannot be released for privacy or licensing reasons [1–3]. The bar for a useful generator is that it reproduce, simultaneously, the canonical stylized facts of daily returns: heavy-tailed marginals, negligible linear autocorrelation, and slow decay of the autocorrelation function (ACF) of absolute returns [4, 5]. Generalized autoregressive conditional heteroskedasticity (GARCH) models [6, 7] capture volatility

*Cornell University, Robert Frederick Smith School of Chemical and Biomolecular Engineering, Ithaca, NY, USA.
Email: aa2725@cornell.edu.

†Cornell Ann S. Bowers College of Computing and Information Science, Cornell University, Ithaca, NY, USA.
Email: cj383@cornell.edu.

‡Cornell University, Robert Frederick Smith School of Chemical and Biomolecular Engineering, Ithaca, NY, USA.
Email: jdv27@cornell.edu.

clustering but not the multi-regime structure of the marginal; i.i.d. generators match the marginal but destroy temporal persistence; deep generators [8–10] reproduce some stylized facts at the cost of interpretability and reproducibility.

An influential negative result has shaped two decades of work on the regime-switching route. Rydén et al. [11] fitted Gaussian-emission hidden Markov models (HMMs) at $K = 2$ –3 states and concluded that the geometric sojourn-time distribution of a standard Markov chain produces exponential ACF decay too fast to match the observed slow decay; Bulla and Bulla [12] restored ACF fidelity by replacing the Markov chain with a hidden semi-Markov chain at the cost of substantially increased complexity.

We identify the spectral root cause. The absolute-return autocovariance of a stationary continuous HMM admits a closed-form mixture-of-eigenvalues identity $\rho_{|G|}(\tau) = \sum_{k \geq 2} w_k \lambda_k^\tau$ over the non-unit eigenvalues of $\mathbf{T}$ (Section 3.8); the empirical Rydén et al. result is then a K -rank statement on $\mathbf{T}$. At $K = 2$ the ACF is mono-exponential and slow decay can be matched by tuning λ_2 near 1; the binding constraint at $K = 2$ is instead the two-component marginal, whose representational capacity is insufficient for the empirical kurtosis target, failing distributional Kolmogorov–Smirnov (KS). At $K \geq 3$ the model admits multiple non-unit eigenvalues and a richer mixture, and the temporal and distributional constraints decouple. The semi-Markov route injects the missing flexibility through an explicit sojourn distribution; moderate- K continuous HMM achieves the same end through the spectral richness of $\mathbf{T}$ and the additional mixture components, inside the standard Baum-Welch scaffold.

We instantiate this on ten years of daily SPDR S&P 500 exchange-traded fund (ETF) (ticker SPY) data with a 2.5-year held-out window, using a continuous HMM trained by Baum-Welch [13, 14] under quantile-based initialization, in three emission variants sharing identical forward-backward recursions: CHMM-N with Gaussian emissions, CHMM-t with per-state Student-t emissions whose degrees of freedom ν_k are updated inside the M-step by a golden-section search on the expectation conditional maximization (ECM) Q -function in the tradition of Peel and McLachlan [15] and Liu and Rubin [16], and CHMM-L with Laplace emissions whose location is the weighted median and whose scale is the weighted mean absolute deviation (MAD), both in closed form. Against a benchmark panel covering the i.i.d. bootstrap (17), Gaussian and Laplace i.i.d., GARCH(1,1) (7), heavy-tailed GARCH(1,1)- t (18), and the regime-switching Markov-switching GARCH (MS-GARCH) foil of Haas et al. [19], the CHMM at moderate K attains 93.8% in-sample (IS) and 84.2% out-of-sample (OoS) KS pass rates and an absolute-return ACF mean absolute error (MAE) of 0.0507, essentially matching GARCH’s temporal fidelity without sacrificing distributional fit. The three emission variants land at three distinct points on the kurtosis axis: CHMM-N undershoots (5.0 simulated vs. 7.7 observed), CHMM-t overshoots (14.6), and CHMM-L matches cleanly (6.8). The unconditional Value-at-Risk (VaR) and Expected-Shortfall (ES) envelopes bracket observed historical values at both 1% and 5% on the IS and OoS windows. No single benchmark dominates every metric; the i.i.d. bootstrap is the strongest distributional baseline, GARCH is the strongest unconditional volatility-clustering baseline, and the CHMM is the strongest joint-fit row at moderate state resolution. Multi-asset synthesis composes cleanly through a rank-based Student-t copula on top of the CHMM marginals (profile maximum likelihood estimation (MLE) selects $\nu^* = 6$; six-asset off-diagonal correlation MAE 0.027). Returns are assumed locally stationary across the ten-year IS window; substantial OoS degradation on individual single names (NVIDIA, ticker NVDA; JPMorgan Chase, ticker JPM) reflects the well-known stationarity-scope limit on financial time series [20–22] and motivates periodic refit at deployment in the spirit of Pesaran and Timmermann [23].

Our contributions are: (i) a closed-form mixture-of-eigenvalues identity for the absolute-return ACF that recasts the Rydén low- K failure as a K -rank statement on $\mathbf{T}$; (ii) three emission families (Gaussian, Student-t, Laplace) under a unified ECM scaffold, where the M-step is the only

architectural difference; (iii) a single-window IS / OoS empirical study on SPY plus a six-ticker generalisation panel and a six-asset Student-t copula extension.

The paper is organized as follows. Section 2 reviews related work; Section 3 describes the data, the continuous HMM, the EM estimator, and the spectral mechanism for the absolute-return ACF; Section 4 presents the empirical study and the VaR / ES envelope calibration; Sections 5–6 discuss and conclude.

2 Related Work

Hidden Markov models in finance. Hamilton [24] introduced Markov-switching to economics; Schaller and Van Norden [25] showed monthly equity returns admit a two-state Markov-switching specification. The seminal evaluation against the Cont stylized facts is due to Rydén et al. [11]: Gaussian-emission HMMs at $K = 2$ –3 states fail the slow ACF decay of absolute returns because the geometric sojourn-time distribution of a standard Markov chain produces exponential decay too fast to match the observed signature. Bulla and Bulla [12] restored fidelity with hidden semi-Markov models at substantially higher complexity. Nystrup et al. [26] extends the same line to long-memory volatility. Maximum-likelihood fitting of heavy-tailed innovations follows Peel and McLachlan [15] and Liu and Rubin [16]. The discrete-state predecessor of this paper [3] reproduced the same three stylized facts via Laplace quantile binning plus a Poisson-jump duration mechanism; here we replace the discretization step with continuous emissions and recover the same target through the spectral mechanism of Section 3.8.

GARCH and its extensions. GARCH [6, 7] encodes volatility persistence through a single decay parameter but imposes a unimodal innovation distribution and fails distributional Kolmogorov–Smirnov on equity returns. Heavy-tailed innovations [18], asymmetric leverage extensions (EGARCH, 27; GJR-GARCH, 28; threshold GARCH, 29), and regime-switching extensions (MS-GARCH, 19) close some gaps but inherit the unimodal-innovation constraint. Realised-variance frameworks [30, 31] provide a complementary forecasting baseline.

Deep generative models for financial time series. Generative adversarial network (GAN)-based generators reproduce the visual appearance of return series but typically fail volatility-clustering tests unless architectures encode temporal dependence explicitly [32, 33]; the TimeGAN family of Yoon et al. [8], the convolutional Wasserstein GAN (WGAN) QuantGAN of Wiese et al. [9], and the autoregressive-diffusion approach of Rasul et al. [10] are the standard references. We include QuantGAN as the deep-generative row in our extended panel (Appendix A); the present paper’s contribution is grounded against parametric and non-parametric baselines.

Cross-asset dependence and synthetic-data evaluation. For multi-asset synthesis, the Single Index Model [34] propagates a market factor through a rank-one linear decomposition but distorts each non-market marginal; Sklar’s theorem [35, 36] supports a cleaner decomposition in which each asset’s fitted marginal is preserved and only the copula is estimated, with Gaussian and Student-t copulas [37, 38] the standard choices and rank-reordering simulation [39] the standard injection. Synthetic-data evaluation is anchored on the framework of Stenger et al. [40], the proper-scoring-rule literature [41–43], kernel two-sample tests [44], and signature-based metrics [45, 46].

Pipeline A (single-asset, Section 4.3):



Pipeline B (cross-asset, Section 4.5):




Figure 1. Schematic of Pipeline A (single-asset) and Pipeline B (cross-asset). Pipeline A fits a CHMM per ticker independently and evaluates single-asset marginal, tail, and temporal metrics. Pipeline B reuses the per-asset CHMM fits as *marginals* and injects cross-asset *dependence* through the rank-based Student-t copula (profile MLE selects $\nu^* = 6$), preserving each asset’s fitted CHMM distribution exactly while coupling asset ranks to match observed multi-asset correlation structure. Headline: the two pipelines share the single-asset scaffold, so improvements to the CHMM at the marginal step propagate into Pipeline B without any changes to the dependence layer.

3 Methods

3.1 Data and Excess Growth Rates

The single-asset analysis uses daily SPDR S&P 500 ETF (SPY) prices from January 3, 2014 to January 3, 2024 (a ten-year training window, $T_{\text{IS}} = 2,516$), with the period from January 4, 2024 through April 20, 2026 ($T_{\text{OoS}} = 572$) held out. Cross-ticker generalization additionally covers NVIDIA (NVDA), Johnson & Johnson (JNJ), JPMorgan Chase (JPM), Apple (AAPL), and the Invesco QQQ Trust (QQQ). For ticker i and trading day j , the annualized excess growth rate is

$$G_{i,j} \equiv \frac{1}{\Delta t} \ln \left(\frac{P_{i,j}}{P_{i,j-1}} \right) - r_f, \quad \Delta t = 1/252, \quad r_f = 0. \quad (1)$$

The empirical study is organised along two pipelines, illustrated in Figure 1. *Pipeline A* fits a CHMM to each ticker independently and evaluates single-asset marginal, tail, and temporal metrics; this is the scope of the SPY headline panel (Section 4.3), the cross-ticker generalisation (Section 4.4), and the unconditional VaR / ES envelope (Section 4.6). *Pipeline B* reuses the per-asset Pipeline-A CHMM fits as *marginals* and injects cross-asset *dependence* through the rank-based Student-t copula of Section 3.4, evaluated in Section 4.5.

3.2 Continuous Hidden Markov Model

We model the continuous excess growth rate G_t as a continuous HMM $\mathcal{M} = (\mathcal{S}, \mathbf{T}, \mathbf{B}, \boldsymbol{\pi})$ with K latent states, a per-state emission density $b_k(x; \boldsymbol{\theta}_k)$, transition matrix $\mathbf{T} \in \mathbb{R}^{K \times K}$, and initial distribution $\boldsymbol{\pi}$ [14]. The stationary marginal is the K -component mixture

$$f(x) = \sum_{k=1}^K \bar{\pi}_k b_k(x; \boldsymbol{\theta}_k), \quad \bar{\boldsymbol{\pi}} = \bar{\boldsymbol{\pi}} \mathbf{T}, \quad (2)$$




Figure 2. CHMM architecture and training loop. A K -state latent Markov chain with transition matrix $\mathbf{T}$ generates the regime sequence s_t ; each s_t conditionally emits the observed excess growth rate G_t through a state-specific density b_{s_t} . The three CHMM variants (Gaussian, Student-t, Laplace) share everything except the emission family. Joint estimation alternates an E-step (log-space forward-backward) and a family-specific M-step until the log-likelihood change falls below 10^{-4} .

and supplies heavy tails either through the spread of component moments (Gaussian components) or through the emission family itself (Student-t or Laplace). Existence and uniqueness of $\bar{\pi}$ require irreducibility and aperiodicity of $\mathbf{T}$, recorded as Assumption 1 in Section 3.8; the spectral mechanism behind the absolute-return ACF that this mixture induces is derived in Section 3.8. We learn $\mathbf{T}$, $\theta_{1:K}$, and π jointly from continuous observations via EM (Section 3).

CHMM-N: Gaussian emissions. $b_k(x) = \mathcal{N}(x; \mu_k, \sigma_k^2)$ with per-state location μ_k and variance σ_k^2 .

CHMM-t: Student-t emissions. $b_k(x) = t_{\nu_k}(x; \mu_k, \sigma_k)$ with per-state location μ_k , scale σ_k , and degrees of freedom $\nu_k \in [\nu_{\min}, \nu_{\max}]$. Per-state degrees of freedom let heavy-tailed regimes coexist with near-Gaussian regimes in the same model.

CHMM-L: Laplace emissions. $b_k(x) = \text{Laplace}(x; \mu_k, b_k)$ with per-state location μ_k and scale b_k . The heavier-than-Gaussian Laplace tail (per-component excess kurtosis 3) provides an alternative route to tail fidelity that carries no shape parameter and is strictly cheaper than CHMM-t.

3.3 Benchmark Generators

We benchmark the CHMM against representatives from each of the standard generator tiers, all fitted on the same IS window. The *i.i.d. bootstrap* [17] samples T returns uniformly with replacement, preserving the empirical marginal exactly and serving as the non-parametric distributional ceiling. The *Gaussian* i.i.d. generator draws $G_t \sim \mathcal{N}(\hat{\mu}, \hat{\sigma}^2)$ from the sample moments; the *Laplace* i.i.d. generator draws $G_t \sim \text{Laplace}(\hat{m}, \hat{b})$ from the sample median and mean absolute deviation. *GARCH(1,1)* [7] uses $G_t = \mu + \sigma_t z_t$, $\sigma_t^2 = \omega + \alpha(G_{t-1} - \mu)^2 + \beta\sigma_{t-1}^2$ with $z_t \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, 1)$, fitted by Nelder-Mead maximum likelihood; this is the parametric volatility-clustering workhorse. *GARCH(1,1)-t* [18]

replaces the Gaussian innovation with a Student- t , supplying a heavy-tailed volatility-clustering competitor with the same recursion. *MS-GARCH* (Markov-switching GARCH, 19) is the regime-switching foil: a discrete two-state hidden chain modulates the GARCH parameters, providing the closest peer to the CHMM scaffold among classical baselines. The full extended GARCH-family panel (EGARCH, GJR-GARCH, HAR-RV, MS-GARCH $K = 3$), the explicit semi-Markov CHMM foil, and a deep-generative QuantGAN [9] row are reported in Appendix A.

3.4 Cross-Asset Dependence (Pipeline B)

The univariate analysis fits a separate CHMM to each ticker independently (Pipeline A). For multi-asset synthesis we extend Pipeline A through the rank-based Student- t copula [37, 38], which preserves each asset’s fitted CHMM marginal exactly while injecting cross-sectional tail-coupled dependence. Of the three constructions evaluated in our larger panel (Single Index Model (SIM) [34], Gaussian copula, Student- t copula), the Student- t copula is the strongest dependence layer in every relevant metric; we report it here as Pipeline B and relegate the SIM and Gaussian copula constructions to Appendix A. By Sklar’s theorem [35, 36], any joint cumulative distribution function (CDF) with continuous marginals factors uniquely into a copula and the marginals; this lets us fix each marginal at its per-asset CHMM and estimate only the copula. We use the nonparametric probability integral transform $u_{j,t} = \text{rank}(g_{j,t})/(T + 1)$, estimate the copula correlation matrix Σ by the analytic Kendall’s- τ inversion $\rho_{ij} = \sin(\pi\tau_{ij}/2)$ [38, 47], and select the degrees-of-freedom parameter ν by profile maximum likelihood over $\nu \in \{2, 3, 4, 5, 6, 8, 10, 15, 20, 30\}$. Cross-asset simulation follows the rank-reordering construction of Iman and Conover [39]: (i) draw T independent paths from each per-asset CHMM, (ii) sample T copula variates, and (iii) reorder each column of the CHMM matrix to match the ranks of the copula sample. Reordering only permutes simulated values, so every per-asset marginal is preserved *exactly* while cross-asset dependence is injected through rank coupling; this is a standard property of the rank-reordering construction.

3.5 Evaluation Protocol

For each model we simulate $P = 1,000$ independent paths of length T and evaluate four core metrics, following the framework of [40]. All experiments use a deterministic global random seed (20260420); per-experiment sub-seeds are derived additively so that each block is independently reproducible from a single seed root. The two-sample Kolmogorov–Smirnov [48, 49] pass rate at $\alpha = 0.05$ over the 1,000 paths scores distributional fidelity; the mean simulated excess kurtosis scores tail behaviour; and the mean absolute error of the absolute-return ACF over $L = 252$ lags,

$$\text{ACF-MAE} = \frac{1}{L} \sum_{\tau=1}^L \left| \hat{\rho}_{|G|}^{\text{obs}}(\tau) - \bar{\rho}_{|G|}^{\text{sim}}(\tau) \right|, \quad (3)$$

is the direct diagnostic for the Rydén et al. limitation. A separate Value-at-Risk panel (Section 4.6) carries the unconditional-Kupiec [50] and Expected-Shortfall envelope diagnostics. Estimation of $(\mathbf{T}, \{\boldsymbol{\theta}_k\}, \boldsymbol{\pi})$ proceeds by EM in log-space. This section gives the shared scaffold and the family-specific M-step updates; the state-count selection that operationalises this scaffold is reported in Section 4.

3.6 Shared EM Scaffold

All three emission families share the same log-space forward-backward recursions and the same quantile-based initialization [51]: observations are sorted, partitioned into K equal-sized chunks, and

each state's initial location and scale are seeded from the corresponding chunk; $T_{ij}^{(0)} = \pi_k^{(0)} = 1/K$. EM iterates to tolerance $|\mathcal{L}^{(n)} - \mathcal{L}^{(n-1)}| < 10^{-4}$ with `max_iter` = 60; convergence is typically reached in 20–40 iterations. Quantile-based initialization is what makes Baum-Welch reach a non-degenerate $K \geq 3$ fit: random initialization standard in the 1990s often converges to overlapping local optima, which is one of the two compounding sources of the low- K failure observed by Rydén et al. [11] (the other is the spectral one of Section 3.8). Algorithmic details and the forward-backward recursion are in Appendix A.

3.7 Family-Specific M-Step Updates

CHMM-N (Gaussian). The classical Baum-Welch M-step [13] is closed form: $\mu_k \leftarrow \sum_t \gamma_t(k) O_t / \sum_t \gamma_t(k)$ and $\sigma_k^2 \leftarrow \sum_t \gamma_t(k) (O_t - \mu_k)^2 / \sum_t \gamma_t(k)$.

CHMM-t (Student-t via ECM). We adopt the ECM formulation of Peel and McLachlan [15] and Liu and Rubin [16], the same heavy-tailed innovation choice studied by Abanto-Valle et al. [52] for stochastic-volatility-in-mean models. The latent precision

$$u_{t,k} = \frac{\nu_k + 1}{\nu_k + ((O_t - \mu_k)/\sigma_k)^2} \quad (4)$$

is treated as an augmented E-step quantity, and the M-step updates

$$\mu_k \leftarrow \frac{\sum_t \gamma_t(k) u_{t,k} O_t}{\sum_t \gamma_t(k) u_{t,k}}, \quad \sigma_k^2 \leftarrow \frac{\sum_t \gamma_t(k) u_{t,k} (O_t - \mu_k)^2}{\sum_t \gamma_t(k)} \quad (5)$$

have closed form conditional on ν_k . The degrees of freedom ν_k are updated by a golden-section search maximizing the per-state Q -function $Q_k(\nu) = \sum_t \gamma_t(k) \log t_\nu(O_t; \mu_k, \sigma_k)$ over the bracket $(\nu_{\min}, \nu_{\max}) = (2.1, 50)$. The two-block CM step preserves the EM monotonicity guarantee on the compact bracket, so the incomplete-data log-likelihood is non-decreasing across iterations.

CHMM-L (Laplace). The weighted maximum-likelihood estimators are in closed form: μ_k is the posterior-weighted median of the observations (weights $\gamma_t(k)$, computed by cumulative-weight bracketing with linear interpolation between neighbouring order statistics), and $b_k \leftarrow \sum_t \gamma_t(k) |O_t - \mu_k| / \sum_t \gamma_t(k)$.

3.8 Spectral Mechanism for the Absolute-Return ACF

This subsection derives the spectral mechanism behind the absolute-return ACF (the central technical contribution) and states the assumptions under which it holds. Identifiability, MLE consistency, ECM monotonicity, and rank-reordering marginal preservation are each one-line specialisations of standard results; we collect them as cited propositions in Appendix A with one-sentence proof sketches.

Assumptions.

Assumption 1 (Irreducibility and aperiodicity). *The transition matrix $\mathbf{T} \in \mathbb{R}^{K \times K}$ is irreducible and aperiodic on the state space $\{1, \dots, K\}$.*

Assumption 2 (Finite second moment and emission contrast). *Each per-state emission density $b_k(\cdot; \theta_k)$ has finite second moment, $\mathbb{E}_{b_k}[G^2] < \infty$, and at least two states differ in either location μ_k or scale σ_k .*

Under Assumption 1 the chain $\mathbf{T}$ has a unique stationary distribution $\bar{\pi}$ satisfying $\bar{\pi} \mathbf{T} = \bar{\pi}$, and admits a spectral decomposition with $\lambda_1 = 1$ a simple eigenvalue and all other eigenvalues $\lambda_2, \dots, \lambda_K$ strictly inside the unit disk [53].

Spectral mechanism. Let G_t be a stationary CHMM, $\bar{\pi}$ the stationary distribution of $\mathbf{T}$, and define the per-state absolute-return moments

$$m_k = \mathbb{E}[|G_t| \mid s_t = k], \quad M_k = \mathbb{E}[G_t^2 \mid s_t = k], \quad k = 1, \dots, K. \quad (6)$$

Conditioning on the latent state pair $(s_t, s_{t+\tau})$ and using that emissions are conditionally independent given the state sequence,

$$\mathbb{E}[|G_t| |G_{t+\tau}|] = \sum_{i,j} \bar{\pi}_i (\mathbf{T}^\tau)_{ij} m_i m_j = \mathbf{m}^\top \text{diag}(\bar{\pi}) \mathbf{T}^\tau \mathbf{m},$$

with $\mathbf{m} = (m_1, \dots, m_K)^\top$. Under Assumption 1 the spectral decomposition of $\mathbf{T}$ reads $\mathbf{T}^\tau = \mathbf{1} \bar{\pi}^\top + \sum_{k=2}^K \lambda_k^\tau \mathbf{v}_k \mathbf{w}_k^\top$ with $(\mathbf{v}_k, \mathbf{w}_k)$ the right and left eigenvectors of $\mathbf{T}$ associated with λ_k (normalised so that $\mathbf{w}_k^\top \mathbf{v}_k = 1$). The first term contributes $(\bar{\pi}^\top \mathbf{m})^2 = (\mathbb{E}|G_t|)^2$, and subtracting the product of marginals yields the autocovariance

$$\text{Cov}(|G_t|, |G_{t+\tau}|) = \sum_{k=2}^K c_k \lambda_k^\tau, \quad c_k = (\mathbf{m}^\top \text{diag}(\bar{\pi}) \mathbf{v}_k) (\mathbf{w}_k^\top \mathbf{m}), \quad (7)$$

and the autocorrelation

$$\rho_{|G|}(\tau) = \sum_{k=2}^K w_k \lambda_k^\tau, \quad w_k = c_k / \sigma_{|G|}^2. \quad (8)$$

Equation (8) is the closed-form mixture-of-eigenvalues identity for the absolute-return ACF: the autocorrelation at every lag is a weighted sum of geometric decays at the non-unit eigenvalues of $\mathbf{T}$. Setting $\tau = 0$ gives $\rho_{|G|}(0) = \sum_{k \geq 2} w_k = 1$, so the weights w_k form a (signed) decomposition of unity across the non-unit spectrum. The same derivation with $|G_t|$ replaced by G_t^2 and $\mathbf{m}$ by $(M_1, \dots, M_K)^\top$ gives the analogous form for the squared-return ACF.

Rydén separation as a K -rank statement. At $K = 2$, equation (8) reduces to the mono-exponential form $\rho_{|G|}(\tau) = w_2 \lambda_2^\tau$ with $w_2 = 1$. The eigenvalues of a 2×2 stochastic matrix are $\lambda_1 = 1$ and $\lambda_2 = T_{11} + T_{22} - 1$, so a target slow-decay rate $\rho^*(\tau) = \exp(-\tau/\tau^*)$ over horizon $\tau \in \{1, \dots, L\}$ is matched by tuning $T_{11} + T_{22}$ to set $\lambda_2 \approx \exp(-1/\tau^*)$. The temporal constraint at $K = 2$ is therefore satisfiable; the binding constraint is the marginal. A two-component mixture $f(x) = \bar{\pi}_1 b_1(x) + \bar{\pi}_2 b_2(x)$ has bounded representational capacity, and at the empirical equity-return target (excess kurtosis ≈ 7 to 8) it cannot simultaneously match the central peak and the heavy tail under any choice of (μ_k, σ_k, ν_k) within the standard families, irrespective of how the transitions are tuned. At $K \geq 3$ the right-hand side admits at least two non-unit eigenvalues, the ACF can match $\rho^*(\tau)$ over the same horizon with each individual $|\lambda_k|$ bounded away from 1, and the marginal mixture has enough components to match the observed tail. The empirical Rydén et al. low- K failure is therefore a K -rank statement on two coupled axes: K controls both the number of available decay modes in the ACF and the component count of the marginal mixture, and both matter for the joint fit. The empirical correlate is that ACF-MAE is essentially flat across $K \in \{3, 6, \dots, 21\}$ (the temporal constraint is non-binding once $K \geq 2$) while KS pass rates rise with K (the marginal constraint binds at low K and relaxes at moderate K); see Section 4.

The spectral identity (8) is the mechanism behind the moderate- K ACF claim. Under standard EM monotonicity [54, 55], the per-state ν_k ECM update of Section 3.7 preserves the incomplete-data log-likelihood as non-decreasing on the compact bracket $[\nu_{\min}, \nu_{\max}]$; the closed-form weighted-median / weighted-MAD M-step for Laplace emissions [14] preserves the same guarantee. Identifiability up to label permutation in all three emission families (CHMM-N, CHMM-t, CHMM-L) follows from Allman et al. [56] together with Yakowitz and Spragins [57], and MLE consistency from Bickel et al. [58] and Douc et al. [59] under Assumptions 1–2. Formal statements with proof sketches are collected in Appendix A.

4 Empirical Study

4.1 Descriptive Statistics and Stylized Facts

All three canonical stylized facts are present in both the IS ($T = 2,516$) and OoS ($T = 572$) windows. The IS excess growth rates exhibit excess kurtosis of 7.678, strong negative skewness (-0.751), and a Jarque–Bera statistic of 6,416.6, decisively rejecting normality ($p < 0.001$). The Ljung–Box test on absolute returns rejects no autocorrelation ($LB = 2,959.3$, $p < 0.001$), confirming pronounced volatility clustering. The OoS window shares the same qualitative profile (excess kurtosis 5.294, $LB_{|G|} = 250.6$).

4.2 State-Count Selection: Two Reference Operating Points

We sweep $K \in \{3, 6, 9, 12, 15, 18, 21\}$ for CHMM-N on the SPY IS window and score each fit on (i) held-out log-likelihood on a clean pre-OoS validation slice (estimation 2014-01-03 to 2021-12-31, validation 2022-01-03 to 2024-01-03), (ii) the four information criteria of Nguyen [60] (Akaike information criterion (AIC), Bayesian information criterion (BIC), Hannan–Quinn criterion (HQC), consistent AIC (CAIC), $p = K^2 + 2K - 1$), and (iii) the IS and OoS two-sample Kolmogorov–Smirnov pass rates over 1,000 simulated paths.

The held-out criterion picks $K^* = 3$: validation log-likelihood per observation degrades monotonically from -2.5345 at $K = 3$ to -2.5847 at $K = 21$, a difference of 25.2 nats over the $T_{\text{val}} = 502$ -day validation slice. BIC on the estimation window agrees ($K^* = 3$). Distributional fidelity, however, rises across the sweep before plateauing: IS KS rises from 89.8% at $K = 3$ to 94.7% at $K = 18$. ACF-MAE is remarkably stable across the entire sweep (0.047 to 0.053), consistent with the spectral identity (8): the ACF is determined by the eigenvalue spectrum of $\mathbf{T}$, which the quantile-initialized Baum–Welch recovers at every $K \geq 3$.

We therefore report two reference operating points throughout this section. $K = 3$ is the held-out-likelihood and BIC choice, the methodologically conservative default. $K = 18$ is the explicit tail-fidelity operating point, chosen because it sits at the top of the IS KS plateau and inside the CAIC-minimizing band. The cost of moving from $K^* = 3$ to $K = 18$ is the 25.2-nat held-out-likelihood penalty above; the gain is documented in Tables 1 and 2. The complete per- K panel ($K \in \{3, 6, 9, 12, 15, 18, 21\}$) and the full held-out-validation sweep are reported in Appendix A.

4.3 Six-Generator Comparison (Pipeline A)

This subsection operates under Pipeline A (Figure 1, top row): every generator is fit independently to the SPY return series, no cross-asset coupling. Table 1 reports the IS and OoS validation results for the benchmark generators of Section 3.3 together with the three CHMM emission variants at the headline operating point $K = 18$ and a CHMM-N reference row at the held-out-likelihood

Table 1. Headline generator comparison on SPY. 1,000 simulated paths, $\alpha = 0.05$. IS: 2014-01-03 through 2024-01-03 ($T = 2,516$); OoS: 2024-01-04 through 2026-04-20 ($T = 572$). KS = Kolmogorov–Smirnov pass rate (%); Kurt = mean simulated excess kurtosis (observed: 7.68 IS, 5.29 OoS); ACF-MAE = mean absolute error of $\text{ACF}(|G_t|)$ over 252 lags; CRPS = sample continuous-ranked-probability-score on the OoS series, mean over $T_{\text{OoS}} = 572$ days, ensemble of 1,000 paths (lower is better). The CHMM-N $K = 3$ row is the held-out-likelihood reference; CHMM-N CRPS at $K = 3$ is omitted because the CRPS run uses the cached $K = 18$ archive. GARCH(1,1)-t is the Student- t -innovations extension of Bollerslev [18]; MS-GARCH ($K = 2$) is the Markov-switching GARCH foil of Haas et al. [19]; full extended-GARCH panel (EGARCH, GJR-GARCH, HAR-RV, MS-GARCH $K = 3$) and the semi-Markov CHMM foil are reported in Appendix A.

Model	KS (%) $\uparrow$		Kurtosis		ACF-MAE	CRPS OoS
	IS	OoS	IS	OoS	$\downarrow$	$\downarrow$
<i>Observed</i>	–	–	7.68	5.29	–	–
Bootstrap	99.9	91.0	7.54	6.80	0.0628	1.0398
Gaussian i.i.d.	0.0	1.2	−0.01	−0.02	0.0627	1.0611
Laplace i.i.d.	99.1	86.1	2.96	2.83	0.0627	1.0386
GARCH(1,1)	25.9	58.2	6.97	2.88	0.0484	1.0440
GARCH(1,1)-t	57.3	80.8	15.13	–	0.0316	–
MS-GARCH ($K = 2$)	27.7	38.7	4.73	–	0.0367	–
CHMM-N ($K^* = 3$)	89.8	80.1	3.86	3.56	0.0462	–
CHMM-N ($K = 18$)	93.8	84.2	4.99	4.51	0.0507	1.0384
CHMM-t ($K = 18$)	93.7	85.9	14.57	9.74	0.0548	1.0399
CHMM-L ($K = 18$)	93.2	79.3	6.75	6.27	0.0565	1.0400

choice $K^* = 3$. The headline panel covers one representative from each benchmark tier: a non-parametric distributional ceiling (i.i.d. bootstrap, 17), a Gaussian negative control, a Laplace i.i.d. comparator, the workhorse GARCH(1,1) volatility-clustering baseline of Bollerslev [7], the heavy-tailed GARCH(1,1)- t extension of Bollerslev [18], and the regime-switching MS-GARCH foil of Haas et al. [19]. The full extended GARCH-family panel (EGARCH, GJR-GARCH, HAR-RV, MS-GARCH $K = 3$), the semi-Markov CHMM foil, and the deep-generative QuantGAN baseline (9) are reported in Appendix A; the headline conclusions of this section are unchanged under that extended panel.

The headline finding is in the bottom four rows. At $K = 18$, CHMM-N attains 93.8% IS KS and 84.2% OoS KS, comparable to the bootstrap distributional ceiling and substantially above GARCH (25.9% IS KS); ACF-MAE of 0.0507 is essentially on par with GARCH (0.0484), the empirical instance of the spectral identity (8): with $K - 1 = 17$ non-unit eigenvalues at the operating point, the mixture-of-eigenvalues sum reproduces the slow absolute-return ACF that the Rydén et al. [11] $K = 2$ –3 Gaussian setup could not recover, inside a standard HMM at moderate state resolution. The bootstrap (at 0.0628) and the i.i.d. baselines all sit far above the temporal floor.

CHMM-N at the held-out-likelihood reference $K^* = 3$ remains a credible operating point: 89.8% IS KS, 80.1% OoS KS, ACF-MAE 0.0462. The cost of staying at $K^* = 3$ is a roughly 4 percentage point drop in IS KS and a doubling of the kurtosis residual (3.86 vs 4.99 at $K = 18$). The benefit is the 25.2-nat held-out-likelihood gain documented in Section 4.2. For risk-management consumers who require minimum model complexity the $K^* = 3$ row is the recommended default; for synthetic-data consumers who require the strongest joint match on KS and tail moments the $K = 18$ row is the operating point.

The three CHMM emission variants land at three distinct points on the kurtosis axis without retuning K or the transition structure: CHMM-N undershoots (5.0 IS vs. observed 7.7), CHMM-t overshoots (14.6, an artefact discussed in Section 5), and CHMM-L matches cleanly (6.8, within

0.9 units of observed). CHMM-L is the strongest entry on IS kurtosis in the panel and adds no shape parameter to CHMM-N (location is the weighted median, scale is the weighted mean absolute deviation, both in closed form). All three variants match each other on KS pass rates to within 1 percentage point IS and within 7 percentage points OoS, and all three are within 0.006 of each other on ACF-MAE: the emission-family swap moves tail weight cleanly without disturbing the central or temporal fit.

The OoS sample-CRPS column tightens this picture. CHMM-N at $K = 18$ attains the lowest mean OoS CRPS in the panel (1.0384); CHMM-t, CHMM-L, and the Laplace i.i.d. baseline are within 0.002 and statistically indistinguishable from CHMM-N under a Diebold-Mariano (DM) test on the loss differential (Newey-West heteroskedasticity- and autocorrelation-consistent (HAC) variance, two-sided $p > 0.45$ in each pair). CHMM-N is significantly better than the Gaussian i.i.d. baseline ($p \approx 0.002$). On a proper scoring rule the CHMM family ties or beats every benchmark; the joint dominance on KS, ACF-MAE, and CRPS rules out the concern that the headline KS finding is a non-proper-metric artefact. Per-pair DM verdicts and the methods note are in Appendix A.

No single benchmark dominates every metric: the i.i.d. bootstrap is the strongest distributional baseline but the worst on temporal fidelity (ACF-MAE 0.0628); GARCH(1,1) and the GARCH(1,1)- t extension of Bollerslev [18] are the strongest unconditional volatility-clustering baselines (ACF-MAE 0.0316–0.0484) but fail distributional KS decisively, with GARCH- t recovering substantial OoS marginal fidelity (80.8%) at the cost of an explicit shape parameter; the Markov-switching GARCH foil of Haas et al. [19] sits between standard GARCH and the CHMM family on KS but does not match the CHMM on the joint corner. The CHMM at moderate K occupies a distinctive corner of the metric panel where both KS and ACF-MAE simultaneously sit close to their respective best-in-panel values, and the choice between CHMM-N, CHMM-t, and CHMM-L is the choice of where to land the kurtosis residual.

The headline KS pass rates are stable under multi-seed and KS-power-calibration controls. A ten-seed Monte Carlo at $K = 18$ ($N_{\text{paths}} = 250$ per seed) gives CHMM-N OoS KS $82.6 \pm 1.79\%$, CHMM-t $85.6 \pm 1.92\%$, and CHMM-L $85.0 \pm 2.4\%$ (Appendix A.20); the seed-to-seed standard deviation is small relative to the cross-generator gaps in Table 1. A KS power-calibration positive control (i.i.d. resamples of the IS series at length T_{OoS}) passes at 90% asymptotic and 98–99% block-bootstrap (Appendix A.21, A.22); the CHMM rows under a stricter block-bootstrap critical value retain 81–92% pass rates, consistent with the asymptotic numbers above.

4.4 Cross-Ticker Generalization (Pipeline A)

Still under Pipeline A: each ticker is fit independently, no cross-asset coupling. Table 2 reports CHMM-t IS and OoS KS pass rates across the five additional tickers (NVDA, JNJ, JPM, AAPL, QQQ), each fit independently at $K = 18$. CHMM-t is the natural generalization vehicle here because it carries the strongest OoS KS in the SPY headline panel (85.9%, against 84.2% for CHMM-N and 79.3% for CHMM-L) and supplies the per-state heavy tails that the outlier-prone single-name tickers (NVDA, JNJ) require. Per-family K -sensitivity for SPY across all three emission variants and per-ticker price-level metrics across all three families are in Appendix A. The univariate fidelity persists across the panel on IS KS (95.2% to 99.8%). On OoS, JNJ, AAPL, and QQQ generalise cleanly (90–95%), while NVDA (57.2%) and JPM (53.0%) show substantial OoS degradation, a stationarity artefact: the IS-fixed transitions and emissions cannot capture the artificial-intelligence (AI) driven NVDA regime or the rate-repricing cycle for financials in the 2024–2026 holdout. We acknowledge the JPM and NVDA OoS cliffs as stationarity-scope limitations rather than model-selection issues; the pattern persists at every K in the sweep, and we discuss the operational implication (periodic refit at deployment) in Section 5.




Figure 3. IS CHMM-N validation on SPY ($K = 18$, 1,000 simulated paths). Panel (a): marginal density of the annualized excess log return G_t ; gray histogram = observed IS series, blue curve = one simulated path. Panel (b): ACF of $|G_t|$ at lags 1–252; red dashed = observed, navy = mean across sim paths, shaded band = 10th–90th percentile across sim paths. Panel (c): tail Q-Q plot at 200 quantile levels; blue points = simulated pooled over paths, black dashed = identity.

Table 2. Cross-ticker generalization, CHMM-t at $K = 18$. Each ticker fit independently; 1,000 simulated paths, $\alpha = 0.05$. Observed kurtosis is the IS sample value.

Ticker	KS (%) $\uparrow$		Kurtosis		ACF-MAE $\downarrow$
	IS	OoS	Obs	Sim	
SPY	95.2	84.1	7.7	14.68	0.0551
NVDA	99.1	57.2	5.5	46.40	0.0478
JNJ	99.8	94.3	9.0	99.66	0.0332
JPM	99.2	53.0	7.6	18.82	0.0483
AAPL	99.1	94.7	3.2	7.54	0.0424
QQQ	96.1	90.3	3.4	3.30	0.0576




Figure 4. OoS CHMM-N validation on SPY ($K = 18$, 1,000 paths, $T_{\text{OoS}} = 572$). Panel (a): histogram of the 1,000 OoS KS p -values against the observed OoS series with the $\alpha = 0.05$ threshold (red dashed). Panel (b): OoS return-density fan; blue thin curves = simulated densities, red thick curve = observed OoS density. Panel (c): ACF of $|G_t|$ on the OoS window; red dashed = observed, navy = mean across sim paths, shaded band = 10th–90th percentile.

Table 3. Cross-asset Student-t copula on CHMM-N marginals (Pipeline B; $\nu^* = 6$ by profile MLE; 200 paths). Per-asset KS pass rate at $\alpha = 0.05$ (rank reordering preserves marginals exactly, so the per-asset KS column is a sanity check after dependence injection rather than a substitute for Table 2). The lower block reports correlation reproduction averaged over paths: Frobenius norm of $\hat{\Sigma}_{\text{sim}} - \hat{\Sigma}_{\text{obs}}$ and off-diagonal mean absolute error. Comparison against SIM and the Gaussian copula on the same marginals is in Appendix A; the Student-t copula is the best of the three on every metric (off-diagonal MAE 0.027 vs. 0.031 Gaussian, 0.076 SIM).

Ticker	KS (%) $\uparrow$	
	IS	OoS
SPY	95.5	86.0
NVDA	93.0	50.5
JNJ	100.0	93.5
JPM	96.5	49.0
AAPL	99.0	94.5
QQQ	93.0	95.0
Mean	96.2	78.1
Median	96.0	89.8
<i>Correlation reproduction (averaged over 200 paths)</i>		
Frobenius IS	0.184	
Off-diag MAE IS	0.027	
Frobenius OoS	1.499	
Off-diag MAE OoS	0.208	

4.5 Cross-Asset Dependence (Pipeline B): Student-t Copula

This subsection is the only Pipeline B subsection in the main text (Figure 1, bottom row): we reuse the per-asset CHMM-N marginals from the Pipeline-A fits of Section 4.4 and add a cross-asset dependence layer on top. For multi-asset synthesis we inject cross-sectional dependence through the rank-based Student-t copula of Section 3.4, the strongest dependence layer in our larger panel; the SIM and Gaussian copula constructions are reported in Appendix A. The per-asset marginals here are the CHMM-N fits at $K = 18$; rank reordering preserves whichever marginal it receives exactly (Proposition 4), so the dependence-layer ranking is invariant to the choice of CHMM emission family, and the Pipeline-B numbers below transfer unchanged to CHMM-t or CHMM-L marginals up to the per-asset KS sanity-check column. Profile maximum likelihood over $\nu \in \{2, 3, 4, 5, 6, 8, 10, 15, 20, 30\}$ on the six-asset universe (SPY, NVDA, JNJ, JPM, AAPL, QQQ; $T = 2,516$ IS, 200 simulated paths) selects $\nu^* = 6$, consistent with values typically reported for equity indices [37, 38].

The Student-t copula preserves each asset’s fitted CHMM marginal exactly by the rank-reordering construction, so the per-asset KS column (median 96.0% IS, 89.8% OoS) is a sanity check on the copula sample rather than a substitute for Table 2. The central dependence-layer claim is the off-diagonal correlation MAE: 0.027 on IS, against 0.031 for the Gaussian copula and 0.076 for the SIM factor baseline on the same marginals (Appendix A). The copula advantage is intrinsic to the rank-reordering construction rather than specific to the marginal model.

The OoS off-diagonal MAE of 0.208 is roughly an order of magnitude larger than the 0.027 IS value, a degradation we attribute to the same stationarity-scope mechanism as the NVDA / JPM single-name OoS cliffs in Table 2: the IS-fixed copula has no mechanism to track regime shifts in the cross-asset correlation structure that lie outside the IS support. Two appendix diagnostics localise the cause. First, extending the universe with GLD (a non-equity commodity asset class) reduces

Table 4. VaR and Expected-Shortfall envelope for SPY ($N_{\text{paths}} = 1,000$). Each entry is the median and [5%, 95%] envelope across paths of the left-tail VaR / ES at $\alpha \in \{0.01, 0.05\}$; observed historical values are in the first row. All three CHMM variants bracket the observed historical VaR and ES at both levels on both windows.

Model	IS VaR _{0.01} median [5%, 95%]	IS ES _{0.01} median [5%, 95%]	IS VaR _{0.05} median [5%, 95%]	IS ES _{0.05} median [5%, 95%]
<i>Observed</i>	−6.38	−9.00	−3.43	−5.50
Bootstrap	−6.38 [−7.04, −5.99]	−8.86 [−10.37, −7.78]	−3.43 [−3.68, −3.20]	−5.46 [−5.99, −5.05]
GARCH	−5.80 [−8.62, −4.52]	−7.86 [−13.54, −5.77]	−3.19 [−3.92, −2.71]	−4.89 [−7.11, −3.91]
CHMM-N	−6.64 [−8.01, −5.47]	−8.80 [−10.16, −7.34]	−3.45 [−4.05, −2.95]	−5.45 [−6.33, −4.61]
CHMM-t	−6.58 [−7.51, −5.83]	−9.06 [−10.70, −7.79]	−3.40 [−3.90, −2.90]	−5.51 [−6.18, −4.85]
CHMM-L	−6.89 [−7.85, −6.01]	−9.15 [−10.42, −7.97]	−3.30 [−3.77, −2.88]	−5.53 [−6.17, −4.87]
	OoS VaR _{0.01}	OoS ES _{0.01}	OoS VaR _{0.05}	OoS ES _{0.05}
<i>Observed</i>	−5.51	−7.94	−2.77	−4.67
Bootstrap	−6.33 [−7.60, −5.28]	−8.43 [−11.76, −6.59]	−3.39 [−3.91, −2.88]	−5.34 [−6.38, −4.50]
GARCH	−5.35 [−9.98, −3.61]	−6.72 [−13.51, −4.38]	−3.14 [−4.94, −2.36]	−4.59 [−7.92, −3.23]
CHMM-N	−6.30 [−9.02, −4.30]	−8.44 [−11.36, −5.42]	−3.39 [−4.74, −2.47]	−5.27 [−7.23, −3.74]
CHMM-t	−6.32 [−8.13, −4.80]	−8.60 [−12.20, −6.29]	−3.29 [−4.27, −2.45]	−5.34 [−6.72, −4.08]
CHMM-L	−6.69 [−8.59, −4.91]	−8.88 [−11.48, −6.73]	−3.27 [−4.30, −2.51]	−5.47 [−6.82, −4.20]

the OoS off-diag MAE to 0.179 (Appendix A.32), so the OoS gap is not driven by equity-cluster size. Second, a per-pair breakdown (Table 20) shows the gap concentrates in three pairs involving JNJ (JNJ-QQQ, SPY-JNJ, NVDA-JNJ), each with $|\Delta| > 0.48$ on OoS, while gold-equity pairs sit at $|\Delta| \in [0.10, 0.19]$. The dominant cause is a single-name regime shift in the JNJ-equity-factor relationship rather than copula-degree mis-specification.

4.6 Value-at-Risk and Expected-Shortfall Back-Test (Pipeline A)

Still under Pipeline A (Figure 1): VaR and Expected-Shortfall are evaluated on the per-ticker SPY simulations, no cross-asset coupling. For a financial-risk consumer the canonical diagnostics are one-tail Value-at-Risk (VaR) and Expected-Shortfall (ES) at daily probability levels of 1% and 5%. A synthetic-data generator is well calibrated if the envelope of per-path simulated VaR/ES brackets the observed historical value. We simulate 1,000 independent paths from each of the five most relevant generators (Bootstrap, GARCH(1,1), CHMM-N, CHMM-t, CHMM-L) and report median and [5%, 95%] envelope across paths in Table 4.

All three CHMM variants bracket the observed historical VaR and ES at both 1% and 5% levels within their 5%–95% envelopes on both IS and OoS windows; the IS CHMM-N median VaR_{0.01} at −6.64 (observed −6.38) and ES_{0.01} at −8.80 (observed −9.00) are the closest point estimates across the five generators. The CHMM envelopes are roughly $1.7\times$ tighter than GARCH’s and pass the unconditional Kupiec test cleanly at both levels (CHMM-N $LR_{uc} = 3.83$ at 5% against the $\chi^2_1 = 3.84$ critical value; CHMM-t at 3.83; CHMM-L at 3.03). CHMM-t’s IS kurtosis overshoot does *not* translate into a systematically more conservative VaR/ES point estimate; the IS median VaR_{0.01} of −6.58 is almost identical to CHMM-N’s −6.64 and the envelope width is comparable. CHMM-L’s IS median VaR_{0.01} of −6.89 is mildly more conservative than CHMM-N’s, consistent with its closer-to-observed kurtosis.

5 Discussion

A continuous HMM trained by Baum-Welch at moderate state resolution ($K = 18$) reproduces the three Cont stylized facts simultaneously on SPY (93.8% IS KS, 84.2% OoS KS, ACF-MAE within 0.003 of GARCH). Three emission variants share the same scaffold and land at three distinct points on the kurtosis axis without retuning K or the transition structure: CHMM-N undershoots, CHMM-L matches observed kurtosis cleanly without a shape parameter, and CHMM-t overshoots while supplying per-state heavy tails for tail-conditional consumers. The mechanism behind the simultaneous KS / ACF fit is the spectral identity of equation (8): at $K \geq 3$ with quantile-initialised Baum-Welch, the absolute-return ACF is a population-level mixture of geometric decays at the non-unit eigenvalues of $\mathbf{T}$.

Replicating the slow-ACF behaviour at moderate state resolution. Two issues compound in the Rydén et al. [11] low- K setup: a single self-transition rate produces a single exponential mode that cannot be simultaneously fast (good marginal) and slow (good ACF), and the random initialisations standard in the 1990s often converge to degenerate local optima with overlapping states. Our $K = 2$ replication on SPY (Appendix A) clarifies the mechanism: ACF-MAE is essentially constant at ≈ 0.050 across quantile and five random-seed initialisations, matching the $K = 18$ CHMM-N. The binding low- K constraint is therefore distributional rather than temporal (KS at $K = 2$ drops to 67–72%): with modern initialisation the ACF is reproduced at any $K \geq 2$, while distributional fidelity is what actually requires moderate state resolution.

OoS KS pass rates against the test-power ceiling. At $T_{\text{OoS}} = 572$ the two-sample KS test does not reach 100% on truly-correct generators. A length-matched i.i.d. resample of the IS distribution itself passes at 90.0%, while the Gaussian i.i.d. negative control is rejected at 0%, so the test retains ample power to refuse clearly-wrong generators. The CHMM OoS KS pass rates of 79–86% sit just below the 90% ceiling (a genuine fidelity claim, not a low-power artefact); GARCH’s 58% is correctly rejected.

Closing the kurtosis gap with heavy-tailed emissions. CHMM-N produces excess kurtosis of 5.02 against an observed 7.68, a gap inherent to Gaussian emissions: a Gaussian mixture’s kurtosis is bounded by the variance of component means relative to their variances, and pushing tail states out degrades the central fit. CHMM-t addresses this directly through per-state Student-t emissions whose ν_k is updated by ECM one-dimensional search, allowing tail states to carry arbitrarily high individual kurtosis at the cost of an IS overshoot to 14.57. The unpenalised overshoot is a single-state artefact: a bracket sweep $\nu_{\min} \in \{2.1, 2.5, 3.0, 4.0\}$ shows the fit does not pile against the lower bracket (median $\nu_k = 50$, only 2/18 states near the lower edge), and an exponential $1/\nu_k$ shrinkage prior at rate $\lambda = 20$ brings simulated kurtosis to ≈ 8 with a 1pp KS cost (full rate sweep in Appendix A.34). CHMM-L sits between CHMM-N and CHMM-t on the kurtosis axis without any shape parameter at all: per-component excess kurtosis of 3 (versus 0 Gaussian and ν -dependent Student-t) brings simulated IS kurtosis to 6.75, the closest match in the panel to observed 7.68. The three-way decision is concrete and is summarised in Table 5; all three variants’ 5–95% envelopes bracket observed VaR/ES on both windows.

Stationarity scope and operational deployment. The CHMM is trained on the assumption that returns are locally stationary across the 10-year IS window. Single-name equity returns are well known to carry sector-specific structural breaks (20; 21; surveyed by 22), and Cross-ticker

Table 5. Variant decision guide. Recommended CHMM emission family by use-case priority. All three variants share the forward-backward scaffold and the quantile-based initialisation; the M-step is the only architectural difference.

Use-case priority	Recommended variant
Simplest closed-form fit; kurtosis fidelity secondary	CHMM-N
Cleanest IS kurtosis match; no shape parameter	CHMM-L
Per-state heavy tails for tail-conditional consumers	CHMM-t

generalization (Section 4.4) reflects this directly: at $K = 18$, OoS distributional fidelity is preserved on SPY, JNJ, AAPL, and QQQ but degrades on NVDA (57.2%) and JPM (53.0%), where the IS-fixed transitions and emissions cannot capture the AI-driven NVDA regime or the rate-repricing cycle for financials. The pattern persists at every K in the sweep, so it is not a model-selection issue. Operational deployment of the CHMM should follow standard practice for time-series generators on financial data, with periodic refit at quarterly or finer cadence (23; an online-EM formulation is given by 61). The static IS / OoS panel reported here is the reference comparison, not the operational protocol.

Why the Student-t copula beats SIM on cross-asset fidelity. The Pipeline B Student-t copula on top of the CHMM-N marginals attains an off-diagonal correlation MAE of 0.027 on the six-asset universe, against 0.076 for a Single Index Model factor baseline on the same marginals (Section 4.5; full SIM and Gaussian copula panels in Appendix A). Two structural reasons. First, rank reordering preserves marginals by construction: the simulated path is a permutation of the stand-alone CHMM path, so its empirical CDF coincides with the CHMM marginal exactly, while SIM’s affine combination $\alpha_j + \beta_j G_{m,t} + \eta_{j,t}$ convolves the market factor with idiosyncratic residuals and distorts the marginal whenever $\beta_j \neq 1$ or residuals are heavy. Second, the Student-t copula’s $\nu^* = 6$ implies non-trivial joint tail dependence that injects rare concordant extreme moves the Gaussian copula systematically understates. The copula advantage is intrinsic to the rank-reordering construction rather than specific to the marginal family.

Limitations. The cross-ticker study covers six tickers; scaling to a larger universe via vine or factor copulas is a natural next step. State selection is via information-criterion plus distributional pass-rate rather than held-out log-likelihood; a clean re-selection on a non-OoS validation slice picks $K^* = 3$ by held-out log-lik (Appendix A), so the $K = 18$ operating point is best read as a multi-metric choice that weights tail fidelity alongside likelihood. Skew-heavy-tailed emissions (skew- t or skew-Laplace), explicit-duration semi-Markov sojourns, and operational regime-conditional VaR are deferred as companion-paper directions.

6 Conclusion

A continuous hidden Markov model trained by Baum-Welch with quantile-based initialization reproduces the three Cont stylized facts of daily equity returns simultaneously at moderate state resolution; the closed-form mixture-of-eigenvalues identity for the absolute-return ACF recasts the Rydén et al. low- K failure as a K -rank statement on $\mathbf{T}$. Three emission variants share the same scaffold and land at three distinct points on the kurtosis axis: the Gaussian CHMM-N undershoots, the Student-t CHMM-t overshoots and is the per-state heavy-tailed choice for tail-conditional

consumers, and the Laplace CHMM-L matches observed kurtosis cleanly without a shape parameter. The M-step swap is the only architectural difference across the three.

The headline quantitative findings on SPY are 93.8% in-sample and 84.2% out-of-sample Kolmogorov–Smirnov pass rates, an absolute-return ACF mean absolute error within 0.003 of GARCH(1,1), a clean Kupiec pass on Value-at-Risk at both 1% and 5%, and Value-at-Risk and Expected-Shortfall envelopes that bracket the observed historical values on both windows. Substantial OoS degradation on NVDA and JPM reflects the well-known stationarity scope limit on individual financial time series and motivates periodic refit in operational deployment, in line with standard practice for time-series generators in this domain.

Multi-asset synthesis composes cleanly with the rank-based Student-t copula on top of the CHMM-N marginals: profile MLE selects $\nu^* = 6$, and the resulting cross-asset construction reproduces the off-diagonal correlation matrix on six SPY-member assets with a mean absolute error of 0.027, against 0.076 for a Single Index Model factor baseline on the same marginals.

Skew-heavy-tailed emissions, explicit-duration semi-Markov sojourns, and operational regime-conditional VaR are companion-paper directions; scaling the cross-asset construction to a larger universe via vine or factor copulas is the natural multi-asset follow-up.

Data and code availability. The simulation scripts and data to reproduce the paper are available under an MIT license at <https://github.com/altashly1/CHMM-paper>. The companion Julia package `CHMM-Model.jl` is released alongside this paper at <https://github.com/altashly1/CHMM-Model>; the public API listed in Appendix A regenerates every table and figure from a single seed root under Julia ≥ 1.12 .

Conflict of interest. The authors declare no competing interests.

Author contributions. J.V. directed the study. A.A. developed the continuous model, implemented the Baum-Welch algorithm, conducted the empirical analysis, and generated all figures and tables. C.J. contributed to methodology review and validation. All authors edited and reviewed the final manuscript.

References

- [1] James Jordon, Lukasz Szpruch, Florimond Houssiau, Mirko Bottarelli, Giovanni Cherubin, Carsten Maple, Samuel N. Cohen, and Adrian Weller. Synthetic data: what, why and how? *arXiv preprint arXiv:2205.03257*, 2022.
- [2] Samuel A. Assefa, Danial Dervovic, Mahmoud Mahfouz, Robert E. Tillman, Prashant Reddy, and Manuela Veloso. Generating synthetic data in finance: opportunities, challenges and pitfalls. In *Proceedings of the First ACM International Conference on AI in Finance*, pages 1–8, 2020.
- [3] Abdulrahman Alswaidan and Jeffrey D. Varner. Hybrid hidden markov model for modeling equity excess growth rate dynamics: A discrete-state approach with jump-diffusion. *arXiv preprint arXiv:2603.10202*, 2026.
- [4] Benoit Mandelbrot. The variation of certain speculative prices. *The Journal of Business*, 36(4): 394–419, 1963.
- [5] R. Cont. Empirical properties of asset returns: stylized facts and statistical issues. *Quantitative Finance*, 1(2):223–236, 2001.

- [6] Robert F. Engle. Autoregressive conditional heteroscedasticity with estimates of the variance of united kingdom inflation. *Econometrica*, 50(4):987–1007, 1982.
- [7] Tim Bollerslev. Generalized autoregressive conditional heteroscedasticity. *Journal of Econometrics*, 31(3):307–327, 1986.
- [8] Jinsung Yoon, Daniel Jarrett, and Mihaela van der Schaar. Time-series generative adversarial networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 32, 2019.
- [9] Magnus Wiese, Robert Knobloch, Ralf Korn, and Peter Kretschmer. Quant GANs: Deep generation of financial time series. *Quantitative Finance*, 20(9):1419–1440, 2020.
- [10] Kashif Rasul, Calvin Seward, Ingmar Schuster, and Roland Vollgraf. Autoregressive denoising diffusion models for multivariate probabilistic time series forecasting. In *Proceedings of the 38th International Conference on Machine Learning*, pages 8857–8868, 2021.
- [11] Tobias Rydén, Timo Teräsvirta, and Stefan Åsbrink. Stylized facts of daily return series and the hidden Markov model. *Journal of Applied Econometrics*, 13(3):217–244, 1998.
- [12] Jan Bulla and Ingo Bulla. Stylized facts of financial time series and hidden semi-Markov models. *Computational Statistics & Data Analysis*, 51(4):2192–2209, 2006.
- [13] Leonard E. Baum, Ted Petrie, George Soules, and Norman Weiss. A maximization technique occurring in the statistical analysis of probabilistic functions of markov chains. *The Annals of Mathematical Statistics*, 41(1):164–171, 1970.
- [14] L. Rabiner and B. Juang. An introduction to hidden Markov models. *IEEE ASSP Magazine*, 3(1):4–16, 1986.
- [15] David Peel and Geoffrey J. McLachlan. Robust mixture modelling using the t distribution. *Statistics and Computing*, 10(4):339–348, 2000.
- [16] Chuanhai Liu and Donald B. Rubin. ML estimation of the t distribution using EM and its extensions, ECM and ECME. *Statistica Sinica*, 5(1):19–39, 1995.
- [17] Dimitris N. Politis and Joseph P. Romano. The stationary bootstrap. *Journal of the American Statistical Association*, 89(428):1303–1313, 1994.
- [18] Tim Bollerslev. A conditionally heteroskedastic time series model for speculative prices and rates of return. *Review of Economics and Statistics*, 69(3):542–547, 1987.
- [19] Markus Haas, Stefan Mittnik, and Marc S. Paoletta. A new approach to Markov-switching GARCH models. *Journal of Financial Econometrics*, 2(4):493–530, 2004.
- [20] Ľuboš Pástor and Robert F. Stambaugh. The equity premium and structural breaks. *Journal of Finance*, 56(4):1207–1239, 2001.
- [21] Elena Andreou and Eric Ghysels. Detecting multiple breaks in financial market volatility dynamics. *Journal of Applied Econometrics*, 17(5):579–600, 2002.
- [22] Andrew Ang and Allan Timmermann. Regime changes and financial markets. *Annual Review of Financial Economics*, 4(1):313–337, 2012.

- [23] M. Hashem Pesaran and Allan Timmermann. Selection of estimation window in the presence of breaks. *Journal of Econometrics*, 137(1):134–161, 2007.
- [24] James D. Hamilton. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2):357–384, 1989.
- [25] Huntley Schaller and Simon Van Norden. Regime switching in stock market returns. *Applied Financial Economics*, 7(2):177–191, 1997.
- [26] Peter Nystrup, Henrik Madsen, and Erik Lindström. Long memory of financial time series and hidden markov models with time-varying parameters. *Journal of Forecasting*, 36(8):989–1002, 2017.
- [27] Daniel B. Nelson. Conditional heteroskedasticity in asset returns: A new approach. *Econometrica*, 59(2):347–370, 1991.
- [28] Lawrence R. Glosten, Ravi Jagannathan, and David E. Runkle. On the relation between the expected value and the volatility of the nominal excess return on stocks. *The Journal of Finance*, 48(5):1779–1801, 1993.
- [29] Jean-Michel Zakoïan. Threshold heteroskedastic models. *Journal of Economic Dynamics and Control*, 18(5):931–955, 1994.
- [30] Torben G. Andersen and Tim Bollerslev. Heterogeneous information arrivals and return volatility dynamics: Uncovering the long-run in high frequency returns. *The Journal of Finance*, 52(3):975–1005, 1997.
- [31] Fulvio Corsi. A simple approximate long-memory model of realized volatility. *Journal of Financial Econometrics*, 7(2):174–196, 2009.
- [32] Shuntaro Takahashi, Yu Chen, and Kumiko Tanaka-Ishii. Modeling financial time-series with generative adversarial networks. *Physica A: Statistical Mechanics and its Applications*, 527:121261, 2019.
- [33] Sohyeon Kwon and Yongjae Lee. Can GANs learn the stylized facts of financial time series? In *Proceedings of the 5th ACM International Conference on AI in Finance*, pages 126–133, 2024.
- [34] William F. Sharpe. A simplified model for portfolio analysis. *Management Science*, 9(2):277–293, 1963.
- [35] A. Sklar. Fonctions de répartition à n dimensions et leurs marges. *Publications de l’Institut de Statistique de l’Université de Paris*, 8:229–231, 1959.
- [36] Roger B. Nelsen. *An Introduction to Copulas*. Springer, 2 edition, 2006.
- [37] Stefano Demarta and Alexander J. McNeil. The t copula and related copulas. *International Statistical Review*, 73(1):111–129, 2005.
- [38] Alexander J. McNeil, Rüdiger Frey, and Paul Embrechts. *Quantitative Risk Management: Concepts, Techniques and Tools*. Princeton University Press, revised edition, 2015.
- [39] Ronald L. Iman and W. J. Conover. A distribution-free approach to inducing rank correlation among input variables. *Communications in Statistics–Simulation and Computation*, 11(3):311–334, 1982.

- [40] Michael Stenger, André Bauer, Thomas Prantl, Robert Leppich, Nathaniel Hudson, Kyle Chard, Ian Foster, and Samuel Kounev. Thinking in categories: A survey on assessing the quality for time series synthesis. *Journal of Data and Information Quality*, 16(2):1–32, 2024.
- [41] Tilmann Gneiting and Adrian E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007.
- [42] Tilmann Gneiting and Matthias Katzfuss. Probabilistic forecasting. *Annual Review of Statistics and Its Application*, 1:125–151, 2014.
- [43] Francis X. Diebold and Roberto S. Mariano. Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3):253–263, 1995.
- [44] Arthur Gretton, Karsten M. Borgwardt, Malte J. Rasch, Bernhard Schölkopf, and Alexander Smola. A kernel two-sample test. *Journal of Machine Learning Research*, 13(1):723–773, 2012.
- [45] Ilya Chevyrev and Andrey Kormilitzin. A primer on the signature method in machine learning. *arXiv preprint arXiv:1603.03788*, 2016.
- [46] Hao Ni, Lukasz Szpruch, Magnus Wiese, Shujian Liao, and Baoren Xiao. Conditional Sig-Wasserstein GANs for time series generation. *arXiv preprint arXiv:2006.05421*, 2020.
- [47] Paul Embrechts, Alexander McNeil, and Daniel Straumann. Correlation and dependence in risk management: Properties and pitfalls. *Risk Management: Value at Risk and Beyond*, pages 176–223, 2002.
- [48] A. N. Kolmogorov. Sulla determinazione empirica di una legge di distribuzione. *Giornale dell’Istituto Italiano degli Attuari*, 4:83–91, 1933.
- [49] N. Smirnov. Table for estimating the goodness of fit of empirical distributions. *The Annals of Mathematical Statistics*, 19(2):279–281, 1948.
- [50] Paul H. Kupiec. Techniques for verifying the accuracy of risk measurement models. *Journal of Derivatives*, 3(2):73–84, 1995.
- [51] Jeff A. Bilmes. A gentle tutorial of the em algorithm and its application to parameter estimation for gaussian mixture and hidden markov models. *Technical Report TR-97-021, International Computer Science Institute*, 1998.
- [52] Carlos A. Abanto-Valle, Roland Langrock, Ming-Hui Chen, and Marcos V. Cardoso. Maximum likelihood estimation for stochastic volatility in mean models with heavy-tailed distributions. *Applied Stochastic Models in Business and Industry*, 33(4):394–408, 2017.
- [53] David A. Levin and Yuval Peres. *Markov Chains and Mixing Times*. American Mathematical Society, 2 edition, 2017.
- [54] C. F. Jeff Wu. On the convergence properties of the EM algorithm. *The Annals of Statistics*, 11(1):95–103, 1983.
- [55] Xiao-Li Meng and Donald B. Rubin. Maximum likelihood estimation via the ECM algorithm: A general framework. *Biometrika*, 80(2):267–278, 1993.

- [56] Elizabeth S. Allman, Catherine Matias, and John A. Rhodes. Identifiability of parameters in latent structure models with many observed variables. *The Annals of Statistics*, 37(6A):3099–3132, 2009.
- [57] Sidney J. Yakowitz and John D. Spragins. On the identifiability of finite mixtures. *The Annals of Mathematical Statistics*, 39(1):209–214, 1968.
- [58] Peter J. Bickel, Ya’acov Ritov, and Tobias Rydén. Asymptotic normality of the maximum-likelihood estimator for general hidden Markov models. *The Annals of Statistics*, 26(4):1614–1635, 1998.
- [59] Randal Douc, Eric Moulines, and Tobias Rydén. Asymptotic properties of the maximum likelihood estimator in autoregressive models with Markov regime. *The Annals of Statistics*, 32(5):2254–2304, 2004.
- [60] Nguyet Nguyen. Hidden Markov model for stock trading. *International Journal of Financial Studies*, 6(2):36, 2018.
- [61] Olivier Cappé. Online EM algorithm for hidden Markov models. *Journal of Computational and Graphical Statistics*, 20(3):728–749, 2011.
- [62] Paul Glasserman. *Monte Carlo Methods in Financial Engineering*. Springer, 2003.
- [63] Martin Arjovsky, Soumith Chintala, and Léon Bottou. Wasserstein generative adversarial networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 214–223, 2017.
- [64] Mike Innes. Flux: Elegant machine learning with Julia. *Journal of Open Source Software*, 3(25):602, 2018.

A Supplementary Material

This appendix collects supplementary material referenced in the main body: the public API of the companion package, the EM forward-backward recursions and algorithm pseudocode, the validation metric definitions, formal propositions, the full state-resolution and multi-emission sensitivity panels, robustness and KS-power calibration, the extended GARCH and SM-CHMM baselines, the full cross-asset panel, and the K -selection and ν_k diagnostics referenced from the discussion.

A.1 CHMM-Model.jl Public API and Reproducer

The companion Julia package `CHMM-Model.jl` (<https://github.com/altashly1/CHMM-Model>) is the entry point used to regenerate every table and figure in this paper from a single seed root under Julia ≥ 1.12 . The public surface is a single `CHMM` type, three emission constructors, a unified `fit!` method (log-space Baum-Welch / ECM), and a unified `simulate` method.

```
using CHMM
chmm_n = make_chmm_n(returns; K = 18)           # Gaussian emissions
chmm_t = make_chmm_t(returns; K = 18)           # Student-t emissions
chmm_l = make_chmm_l(returns; K = 18)           # Laplace emissions
fit!(chmm_n; max_iter = 60, tol = 1e-4)         # log-space Baum-Welch
paths = simulate(chmm_n; T = 2516, n = 1000)    # 1000 paths of length T
```

The same `simulate` interface is used for every benchmark generator (Bootstrap, Gaussian i.i.d., Laplace i.i.d., GARCH(1,1)) so that the seven-metric panel evaluates each row on byte-identical windows. The Pipeline B Student-t copula sampler is exposed as `simulate_copula(...; nu = 6)` on the per-asset CHMM marginals; the rank-reordering construction is one function call. Running `julia -project=. run_full_rebuild.jl` from the paper repository (<https://github.com/altashly1/CHMM-paper>) reproduces every numerical artifact under the global seed policy of Section 3.5; `Manifest.toml` pins the dependency graph.

A.2 CHMM Algorithms and Derivations

A.3 Forward-Backward Recursions and M-Step Updates

For reference, the in-text Baum-Welch description of Section 3.2 is underpinned by the standard log-space recursions [14, 51]. Define the forward variable $\alpha_t(k) = \mathbb{P}(O_{1:t}, S_t = k \mid \mathcal{M})$ and the backward variable $\beta_t(k) = \mathbb{P}(O_{t+1:T} \mid S_t = k, \mathcal{M})$. The log-space recursions are

$$\log \alpha_1(k) = \log \pi_k + \log b_k(O_1), \quad (9)$$

$$\log \alpha_t(j) = \log \sum_i \exp(\log \alpha_{t-1}(i) + \log T_{ij}) + \log b_j(O_t), \quad t = 2, \dots, T, \quad (10)$$

$$\log \beta_T(k) = 0, \quad (11)$$

$$\log \beta_t(i) = \log \sum_j \exp(\log T_{ij} + \log b_j(O_{t+1}) + \log \beta_{t+1}(j)), \quad t = T-1, \dots, 1, \quad (12)$$

where $\log \text{sumexp}_i(x_i) = x_{\max} + \log \sum_i \exp(x_i - x_{\max})$. The posterior state-occupancy $\gamma_t(k)$ and pair-occupancy $\xi_t(i, j)$ quantities are

$$\gamma_t(k) = \frac{\alpha_t(k) \beta_t(k)}{\sum_j \alpha_t(j) \beta_t(j)}, \quad \xi_t(i, j) = \frac{\alpha_t(i) T_{ij} b_j(O_{t+1}) \beta_{t+1}(j)}{\sum_{m,n} \alpha_t(m) T_{mn} b_n(O_{t+1}) \beta_{t+1}(n)}, \quad (13)$$

and the closed-form M-step updates are

$$\pi_k^{\text{new}} = \gamma_1(k), \quad \mu_k^{\text{new}} = \frac{\sum_t \gamma_t(k) O_t}{\sum_t \gamma_t(k)}, \quad (\sigma_k^{\text{new}})^2 = \frac{\sum_t \gamma_t(k) (O_t - \mu_k^{\text{new}})^2}{\sum_t \gamma_t(k)}, \quad T_{ij}^{\text{new}} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \gamma_t(i)}. \quad (14)$$

The convergence criterion is $|\mathcal{L}^{(n)} - \mathcal{L}^{(n-1)}| < 10^{-4}$ with $\mathcal{L}^{(n)} = \log \text{sumexp}_k \log \alpha_T^{(n)}(k)$, and simulation draws $S_1 \sim \text{Categorical}(\boldsymbol{\pi})$ and then iterates $G_t \sim \mathcal{N}(\mu_{S_t}, \sigma_{S_t}^2)$, $S_{t+1} \sim \text{Categorical}(\mathbf{T}_{S_t, \cdot})$ with no jump process.

CHMM-t M-step (per-state ECM, closed form given $u_{t,k}$). For Student-t emissions $b_k(x) = t_{\nu_k}(x; \mu_k, \sigma_k)$ the ECM formulation of Peel and McLachlan [15], Liu and Rubin [16] treats the latent precision as an augmented E-step quantity,

$$u_{t,k} = \frac{\nu_k + 1}{\nu_k + ((O_t - \mu_k)/\sigma_k)^2}, \quad (15)$$

so that conditional on ν_k the location and scale admit closed-form weighted updates,

$$\mu_k^{\text{new}} = \frac{\sum_t \gamma_t(k) u_{t,k} O_t}{\sum_t \gamma_t(k) u_{t,k}}, \quad (\sigma_k^{\text{new}})^2 = \frac{\sum_t \gamma_t(k) u_{t,k} (O_t - \mu_k^{\text{new}})^2}{\sum_t \gamma_t(k)}. \quad (16)$$

The degrees-of-freedom parameter ν_k is then refined by a one-dimensional golden-section search on the per-state Q -function $Q_k(\nu) = \sum_t \gamma_t(k) \log t_{\nu}(O_t; \mu_k, \sigma_k)$ over the bracket $(\nu_{\min}, \nu_{\max}) = (2.1, 50)$.

This two-block ECM preserves the monotonicity of the standard EM guarantee on the compact bracket $[\nu_{\min}, \nu_{\max}]$, with the search invariant $Q_k(\nu^{(n+1)}) \geq Q_k(\nu^{(n)})$ at every iteration; Proposition 1 states the guarantee in full. We extend this M-step with an optional exponential shrinkage prior on $1/\nu_k$, corresponding to the penalised objective

$$Q_k^{\text{pen}}(\nu) = \sum_t \gamma_t(k) \log t_\nu(O_t; \mu_k, \sigma_k) - \lambda / \nu, \quad (17)$$

maximised by the same golden-section search ($\lambda = 0$ recovers the unpenalised Peel and McLachlan [15], Liu and Rubin [16] ECM; the rate sweep and its effect on simulated IS kurtosis are in Section 5).

CHMM-L M-step (closed-form weighted Laplace MLE). For Laplace emissions $b_k(x) = \frac{1}{2b_k} \exp(-|x - \mu_k|/b_k)$ the weighted-MLE estimators are available in closed form:

$$\mu_k^{\text{new}} = \text{WeightedMedian}(O_{1:T}; \gamma_{1:T}(k)), \quad b_k^{\text{new}} = \frac{\sum_t \gamma_t(k) |O_t - \mu_k^{\text{new}}|}{\sum_t \gamma_t(k)}. \quad (18)$$

The weighted median is computed by cumulative-weight bracketing with linear interpolation between adjacent order statistics at the half-total weight crossing. The Laplace M-step carries no shape parameter, so CHMM-L is strictly cheaper per iteration than CHMM-t (which requires the per-state Q_k search).

A.4 Algorithm Descriptions

This section provides detailed pseudocode for the main algorithms used in this study. Algorithm 1 is the unified EM procedure shared across all three emission families. It makes explicit the *shared scaffold* (quantile-based initialization, log-space forward-backward, transition update) and the three *family-specific branches* (Gaussian M-step for CHMM-N, ECM with one-dimensional golden-section ν -search for CHMM-t, closed-form weighted-median + weighted-MAD for CHMM-L); the branch points are lines 3, 12, and 28. Visualizing the three variants as a single algorithm with three M-step branches makes the architectural contribution of the paper (one EM engine, three interchangeable emission families) immediate from the pseudocode. Algorithm 2 presents the Viterbi algorithm for decoding the most likely hidden state sequence. Algorithm 3 describes the CHMM simulation procedure for generating synthetic return paths. Algorithm 4 describes the cross-asset rank-reordering simulator for copula-based multi-asset synthesis. Algorithm 5 details the SIM regression-and-resampling pipeline.

A.5 CHMM-t Degrees-of-Freedom Diagnostics

This appendix backs the discussion in Section 5 of the CHMM-t IS kurtosis overshoot. Figure 5 plots the per-state ν_k histogram for the $K = 18$ CHMM-t fit on SPY IS (seed = 20260420). Thirteen of the eighteen states rest at the upper ECM bracket ($\nu_k = 50$), and only two states lie near the lower bracket ($\nu_2 = 2.19$, $\nu_4 = 2.10$). The simulated IS mixture kurtosis of 17.34 is driven by the posterior mass routed to those two very heavy-tailed states, not by a systemic collapse to the lower bracket.

Table 6 reports the seven-metric panel for CHMM-t fits whose ECM bracket lower bound is lifted from $\nu_{\min} = 2.1$ through $\{2.5, 3.0, 4.0\}$, plus the Gaussian limit ($\nu = \infty$, equivalent to CHMM-N). Lifting the bracket does not reduce the IS kurtosis overshoot: the simulated kurtosis moves within the 12.5–14.0 band across bracket floors and only collapses to 5.05 in the full Gaussian limit. The overshoot is therefore a feature of the Student-t emission family at $K = 18$ and not a lower-bracket artefact.

Algorithm 1 Unified EM algorithm for the CHMM family. Lines 3, 12, and 28 branch on the emission family $F \in \{\mathcal{N}, t, L\}$; every other line is identical across families. CHMM-N and CHMM-L instantiate classical EM; CHMM-t instantiates an ECM sequence whose second CM step updates ν_k by a one-dimensional golden-section search on the per-state Q -function.

Require: Observations $O_{1:T}$, number of states K , tolerance tol , max_iter , emission family $F \in \{\mathcal{N}, t, L\}$, if $F = t$ bracket $(\nu_{\min}, \nu_{\max})$ and initial $\nu^{(0)}$.

Ensure: Transition matrix $\mathbf{T}$, per-state emission parameters $\theta_{1:K}$, initial distribution π .

```

1: Quantile-Based Initialization.
2: Sort  $O^{(\text{sort})} \leftarrow \text{sort}(O_{1:T})$  and partition into  $K$  equal chunks  $\{C_1, \dots, C_K\}$ .
3: for  $k = 1$  to  $K$  do
4:   switch  $F$ 
5:      $F = \mathcal{N}$ :  $\mu_k \leftarrow \text{mean}(C_k)$ ,  $\sigma_k \leftarrow \max(\text{std}(C_k), 10^{-6})$ .
6:      $F = t$ :  $\mu_k \leftarrow \text{mean}(C_k)$ ,  $\sigma_k \leftarrow \max(\text{std}(C_k), 10^{-6})$ ,  $\nu_k \leftarrow \nu^{(0)}$ .
7:      $F = L$ :  $\mu_k \leftarrow \text{median}(C_k)$ ,  $b_k \leftarrow \max(\text{mean}(|C_k - \mu_k|), 10^{-6})$ .
8:   end for
9:  $T_{ij} \leftarrow 1/K$ ;  $\pi_k \leftarrow 1/K$ ;  $\mathcal{L}_{\text{prev}} \leftarrow -\infty$ .
10: EM Loop.
11: for  $n = 1$  to  $\text{max\_iter}$  do
12:   E-Step, per-family emission log-likelihood.
13:   switch  $F$ 
14:      $F = \mathcal{N}$ :  $\log B_{t,k} \leftarrow \log \mathcal{N}(O_t | \mu_k, \sigma_k^2)$ .
15:      $F = t$ :  $\log B_{t,k} \leftarrow \log t_{\nu_k}(O_t; \mu_k, \sigma_k)$ .
16:      $F = L$ :  $\log B_{t,k} \leftarrow -\log(2b_k) - |O_t - \mu_k|/b_k$ .
17:   E-Step, shared log-space forward-backward.
18:    $\log \alpha_1(k) \leftarrow \log \pi_k + \log B_{1,k}$ .
19:   for  $t = 2$  to  $T$ ,  $j = 1$  to  $K$  do
20:      $\log \alpha_t(j) \leftarrow \log \sum_i \exp(\log \alpha_{t-1}(i) + \log T_{ij}) + \log B_{t,j}$ .
21:   end for
22:    $\log \beta_T(k) \leftarrow 0$ .
23:   for  $t = T - 1$  down to  $1$ ,  $i = 1$  to  $K$  do
24:      $\log \beta_t(i) \leftarrow \log \sum_j \exp(\log T_{ij} + \log B_{t+1,j} + \log \beta_{t+1}(j))$ .
25:   end for
26:   Compute  $\gamma_t(k)$  and  $\xi_t(i, j)$  from equation (13).
27:   Shared transition update.  $T_{ij} \leftarrow \sum_t \xi_t(i, j) / \sum_t \gamma_t(i)$ ;  $\pi_k \leftarrow \gamma_1(k)$ .
28:   M-Step, per-family emission update.
29:   switch  $F$ 
30:      $F = \mathcal{N}$  {classical Baum-Welch [13].}
31:      $\mu_k \leftarrow \sum_t \gamma_t(k) O_t / \sum_t \gamma_t(k)$ .
32:      $\sigma_k \leftarrow \max(\sqrt{\sum_t \gamma_t(k) (O_t - \mu_k)^2 / \sum_t \gamma_t(k)}, 10^{-6})$ .
33:      $F = t$  {ECM of Peel and McLachlan [15], Liu and Rubin [16]; closed-form  $(\mu_k, \sigma_k)$  given  $\nu_k$ , then 1D search on  $\nu_k$ .}
34:      $u_{t,k} \leftarrow (\nu_k + 1) / (\nu_k + ((O_t - \mu_k) / \sigma_k)^2)$  for all  $t, k$ .
35:      $\mu_k \leftarrow \sum_t \gamma_t(k) u_{t,k} O_t / \sum_t \gamma_t(k) u_{t,k}$ .
36:      $\sigma_k \leftarrow \max(\sqrt{\sum_t \gamma_t(k) u_{t,k} (O_t - \mu_k)^2 / \sum_t \gamma_t(k)}, 10^{-6})$ .
37:      $\nu_k \leftarrow \arg \max_{\nu \in [\nu_{\min}, \nu_{\max}]} \sum_t \gamma_t(k) \log t_{\nu}(O_t; \mu_k, \sigma_k)$  {40-iter golden-section search.}
38:      $F = L$  {closed-form weighted Laplace MLE.}
39:      $\mu_k \leftarrow \text{WeightedMedian}(O_{1:T}, \gamma_{1:T}(k))$  {cumulative-weight bracketing with interpolation at the  $W_k/2$  crossing.}
40:      $b_k \leftarrow \max(\sum_t \gamma_t(k) |O_t - \mu_k| / \sum_t \gamma_t(k), 10^{-6})$ .
41:   Convergence check.  $\mathcal{L}^{(n)} \leftarrow \log \sum_k \exp(\log \alpha_T(k))$ .
42:   if  $|\mathcal{L}^{(n)} - \mathcal{L}_{\text{prev}}| < \text{tol}$  then
43:     break
44:   end if
45:    $\mathcal{L}_{\text{prev}} \leftarrow \mathcal{L}^{(n)}$ .
46: end for
47: return  $\mathbf{T}, \theta_{1:K}, \pi$  { $\theta_k = (\mu_k, \sigma_k)$  for  $F = \mathcal{N}$ ;  $(\mu_k, \sigma_k, \nu_k)$  for  $F = t$ ;  $(\mu_k, b_k)$  for  $F = L$ .}

```

Algorithm 2 Viterbi decoding for the most likely hidden state sequence

Require: Observations $O_{1:T}$, Trained CHMM $\mathcal{M} = (\mathbf{T}, \{\mu_k, \sigma_k\}, \boldsymbol{\pi})$ **Ensure:** Most probable state sequence $S_{1:T}^*$

```
1: for  $k = 1$  to  $K$  do
2:    $\log \delta_1(k) \leftarrow \log(1/K) + \log \mathcal{N}(O_1 \mid \mu_k, \sigma_k^2)$ .
3: end for
4: for  $t = 2$  to  $T$  do
5:   for  $j = 1$  to  $K$  do
6:      $\text{vals}(j) \leftarrow \log \delta_{t-1}(j) + \log T_{ij}$  for all  $i$ .
7:      $\log \delta_t(j) \leftarrow \max_i \text{vals}(i) + \log \mathcal{N}(O_t \mid \mu_j, \sigma_j^2)$ .
8:      $\psi_t(j) \leftarrow \arg \max_i \text{vals}(i)$ .
9:   end for
10: end for
11:  $S_T^* \leftarrow \arg \max_k \log \delta_T(k)$ .
12: for  $t = T - 1$  down to  $1$  do
13:    $S_t^* \leftarrow \psi_{t+1}(S_{t+1}^*)$ .
14: end for
15: return  $S_{1:T}^*$ 
```

Algorithm 3 CHMM simulation of synthetic excess growth rate paths

Require: Trained CHMM $\mathcal{M} = (\mathbf{T}, \{\mu_k, \sigma_k\}, \boldsymbol{\pi})$, Number of steps M , Number of paths P **Ensure:** Simulated return paths $\{\hat{G}_{1:M}^{(p)}\}_{p=1}^P$

```
1: for  $p = 1$  to  $P$  do
2:   Sample initial state  $S_1 \sim \text{Categorical}(\boldsymbol{\pi})$ .
3:   for  $t = 1$  to  $M$  do
4:     Emit observation  $\hat{G}_t^{(p)} \sim \mathcal{N}(\mu_{S_t}, \sigma_{S_t}^2)$ .
5:     if  $t < M$  then
6:       Transition to next state  $S_{t+1} \sim \text{Categorical}(\mathbf{T}_{S_t, \cdot})$ .
7:     end if
8:   end for
9: end for
10: return  $\{\hat{G}_{1:M}^{(p)}\}_{p=1}^P$ 
```

Table 6. CHMM-t $\nu_{\min}$ bracket sensitivity at $K = 18$ on SPY (seed = 20260420, 1,000 paths). AD = Anderson–Darling pass rate; W_1 = Wasserstein-1; H = Hellinger distance (definitions in Appendix A.7).

$\nu_{\min}$	KS IS	AD IS	KS OoS	AD OoS	Kurt IS	Kurt OoS	ACF-MAE	W_1	H
2.1	95.7	90.2	86.2	82.0	13.95	10.08	0.0546	0.100	0.0759
2.5	95.8	90.6	84.2	79.1	15.12	10.34	0.0547	0.100	0.0759
3.0	93.8	89.5	85.3	80.4	12.25	9.65	0.0548	0.099	0.0760
4.0	95.5	89.5	85.0	80.3	13.25	9.39	0.0550	0.101	0.0759
∞ (Gauss)	94.1	86.7	84.2	75.5	5.03	4.46	0.0510	0.110	0.0809
<i>Observed</i>	–	–	–	–	7.68	5.29	–	–	–

A.6 Validation Metric Details

A.7 Validation Metric Details

This appendix documents the precise form of the seven per-path metrics referenced in Section 3.5. For each model, 1,000 independent paths of length T are simulated (200 paths for the cross-asset

Algorithm 4 Cross-asset copula simulation via rank reordering

Require: Fitted per-asset CHMMs $\{\mathcal{M}_j\}_{j=1}^d$, correlation matrix Σ (positive semi-definite (PSD), unit diagonal), copula family $\mathcal{C} \in \{\text{Gaussian}, t_\nu\}$, path length T , number of paths P

Ensure: Cross-asset path tensor $\{\hat{G}_{j,t}^{(p)}\}$ of shape $T \times d \times P$

```
1: for  $p = 1$  to  $P$  do
2:   Draw copula sample  $\mathbf{U}^{(p)} \in [0, 1]^{T \times d}$  from  $\mathcal{C}(\Sigma)$ :
3:    $L \leftarrow \text{Cholesky}(\Sigma)$ ,  $Z \in \mathbb{R}^{T \times d} \sim \mathcal{N}(0, I_d)$  i.i.d.
4:    $Y \leftarrow ZL^\top$ 
5:   if  $\mathcal{C} = \text{Gaussian}$  then
6:      $\mathbf{U}^{(p)} \leftarrow \Phi(Y)$  componentwise
7:   else
8:     Draw  $W \sim \chi_\nu^2/\nu$  i.i.d. for  $t = 1, \dots, T$ ;  $X_{t,\cdot} \leftarrow Y_{t,\cdot}/\sqrt{W_t}$ 
9:      $\mathbf{U}^{(p)} \leftarrow t_\nu(X)$  componentwise
10:  end if
11:  for  $j = 1$  to  $d$  do
12:    Simulate  $\tilde{g}_{j,1:T}^{(p)}$  from  $\mathcal{M}_j$  via Algorithm 3.
13:    Compute order statistics  $\tilde{g}_{j,(1)}^{(p)} \leq \dots \leq \tilde{g}_{j,(T)}^{(p)}$ .
14:    Compute ranks  $r_{j,t}^{(p)} \leftarrow \text{rank}(U_{j,t}^{(p)})$  for  $t = 1, \dots, T$ .
15:     $\hat{G}_{j,t}^{(p)} \leftarrow \tilde{g}_{j,(r_{j,t}^{(p)})}^{(p)}$  for  $t = 1, \dots, T$ .
16:  end for
17: end for
18: return  $\{\hat{G}_{j,t}^{(p)}\}$ 
```

Algorithm 5 Single Index Model (SIM) construction and simulation

Require: IS return matrix $\mathbf{R} \in \mathbb{R}^{T \times d}$, market-column index m , market CHMM $\mathcal{M}_m$, path length T , number of paths P

Ensure: Cross-asset path tensor $\{\hat{G}_{j,t}^{(p)}\}$ of shape $T \times d \times P$

```
1: Ordinary least squares (OLS) fit. For  $j \neq m$ :
2:    $\hat{\beta}_j \leftarrow \text{Cov}(\mathbf{R}_{:,j}, \mathbf{R}_{:,m})/\text{Var}(\mathbf{R}_{:,m})$ 
3:    $\hat{\alpha}_j \leftarrow \overline{\mathbf{R}_{:,j}} - \hat{\beta}_j \overline{\mathbf{R}_{:,m}}$ 
4:    $\hat{\eta}_{j,t} \leftarrow R_{t,j} - \hat{\alpha}_j - \hat{\beta}_j R_{t,m}$ ,  $R_j^2 \leftarrow 1 - \sum_t \hat{\eta}_{j,t}^2 / \sum_t (R_{t,j} - \overline{\mathbf{R}_{:,j}})^2$ 
5: for  $p = 1$  to  $P$  do
6:   Simulate market path  $G_{m,1:T}^{(p)}$  from  $\mathcal{M}_m$  via Algorithm 3.
7:    $\hat{G}_{m,t}^{(p)} \leftarrow G_{m,t}^{(p)}$ .
8:   for  $j \neq m$  do
9:     Draw residual indices  $\{\iota_t\}_{t=1}^T \sim \text{Uniform}(\{1, \dots, T\})$ .
10:     $\hat{G}_{j,t}^{(p)} \leftarrow \hat{\alpha}_j + \hat{\beta}_j G_{m,t}^{(p)} + \hat{\eta}_{j,\iota_t}$ .
11:  end for
12: end for
13: return  $\{\hat{G}_{j,t}^{(p)}\}$ 
```




Figure 5. Per-state degrees-of-freedom ν_k for the CHMM-t fit at $K = 18$ on SPY IS. The dashed red line marks the lower ECM bracket $\nu_{\min} = 2.1$. Only two states ($\nu_2 = 2.19$, $\nu_4 = 2.10$) sit near the lower bracket; the median ν_k is at the upper bracket.

extension) and each metric is computed per path before aggregation.

Kolmogorov–Smirnov (KS) pass rate. The two-sample KS statistic [48, 49] measures the maximum vertical distance between two empirical CDFs,

$$D_{n,m} = \sup_x |F_n(x) - G_m(x)|. \quad (19)$$

The KS pass rate is the fraction of paths for which the test fails to reject distributional equivalence against the reference data at $\alpha = 0.05$.

Anderson–Darling (AD) pass rate. The AD test is the integrated squared CDF deviation with a quadratic tail weighting, making it more sensitive than KS to tail discrepancies. The AD pass rate is the fraction of paths for which the two-sample AD test fails to reject at $\alpha = 0.05$.

Excess kurtosis. Mean simulated excess kurtosis across paths, compared to the observed value. Given Gaussian emissions, the CHMM is expected to underestimate kurtosis relative to the heavy-tailed empirical distribution, and the shortfall is one of the reported diagnostics.

ACF-MAE. Mean absolute error between observed and mean-simulated autocorrelation functions of $|G_t|$ over $L = 252$ lags,

$$\text{ACF-MAE} = \frac{1}{L} \sum_{\tau=1}^L \left| \hat{\rho}_{|G|}^{\text{obs}}(\tau) - \bar{\hat{\rho}}_{|G|}^{\text{sim}}(\tau) \right|. \quad (20)$$

This is the direct diagnostic for the Rydén et al. limitation.

Wasserstein-1 distance. The mean absolute difference of sorted empirical quantiles, equivalent to the L^1 distance between the empirical CDFs [62],

$$W_1(F, G) = \int_0^1 |F^{-1}(u) - G^{-1}(u)| du. \quad (21)$$

Hellinger distance. A histogram-based overlap distance in $[0, 1]$,

$$H(P, Q) = \frac{1}{\sqrt{2}} \sqrt{\sum_{i=1}^B (\sqrt{p_i} - \sqrt{q_i})^2}, \quad (22)$$

with B equal-width bins spanning the joint support and p_i, q_i the observed and simulated bin probabilities.

Quantile coverage. The fraction of 99 equally-spaced observed quantiles (1st through 99th percentile) falling inside the 5th–95th percentile envelope of the corresponding simulated quantile across 1,000 paths. A coverage of 100% indicates that the simulation envelope fully brackets the observed distribution at every percentile.

Cross-asset metrics. For the cross-asset extension we additionally track the mean Frobenius norm $\|\hat{\Sigma}_{\text{sim}}^{(p)} - \hat{\Sigma}_{\text{obs}}\|_F$ and the off-diagonal Pearson correlation MAE between simulated and observed return matrices, averaged over paths. The off-diagonal MAE is averaged over the $d(d-1)/2$ unique upper-triangular entries of the symmetric correlation matrix (so 15 entries for the six-asset cross-section of Section 4.5), not double-counted across both triangles:

$$\text{MAE}_{\text{off-diag}} = \frac{1}{P} \sum_{p=1}^P \frac{2}{d(d-1)} \sum_{1 \leq i < j \leq d} \left| \hat{\Sigma}_{\text{sim}, ij}^{(p)} - \hat{\Sigma}_{\text{obs}, ij} \right|.$$

A.8 Extended Evaluation: Detail Panels

Per-generator detail for the leverage, aggregational-Gaussianity, and joint p -value diagnostics on the SPY IS and OoS windows.

Table 7. Leverage effect (Pipeline A, SPY). Observed IS $\text{avg}_{k=1}^{20} \text{corr}(r_t^2, r_{t-k}) = -0.088$ with down-vs-up asymmetry 0.358; observed OoS averages are -0.066 and 0.478 . p_{avg} is the two-sided simulation-based coverage p -value for the observed average over 1,000 paths.

Model	avg IS	asym IS	p_{avg} IS	avg OoS	asym OoS	p_{avg} OoS
Bootstrap	0.000	0.001	0.000	0.001	−0.006	0.000
Gaussian	−0.000	−0.001	0.000	−0.000	0.001	0.000
Laplace	−0.000	−0.000	0.000	−0.000	0.005	0.000
GARCH(1,1)	−0.001	0.064	0.014	−0.001	0.065	0.061
CHMM-N	−0.038	0.292	0.000	−0.033	0.287	0.205
CHMM-t	−0.025	0.281	0.000	−0.029	0.274	0.048
CHMM-L	−0.024	0.322	0.000	−0.025	0.310	0.019

Table 8. Aggregational Gaussianity (Pipeline A, SPY). Median simulated kurtosis at aggregation horizons $h \in \{1, 5, 10, 21\}$ over 1,000 paths. Observed IS kurtosis is $\{7.68, 4.34, 6.82, 2.63\}$; observed OoS kurtosis is $\{5.29, 1.98, 0.09, -0.16\}$.

Model	IS kurtosis				OoS kurtosis			
	$h = 1$	$h = 5$	$h = 10$	$h = 21$	$h = 1$	$h = 5$	$h = 10$	$h = 21$
Observed	7.68	4.34	6.82	2.63	5.29	1.98	0.09	-0.16
Bootstrap	7.44	1.31	0.54	0.13	5.24	0.77	0.04	-0.39
Gaussian	-0.01	-0.03	-0.09	-0.17	-0.05	-0.17	-0.30	-0.54
Laplace	2.86	0.51	0.16	-0.04	2.55	0.23	-0.15	-0.49
GARCH(1,1)	4.35	4.79	3.74	2.55	1.79	1.87	1.26	0.34
CHMM-N	4.88	3.70	3.27	2.38	4.08	2.38	1.60	0.40
CHMM-t	8.46	2.65	2.06	1.35	5.55	1.89	1.15	0.21
CHMM-L	6.33	2.45	1.80	1.02	5.49	1.97	1.04	0.07

Table 9. Joint simulation-based p -value coverage (Pipeline A, SPY). Each cell is the two-sided simulation-based p -value for the observed statistic under the generator. $\bar{p}v$ is the mean across $\{\text{kurt}, \text{avg leverage}, \text{agg kurt } h = 5, \text{agg kurt } h = 21\}$; larger $\bar{p}v$ means better joint coverage.

Model	p_{kurt}	p_{acf}	p_{lev}	p_{ak5}	p_{ak21}	$\bar{p}v$ IS	$\bar{p}v$ OoS
Bootstrap	0.894	0.000	0.000	0.012	0.005	0.228	0.466
Gaussian	0.000	0.000	0.000	0.000	0.002	0.000	0.116
Laplace	0.002	0.000	0.000	0.000	0.002	0.001	0.156
GARCH(1,1)	0.271	0.021	0.014	0.903	0.974	0.540	0.468
CHMM-N	0.009	0.000	0.000	0.659	0.894	0.390	0.539
CHMM-t	0.872	0.000	0.000	0.196	0.228	0.324	0.661
CHMM-L	0.369	0.000	0.000	0.061	0.130	0.140	0.692

A.9 Continuous Ranked Probability Score and Diebold-Mariano

The OoS CRPS column of Table 1 reports a proper-scoring-rule complement to the KS pass rate. We use the unbiased sample CRPS estimator: for an ensemble of N simulated paths $\{x_i\}_{i=1}^N$ at time t and observed value y_t ,

$$\widehat{\text{CRPS}}_t = \frac{1}{N} \sum_{i=1}^N |x_i - y_t| - \frac{1}{N(N-1)} \sum_{1 \leq i < j \leq N} |x_i - x_j|. \quad (23)$$

The double sum is computed in $O(N \log N)$ via the sorted-ensemble identity $\sum_{i < j} (x_{(j)} - x_{(i)}) = \sum_i x_{(i)} (2i - N - 1)$. The ensemble at each t is the cross-section of the $N = 1,000$ unconditional simulated paths used throughout Section 4; this yields a marginal-predictive CRPS suitable for an unconditional generative-fidelity assessment, distinct from the conditional one-step-ahead CRPS used in forecasting.

For pairwise comparison we report a Diebold-Mariano test on the per- t loss differential $d_t = \widehat{\text{CRPS}}_t^A - \widehat{\text{CRPS}}_t^B$ with a Newey-West HAC variance under a Bartlett kernel and bandwidth $h = \lfloor T_{\text{OoS}}^{1/3} \rfloor = 8$. Two-sided p -values are reported under the standard-normal null. The body claim of Section 4.3 is: CHMM ties or beats every benchmark on mean OoS CRPS, with Gaussian i.i.d. the only significantly worse comparator.

A.10 VaR / ES Envelope Visualization (Companion to Table 4)

Figure 6 visualizes the median and [5%, 95%] envelopes across simulated paths whose numeric values are reported in Table 4.




Figure 6. VaR / ES envelope diagnostic (SPY, $N_{\text{paths}} = 1,000$; global seed policy in Section 3.5). Each subplot shows the simulated median (dot) and [5%, 95%] envelope (error bar) across paths against the observed historical value (dashed red line); rows are IS and OoS windows, columns are VaR and ES, and the x-axis lists the five generators (Bootstrap, GARCH(1,1), CHMM-N, CHMM-t, CHMM-L). The y-axis is the annualized log excess growth rate (daily log return scaled by $1/\Delta t$ with $\Delta t = 1/252$), consistent with the $R_{\text{IS}}/R_{\text{OoS}}$ convention used throughout the paper. Headline: every observed red line falls inside the corresponding simulated envelope, confirming that the CHMM family passes the coverage leg of the back-test complementary to the Kupiec / Christoffersen likelihood-ratio leg in Section 4.6.

A.11 Formal Theoretical Statements

The following propositions are referenced in Section 3.8 and stated formally here for reference; each is a one-line specialisation of a standard result and the proof sketch is one sentence.

Proposition 1 (ECM monotonicity for CHMM-t (and EM monotonicity for CHMM-N, CHMM-L)). *Under compactness of the parameter space (μ_k, σ_k or b_k, ν_k each in a closed bounded interval, $T_{ij} \geq \epsilon_T > 0$), the two-block CM step for CHMM-t of Section 3.7, namely the closed-form weighted updates of $(\mathbf{T}, \boldsymbol{\pi}, \mu_k, \sigma_k)$ at fixed ν_k , followed by a golden-section maximisation of the per-state Q -function over $\nu_k \in [\nu_{\min}, \nu_{\max}]$, satisfies $Q(\theta^{(n+1)} \mid \theta^{(n)}) \geq Q(\theta^{(n)} \mid \theta^{(n)})$, and consequently*

$\mathcal{L}(\theta^{(n+1)}) \geq \mathcal{L}(\theta^{(n)})$ where $\mathcal{L}$ is the incomplete-data log-likelihood. The same monotonicity holds for the closed-form CHMM-N and CHMM-L M-steps under the standard EM guarantee.

Proof sketch. Each CM step does not decrease Q by definition (CM-1 is closed-form maximum, CM-2 is golden-section maximum); the standard EM inequality $\mathcal{L}(\theta^{(n+1)}) - \mathcal{L}(\theta^{(n)}) \geq Q(\theta^{(n+1)} | \theta^{(n)}) - Q(\theta^{(n)} | \theta^{(n)})$ [54, 55] converts Q -monotonicity into log-likelihood monotonicity. Compactness gives boundedness above, hence convergence.

Proposition 2 (Identifiability up to label permutation). *Under irreducibility and aperiodicity of $\mathbf{T}$ (Assumption 1) and the emission contrast condition (Assumption 2), the CHMM parameter is identifiable up to label permutation in all three emission families considered (CHMM-N, CHMM-t with $\nu_k \in [\nu_{\min}, \nu_{\max}]$, and CHMM-L).*

Proof sketch. The HMM identifiability theorem of Allman et al. [56] requires $\mathbf{T}$ of full rank and per-state emissions linearly independent in the space of densities on $\mathbb{R}$. Full-rank $\mathbf{T}$ follows from irreducibility [53]; linear independence of Gaussian, bracketed- ν Student-t, and Laplace emissions with distinct location-scale parameters follows from Yakowitz and Spragins [57].

Proposition 3 (MLE consistency). *Under Assumptions 1–2 and the assumption that the true parameter lies in the interior of a compact Θ , the MLE $\hat{\theta}_T = \arg \max_{\theta \in \Theta} \mathcal{L}_T(\theta)$ is consistent: $\hat{\theta}_T \xrightarrow{p} \theta_0$ as $T \rightarrow \infty$, modulo label permutation.*

Proof reference. Bickel et al. [58] (Theorem 1) and Douc et al. [59] apply directly under our assumption set: irreducibility supplies the mixing condition, finite second moment the moment condition, compactness the parameter-space condition.

Proposition 4 (Marginal preservation under rank reordering). *Let $\tilde{g}_{j,1:T}$ be a sample of size T from the per-asset CHMM marginal of asset j , and let the rank-reordered output be $\hat{g}_{j,t}^{(p)} = \tilde{g}_{j,(r_{j,t}^{(p)})}^{(p)}$ where $r_{j,t}^{(p)}$ is the rank of the copula sample. Then for every asset j the empirical CDF of $\{\hat{g}_{j,t}^{(p)}\}_{t=1}^T$ equals that of $\{\tilde{g}_{j,t}^{(p)}\}_{t=1}^T$.*

Proof sketch. The map $t \mapsto r_{j,t}^{(p)}$ is a permutation almost surely, so the multiset of values is preserved; the empirical CDF depends only on the multiset.

A.12 State-Resolution and Multi-Emission Sensitivity

A.13 Full State-Resolution Sensitivity Table

Table 10 reports the full K -sweep referenced in Section 4.2 of the main text. Observed excess kurtosis is 7.68 IS and 5.29 OoS at all K .

Table 10. State resolution sensitivity for SPY CHMM-N (1,000 paths, $\alpha = 0.05$). Every metric is computed on the identical simulation pipeline used for Table 1 in the main text. IS coverage was 100% at every K ; OoS coverage is reported separately.

K	KS (%)		AD (%)		Kurtosis		ACF-MAE	W_1	H	Cov	OoS (%)
	IS	OoS	IS	OoS	Obs	Sim					
3	88.8	77.1	79.2	67.9	7.68	3.80	0.0466	0.130	0.0836		94.9
6	93.1	80.0	87.1	74.6	7.68	5.26	0.0499	0.113	0.0828		90.9
9	94.6	82.4	85.8	73.6	7.68	4.56	0.0494	0.116	0.0806		93.9
12	94.6	82.8	86.0	74.9	7.68	4.57	0.0495	0.112	0.0802		88.9
15	93.9	82.7	86.9	75.6	7.68	4.89	0.0505	0.116	0.0806		90.9
18	95.6	81.7	88.4	73.2	7.68	4.98	0.0509	0.110	0.0805		90.9
21	96.1	86.6	89.0	78.4	7.68	3.83	0.0532	0.111	0.0798		100.0

A.14 Multi-Emission State-Resolution Sensitivity

Table 11 extends Table 10 across the three emission families, confirming that the qualitative operating point $K = 18$ is stable when the Gaussian component is replaced with Student-t or Laplace emissions. For CHMM-t the per-state degrees of freedom ν_k are refit by ECM-plus-golden-section at every K ; for CHMM-L the weighted-median location and weighted MAD scale are closed form and impose no additional cost. Three patterns extend from Table 10 to the new families. First, the distributional pass rates plateau in the $K \in \{12, 15, 18, 21\}$ band for all three families: CHMM-t reaches its highest IS KS at $K = 18$ (95.2%) and its highest OoS KS at $K = 18$ (85.1%), and CHMM-L reaches its highest IS KS at $K = 18$ (96.2%) and its highest IS AD at $K = 18$ (93.0%). Once a family enters the plateau the operating point is flat with respect to K , so we report the main-text results at the same $K = 18$ used for CHMM-N. Second, the ACF-MAE remains within the narrow 0.047–0.056 band observed for the Gaussian sweep, indicating that the volatility-clustering mechanism is a property of the regime structure rather than the emission family. Third, simulated kurtosis behaves as expected from the component families: the Gaussian sweep peaks around 5.0 at the high- K end, the Laplace sweep lands squarely in the 5.2–6.6 range (closest to the observed 7.68), and the Student-t sweep climbs sharply above 13 across the entire sweep because the ECM search assigns small ν_k to a subset of tail states, over-correcting IS kurtosis.

Table 11. Multi-emission state resolution sensitivity for SPY (1,000 paths, $\alpha = 0.05$). CHMM-N: Gaussian; CHMM-t: Student-t with per-state ν_k refit by ECM-plus-golden-section; CHMM-L: Laplace with weighted-median location and weighted-MAD scale. Every metric is computed on the identical simulation pipeline used for Table 1. Observed excess kurtosis is 7.68 IS and 5.29 OoS.

Family	K	KS (%)		AD (%)		Kurt (IS sim)	ACF-MAE	W_1 IS	H IS
		IS	OoS	IS	OoS				
CHMM-N	3	89.8	76.3	81.3	67.7	3.81	0.0468	0.128	0.0835
	6	91.7	79.6	85.9	73.6	5.30	0.0497	0.115	0.0831
	9	92.8	82.2	84.7	75.5	4.55	0.0497	0.118	0.0810
	12	95.0	83.7	86.1	76.6	4.55	0.0497	0.114	0.0805
	15	95.2	83.5	86.9	76.3	4.91	0.0508	0.111	0.0802
	18	94.7	82.4	85.8	76.5	5.02	0.0509	0.111	0.0809
	21	95.2	87.6	89.1	80.9	3.80	0.0530	0.112	0.0802
CHMM-t	3	91.1	81.6	85.2	77.1	27.12	0.0540	0.109	0.0764
	6	92.2	82.1	86.1	76.6	20.80	0.0533	0.108	0.0768
	9	92.7	84.8	86.9	80.1	15.96	0.0537	0.104	0.0756
	12	94.1	83.1	88.3	80.6	14.68	0.0542	0.102	0.0757
	15	93.5	82.9	88.0	76.5	15.85	0.0537	0.103	0.0757
	18	95.2	85.1	90.9	80.7	14.40	0.0550	0.101	0.0756
	21	94.9	82.2	89.3	77.8	13.28	0.0540	0.103	0.0758
CHMM-L	3	82.4	63.6	82.9	66.9	5.24	0.0533	0.112	0.0853
	6	88.8	77.0	83.3	72.5	5.74	0.0511	0.113	0.0823
	9	90.0	77.9	84.1	74.8	6.05	0.0519	0.113	0.0823
	12	92.4	77.8	87.6	77.9	6.21	0.0540	0.106	0.0797
	15	92.6	79.4	87.0	77.6	6.06	0.0542	0.106	0.0795
	18	96.2	80.9	93.0	80.6	6.60	0.0564	0.101	0.0793
	21	90.7	80.7	87.3	78.8	5.99	0.0544	0.107	0.0809

A.15 Baum-Welch Convergence

Figure 7 shows the Baum-Welch convergence curves for representative state resolutions. All models converged within the 60-iteration budget, with the log-likelihood monotonically increasing (as guaranteed by the EM algorithm). Lower K values typically converged faster (15–25 iterations), while higher K values required 25–40 iterations to reach the convergence tolerance of $|\Delta\mathcal{L}| < 10^{-4}$. The monotonic increase confirms that the quantile-based initialization provided a valid starting point from which the EM algorithm could improve without encountering degenerate solutions.




Figure 7. Baum-Welch convergence for the SPY CHMM-N fit (Pipeline A, IS) at selected state resolutions. Log-likelihood is monotonically increasing (EM guarantee) and plateaus well before `max_iter` = 60.

A.16 IS Comparison at Non-Selected State Resolutions

Figures 8–11 show the IS model comparison panels for $K \in \{3, 6, 12, 21\}$, complementing the $K = 18$ comparison in Figure 3 of the main text. Across all resolutions, the CHMM’s simulated density closely overlays the observed density, and the ACF of absolute returns tracks the slow decay pattern. The visible improvement from $K = 3$ to $K = 18$ is concentrated in the tails of the Q-Q panel and the steady-state level of the ACF panel. $K = 9$ and $K = 15$ panels are omitted because they interpolate smoothly between the shown resolutions; their numerical metrics are in Table 10.




Figure 8. IS SPY CHMM-N comparison at $K = 3$ (Pipeline A, 1,000 simulated paths; panels as in Figure 3). The $K = 18$ counterpart is Figure 3 of the main text.




Figure 9. IS SPY CHMM-N comparison at $K = 6$ (Pipeline A, 1,000 simulated paths; panels as in Figure 3).




Figure 10. IS SPY CHMM-N comparison at $K = 12$ (Pipeline A, 1,000 simulated paths; panels as in Figure 3).




Figure 11. IS SPY CHMM-N comparison at $K = 21$ (Pipeline A, 1,000 simulated paths; panels as in Figure 3).

A.17 OoS Validation at Non-Selected State Resolutions

Figures 12–15 show the OoS validation panels for $K \in \{3, 6, 12, 21\}$. KS pass rates remain above 91% at all resolutions; the $K = 18$ panel is reproduced in Figure 4 of the main text. As with the IS panels, $K = 9$ and $K = 15$ are omitted; numerical OoS KS pass rates are in Table 10.




Figure 12. OoS SPY CHMM-N validation at $K = 3$ (Pipeline A, 1,000 simulated paths; panels as in Figure 4). The $K = 18$ counterpart is Figure 4 of the main text.




Figure 13. OoS SPY CHMM-N validation at $K = 6$ (Pipeline A, 1,000 simulated paths; panels as in Figure 4).




Figure 14. OoS SPY CHMM-N validation at $K = 12$ (Pipeline A, 1,000 simulated paths; panels as in Figure 4).




Figure 15. OoS SPY CHMM-N validation at $K = 21$ (Pipeline A, 1,000 simulated paths; panels as in Figure 4).

A.18 Multi-Emission Figures at $K = 18$

This appendix reproduces the standard diagnostic suite (emission PDFs, residence times, IS comparison, OoS validation, and convergence) for the two heavy-tailed emission families (CHMM-t and CHMM-L) at the selected state resolution $K = 18$; transition matrices and stationary distributions

are discussed in prose only below, since they are visually indistinguishable across families at $K = 18$. The Gaussian counterparts are already shown as Figures 3 and 4 in the main text and as Figure 7 above in this appendix. The figures confirm that the shared forward-backward scaffold produces internally-similar regime structure across the three families: the near-diagonal transition band, the K -component dense emission cover, the longer central-state residence times, and the monotone log-likelihood convergence are features of the CHMM framework and not of the Gaussian component specifically.

Emission PDFs. Figure 16 compares the learned $K = 18$ emissions for the three families. The CHMM-t panel's per-state ν_k allocation is visible in the narrower modes of certain tail states; the CHMM-L panel's Laplace components have sharper peaks and slower exponential tail decay than the Gaussians.




Figure 16. Learned emission distributions at $K = 18$ across the three emission families (CHMM-N Gaussian, CHMM-t Student-t with per-state ν_k , CHMM-L Laplace), overlaid on the observed SPY IS return histogram (Pipeline A). Heavier-tailed emission families (t, L) give individual components more tail mass, which is what closes the kurtosis gap of the Gaussian CHMM.

Transition matrices and stationary distributions. The $K = 18$ transition matrices and stationary distributions learned under CHMM-t and CHMM-L are visually indistinguishable from their CHMM-N counterparts: the near-diagonal band, the magnitude of self-transition probabilities, and the mass concentration at central states are features of the regime structure, not the emission family. This is consistent with the ACF-MAE being stable across the family axis (Table 11). The per-family residence times $\tau_k = 1/(1 - T_{kk})$, which magnify small differences in the self-transition probabilities each family recovers, are compared in Figure 17.




Figure 17. Natural residence times $\tau_k = 1/(1 - T_{kk})$ at $K = 18$ for the SPY IS fit (Pipeline A) across the three emission families (CHMM-N Gaussian, CHMM-t Student-t, CHMM-L Laplace).

IS comparison. Figures 18–19 reproduce the three-panel IS comparison (density, ACF of $|G_t|$, tail Q-Q) at $K = 18$ for the two heavy-tailed families; the CHMM-N panel is Figure 3 in the main text.




Figure 18. IS SPY model comparison at $K = 18$ for CHMM-t Student-t emissions (Pipeline A, 1,000 simulated paths). Panels show (a) marginal density with KS pass rate, (b) ACF of $|G_t|$, and (c) tail Q-Q plot. The CHMM-N counterpart is Figure 3; the CHMM-L counterpart is Figure 19.




Figure 19. IS SPY model comparison at $K = 18$ for CHMM-L Laplace emissions (Pipeline A, 1,000 simulated paths). Panels as in Figure 18.

OoS validation. Figures 20–21 show the OoS validation panels for CHMM-t and CHMM-L at $K = 18$. The CHMM-t panel demonstrates the OoS kurtosis alignment noted in the main text (simulated 6.40 vs. observed 6.24), and the CHMM-L panel shows the tighter density fan produced by the heavier Laplace tails relative to the Gaussian.




Figure 20. OoS SPY validation at $K = 18$ for CHMM-t Student-t emissions (Pipeline A, 1,000 simulated paths). Panels show (a) KS p -values, (b) OoS density fan, and (c) OoS ACF of $|G_t|$. The CHMM-N counterpart is Figure 4; the CHMM-L counterpart is Figure 21.




Figure 21. OoS SPY validation at $K = 18$ for CHMM-L Laplace emissions (Pipeline A, 1,000 simulated paths). Panels as in Figure 20.

Convergence. Figure 22 shows the Baum-Welch / ECM / weighted-EM log-likelihood histories at $K = 18$ for all three families. The CHMM-t and CHMM-L curves inherit the monotone increase of the classical EM guarantee; the CHMM-t curve has a slightly longer pre-plateau phase because each ECM sub-iteration refines ν_k with a golden-section search.




Figure 22. Log-likelihood histories for the SPY IS fit (Pipeline A) at $K = 18$ across the three emission families (CHMM-N Gaussian classical EM; CHMM-t Student-t ECM with per-state ν_k golden-section sub-step; CHMM-L Laplace weighted-EM). All three curves are monotonically non-decreasing and plateau well before the `max_iter` = 60 budget.

A.19 Rydén $K = 2$ Replication

Table 12 reports the direct replication of Rydén et al.’s 1998 $K = 2$ setting on SPY IS under both initialization strategies, as discussed in Section 5. The quantile-based initialization used in the main body is row 1; five independent random-seed initializations are rows 2–6. Random init here draws per-state (μ_k, σ_k) from a moderate perturbation of the sample moments, which is representative of the random-init practice common in the 1990s HMM literature.

Table 12. Rydén replication: $K = 2$ CHMM-N on SPY IS under quantile vs random initialization (seed sub-index varies the random init).

Init	KS IS (%)	AD IS (%)	ACF-MAE	Kurt IS	KS OoS (%)
Quantile (main body)	79.0	71.5	0.0501	2.57	72.4
Random seed = 1	78.7	70.7	0.0502	2.57	73.2
Random seed = 7	75.3	69.6	0.0498	2.56	74.0
Random seed = 42	77.6	70.0	0.0498	2.59	71.8
Random seed = 100	76.2	69.4	0.0500	2.59	72.3
Random seed = 2024	76.4	69.8	0.0499	2.54	72.9

The key observation is that ACF-MAE is essentially constant at ≈ 0.050 across all six rows, matching the $K = 18$ CHMM-N of the main body. IS KS pass rate remains below the $K = 18$ operating point (75.3–79.0%) because two Gaussian components cannot tile the full multi-regime return distribution. The combination is consistent with the reading that the binding low- K constraint in the Rydén et al. [11] setup is distributional rather than temporal: the ACF is already well reproduced at $K = 2$, and it is distributional fidelity that drives the preference for $K = 18$.

A.20 Robustness, Test Power, Multi-Seed, and Auxiliary Baselines

A.21 OoS KS Test-Power Calibration

Table 13 reports the two-sample KS test-power calibration used to anchor the OoS KS pass rates of Section 4.3. For each of 1,000 Monte Carlo replications (seed = 20260420), a simulated series of length $T_{\text{OoS}} = 572$ is drawn from each of three reference generators and tested against the observed OoS series using the same $\alpha = 0.05$ threshold as the main-body pass-rate metric.

Table 13. OoS KS pass-rate calibration (1,000 replications, seed = 20260420).

Reference generator	Pass rate (%)
I.i.d. resamples of R_{OoS} at length $T_{\text{OoS}} = 572$ (known-correct ceiling)	100.0
I.i.d. resamples of R_{IS} at length $T_{\text{OoS}} = 572$ (nearly correct)	90.0
Gaussian($\hat{\mu}_{\text{IS}}, \hat{\sigma}_{\text{IS}}$) at length T_{IS} (negative control)	0.0

The 90.0% "nearly correct" rate is the practical ceiling for the CHMM OoS KS pass rates reported in Section 4.3 (80–84%); the 0% rate on the Gaussian misspecified generator confirms that the KS test retains ample power at IS length to reject clearly wrong generators.

A.22 Block-Bootstrap KS Recalibration and Mean p-value Distribution

Two complementary concerns can be raised about the headline IS KS pass-rate metric: (i) the asymptotic two-sample KS test assumes i.i.d. samples, but the simulated paths from CHMM and

Table 14. Block-bootstrap KS recalibration and p-value distribution. All numbers under seed 20260422 with $N_{\text{paths}} = 500$. “asyp pass” is the original Section 4.3 metric (two-sample KS p-value ≥ 0.05). “mean pv” is the mean two-sample KS p-value across paths, with q_5 and q_{95} as the lower / upper p-value quantiles. “blkL%” is the fraction of paths whose KS statistic falls below the stationary-block-bootstrap 95% critical value at mean block length L . Block-bootstrap critical values: 0.0306 ($L = 5$), 0.0338 ($L = 10$), 0.0350 ($L = 20$); asymptotic critical value 0.0383.

Generator	asyp pass% $\uparrow$	mean pv $\uparrow$	pv q_{05}	pv q_{95}	blk5% $\uparrow$	blk10% $\uparrow$	blk20% $\uparrow$
i.i.d. bootstrap	99.8	0.812	0.371	0.999	98.2	99.0	99.6
GARCH(1,1)	25.4	0.042	0.000	0.201	5.6	11.2	16.2
CHMM-N	94.8	0.546	0.047	0.976	81.2	89.4	91.6
CHMM-t	95.8	0.571	0.059	0.987	82.4	89.2	91.6
CHMM-L	94.8	0.498	0.048	0.948	81.0	88.2	90.0

GARCH all carry volatility clustering, so the asymptotic null distribution may not apply; and (ii) the binary “pass / fail at $\alpha = 0.05$ ” reduces a continuous p-value to a one-bit summary. Table 14 reports two recalibrations addressing each concern.

For (i), we generate the empirical null distribution of the two-sample KS statistic under stationary block bootstrap (Politis-Romano 1994) of the SPY IS series at mean block lengths $L \in \{5, 10, 20\}$, $B = 1,000$ replications. The resulting 95% block-bootstrap critical values are 0.0306, 0.0338, and 0.0350 respectively, all *smaller* than the asymptotic critical value $1.36\sqrt{2/n_{\text{IS}}} = 0.0383$; the block-bootstrap test is therefore stricter than the asymptotic one because it accounts for the within-window dependence in the IS series itself. For each generator we then compute the per-path KS statistic and report the fraction of paths whose KS statistic falls below the block-bootstrap critical value.

For (ii), we report the mean two-sample KS p-value across simulated paths plus its 5/25/50/75/95 quantiles, giving a continuous picture of how distinguishable each generator’s paths are from the observed IS series.

Two patterns emerge that the asymptotic pass-rate metric of the main panel does not surface. First, the cross-generator ordering on the mean p-value scale (bootstrap 0.81, CHMM-t 0.57, CHMM-N 0.55, CHMM-L 0.50, GARCH 0.04) is much more compressed for CHMM-vs-bootstrap than the 94.8–99.8% asyp pass rates suggest: the CHMM family lands at mean p-value ~ 0.5 , the natural value for paths that are statistically indistinguishable from the observed series. Second, GARCH(1,1)’s drop from 25.4% asyp $\rightarrow$ 5.6–16.2% block-aware confirms concern (i): the asymptotic test is mildly lenient on volatility-clustering generators with the wrong marginal, and the block-aware test catches GARCH’s misspecification more sharply. The CHMM family, by contrast, sees only a small block-bootstrap penalty (94.8% $\rightarrow$ 88–92%), indicating that its asymptotic-test passes are not a clustering-induced artefact.

A.23 Bandwidth, Bin-Count, and Multi-Seed Robustness

Fixed observed-sample MMD bandwidth. An auxiliary MMD metric we report in this appendix (Section A.6) computes Maximum Mean Discrepancy with median-heuristic bandwidth on the pooled (observed + simulated) window matrix, so the kernel scale γ varies with the generator under evaluation. We refit MMD across the full extended generator panel under a single bandwidth γ_{obs} computed once on 500 observed standardised 20-day windows from SPY IS. The legacy and fixed MMDs differ only in the third or fourth decimal across all ten generators in our reproduction; the cross-generator ranking under the fixed bandwidth is preserved. The bandwidth-heuristic concern is therefore real in principle but small in practice for this panel.

Multi-seed Monte Carlo on headline metrics. The paper reports headline statistics under a single global seed. We re-run the CHMM-N / -t / -L pipeline at $K = 18$ across ten alternative global seeds (base 20260422, increment 100) with $N_{\text{paths}} = 250$ per seed and report mean $\pm$ std for the headline metrics in Table 15. Seed-to-seed variability is small: IS KS std is 1.2–1.8%; OoS KS std is 1.8–3.6%; simulated IS kurtosis std is 0.06 for CHMM-N, 0.05 for CHMM-L, and 1.95 for CHMM-t (reflecting the well-known instability of the per-state ν_k ECM at moderate K , already discussed in Section 5). Kupiec $\text{LR}_{\text{uc}}^{1\%}$ has std up to 1.25 for CHMM-N because the breach count moves on a coarse 0–10 grid at $T_{\text{OoS}} = 572$. The single-seed headline numbers in Table 1 (e.g., CHMM-N IS KS 93.8%, simulated kurt 4.99) sit comfortably inside the multi-seed bands.

Table 15. Multi-seed Monte Carlo on headline metrics. Ten global seeds; $N_{\text{paths}} = 250$ per seed (reduced from the headline panel’s 1,000 for tractability across 30 fits, so the seed-to-seed std reported here is an upper bound on the std the headline panel would show at $N_{\text{paths}} = 1,000$).

Model	IS KS %	OoS KS %	sim Kurt	ACF-MAE	$\text{LR}_{\text{uc}}^{1\%}$	$\text{LR}_{\text{uc}}^{5\%}$
CHMM-N	94.7 ± 1.2	82.6 ± 1.8	5.03 ± 0.06	0.0467 ± 0.0006	2.12 ± 1.25	3.83 ± 0.00
CHMM-t	94.6 ± 1.2	85.6 ± 1.9	13.95 ± 1.95	0.0518 ± 0.0003	0.54 ± 0.15	3.19 ± 0.34
CHMM-L	94.4 ± 1.8	80.8 ± 3.6	6.71 ± 0.05	0.0537 ± 0.0003	3.26 ± 0.00	2.71 ± 0.59

A.24 Block-Bootstrap Baseline

Table 16 reports the full seven-metric panel for the stationary block bootstrap at block length $b = 10$ trading days, the temporally-aware non-parametric baseline added to Table 1.

Table 16. Stationary block bootstrap (block length $b = 10$ trading days) on SPY (1,000 paths, seed = 20260420).

Window	KS	AD	Kurt	ACF-MAE	W_1	H	Cov
IS	99.2	96.4	7.07	0.0568	0.092	0.0534	100.0
OoS	87.5	86.4	5.34	0.0512	0.225	0.1570	86.9

A.25 QuantGAN Deep-Generative Baseline

Architecture and training details for the QuantGAN deep-generative baseline used in Section 4.3 are summarised here. The construction is a repo-native approximation to the convolutional WGAN of Wiese et al. [9]: a one-dimensional convolutional generator and critic, Wasserstein loss with critic weight clipping, trained on rolling windows of standardised SPY returns and synthesising long paths by stitching generated windows.

- **Generator.** Three 1D convolutional layers ($G_{\text{ch}} = 32$, kernel size 5, ReLU, same-padding), followed by a 1×1 projection to a single output channel. Input is i.i.d. Gaussian noise of latent dimension $L = 8$ at every position of the rolling window.
- **Critic.** Three 1D convolutional layers ($D_{\text{ch}} = 32$, kernel size 5, leaky-ReLU slope 0.2, same-padding), global mean pooling along the time axis, and a single linear scalar output.
- **Loss.** Wasserstein loss with critic weight clipping at ± 0.01 [63]. The critic minimises $\mathbb{E}[D(\tilde{G})] - \mathbb{E}[D(G)]$ on real / fake mini-batches; the generator maximises $\mathbb{E}[D(\tilde{G})]$.

- **Optimiser.** Adam, learning rate 1.0×10^{-4} for both networks, batch size 64.
- **Training schedule.** 15 epochs, 120 outer steps per epoch, 4 critic updates per generator update, training on rolling windows of length $W = 64$ standardised by the IS mean and standard deviation.
- **Synthesis.** Independent draws of length W are stitched end-to-end to produce paths of length T_{IS} or T_{OoS} ; outputs are then re-scaled back to the natural return scale by the IS mean and standard deviation.
- **Implementation.** Julia with Flux.jl [64] and Optimisers.Adam state. Random seed for weight initialisation, training shuffling, and per-path sampling: 20260422, with deterministic per-path sub-seeds.

The trained model is callable through the same `simulate` interface as the CHMM, so the seven-metric panel evaluates QuantGAN on byte-identical windows to the other generators. The QuantGAN row is reported alongside the extended-GARCH and SM-CHMM rows in Table 18; the deep-generative row reaches 0% IS and OoS KS pass rates and undershoots IS kurtosis decisively (2.1 vs. observed 7.7), consistent with the documented difficulty of GAN-based generators on financial-time-series volatility-clustering benchmarks [32, 33].

A.26 Per-Ticker OoS Price-Simulation Figures and Full Path-Level Metrics (Unconditional CHMM)

This appendix collects (i) the full per-family path-level probabilistic-forecast metrics for an auxiliary OoS price-simulation diagnostic on the six SPY-member tickers and (ii) the per-ticker price-fan and terminal-distribution figures for NVDA, JNJ, JPM, AAPL, and QQQ. Plotting conventions: blue shaded bands are the 5–95% and 25–75% quantile envelopes of the 500 simulated paths, blue solid line is the simulated median, red solid line is the observed OoS price track. Terminal-histogram panels show the cross-path distribution of $\tilde{P}_T$ at T_{OoS} in gray, with the observed terminal P_T^{obs} as a red vertical line and the simulated median as a blue dashed line. For brevity, the figures show CHMM-N only (the representative family; differences across the three emission families in the price-level metrics are small, see Table 17).

A.27 Expanded Baseline Panel: GARCH Family, Semi-Markov CHMM, and QuantGAN

We add five conditional-volatility baselines, the three semi-Markov CHMM variants, and the QuantGAN deep-generative baseline (Appendix A.25) to the single GARCH(1,1) Gaussian row of Table 1. The conditional-volatility models are fit by grid-initialised Nelder-Mead maximum likelihood on the SPY IS window ($T_{\text{IS}} = 2,516$); simulation follows the same recursion used in estimation. Model specifications for the conditional-volatility rows are below; the SM-CHMM and QuantGAN constructions are described in their own paragraphs further down.

EGARCH(1,1) Gaussian. The exponential GARCH of Nelson [27] evolves $\log \sigma_t^2$ in levels:

$$\log \sigma_t^2 = \omega + \beta \log \sigma_{t-1}^2 + \alpha \left(|z_{t-1}| - \sqrt{2/\pi} \right) + \gamma z_{t-1}, \quad z_t = (r_t - \mu)/\sigma_t, \quad z_t \sim \mathcal{N}(0, 1). \quad (24)$$

The log-variance formulation removes the GARCH positivity constraints, and the leverage term γz_{t-1} captures the asymmetric response to negative returns.

Table 17. Full path-level probabilistic-forecast metrics for OoS equity-price simulation, all (ticker, family) cells ($K = 18$, 500 paths, $T_{\text{OoS}} = 572$; global seed policy in Section 3.5). $\overline{C}_{0.90}$: fraction of OoS steps at which the observed price falls inside the simulated 5–95% band (calibration); $\text{MAPE}_{0.50}$: mean absolute percentage error between observed path and simulated median (sharpness); $\overline{\text{CRPS}}$: horizon-average bias-corrected sample CRPS of Gneiting and Raftery [41]; $\overline{\text{CRPS}}/S_0$: CRPS scaled by S_0 to remove price-level effects. Lower MAPE and CRPS, higher $\overline{C}_{0.90}$, are better. **Headline:** no CHMM emission family dominates on all four metrics across all six tickers; differences between CHMM-N, CHMM-t and CHMM-L live mostly in the return tails and propagate only weakly into aggregated price-level scores.

Ticker	Family	$\overline{C}_{0.90}$	$\text{MAPE}_{0.50}$	$\overline{\text{CRPS}}$	$\overline{\text{CRPS}}/S_0$
SPY	CHMM-N	1.000	0.100	39.25	0.0835
	CHMM-t	1.000	0.112	42.66	0.0908
	CHMM-L	1.000	0.102	38.68	0.0823
NVDA	CHMM-N	0.698	0.339	30.62	0.6424
	CHMM-t	0.636	0.383	34.77	0.7293
	CHMM-L	0.666	0.411	36.58	0.7674
JNJ	CHMM-N	1.000	0.101	12.55	0.0780
	CHMM-t	1.000	0.098	12.62	0.0784
	CHMM-L	1.000	0.101	12.65	0.0786
JPM	CHMM-N	1.000	0.195	32.23	0.1881
	CHMM-t	1.000	0.241	40.38	0.2357
	CHMM-L	1.000	0.211	35.25	0.2057
AAPL	CHMM-N	1.000	0.108	21.30	0.1156
	CHMM-t	1.000	0.075	18.90	0.1025
	CHMM-L	1.000	0.077	18.82	0.1021
QQQ	CHMM-N	0.998	0.073	30.48	0.0763
	CHMM-t	0.995	0.079	31.63	0.0792
	CHMM-L	1.000	0.118	41.12	0.1029

GJR-GARCH(1,1) Gaussian. The threshold GARCH of Glosten et al. [28] / Zakoïan [29] augments GARCH(1,1) with an indicator on negative innovations:

$$\sigma_t^2 = \omega + \alpha \varepsilon_{t-1}^2 + \gamma \varepsilon_{t-1}^2 \mathbf{1}\{\varepsilon_{t-1} < 0\} + \beta \sigma_{t-1}^2, \quad (25)$$

with $\varepsilon_t = r_t - \mu$. Stationarity requires $\alpha + \beta + \gamma/2 < 1$; positive γ indicates leverage.

GARCH(1,1) Student-t innovations. A direct tail-aware competitor to CHMM-t: the GARCH(1,1) variance recursion is unchanged, but innovations follow a standardised Student-t with $\nu > 2$ degrees of freedom rescaled to unit variance. The log-likelihood uses the Student-t density:

$$r_t - \mu = \sigma_t \cdot z_t, \quad z_t = \eta_t \sqrt{(\nu - 2)/\nu}, \quad \eta_t \sim t_\nu. \quad (26)$$

The fitted $\hat{\nu} = 6.89$ on SPY IS sits inside the moderate-tail regime the equity-risk literature typically reports.

HAR-RV on daily squared returns. The heterogeneous-autoregressive model of Corsi [31] models realised variance as a mixture of daily, weekly, and monthly components:

$$\log \text{RV}_t = \beta_0 + \beta_d \log \text{RV}_{t-1} + \beta_w \log \overline{\text{RV}}_{t-5:t-1} + \beta_m \log \overline{\text{RV}}_{t-22:t-1} + \eta_t. \quad (27)$$




Figure 23. NVDA OoS price fan, CHMM-N ($S_0 = 47.67 \rightarrow P_T^{\text{obs}} = 198.07$, a $4.15\times$ rally). The observed path exits the 25–75% interquartile band by mid-horizon but remains inside the 5–95% envelope; horizon-wide coverage under CHMM-N is 0.70 (CHMM-t: 0.64, CHMM-L: 0.67; Table 17), the lone calibration outlier in the panel.




Figure 24. OoS terminal-price distributions under CHMM-N for SPY and NVDA. The right-skew is characteristic of log-returns aggregated to prices; the SPY terminal sits close to the simulated median, while the NVDA terminal sits in the upper-right of the simulated distribution (the realised rally is in the upper tail of the CHMM-N terminal density).

In the absence of intraday bars we use daily squared demeaned returns as the RV proxy; simulation samples $\eta_t \sim \mathcal{N}(0, \hat{\sigma}_\eta^2)$ and rolls forward the recursion, then draws $r_t - \mu = \sqrt{\text{RV}_t} z_t$ with $z_t \sim \mathcal{N}(0, 1)$.

MS-GARCH $K = 3$. A three-regime extension of the Markov-switching GARCH(1,1) ($G_t = \mu + \sigma_t z_t$ with regime-specific recursion $\sigma_t^2(k) = \omega_k + \alpha_k (G_{t-1} - \mu)^2 + \beta_k \sigma_{t-1}^2(k)$). Parameter layout is three regime triples $(\omega_k, \alpha_k, \beta_k)$ plus six transition-matrix logits (softmax-normalised per row, enforcing a non-trivial diagonal); fit by grid-initialised Nelder-Mead. States are canonicalised by ascending unconditional variance, producing the calm / medium / stress decomposition $\sigma = [0.10, 1.77, 10.04]$ on SPY IS.

MS-GARCH $K = 2$ (Haas-Mittnik-Paoletta 2004). The path-independent two-regime variant of Haas et al. [19]. Regime stickiness is high: $p_{11} = 0.915$ (calm) and $p_{22} = 0.547$ (stress), with




Figure 25. JNJ OoS price fan and terminal distribution, CHMM-N (Pipeline A, $K = 18$, 500 simulated paths, $T_{\text{OoS}} = 572$; plotting conventions as above).




Figure 26. JPM OoS price fan and terminal distribution, CHMM-N (Pipeline A, $K = 18$, 500 simulated paths, $T_{\text{OoS}} = 572$; plotting conventions as above).

unconditional standard deviations 0.97 (calm) and 12.67 (stress); the calm regime is near-GARCH behaviour and the stress regime is a rare, transient high-variance mode. Common $\mu = 0.094$, log-likelihood $-4,965$.

Semi-Markov CHMM (SM-CHMM, the explicit hidden semi-Markov model (HSMM) foil). The conceptual foil of the present paper: a hidden semi-Markov extension of the CHMM in which the geometric sojourn distribution of the Markov chain is replaced by a per-state explicit-duration family chosen by AIC from {Pareto, negative binomial, geometric}. Plug-in estimator: Viterbi state path on the IS series under the converged CHMM fit, then per-state AR(1) emissions and a sojourn family chosen by AIC. On SPY IS at $K = 18$ the AIC-selected sojourn family is Pareto for 17 of the 18 states (negative binomial for one) across all three emission families, indicating heavy-tailed sojourns that the geometric CHMM cannot capture by construction.

Table 18 reports the IS / OoS distributional fidelity, kurtosis, ACF-MAE, and Kupiec breach diagnostics for the extended GARCH family, the SM-CHMM rows, and the QuantGAN deep-generative row, alongside the headline CHMM panel. Three structural observations.

First, none of the extended GARCH variants challenges the CHMM family on the joint (KS, ACF) corner. GARCH(1,1)- t is the strongest extended row on KS (57.3% IS, 80.8% OoS) but its IS




Figure 27. AAPL OoS price fan and terminal distribution, CHMM-N (Pipeline A, $K = 18$, 500 simulated paths, $T_{\text{OoS}} = 572$; plotting conventions as above).




Figure 28. QQQ OoS price fan and terminal distribution, CHMM-N (Pipeline A, $K = 18$, 500 simulated paths, $T_{\text{OoS}} = 572$; plotting conventions as above).

excess kurtosis of 15.1 overshoots the observed 7.7 by a factor of two and its ACF-MAE of 0.0316 is below the CHMM's only because the GARCH variance recursion is hard-wired to track conditional volatility; CHMM at $K = 18$ matches GARCH on ACF-MAE within 0.02 while gaining 36 to 36.5 percentage points on IS KS. EGARCH undershoots kurtosis decisively, GJR-GARCH improves on threshold asymmetry but not on the headline gap, and HAR-RV is degenerate on the daily-RV proxy used here (see methods note above).

Second, the SM-CHMM rows are weaker than the headline CHMM rows on every metric in the panel: IS KS 77–80% vs CHMM 93–94%, OoS KS 63–69% vs CHMM 79–86%. The geometric-sojourn restriction of the standard CHMM is therefore not the bottleneck for the headline KS / ACF / kurtosis story. The Viterbi-AR(1) plug-in estimator used here is one specific HSMM construction; a maximum-likelihood explicit-duration HSMM under a Pareto sojourn family may close part of the gap, but the present empirical evidence is that the spectral mechanism of Section 3.8 dominates the sojourn-family extension on this universe and these horizons.

Third, the QuantGAN deep-generative row (Appendix A.25) is the weakest entry in the panel on every distributional metric. IS and OoS KS pass rates collapse to 0% at $\alpha = 0.05$ across all 1,000 simulated paths, IS and OoS AD pass rates are also 0%, and simulated IS excess kurtosis of 2.1 undershoots the observed 7.7 by more than a factor of three. ACF-MAE of 0.0591 is comparable to

Table 18. Extended baseline panel. SPY IS / OoS, $T_{\text{IS}} = 2,516$, $T_{\text{OoS}} = 572$, $N_{\text{paths}} = 1,000$, $\alpha = 0.05$ for KS. Columns: IS / OoS KS pass rate (%); IS / OoS AD pass rate (%); simulated IS excess kurtosis; ACF-MAE on $|G_t|$; pooled-archive Kupiec breach rate (%) and unconditional likelihood-ratio statistic at the 1% and 5% tails.

Model	KS (%) $\uparrow$		AD (%) $\uparrow$		Kurt IS	ACF MAE	1% tail		5% tail	
	IS	OoS	IS	OoS			br%	LR _{uc}	br%	LR _{uc}
GARCH(1,1) Gaussian	28.0	59.4	12.5	59.5	7.0	0.0309	0.9	0.10	4.0	1.23
EGARCH(1,1)	6.3	34.9	4.1	41.7	2.8	0.0430	1.0	0.01	3.5	3.03
GJR-GARCH(1,1)	40.1	57.7	28.7	61.4	7.6	0.0330	1.2	0.27	4.7	0.10
GARCH(1,1)- t	57.3	80.8	36.6	65.3	15.1	0.0316	1.0	0.01	4.9	0.01
HAR-RV	0.0	0.0	0.0	0.0	189	0.0566	0.2	5.99	3.5	3.03
MS-GARCH $K = 2$	27.7	38.7	24.0	44.4	4.7	0.0367	1.0	0.01	4.2	0.82
MS-GARCH $K = 3$	36.1	33.1	28.8	38.0	4.1	0.0284	1.0	0.01	4.2	0.82
CHMM-N ($K = 18$)	93.8	84.2	36.6	73.0	5.0	0.0507	0.7	1.58	4.0	3.83
CHMM-t ($K = 18$)	93.7	85.9	33.4	69.8	14.6	0.0548	1.2	0.58	5.0	3.83
CHMM-L ($K = 18$)	93.2	79.3	30.1	65.4	6.8	0.0565	1.4	3.26	5.6	3.03
SM-CHMM-N ($K = 18$)	79.2	66.7	19.5	57.1	4.7	0.0598	1.1	0.01	5.4	0.82
SM-CHMM-t ($K = 18$)	77.1	63.4	17.2	52.3	13.9	0.0625	1.4	0.10	5.7	1.74
SM-CHMM-L ($K = 18$)	80.4	69.1	21.0	58.9	6.4	0.0581	1.4	3.26	5.5	3.03
QuantGAN	0.0	0.0	0.0	0.0	2.1	0.0591	1.0	0.01	3.3	3.83

the i.i.d. baselines of Table 1 (0.0628), so the GAN does not recover absolute-return autocorrelation either. The convolutional WGAN reproduces some surface-level visual properties of the return series but the stitched-window simulation scheme has no mechanism for modelling either the heavy-tailed marginal or the persistent absolute-return autocorrelation; this matches the documented failure mode of GAN-based generators on financial-time-series volatility-clustering benchmarks [32, 33]. The runner is a repo-native convolutional WGAN with critic weight clipping (Appendix A.25), materially smaller than the original temporal convolutional network of Wiese et al. [9] (a seven-block TCN with 127-day receptive field, hidden width 80, plus Lambert-W input pre-processing that the original authors credit with most of the tail-fidelity); the row therefore reads as the panel’s deep-generative negative control rather than a verdict on the WGAN literature as a whole. The QuantGAN row complements the i.i.d. Gaussian negative control of the headline Table 1.

A.28 Cross-Asset Dependence: Full Pipeline B Panel

A.29 Cross-Asset Dependence Models: Derivations and Diagnostics (Pipeline B)

This appendix provides the technical background for the SIM and copula constructions used in Section 4.5 of the main text. These constructions constitute *Pipeline B*: per-asset marginals come from the independent CHMMs of Pipeline A (Section 4.4) and joint dependence is overlaid on top. Pipeline A stays purely univariate; Pipeline B is the sole multi-asset pipeline in the paper.

Sklar’s theorem and the rank-reordering simulator. Sklar’s theorem [35, 36] states that any d -dimensional joint CDF $H(g_1, \dots, g_d)$ with continuous marginals $F_1, \dots, F_d$ can be written as $H = C(F_1, \dots, F_d)$ for a unique copula C . The rank-reordering simulator exploits this uniqueness. If $\tilde{\mathbf{G}}$ is a matrix whose j -th column is an i.i.d. sample from F_j , and $\mathbf{U}$ is an independent sample from C , then reordering each column of $\tilde{\mathbf{G}}$ so that its ranks match $\mathbf{U}$ ’s produces a sample from H ,

up to discrete approximation error in the empirical ranks. Crucially, reordering never creates new values: each column of the output is a permutation of the corresponding column of $\tilde{\mathbf{G}}$, so marginal distributions are preserved exactly [39].

Kendall’s τ inversion. Given IS returns $\mathbf{R} \in \mathbb{R}^{T \times d}$, we estimate the pairwise Kendall’s τ_{ij} from concordant/discordant pair counts and invert to the correlation scale via the analytic relationship $\rho_{ij} = \sin(\pi\tau_{ij}/2)$, which holds exactly for the Gaussian copula family and is commonly used as a robust correlation estimator for the Student-t copula [38, 47]. The resulting matrix $\hat{\Sigma}$ is symmetric but not guaranteed positive semi-definite; we project to the nearest PSD matrix by clipping negative eigenvalues to 10^{-8} and rescaling the diagonal to unity.

Profile MLE for ν . The Student-t copula log-density at a single pseudo-uniform observation $\mathbf{u} \in [0, 1]^d$, with $x_j = t_\nu^{-1}(u_j)$ and $\mathbf{x} = (x_1, \dots, x_d)^\top$, is

$$\begin{aligned} \log c_t(\mathbf{u}; \Sigma, \nu) = & \log \Gamma\left(\frac{\nu+d}{2}\right) + (d-1) \log \Gamma\left(\frac{\nu}{2}\right) - d \log \Gamma\left(\frac{\nu+1}{2}\right) - \frac{1}{2} \log |\Sigma| \\ & - \frac{\nu+d}{2} \log \left(1 + \frac{\mathbf{x}^\top \Sigma^{-1} \mathbf{x}}{\nu}\right) + \sum_{j=1}^d \frac{\nu+1}{2} \log \left(1 + \frac{x_j^2}{\nu}\right). \end{aligned} \quad (28)$$

Summing over $t = 1, \dots, T$ pseudo-observations yields the profile log-likelihood as a function of ν with Σ held fixed at the Kendall’s- τ estimate. We evaluate equation (28) on the discrete grid $\nu \in \{2, 3, 4, 5, 6, 8, 10, 15, 20, 30\}$ and select $\hat{\nu}$ as the grid point with the largest total log-likelihood. In Section 4.5 the selected value was $\hat{\nu} = 6$ on the six-asset universe.

Sampling the copula. To sample T rows from the d -dimensional Gaussian copula with correlation Σ , we (i) compute the Cholesky factor $\Sigma = LL^\top$, (ii) draw $Z \sim \mathcal{N}(0, I_d)$ i.i.d. and form $Y = LZ$, and (iii) return $U = \Phi(Y)$ componentwise. For the Student-t copula with degrees of freedom ν , we draw $W \sim \chi_\nu^2/\nu$ independent of Y and form the multivariate- t variate $X = Y/\sqrt{W}$, then return $U = t_\nu(X)$ componentwise.

SIM regression summary. Table 19 reports the fitted SIM regression coefficients and goodness-of-fit statistics for the non-market assets in Section 4.5. QQQ’s high $R^2 = 0.840$ reflects that the Nasdaq-100 ETF is essentially a linear combination of large-cap technology names that also drive the S&P 500 index, while JNJ’s lower $R^2 = 0.260$ reflects the limited systematic exposure of low-beta healthcare stocks. These factor loadings and residual magnitudes are what the SIM simulation re-injects at simulation time, and they explain the uneven per-asset KS pass rates reported in Table 21.

Table 19. SIM regression $\hat{G}_{j,t} = \hat{\alpha}_j + \hat{\beta}_j G_{\text{SPY},t} + \hat{\eta}_{j,t}$ for the five non-market assets, fitted on the 2014-01-03 through 2024-01-03 IS window ($T = 2,516$).

Ticker	$\hat{\alpha}$	$\hat{\beta}$	R^2
NVDA	0.321	1.694	0.371
JNJ	0.002	0.576	0.260
JPM	−0.008	1.223	0.507
AAPL	0.111	1.205	0.492
QQQ	0.047	1.122	0.840

A.30 Student-t Copula Profile Log-Likelihood (Pipeline B)

Figure 29 plots the profile log-likelihood of the Student-t copula used inside Pipeline B, computed on the six-asset SPY cross-section over the grid $\nu \in \{2, 3, 4, 5, 6, 8, 10, 15, 20, 30\}$ used in Section 4.5. The curve is smooth and single-peaked, with $\nu^* = 6$ attaining the maximum profile log-likelihood (6157.47), and the neighbouring grid points $\nu = 5$ (6143.37) and $\nu = 8$ (6148.37) each within 15 profile-log-likelihood units of the peak, indicating a shallow-but-clean optimum that agrees with the 4–10 range typically reported for equity-portfolio Student-t copulas in the risk-management literature.




Figure 29. Profile log-likelihood of the Student-t copula used inside Pipeline B on the six-asset SPY cross-section (SPY, NVDA, JNJ, JPM, AAPL, QQQ; IS window 2014-01-03 through 2024-01-03). The vertical axis is the profile log-likelihood evaluated at each degrees-of-freedom value ν on the grid $\{2, 3, 4, 5, 6, 8, 10, 15, 20, 30\}$; the dashed red line marks the profile MLE $\nu^* = 6$ (peak value 6157.47). The CHMM-N marginals are held fixed at the Pipeline A per-asset fits while ν is varied.

A.31 Cross-Asset Correlation Heat Maps (Pipeline B, Companion to Table 21)

Figure 30 visualises the observed and path-averaged simulated correlation matrices for the dependence models of Section 4.5. The panel ordering is SIM, Gaussian copula, Student-t copula (each reported quantitatively in Table 21), plus a truncated C-vine variant with SPY as root for visual comparison. All five panels share the six-asset SPY cross-section and the same CHMM-N marginals at $K = 18$, so any visible differences between the simulated panels reflect only the dependence construction.




Figure 30. Cross-asset IS correlation reproduction across dependence models (Pipeline B, companion to Table 21). Panel (a), top-left: *observed* correlation matrix on 2014-01-03 through 2024-01-03 daily excess log returns for SPY, NVDA, JNJ, JPM, AAPL, QQQ (IS window, $T_{IS} = 2,516$). Panel (b), top-right: path-averaged correlation under the Single Index Model with SPY as market factor. Panel (c), middle-left: path-averaged correlation under the Gaussian copula on CHMM marginals. Panel (d), middle-right: path-averaged correlation under the Student-t copula on CHMM marginals, $\nu^* = 6$. Panel (e), bottom-left: path-averaged correlation under a truncated C-vine copula with SPY as root. Each simulated panel averages over 200 paths of length T_{IS} . All marginals are the per-asset CHMM-N fits at $K = 18$ from Table 3; color scale in $[-1, 1]$, red-white-blue (RdBu) diverging. Closer visual match to panel (a) is better.

A.32 Non-US Asset Class Extension and Per-Pair OoS Off-Diagonal Breakdown

The body universe (SPY, NVDA, JNJ, JPM, AAPL, QQQ) is six US-listed equity tickers from one country-index family. To probe whether the cross-asset construction generalises off the equity cluster we extend the universe with GLD (SPDR Gold Trust), a US-listed ETF whose underlying is physically-backed gold, i.e. a non-equity commodity asset class.

GLD Pipeline-A univariate fidelity. Fitting CHMM-N, CHMM-t, and CHMM-L on the GLD return series at $K = 18$ produces 100% IS KS pass rate for all three families and observed IS excess kurtosis 3.6 matched closely by CHMM-L (2.2); on OoS, however, the KS pass rate collapses to 0% for all three families against an observed OoS excess kurtosis of 6.5. Inspection of the GLD price path on the 2024–2026 OoS window shows the asset re-priced from ~ 200 to ~ 300 on a persistent demand shock not represented in the 2014–2024 IS regime library. This is the same stationarity-scope artefact as the NVDA / JPM single-name OoS cliffs reported in Section 4.4, on a non-equity asset; the limitation is not equity-specific.

Pipeline-B: 7-ticker Student-t copula off-diagonal MAE. Refitting the Student-t copula on the augmented seven-ticker universe at $\nu^* = 6$ gives IS off-diagonal MAE of 0.029 (six-ticker baseline: 0.027) and OoS off-diagonal MAE of 0.179 (six-ticker baseline: 0.211). Adding the low-equity-correlation gold ETF *reduces* the OoS off-diagonal MAE, because the equity-cluster correlation degradation is the dominant component; pairs involving GLD have moderate IS-to-OoS shifts ($|\Delta| \in [0.04, 0.19]$) while the largest equity-cluster pairs (JNJ-QQQ, SPY-JNJ, NVDA-JNJ) have $|\Delta| > 0.48$ on OoS.

Per-pair OoS gap localisation. Table 20 reports the per-pair off-diagonal absolute deviation between observed and path-averaged simulated correlation matrices on IS and OoS, sorted by OoS $|\Delta|$. The OoS gap concentrates in pairs involving JNJ (the defensive single-name proxy): JNJ’s IS correlation with the broad equity factor (QQQ, SPY) collapses on OoS, in some pairs flipping sign, and the IS-fixed copula has no mechanism for tracking that shift. The gold-ETF pairs (SPY-GLD, JPM-GLD, JNJ-GLD, QQQ-GLD) sit in the middle of the ranking with $|\Delta|$ between 0.10 and 0.19 on OoS. The dominant cause of the OoS off-diagonal gap is therefore neither equity-cluster size nor copula-degree mis-specification but a single-name regime shift in the JNJ-equity-factor relationship that the IS-fitted dependence layer cannot anticipate.

A.33 Per-Asset KS Bar Chart (Pipeline B, Sanity Check for Table 21)

Visualisation of the per-asset KS pass rates from Table 21; the SIM-vs-copula gap on JPM and AAPL is the visual headline.

Table 20. Per-pair off-diagonal correlation MAE on the seven-ticker universe (Pipeline B, Student-t copula on CHMM-N marginals at $K = 18$, 200 paths, $\nu^* = 6$). For each ticker pair (i, j) the table reports $|\Sigma_{\text{sim,avg}}[i, j] - \Sigma_{\text{obs}}[i, j]|$ on IS and OoS, ranked by OoS $|\Delta|$. Only the ten largest-OoS-gap pairs are shown for compactness.

Pair	IS $ \Delta $	OoS $ \Delta $	OoS – IS
JNJ-QQQ	0.027	0.544	0.517
SPY-JNJ	0.037	0.509	0.473
NVDA-JNJ	0.017	0.489	0.473
NVDA-AAPL	0.011	0.278	0.268
JNJ-JPM	0.046	0.261	0.216
JNJ-AAPL	0.042	0.244	0.202
AAPL-QQQ	0.003	0.227	0.224
JPM-GLD	0.027	0.187	0.160
SPY-GLD	0.038	0.127	0.089
JNJ-GLD	0.015	0.117	0.101




Figure 31. Per-asset IS KS pass rates after dependence injection (Pipeline B sanity check for Table 21). For each of the six tickers and each of the three dependence constructions (SIM, Gaussian copula, Student-t copula with $\nu^* = 6$), the bar reports the percentage of 200 simulated paths whose marginal distribution passes the two-sample KS test against the observed IS series at $\alpha = 0.05$. The market-factor asset SPY is approximately tied across constructions (it is the SIM engine). For non-market assets the copulas dominate because rank reordering preserves each CHMM marginal exactly; the SIM's affine propagation visibly degrades JPM and AAPL.

A.34 Pre-OoS Validation K -Selection, Walk-Forward Summary, and ν_k Diagnostics

Pre-OoS validation K -selection. The $K = 18$ operating point in the main panel was chosen via a multi-objective criterion combining information-criterion (IC) rank with IS and OoS distributional pass rates, which uses the same 2024–2026 window later used for VaR. As a clean re-selection that decouples the two, we fit CHMM-N on 2014-01-03 through 2021-12-31 ($n = 2,013$) and score every K in the sweep by the forward-algorithm held-out log-likelihood on 2022-01-03 through 2024-01-03 ($n = 502$); the 2024–2026 window is not touched. Held-out log-lik, BIC, HQC, and CAIC all select $K^* = 3$ (per-observation log-lik -2.5345 at $K = 3$ vs. -2.5694 at $K = 18$); AIC selects $K^* = 6$; held-out two-sample KS selects $K^* = 9$. None of the six criteria that avoid the OoS window select $K = 18$. The $K = 18$ choice is best read as a multi-metric one that weights tail fidelity (kurtosis, AD) and downstream VaR behaviour alongside likelihood; a parallel main panel at $K = 3$ would re-rank some rows in CHMM’s favour on ACF-MAE and against it on tail fidelity. The 2022–2023 validation slice is itself partly a structural-break period (the rate-hike cycle), so any model with substantial IS-specific structure generalises worse on it.

ν_k bracket and penalised-ECM rate sweep. A bracket sweep $\nu_{\min} \in \{2.1, 2.5, 3.0, 4.0\}$ shows the CHMM-t fit does not pile against the lower bracket: median $\nu_k = 50$, only 2/18 states near the lower edge. Lifting $\nu_{\min}$ does not reduce the $14.30 \rightarrow 7.69$ overshoot because the fit shifts posterior mass into the newly-recentred tail states. An exponential $1/\nu_k$ shrinkage prior at rate $\lambda \in \{0, 5, 20, 50, 100, 200\}$ reduces simulated excess kurtosis monotonically: 14.30, 11.19, 8.43, 5.41, 3.96, 3.67. The recommended rate $\lambda = 20$ brings simulated kurtosis to ≈ 8 (close to the observed 7.69) at a 1pp IS KS cost. Only the $\nu_{\min}$ state parameter actually moves ($2.1 \rightarrow 4.89$); upper bracket and median remain at 50. The overshoot is therefore a single-state artefact of the two tail regimes hitting the lower bracket under unpenalised ECM.

Cross-asset summary. The full per-construction Pipeline B panel is in Section A.28; for reference, the headline ordering is summarised in Table 21.

Table 21. Cross-asset dependence summary (Pipeline B; CHMM-N marginals at $K = 18$; 200 paths). Median per-asset KS pass rate (%) and off-diagonal MAE of the simulated correlation matrix vs. observed. The Student-t copula at $\nu^* = 6$ is the strongest dependence layer on every IS metric; full per-asset detail is in Section A.28.

	Median IS KS	Median OoS KS	Off-diag MAE IS	Off-diag MAE OoS
Single Index Model	78.5	85.0	0.076	0.261
Gaussian copula	97.8	87.2	0.031	0.202
Student-t copula ($\nu^* = 6$)	96.0	89.8	0.027	0.208