

A Continuous Hidden Markov Model for Daily US-Equity Symmetric Stylized-Fact Reproduction

Abdulrahman Alswaidan*

Cade Jin[†]

Jeffrey D. Varner[‡]

April 30, 2026

Abstract

We evaluate a continuous hidden Markov model (CHMM) trained by Baum–Welch as a synthetic-data generator for daily US-equity returns, on SPY 2014–2026 (body window) and 1994–2006 (cross-decade validation via CRSP), a sector-balanced 30-ticker panel, and a six-asset US-equity copula extension. The empirical scope is the three *symmetric* Cont (2001) stylized facts (heavy-tailed marginals, negligible linear ACF, slow $|G_t|$ ACF); the leverage effect and gain-loss asymmetry are out of scope under the body’s symmetric per-state emissions and only borderline-bracketed via state-mixing on the IS window (Appendix A.7.9). A unified scaffold supports four emission families (Gaussian, Student- t , Laplace, Generalised Error Distribution) sharing identical forward-backward recursions and quantile-based initialisation; the M-step is the only architectural difference. The four families are statistically indistinguishable on OoS sample-CRPS at $T_{\text{OoS}} = 572$ (within-CHMM Diebold–Mariano $p > 0.45$ across all three pairs), so the scaffold contribution is the unified swap interface itself rather than four substantively-distinct points on the OoS distributional axis; the variant choice is driven by per-row IS kurtosis match.

At the state-resolution-robust held-out operating point $K^* = 3$ (selected by held-out per-observation log-likelihood and held-out KS on the strictly pre-2020 slice under both single-fold and four-fold / six-fold rolling-origin CV), the four CHMM variants attain 79.8–90.6% IS / 63.1–83.2% OoS Kolmogorov–Smirnov pass rates on the single SPY OoS window, with CHMM-N at 89.7% IS / 80.5% OoS and the penalised CHMM-t at $\lambda = 20$ providing the cleanest IS / OoS heavy-tail match at simulated excess kurtosis 14.91 / 8.50 against observed 7.68 / 5.29. A $K^* = 6$ block is reported as a sensitivity reference (single-fold pre-2020 selection; not preferred over $K^* = 3$ outside sampling error). The OoS window sits at the upper end of a six-fold rolling-origin walk-forward whose CHMM-N median OoS KS is 67.7% at $K = 18$ and 62.1% at $K^* = 3$ (two stress folds, COVID and 2022 rate-hike, below 10%); single-window panel rankings are walk-forward-fragile and the walk-forward median is the operational summary. ML HSMM dominates CHMM on the single-window distributional axis at every operating point tried: HSMM-N at $K^* = 3$ under Pareto sojourns leads OoS KS at 91.0% (above every CHMM operating point), and the Gamma-sojourn variant at $K = 18$ attains the lowest panel $|G_t|$ ACF-MAE (0.0462 on IS). The CHMM positioning is therefore on the structural use cases (regime-conditional VaR, copula composition, parametric privacy) rather than on raw stylized-fact reproduction.

A regime-conditional Value-at-Risk that propagates the one-step-ahead state forecast through the CHMM mixture passes the Christoffersen joint conditional-coverage test at $\alpha = 0.05$ on the headline OoS window and on 19/24 walk-forward folds; the three persistent walk-forward

*Cornell University, Robert Frederick Smith School of Chemical and Biomolecular Engineering, Ithaca, NY, USA.
Email: aa2725@cornell.edu.

[†]Cornell Ann S. Bowers College of Computing and Information Science, Cornell University, Ithaca, NY, USA.
Email: cj383@cornell.edu.

[‡]Cornell University, Robert Frederick Smith School of Chemical and Biomolecular Engineering, Ithaca, NY, USA.
Email: jdv27@cornell.edu.

rejections concentrate on the W2 (COVID) and W4 (2022 rate-hike onset) stress folds at $p < 10^{-3}$ in each case, so the conditional-VaR value proposition is differentiating only under stationary OoS conditions. At $\alpha = 0.01$ the higher-power Engle–Manganelli (2004) Dynamic Quantile test rejects the conditional-coverage null at $K = 18$ ($p = 0.017$) where the lower-power Christoffersen-cc does not; tighter α levels therefore require either larger T_{OoS} or a model that does not over-couple the high-volatility states at the strict tail. The textbook bilinear ACF identity recasts the SPY instance of the Rydén et al. (1998) low- K failure as an effective-rank statement on the deflated transition matrix: across the 30-ticker cross-section the dominant-mode share has median 75.6% (interquartile range [66%, 86%]) with minimum 32.6% (NEM); the SPY-specific value of 93.6% sits in the right tail of this distribution, so the rank-non-binding statement holds at the cross-ticker median rather than as a universal property of equity-return data. Cross-ticker generalisation gives an OoS KS median of 69.1% at the body headline $K^* = 3$ (75.1% at the $K^* = 6$ sensitivity reference, 73.4% at $K = 18$; identical 11/30 failure count and the same regime-introduction tickers driving the failures across all three operating points), lifted to 83.0% by quarterly refit at $K = 18$; a non-equity GLD/SLV stress test collapses to 0% OoS KS under static IS-fitted CHMM.

Keywords: Continuous Hidden Markov Model, Baum-Welch, Student- t Emissions, Stylized Facts, Volatility Clustering, Value-at-Risk, Synthetic-Data Generator.

1 Introduction

Synthetic generators of equity-return time series support stress testing, regulatory backtesting, scenario design, and the training of downstream machine-learning models on data that cannot be released for privacy or licensing reasons [1–3]. The bar for a useful generator is that it reproduce, simultaneously, the canonical stylized facts of daily returns: heavy-tailed marginals, negligible linear autocorrelation, and slow decay of the autocorrelation function (ACF) of absolute returns [4, 5]. Generalized autoregressive conditional heteroskedasticity (GARCH) models [6, 7] capture volatility clustering but not the multi-regime structure of the marginal; i.i.d. generators match the marginal but destroy temporal persistence; deep generators [8–10] reproduce some stylized facts at the cost of interpretability and reproducibility.

An influential negative result has shaped two decades of work on the regime-switching route. Rydén et al. [11] fitted Gaussian-emission hidden Markov models (HMMs) at $K = 2$ –3 states and concluded that the geometric sojourn-time distribution of a standard Markov chain produces exponential ACF decay too fast to match the observed slow decay; Bulla and Bulla [12] restored ACF fidelity by replacing the Markov chain with a hidden semi-Markov chain at the cost of substantially increased complexity.

We identify the binding axis. The empirical Rydén et al. low- K failure compounds two constraints: a temporal one (the absolute-return ACF must admit slow-decay modes) and a distributional one (the marginal mixture must carry enough components to match observed kurtosis). We apply the textbook closed-form mixture-of-eigenvalues identity $\rho_{|G|}(\tau) = \sum_{k \geq 2} c_k \lambda_k^\tau / \sigma_{|G|}^2$ over the non-unit eigenvalues of $\mathbf{T}$ (folklore in the regime-switching literature; 13–15; restated in Section 3.8 for self-containedness only). The identity gives an algebraic upper bound of $K - 1$ on the number of decay modes through the rank of $\mathbf{T} - \mathbf{1}\bar{\pi}^\top$. An empirical effective-rank diagnostic on the fitted $\hat{\mathbf{T}}$ (Appendix A.6.7) shows this bound is non-binding at $K \geq 3$ at the cross-ticker median on equity-return data: across a 30-ticker panel the dominant-mode share has median 75.6% with interquartile range [66%, 86%] and minimum 32.6% (NEM); the SPY value at $K = 18$ is 93.6%, in the right tail of this distribution. We therefore present the rank-non-binding claim at the cross-ticker median rather than as a universal property; SPY’s lag-1 ACF concentration on a single mode is an asset-specific

extreme rather than a generic feature of equity-return data. The binding axis at $K \geq 3$ is therefore the marginal mixture, whose representational capacity must match the empirical kurtosis target. At $K = 2$ both axes bind simultaneously (mono-exponential ACF that can be tuned, two-component marginal that cannot), and the empirical failure is on the distributional axis; at $K \geq 3$ the temporal axis decouples and the marginal-mixture component count drives the joint fit. The semi-Markov route adds an explicit sojourn distribution to relax the temporal axis, but on equity-return data the temporal axis is already non-binding under standard Baum-Welch at $K \geq 3$, so the operationally informative comparison between the two scaffolds runs on the distributional axis, where the four CHMM emission families and the explicit-duration HSMM at $K^* = 3$ trade KS, kurtosis, and ACF fidelity differently (Section 4).

We instantiate this on ten years of daily SPY ETF data with a 2.5-year held-out window, using a continuous HMM trained by Baum-Welch [16, 17] under quantile-based initialization in four emission variants sharing identical forward-backward recursions: CHMM-N (Gaussian), CHMM-t (per-state Student- t with ν_k updated by ECM golden-section in the tradition of [18, 19]), CHMM-L (Laplace, weighted median and weighted MAD in closed form), and CHMM-GED (Generalized Error Distribution with per-state shape p_k via three-stage ECM, nesting CHMM-N at $p_k = 2$ and CHMM-L at $p_k = 1$). The empirical scope is daily US-equity stylized-fact reproduction: SPY headline, a sector-balanced 30-ticker generalisation panel, and a six-asset US-equity copula extension. A non-equity stress test on GLD and SLV ETFs (Appendix A.8.6) collapses to 0% OoS KS under static IS-fitted CHMM, so the scope claim does not extend to other asset classes without periodic refit; we report this as a documented limitation rather than a deferred follow-up.

The body headline is $K^* = 3$, the state-resolution-robust default selected by held-out per-observation log-likelihood and held-out KS on the strictly pre-2020 slice under both single-fold and four-fold / six-fold rolling-origin CV (Appendix A.9: the k -fold mean of $K^* = 6$ vs $K^* = 3$ has $|z| \leq 0.07$ on held-out per-observation log-likelihood with the sign flipping between fold designs, indistinguishable from sampling noise). At this operating point the four emission variants attain 79.8–90.6% IS / 63.1–83.2% OoS KS on the single OoS window, with CHMM-N at 89.7% IS / 80.5% OoS, simulated kurtosis 3.83/3.53, and $|G_t|$ ACF-MAE 0.0460; the penalised CHMM-t at $\lambda = 20$ provides the cleanest IS / OoS heavy-tail match in the headline block at 14.91/8.50 against observed 7.68/5.29 (the IS overshoot is the per-state ν_k ECM lower-bracket pinning artefact at low K , discussed in Section 5). A $K^* = 6$ block is reported as a sensitivity reference (single-fold pre-2020 selection; same KS / kurtosis trade-off as the headline at one state-resolution step up). The single OoS window sits at the upper end of a six-fold rolling-origin walk-forward whose median CHMM-N OoS KS is 67.7% at $K = 18$ and 62.1% at $K^* = 3$ (Table 70), with two stress folds (W2 COVID 2020 and W4 2022 rate-hike onset) below 10% that every generator in the panel rejects; the walk-forward distribution, not the single-window pair, is the operationally informative summary for production deployment. An extended-state-resolution sensitivity reference at $K = 18$ adds ~ 5 pp IS / ~ 5 pp OoS KS at the cost of moving from a strictly held-out criterion onto a multi-objective rule. Against a benchmark panel (i.i.d. bootstrap of [20; Gaussian and Laplace i.i.d.; GARCH(1,1) of [7 and the heavy-tailed extension of [21; MS-GARCH of [22]), the four CHMM variants at $K = 18$ attain 93.6–95.6% IS and 80.8–85.7% OoS KS, with CHMM-N’s $|G_t|$ ACF-MAE of 0.0509 matching GARCH(1,1)’s 0.0485 to within 0.003 (Table 3). The variants land at four kurtosis points (CHMM-N 5.0, CHMM-L 6.6, CHMM-GED 5.2 with a bimodal Gaussian-bulk / Laplace-tail $\hat{p}_k$ partition, penalised CHMM-t at $\lambda = 20$ at 8.56 IS / 7.07 OoS as the cleanest single-row joint match to observed 7.68 / 5.29; the unpenalised CHMM-t IS kurtosis of 14.4 is the per-state ν_k ECM lower-bracket pinning artefact discussed in Section 5). VaR and Expected-Shortfall envelopes bracket observed values at 1% and 5% on both windows.

Multi-asset synthesis on a US-equity universe composes through a rank-based Student- t copula on

the CHMM marginals: profile MLE selects $\nu^* = 6$ at IS off-diagonal correlation MAE 0.027, but on OoS the Gaussian and Student- t copulas are statistically indistinguishable (0.202 versus 0.209, both above the 0.030–0.027 IS spread by an order of magnitude). The dependence-family choice is therefore an IS-calibration distinction, not an OoS-deployment one; quarterly rolling refit reduces the Student- t OoS off-diagonal MAE to 0.185. On the sector-balanced 30-ticker panel (the sample design is 10 GICS sectors $\times$ 3 large-cap representatives, itself one design choice among many) the IS distribution is tight (median KS 99.5%) and the OoS distribution is sector-stratified (Utilities, Information Technology, and Consumer Staples cluster above 90% median; Health Care collapses to 14.5% on weight-loss-drug regime introductions in LLY and 2024 healthcare-policy compression in UNH; 11/30 tickers below 60% OoS KS at the universe scale), reflecting the well-known stationarity-scope limit on financial time series [23–25] and motivating periodic refit at deployment [26].

Our contributions are empirical and architectural rather than theoretical. (i) We apply the textbook bilinear identity $\rho_{|G|}(\tau) = \sum_{k \geq 2} c_k \lambda_k^T / \sigma_{|G|}^2$ (Hamilton 1994 §22.2; Krolzig 1997 Ch. 3; Timmermann 2000) to the SPY 2014–2024 instance of the empirical Rydén et al. [11] low- K failure and show empirically that the algebraic rank bound on $\mathbf{T} - \mathbf{1}\bar{\pi}^\top$ is non-binding at $K \geq 3$ on this instance: a single non-unit eigenvalue carries 93.6% of the lag-1 absolute-return ACF at $K = 18$ on SPY (Appendix A.6.7). On the 30-ticker cross-section the dominant-mode share has median 0.76 and minimum 0.33 (Appendix A.6.8), so the SPY headline overstates the universality and the body framing is at the cross-ticker median. The substantive content is the empirical effective-rank statement applied to a known instance, not a re-derivation of the textbook identity. (ii) Four emission families (Gaussian, Student- t , Laplace, Generalized Error Distribution) under a unified ECM scaffold, where the M-step is the only architectural difference and CHMM-GED’s per-state shape p_k adapts each state on the Gaussian-to-Laplace axis. (iii) A single-window IS / OoS empirical study on SPY plus a six-fold rolling-origin walk-forward, a sector-balanced 30-ticker generalisation panel (10 GICS sectors $\times$ 3 large-cap representatives), and a six-asset Student- t copula extension; the headline-OoS-KS rankings are walk-forward-fragile and the operationally informative metric is the walk-forward median.

The paper is organized as follows. Section 2 reviews related work; Section 3 describes the data, the continuous HMM, the EM estimator, and the spectral mechanism for the absolute-return ACF; Section 4 presents the empirical study and the VaR / ES envelope calibration; Sections 5–6 discuss and conclude.

2 Related Work

Hidden Markov models in finance. Hamilton [27] introduced Markov-switching to economics; Schaller and Van Norden [28] estimated regime-switching specifications on US monthly stock returns and documented the predictive content of the latent regime for return forecasting. The seminal evaluation against the Cont stylized facts is due to Rydén et al. [11]: Gaussian-emission HMMs at $K = 2$ –3 states fail the slow ACF decay of absolute returns because the geometric sojourn-time distribution of a standard Markov chain produces exponential decay too fast to match the observed signature. Bulla and Bulla [12] restored fidelity with hidden semi-Markov models at substantially higher complexity. Nystrup et al. [29] extends the same line to long-memory volatility. Maximum-likelihood fitting of heavy-tailed innovations follows Peel and McLachlan [18] and Liu and Rubin [19]; the Generalized Error Distribution [30, 31] provides a continuous shape parameter that interpolates between Gaussian ($p = 2$) and Laplace ($p = 1$) and underlies our CHMM-GED variant in the same role that the per-state degrees of freedom ν_k plays in CHMM- t . The discrete-state predecessor of this paper, Alswaidan and Varner [3], reproduced the same three stylized facts via Laplace quantile

binning plus a Poisson-jump duration mechanism; here we replace the discretization step with continuous emissions and recover the same target through the spectral mechanism of Section 3.8.

GARCH and its extensions. GARCH [6, 7] encodes volatility persistence through a single decay parameter but imposes a unimodal innovation distribution and fails distributional Kolmogorov–Smirnov on equity returns. Heavy-tailed innovations [21], asymmetric leverage extensions (EGARCH, 32; GJR-GARCH, 33; threshold GARCH, 34), and regime-switching extensions (MS-GARCH, 22, with the **MSGARCH** R package of 35 the standard reference implementation) close some gaps but inherit the unimodal-innovation constraint. Realised-variance frameworks [36, 37] provide a complementary forecasting baseline.

Deep generative models for financial time series. Generative adversarial network (GAN)-based generators reproduce the visual appearance of return series but typically fail volatility-clustering tests unless architectures encode temporal dependence explicitly [38, 39]; the TimeGAN family of Yoon et al. [8], the convolutional Wasserstein GAN (WGAN) QuantGAN of Wiese et al. [9], and the autoregressive-diffusion approach of Rasul et al. [10] are the standard references. We include QuantGAN as the deep-generative row in our extended panel (Appendix A.7.5); the present paper’s contribution is grounded against parametric and non-parametric baselines.

Cross-asset dependence and synthetic-data evaluation. For multi-asset synthesis, the Single Index Model [40] propagates a market factor through a rank-one linear decomposition but distorts each non-market marginal; Sklar’s theorem [41, 42] supports a cleaner decomposition in which each asset’s fitted marginal is preserved and only the copula is estimated, with Gaussian and Student- t copulas [43, 44] the standard choices and rank-reordering simulation [45] the standard injection. Synthetic-data evaluation is anchored on the framework of Stenger et al. [46], the proper-scoring-rule literature [47–49], kernel two-sample tests [50], and signature-based metrics [51, 52].

3 Methods

3.1 Data and Excess Growth Rates

The single-asset analysis uses daily SPDR S&P 500 ETF (SPY) prices from January 3, 2014 to January 3, 2024 (a ten-year training window, $T_{\text{IS}} = 2,516$), with the period from January 4, 2024 through April 20, 2026 ($T_{\text{OoS}} = 572$) held out. Prices are split-adjusted but not dividend-adjusted; the data sources are described in the acknowledgments. For ticker i and trading day j the annualised excess growth rate is

$$G_{i,j} \equiv \frac{1}{\Delta t} \ln \left(\frac{P_{i,j}}{P_{i,j-1}} \right) - r_f, \quad \Delta t = 1/252, \quad r_f = 0, \quad (1)$$

with $P_{i,j}$ the session volume-weighted average price (VWAP). The factor $1/\Delta t$ is a constant rescaling; all simulations and metrics are evaluated on the daily-scale series. The empirical study is organised along two pipelines (full schematic in Appendix A.2.4). *Pipeline A* fits a CHMM to each ticker independently and evaluates single-asset metrics; *Pipeline B* reuses the per-asset Pipeline-A fits as marginals and injects cross-asset dependence through a rank-based Student- t copula (§3.4).

3.2 Continuous Hidden Markov Model

We model G_t as a continuous HMM $\mathcal{M} = (\mathcal{S}, \mathbf{T}, \mathbf{F}, \boldsymbol{\pi})$ with K latent states, per-state emission density $f_k(x; \boldsymbol{\theta}_k)$ (with corresponding per-state CDF $F_k(x; \boldsymbol{\theta}_k) = \int_{-\infty}^x f_k(u; \boldsymbol{\theta}_k) du$, used in the regime-conditional VaR construction of §4.6), transition matrix $\mathbf{T} \in \mathbb{R}^{K \times K}$, and initial distribution $\boldsymbol{\pi}$ [17]. The stationary marginal is the K -component mixture

$$f(x) = \sum_{k=1}^K \bar{\pi}_k f_k(x; \boldsymbol{\theta}_k), \quad \bar{\boldsymbol{\pi}} = \bar{\boldsymbol{\pi}} \mathbf{T}, \quad (2)$$

with $\bar{\boldsymbol{\pi}}$ the stationary distribution. Existence and uniqueness require irreducibility and aperiodicity of $\mathbf{T}$ (Assumption 1 in §3.8). We learn $\mathbf{T}$, $\boldsymbol{\theta}_{1:K}$, and $\boldsymbol{\pi}$ jointly by EM (§3).

The four emission families share a single forward-backward scaffold and quantile-based initialization; the M-step is the only architectural difference. *CHMM-N*: $f_k = \mathcal{N}(\mu_k, \sigma_k^2)$. *CHMM-t*: $f_k = t_{\nu_k}(\mu_k, \sigma_k)$ with per-state degrees of freedom $\nu_k \in [\nu_{\min}, \nu_{\max}]$. *CHMM-L*: $f_k = \text{Laplace}(\mu_k, \beta_k)$ with per-state Laplace scale denoted β_k to avoid collision with the canonical Laplace-scale symbol b . *CHMM-GED*: $f_k = \text{GED}(\mu_k, \alpha_k, p_k)$ with per-state shape $p_k \in [p_{\min}, p_{\max}]$ that nests Gaussian ($p_k = 2$) and Laplace ($p_k = 1$). The architectural diagram is in Appendix A.2.5.

3.3 Benchmark Generators

We benchmark against representatives from each generator tier, all fit on the same IS window: i.i.d. bootstrap [20] (non-parametric distributional ceiling), Gaussian and Laplace i.i.d., GARCH(1,1) Gaussian [7] and Student- t [21], and Markov-switching GARCH (MS-GARCH; 22). The extended GARCH-family panel (EGARCH, GJR-GARCH, HAR-RV, MS-GARCH at $K \in \{3, 4, 6\}$), an explicit-duration semi-Markov foil, a deep-generative QuantGAN [9] row, plus stochastic-volatility (Taylor-Harvey-Ruiz-Shephard), Markov-switching multifractal [53], and Merton-jump-diffusion [54] baselines are reported in Appendix A.7.7 and A.7.8.

3.4 Cross-Asset Dependence (Pipeline B)

Pipeline B reuses each per-asset CHMM marginal and injects cross-asset dependence through a rank-based Student- t copula [43, 44]. Pseudo-uniform observations $u_{j,t} = \text{rank}(g_{j,t})/(T+1)$ feed an analytic Kendall's- τ correlation estimate $\rho_{ij} = \sin(\pi\tau_{ij}/2)$ [55]; profile MLE selects ν^* . Sklar's theorem [41] guarantees that the rank-reordering simulator of Iman and Conover [45] preserves each fitted CHMM marginal exactly while coupling the asset ranks to the copula sample. SIM, Gaussian-copula, and truncated C-vine comparators are in Appendix A.8.1.

3.5 Evaluation Protocol

For each model we simulate $P = 1,000$ independent paths of length T following [46], under a deterministic global seed root 20260420. Headline metrics are the two-sample Kolmogorov-Smirnov [56, 57] pass rate at $\alpha = 0.05$, mean simulated excess kurtosis, and the absolute-return ACF mean absolute error (MAE) over $L = 252$ lags,

$$\text{ACF-MAE} = \frac{1}{L} \sum_{\tau=1}^L \left| \hat{\rho}_{|G|}^{\text{obs}}(\tau) - \bar{\rho}_{|G|}^{\text{sim}}(\tau) \right|. \quad (3)$$

A separate Value-at-Risk panel (§4.6) carries the unconditional Kupiec [58] and Expected-Shortfall envelope diagnostics. Block-bootstrap KS recalibration, Anderson-Darling, W_1 , Hellinger, and CRPS

auxiliaries are in Appendix A.3. Estimation of $(\mathbf{T}, \{\boldsymbol{\theta}_k\}, \boldsymbol{\pi})$ proceeds by EM in log-space. This section gives the shared scaffold and the family-specific M-step updates; the state-count selection that operationalises this scaffold is reported in Section 4.

3.6 Shared EM Scaffold

All four emission families share the same log-space forward-backward recursions and the same quantile-based initialization [59]: observations are sorted, partitioned into K equal-sized chunks, and each state’s initial location and scale are seeded from the corresponding chunk; $T_{ij}^{(0)} = \pi_k^{(0)} = 1/K$. EM iterates to tolerance $|\mathcal{L}^{(n)} - \mathcal{L}^{(n-1)}| < 10^{-4}$ with `max_iter` = 60; convergence is typically reached in 20–40 iterations. Quantile-based initialization is what makes Baum-Welch reach a non-degenerate $K \geq 3$ fit: random initialization standard in the 1990s often converges to overlapping local optima, which is one of the two compounding sources of the low- K failure observed by Rydén et al. [11] (the other is the spectral one of Section 3.8). Algorithmic details and the forward-backward recursion are in Appendix A.2.

3.7 Family-Specific M-Step Updates

CHMM-N (Gaussian). The classical Baum-Welch M-step [16] is closed form: $\mu_k \leftarrow \sum_t \gamma_t(k) O_t / \sum_t \gamma_t(k)$ and $\sigma_k^2 \leftarrow \sum_t \gamma_t(k) (O_t - \mu_k)^2 / \sum_t \gamma_t(k)$.

CHMM-t (Student-t via ECM). We adopt the ECM formulation of Peel and McLachlan [18] and Liu and Rubin [19], the same heavy-tailed innovation choice studied by Abanto-Valle et al. [60] for stochastic-volatility-in-mean models. The latent precision

$$u_{t,k} = \frac{\nu_k + 1}{\nu_k + ((O_t - \mu_k)/\sigma_k)^2} \quad (4)$$

is treated as an augmented E-step quantity, and the M-step updates

$$\mu_k \leftarrow \frac{\sum_t \gamma_t(k) u_{t,k} O_t}{\sum_t \gamma_t(k) u_{t,k}}, \quad \sigma_k^2 \leftarrow \frac{\sum_t \gamma_t(k) u_{t,k} (O_t - \mu_k)^2}{\sum_t \gamma_t(k)} \quad (5)$$

have closed form conditional on ν_k . The degrees of freedom ν_k are updated by a golden-section search maximizing the per-state Q -function $Q_k(\nu) = \sum_t \gamma_t(k) \log t_\nu(O_t; \mu_k, \sigma_k)$ over the bracket $(\nu_{\min}, \nu_{\max}) = (2.1, 50)$. The two-block CM step preserves the EM monotonicity guarantee on the compact bracket, so the incomplete-data log-likelihood is non-decreasing across iterations.

CHMM-L (Laplace). The weighted maximum-likelihood estimators are in closed form: μ_k is the posterior-weighted median of the observations (weights $\gamma_t(k)$, computed by cumulative-weight bracketing with linear interpolation between neighbouring order statistics), and $\beta_k \leftarrow \sum_t \gamma_t(k) |O_t - \mu_k| / \sum_t \gamma_t(k)$.

CHMM-GED (Generalized Error Distribution via three-stage ECM). For $f_k(x) = \text{GED}(x; \mu_k, \alpha_k, p_k)$ the M-step decomposes into three CM updates per state: a bracketed $\mu_k \leftarrow \arg \min_\mu \sum_t \gamma_t(k) |O_t - \mu|^{p_k}$ via golden-section over $[\hat{\mu}_k - 5\alpha_k, \hat{\mu}_k + 5\alpha_k]$, then the closed-form scale $\alpha_k \leftarrow [(p_k/W_k) \sum_t \gamma_t(k) |O_t - \mu_k|^{p_k}]^{1/p_k}$ with $W_k = \sum_t \gamma_t(k)$, and finally a bracketed $p_k \leftarrow \arg \max_{p \in [p_{\min}, p_{\max}]} \sum_t \gamma_t(k) \log f_k(O_t; \mu_k, \alpha_k, p)$ via golden-section over $[p_{\min}, p_{\max}] = [0.5, 3.0]$. Each CM step is a partial maximization, so EM monotonicity holds on the compact bracket. CHMM-N (all $p_k = 2$) and CHMM-L (all $p_k = 1$) are the boundary cases.

3.8 Spectral Mechanism for the Absolute-Return ACF

The closed-form bilinear identity for the absolute-return ACF of a stationary CHMM, $\mathbb{E}[|G_t| |G_{t+\tau}|] = \mathbf{m}^\top \text{diag}(\bar{\pi}) \mathbf{T}^\tau \mathbf{m}$ with $\mathbf{m}_k = \mathbb{E}[|G_t| \mid s_t = k]$, and its eigendecomposition over the non-unit eigenvalues of $\mathbf{T}$ are folklore in the regime-switching literature (13, §22.2; 14, Ch. 3; 15). Under irreducibility and aperiodicity of $\mathbf{T}$ (Assumption 1, $\bar{\pi}$ unique stationary; 61) and finite per-state second moments with non-trivial location/scale contrast (Assumption 2), the spectral decomposition $\mathbf{T}^\tau = \mathbf{1}\bar{\pi}^\top + \sum_{k=2}^K \lambda_k^\tau \mathbf{v}_k \mathbf{w}_k^\top$ subtracts the marginal product to give the ACF

$$\rho_{|G|}(\tau) = \sum_{k=2}^K w_k \lambda_k^\tau, \quad w_k = (\mathbf{m}^\top \text{diag}(\bar{\pi}) \mathbf{v}_k)(\mathbf{w}_k^\top \mathbf{m}) / \sigma_{|G|}^2. \quad (6)$$

Equation (6) is a real-coefficient sum of geometric decays at the non-unit eigenvalues of $\mathbf{T}$; complex-conjugate λ_k pairs combine into damped oscillatory modes and non-diagonalisable $\mathbf{T}$ adds polynomial-times-geometric prefactors via Jordan blocks. A self-contained derivation, the lag-zero behaviour, and the squared-return ACF are in Appendix A.5; the formal Assumptions 1–2, and the cited propositions on identifiability, MLE consistency, and ECM monotonicity, are in Appendix A.4. Our contribution on this axis is exclusively empirical: applying (6) to recast the Rydén et al. [11] low- K failure on the SPY 2014–2024 instance as a rank statement on $\mathbf{T} - \mathbf{1}\bar{\pi}^\top$ and demonstrating empirically that the implied $K - 1$ rank bound is non-binding at $K \geq 3$.

Rydén separation as a K -rank statement. At $K = 2$, equation (6) reduces to the mono-exponential form $\rho_{|G|}(\tau) = \lambda_2^\tau$ with $\lambda_2 = T_{11} + T_{22} - 1$; slow decay is matched by tuning $T_{11} + T_{22}$. The temporal axis is therefore satisfiable at $K = 2$; the binding axis is the two-component marginal, which cannot match the empirical kurtosis target. At $K \geq 3$ the matrix admits multiple non-unit eigenvalues and the marginal mixture has enough components to match the observed tail. Empirically, ACF-MAE is essentially flat across $K \in \{3, 6, \dots, 21\}$ while KS rises (Section 4); a direct effective-rank diagnostic (Appendix A.6.7) shows a single non-unit eigenvalue carries 93.6% of the lag-1 ACF at $K = 18$ on SPY, so the K -rank upper bound is non-binding at $K \geq 3$ on this instance. Three modes carry $\geq 95\%$ of the lag-1 ACF and twelve modes carry $\geq 99\%$ (Table 1). Repeating the diagnostic on the 30-ticker cross-section (Appendix A.6.8) gives a wider distribution: cross-ticker median dominant-mode share 0.76, interquartile range [0.66, 0.86], minimum 0.33 on NEM. The SPY value is therefore in the right tail of the cross-ticker distribution; the rank-non-binding claim is supported at the cross-ticker median rather than as a universal property of equity-return data.

Table 1. Leading non-unit-eigenvalue contributions to the lag-1 absolute-return ACF, CHMM-N on SPY IS at $K = 18$ (results/diagnostics/spectral_rank.txt). Modes ordered by $|w_k \lambda_k|$ descending. “cum lag-1” is the cumulative fraction of $\sum_{k=2}^K w_k \lambda_k$ accounted for through that rank. The dominant non-unit mode at $|\lambda_2| = 0.929$ carries 93.6% of the lag-1 ACF; the remainder is spread thinly across the lower-magnitude modes.

rank k	$ \lambda_k $	$ w_k $	$ w_k \lambda_k $	cum. lag-1
2	0.929	0.324	0.301	0.936
3	0.335	0.012	0.0039	0.948
4	0.854	0.003	0.0022	0.955
5	0.231	0.008	0.0019	0.961
6	0.231	0.008	0.0019	0.967
$\vdots$				
13	0.119	0.006	0.0007	0.995

4 Empirical Study

Reminder. Subsections 4.1–4.4 and 4.6 operate under **Pipeline A** (single-asset CHMM fit per ticker, no cross-asset coupling); only §4.5 operates under **Pipeline B** (Pipeline A marginals plus a rank-based copula).

4.1 Descriptive Statistics and Stylized Facts

All three Cont stylized facts are present in both windows (Figure 25 in Appendix A.6.31): heavy-tailed marginal (IS excess kurtosis 7.68, Jarque–Bera 6,416.6, $p < 0.001$), negligible linear autocorrelation in raw G_t (Ljung–Box $LB_G \approx 5.5$ at lag 20 vs. $\chi^2_{20}(0.05) = 31.4$), and persistent volatility clustering in $|G_t|$ ($LB_{|G|} = 2,959.3$, $p < 0.001$). The OoS window shares the same qualitative profile (excess kurtosis 5.29). A stationary block bootstrap of Politis and Romano [20] at mean block length $L = 20$ gives 95% CIs of $[2.17, 12.40]$ on the IS excess kurtosis and $[0.90, 8.26]$ on the OoS excess kurtosis (Appendix A.6.20, Table 28); the IS and OoS CIs overlap broadly and the IS-vs-OoS kurtosis difference of ~ 2.4 units lies within sampling error at this block length, so kurtosis-fidelity targets reported throughout this paper should be read against the joint $[0.90, 12.40]$ envelope rather than against the 7.68/5.29 point estimates.

4.2 State-Count Selection

We sweep $K \in \{3, 6, 9, 12, 15, 18, 21\}$ for CHMM-N on the IS window and score each fit on held-out log-likelihood (two slices: estimation 2014–2021 / validation 2022–2023, and the strictly pre-2020 slice estimation 2014-01–2018-06 / validation 2018-07–2019-12), the four information criteria of Nguyen [62] (AIC, BIC, HQC, CAIC), and IS / OoS KS pass rates over 1,000 simulated paths. Table 2 consolidates the chosen K^* across criteria and slices. Absolute-return ACF-MAE is essentially flat (0.047–0.053) across the entire sweep, consistent with the spectral identity (6).

The body headline is $K^* = 3$, the state-resolution-robust held-out-clean default. Selection rule: held-out per-observation log-likelihood and held-out KS on the strictly pre-2020 slice, evaluated under both single-fold (Appendix A.9, “Pre-2020 held-out K -selection”) and four-fold / six-fold rolling-origin CV (Appendix A.9). The k -fold result is the deciding criterion: $K = 6$ vs $K = 3$ has mean held-out log-likelihood difference $|z| = 0.07$ at four-fold full-year cadence and $|z| = 0.04$ at six-fold half-year cadence under independent-fold pooled SEs, and $|z| = 0.90$ / 0.57 under a Diebold–Mariano-style

Table 2. Pre-registered K -selection across criteria and slices. Each cell is the K^* chosen by the row criterion on the column slice. “Pre-2020” is the strictly pre-COVID estimation 2014-01–2018-06 / validation 2018-07–2019-12 split (Appendix A.9); “2022–2023” is the rate-hike-onset slice (Appendix A.6). The single-fold pre-2020 slice’s held-out criteria select $K^* = 6$, but the four-fold and six-fold rolling-origin CV on the same pre-2020 estimation slice (Appendix A.9) cannot distinguish $K^* = 6$ from $K^* = 3$ at conventional levels (mean held-out per-observation log-likelihood $|z| \leq 0.07$ with the sign flipping between fold designs); the body headline is therefore the state-resolution-robust $\mathbf{K}^* = \mathbf{3}$, with $K^* = 6$ retained as a single-fold sensitivity reference. Information criteria favour $K^* = 3$ on the pre-2020 slice, reflecting the parameter-count penalty on a 4.5-year estimation window. No held-out criterion selects $K = 18$.

Criterion	2022–2023 slice	Pre-2020 slice (2018-07 / 2019-12)
Held-out per-obs log-lik	3	6
Held-out KS	9	6
AIC	6	3
BIC	3	3
HQC	3	3
CAIC	3	3

Newey–West HAC variance on the paired diff series across folds (Appendix A.9). Both inferences keep the $K = 6$ vs $K = 3$ comparison well inside conventional sampling noise ($|z| \ll 1.96$), and the sign flips between fold designs. Under that contingency the body is re-headlined at $K^* = 3$; the $K^* = 6$ block is retained as a sensitivity reference. The same k -fold CV reports $K = 18$ as worse than $K = 6$ on mean held-out log-likelihood: the independent-fold $|z| = 1.92$ full-year / 1.70 half-year is borderline, but **under HAC the same statistic moves to $|z_{\text{HAC}}| = 3.56$ (4-fold) and 5.00 (6-fold)**, decisively below $K = 6$. The $K = 18$ block is therefore retained as an *extended-state-resolution sensitivity reference* chosen by a multi-objective rule that uses the OoS window: it improves IS / OoS KS by 1.5–5pp at the cost of being decisively worse than $K = 6$ on held-out per-observation log-likelihood under HAC. The held-out-likelihood / BIC choice $K^* = 3$ on the 2022–2023 slice (Table 8) provides an additional independent corroboration of the $K^* = 3$ headline on a slice that does not overlap the pre-2020 selection. The cross-ticker panel (Table 6) is constructed at $K = 18$ for kurtosis fidelity; the $K^* = 6$ rebuild matches the $K = 18$ KS distribution within sampling error (Appendix A.6.19), and a $K^* = 3$ rebuild of the cross-ticker panel is logged as a follow-up item. Full per- K panel and held-out details are in Appendix A.6 and A.9.

K_{eff} -corrected information-criterion re-rank. The four ICs in Table 2 use the textbook parameter count $p_{\text{nom}} = K(K - 1) + 2K$ in the penalty term. The standardized-distance state-distinctness diagnostic of Appendix A.6.38 reports $K_{\text{eff}} = 11/18$ effective states for CHMM-N at $K = 18$ on SPY IS, so p_{nom} overstates the operationally distinct parameter count whenever multiple states collapse under the single-linkage merge. Sweeping the diagnostic across the entire K -grid on the pre-2020 estimation slice gives $K_{\text{eff}} = K_{\text{nom}}$ at $K \in \{3, 6, 9\}$, $K_{\text{eff}} = 9$ at $K \in \{12, 15, 21\}$, and $K_{\text{eff}} = 11$ at $K = 18$ (Appendix A.6.39, Table 44); the sweep’s apparent K_{nom} blow-up at $K \geq 12$ is therefore largely a parameter-counting artefact. Re-evaluating the four ICs with $p_{\text{eff}} = K_{\text{eff}}(K_{\text{eff}} - 1) + 2K_{\text{eff}}$ in place of p_{nom} in the penalty term re-ranks the lighter-penalty criteria but not the heavier ones: AIC and HQC select $K^* = 21$ under the corrected penalty (vs. 3 under p_{nom}), while BIC and CAIC continue to select $K^* = 3$ under both penalties. The held-out per-observation log-likelihood and held-out KS criteria do not use a parameter-count penalty and are unchanged. The body headline $K^* = 3$ (selected by held-out criteria under k -fold CV, see paragraph above) is unaffected by the IC penalty correction; at $K \in \{3, 6, 9\}$ there is no K_{nom}

vs. K_{eff} distinction in the first place. The IC re-rank is therefore a finding about the $K \geq 12$ parameter-counting regime, not about the body operating point.

4.3 Six-Generator Comparison (Pipeline A)

Table 3 reports IS / OoS validation for the benchmark generators of §3.3 and the four CHMM emission variants at three operating points, against an explicit-duration semi-Markov reference.

At the body-headline operating point $K^* = 3$ (state-resolution-robust under k -fold CV), the four CHMM variants attain 79.8–90.6% IS / 63.1–83.2% OoS KS, with CHMM-N at 89.7% IS / 80.5% OoS KS, simulated kurtosis 3.83 IS / 3.53 OoS, and $|G_t|$ ACF-MAE 0.0460, on par with GARCH (0.0485). The penalised CHMM-t at $\lambda = 20$ provides the cleanest IS / OoS heavy-tail match in the $K^* = 3$ block at 14.91/8.50 against observed 7.68/5.29 (the aggregate IS overshoot reflects the bracket-pinning artefact of the per-state ν_k ECM at low K , discussed in §5; the per-path simulated IS kurtosis distribution has median 7.77 at $K^* = 3$ and 76.6% of paths inside the bootstrap CI on observed (Appendix A.6.20, Table 29), so the headline cleanliness claim is supported as a per-path-median statement rather than as an aggregate-mean statement); CHMM-GED’s per-state shape adapts to 5.48/5.12 as a moderate-tail compromise, and CHMM-L sits at 5.30/4.85 but loses $\sim 10\text{pp}$ on OoS KS relative to the other three families because at $K = 3$ the Laplace per-state shape cannot smooth the low-volatility regime as effectively as the Gaussian. The $K^* = 6$ sensitivity reference attains 88.3–92.6% IS / 76.3–79.0% OoS KS at simulated kurtosis 5.22–5.86 IS, with CHMM-N at $K^* = 6$ matching the OoS observed 5.29 within 0.83 units; the trade-off between $K^* = 3$ and $K^* = 6$ is therefore an OoS-distributional gain (+1.7pp on CHMM-N OoS KS at $K^* = 6$) against a kurtosis-fidelity gain at $K^* = 6$ (+1.43 unit closer to OoS observed). With up to $K - 1 = 2$ non-unit eigenvalues at $K^* = 3$ and 17 at $K = 18$, the mixture-of-eigenvalues sum (6) reproduces the slow absolute-return ACF that the Rydén et al. [11] $K = 2$ –3 Gaussian setup could not recover on the SPY 2014–2024 instance of that experiment (Appendix A.6.17); the additional $|G_t|$ ACF-MAE that $K = 18$ buys over $K^* = 3$ is below sampling-error magnitude on this dataset. The raw-return ACF-MAE column closes the third stylized fact: every CHMM row sits at 0.0234–0.0240, indistinguishable from the i.i.d. baselines that have zero linear ACF by construction. In the $K = 18$ sensitivity reference, the penalised CHMM-t with $1/\nu_k$ shrinkage at $\lambda = 20$ provides the cleanest single-row IS / OoS heavy-tail match at 8.56/7.07 against observed 7.68/5.29 at moderate state resolution. CHMM-GED’s per-state shape $\hat{p}_k$ partitions bimodally into a Gaussian-bulk / Laplace-tail structure that replicates across seeds and tickers (Appendix A.6.33). On the OoS sample-CRPS, the four CHMM variants are statistically indistinguishable on two-sided Diebold and Mariano [49] tests at $T_{\text{OoS}} = 572$ with a Newey–West Bartlett-kernel long-run-variance estimator at default bandwidth $h = \lfloor T_{\text{OoS}}^{1/3} \rfloor = 8$: the within-CHMM pair statistics are $\text{DM}_{\text{N vs t}} = -0.726$ ($p = 0.468$), $\text{DM}_{\text{N vs L}} = -0.750$ ($p = 0.453$), and $\text{DM}_{\text{t vs L}} = -0.044$ ($p = 0.965$), all $p > 0.45$; bandwidth sensitivity over $h \in \{4, 8, 16, 32\}$ leaves the conclusion intact (Appendix A.6.41, Table 47). We read this as evidence that the four families are interchangeable on this evaluation rather than as a per-family strength claim, with the family choice driven instead by the per-row kurtosis match (Table 9, §5). The MS-GARCH benchmark plateaus at 34–36% IS KS across $K \in \{2, 3, 6\}$ in our Nelder-Mead fit, and at 0–0.1% IS / 5–6% OoS KS in the reference Bayesian re-run via the MSGARCH R package of Ardia et al. [35] at $K \in \{2, 3, 4\}$ (posterior-predictive simulation; Table 3 ¶ rows), so the multi-state benefit is specific to the CHMM scaffold rather than to multi-state regime-switching per se on this dataset under either estimator (full higher- K panel and reference-implementation detail: Appendix A.7.7). A six-fold rolling-origin walk-forward (Appendix A.9, Table 70; body summary: Table 4) attains median OoS KS 67.7% at $K = 18$ and 62.1% at $K^* = 3$; the single-window OoS sits at the upper end of these distributions, with two stress folds (W2 COVID, W4 2022 rate-hike

onset) below 10% and explicitly out-of-distribution by KS for every generator in the panel. The walk-forward median, not the single-window OoS pair, is the operationally informative summary for production deployment: the headline single OoS window of 2024–2026 is a relatively low-stress slice (no regime introduction comparable to W2 / W4) and the single-window numbers therefore overstate what a deployed CHMM achieves on average across rolling windows.

ML HSMM as a co-headline result. Replacing the geometric sojourn distribution with an explicit truncated Pareto and refitting through the explicit-duration forward-backward of Yu [63] gives ML HSMM-N as a co-headline scaffold. The OoS KS column at three matched operating points: $K^* = 3$ at 98.4% IS / **91.0%** OoS, $K = 6$ at 95.9% IS / **89.8%** OoS, and $K = 18$ at 0.8% IS / 33.4% OoS (collapse). At $K^* = 3$ and $K = 6$, ML HSMM-N is the strongest single-window OoS KS row in the entire panel, beating every CHMM operating point; the HSMM-vs-CHMM trade-off is therefore real and we frame the headline as “CHMM and ML HSMM are complementary scaffolds on the same Markov backbone, with KS / ACF / kurtosis trade-offs that select between them by use case” rather than as “CHMM is the panel headline”. The HSMM advantage on KS is offset by an absolute-return ACF-MAE that matches the i.i.d. baseline level (0.0629 at $K^* = 3$, 0.0571 at $K = 6$) rather than CHMM-N’s 0.0467 at $K^* = 6$, because the fitted Pareto sojourn concentrates probability mass on a single low-volatility state and the regime-driven slow-ACF mechanism that the CHMM exploits at $K \geq 3$ is muted; the per-state $\bar{\pi}$ concentration is documented in Appendix A.6.34. The Pareto-sojourn HSMM at $K = 18$ collapses to a near-degenerate optimum at 0.8% IS / 33.4% OoS KS; under the Yu (2010) explicit-duration EM at the same $T_{IS} = 2,516$, both KS pass rates collapse and the EM converges in 9 iterations against 19 at $K = 3$. Replacing the truncated Pareto sojourn with a discrete Gamma sojourn (Appendix A.6.35) recovers a non-degenerate fit at $K = 18$ (IS KS 86.0%, OoS KS 80.2%, $|G_t|$ ACF-MAE **0.0462** on IS), and at $K = 3$ trades ~ 20 pp of KS pass rate (98.4 $\rightarrow$ 77.0% IS, 91.0 $\rightarrow$ 80.7% OoS) for an $|G_t|$ ACF-MAE that drops from the i.i.d. baseline level 0.0629 to 0.0528 (Pareto-vs-Gamma comparison: Table 40). The lighter-tailed Gamma sojourn distributes probability mass across sojourn durations more uniformly than the heavy-tailed Pareto, keeping the regime structure non-degenerate at high K and producing the cleanest single-row $|G_t|$ ACF match in the entire HSMM panel. The complementary reading: ML HSMM under truncated Pareto sojourns wins on the marginal-distribution axis at $K \in \{3, 6\}$ (peak 91.0% OoS KS); the Gamma-sojourn variant wins on the joint marginal-plus-volatility-clustering axis at $K = 18$ (the only HSMM operating point where $|G_t|$ ACF-MAE drops below the body CHMM-N $K = 18$ value of 0.0509); CHMM at moderate K remains the right choice when the structural use cases of §4.6–4.5 are in scope. Readers selecting a synthetic-data generator on raw OoS KS alone should prefer ML HSMM-N (Pareto) at $K^* = 3$ or $K = 6$; readers prioritising $|G_t|$ ACF reproduction at high state resolution should prefer ML HSMM-N (Gamma) at $K = 18$. Full HSMM panel and per- K EM-iteration counts in Appendix A.6.34; sojourn-family comparison in Appendix A.6.35.

Bootstrap as a non-parametric synthetic-data benchmark (and how to read Table 3).

The stationary block bootstrap of Politis and Romano [20] attains 99.7% IS / 92.1% OoS KS in Table 3 on the headline single window, comparable to ML HSMM-N at $K^* = 3$ (91.0% OoS) and above every CHMM operating point. The 1-day OoS KS column is therefore the wrong column to read for CHMM-vs-bootstrap differentiation. We report this honestly: on raw 1-day OoS KS at this single window, the simplest non-parametric benchmark dominates the parametric panel, and a reader looking only at that column would not need a CHMM at all. The use-case differentiation the paper actually delivers is therefore not “CHMM beats bootstrap on KS” but rather “CHMM offers parametric controls and structural use cases that the bootstrap cannot serve”. Three differentiating use cases

are exercised in this paper: (i) the regime-conditional Value-at-Risk of §4.6, which propagates the one-step-ahead state forecast through the predictive mixture and passes Christoffersen-cc cleanly on 19/24 walk-forward folds, has no bootstrap analogue at all, since the bootstrap is exchangeable and admits no notion of a forward state filter; (ii) the multi-asset copula composition of §4.5 requires per-asset parametric marginals and is incompatible with a non-parametric bootstrap that only re-emits IS samples; (iii) on the privacy / licensing axis raised in the introduction, every bootstrap path is a permutation of literal IS observations, so the bootstrap is unusable wherever the IS sample is restricted (proprietary feeds, regulatory carve-outs, training-data privacy). The body comparison ordering is therefore: bootstrap is the right choice for raw distributional fidelity on the IS series; ML HSMM-N at $K^* \in \{3, 6\}$ is the right choice when KS is the only objective and parametric structure is not required; CHMM at moderate K is the right choice when one of the three differentiating use cases above is in scope. We avoid framing CHMM as “the headline KS row” on this dataset because Table 3 does not support that framing; the headline is the use-case decision. A natural follow-up question is whether the regime-switching structure adds detectable information on multi-day cumulative-return forecast horizons even when it does not on raw 1-day KS. At $K = 18$, on non-overlapping 20-day cumulative-return blocks ($n = 28$ OoS blocks), CHMM-N attains lower per-block CRPS than the bootstrap with $\Delta\text{CRPS} = -0.180$, $\text{DM} = -2.99$, $p = 0.003$ (two-sided NW-HAC); at $h = 5$ ($n = 114$ blocks) CHMM-L is significant ($\Delta\text{CRPS} = -0.029$, $p = 0.042$); at $h = 1$ neither side dominates ($p > 0.45$). The $h = 20$ result, however, **does not replicate at the body-headline $K^* = 3$ or across the six-asset universe of §4.5**: at $K^* = 3$ on SPY, $\Delta\text{CRPS} = -0.077$, $\text{DM} = -1.16$, $p = 0.244$, and across the six-asset universe at $K^* = 3$ the median ΔCRPS is essentially zero with no per-asset DM significant at $\alpha = 0.05$ (Appendix A.6.11, Table 24). The body’s $K = 18$ multi-day DM result is therefore $K = 18$ -specific on SPY rather than a robust property of the CHMM scaffold across operating points and assets. Full single-asset panel: Appendix A.6.10, Table 23; cross-asset replication: Appendix A.6.11. **The substantive CHMM-vs-bootstrap differentiation is therefore on the structural use cases (regime-conditional VaR, copula composition, parametric privacy), not on multi-day cumulative-return forecast horizons.**

Block-aware OoS KS recalibration (headline reference). The asymptotic two-sample KS critical value $1.36\sqrt{2/T_{\text{OoS}}} = 0.0804$ used in Table 3 assumes i.i.d. samples; under autocorrelated daily returns the appropriate null is the stationary block-bootstrap KS distribution of Politis and Romano [20]. We report the block-aware OoS KS pass rate at mean block length $L = 20$ alongside the asymptotic value for the body headline $K^* = 3$ block, the $K^* = 6$ sensitivity reference, and the $K = 18$ sensitivity reference (Table 5; full $L \in \{5, 10, 20\}$ panel and the OoS-anchored block-bootstrap critical value 0.0647 in Appendix A.7.2). Two takeaways: (i) the cross-generator ordering is preserved (CHMM-t penalised at $\lambda = 20$ at $K = 18$ remains the headline OoS row, the iid bootstrap remains the non-parametric ceiling, GARCH(1,1) the floor); (ii) the absolute level drops by roughly 20–30 percentage points across all rows on OoS, materially larger than the IS recalibration drop of $\sim 5\text{pp}$ (Appendix A.7.2). **Read together with the asymp pass-rates of Table 3, the OoS panel under a temporally-aware null moves CHMM-N from a “passes most of the time” generator to a “passes about half the time” generator at $L = 20$.** The body OoS KS headline at the asymptotic critical value therefore overstates the absolute level relative to a temporally-aware null; the cross-generator ranking is the robust comparison and the absolute-level numbers in the asymp column should be read against the $\sim 25\text{pp}$ OoS recalibration penalty.

Table 3. Headline generator comparison on SPY. 1,000 simulated paths, $\alpha = 0.05$; IS $T = 2,516$, OoS $T = 572$. KS = Kolmogorov–Smirnov pass rate; Kurt = mean simulated excess kurtosis (observed 7.68 IS, 5.29 OoS); ACF-MAE = mean absolute error of the absolute-return / raw-return ACF over 252 lags; CRPS = OoS sample continuous ranked probability score. **Reading guide:** the 1-day OoS KS column is dominated by the simplest non-parametric benchmark (i.i.d. bootstrap at 92.1% OoS); the substantive CHMM-vs-bootstrap differentiation is on the *structural* use cases of §4.6 (regime-conditional VaR Christoffersen-cc) and §4.5 (multi-asset copula composition), not on the per-day or multi-day distributional axis (the body’s $h = 20$ DM result at $K = 18$ does not replicate at the body-headline $K^* = 3$ or across the six-asset universe at $K^* = 3$; see Appendix A.6.11). The $K^* = 3$ block is the body headline (marked with leading $\star$), selected as the state-resolution-robust default by four-fold and six-fold rolling-origin CV on the strictly pre-2020 slice (Appendix A.9: $K^* = 6$ is not preferred over $K^* = 3$ outside sampling error). The $K^* = 6$ block is reported as a sensitivity reference (single-fold pre-2020 selection); the $K = 18$ block is an extended-state-resolution sensitivity reference selected by a multi-objective rule that uses the OoS window. Bold within the $K^* = 3$ headline block marks the cleanest IS / OoS row on each axis; bold within the $K = 18$ sensitivity block marks column-best. [†] Unpenalised CHMM-t reference. ^{††} Bracket-lift ablation at $\nu_{\min} = 4$ and no $1/\nu_k$ penalty; see Remark 1 in §5 for the substantive reading. The OoS-calibrated kurtosis match goes to CHMM-N at $K^* = 6$ (sim 4.46 vs. observed 5.29); the IS-calibrated heavy-tail match goes to penalised CHMM-t at $\lambda = 20$, $K = 18$. Block-bootstrap KS recalibration (the temporally-aware variant) is in Appendix A.7.2; MS-GARCH at $K = 3, 4, 6$ and the extended GARCH family are in Appendix A.7.7. [‡] Reference-implementation re-run of MS-GARCH via the MSGARCH R package of Ardia et al. [35] (version 2.51) under Bayesian estimation (12,500 MCMC draws, 2,500 burn-in, thin 10); 1,000 posterior-predictive paths (one path per retained posterior draw). The lower KS pass rate against our in-house Nelder-Mead rows reflects posterior-predictive integration of parameter uncertainty path-by-path; the frequentist rows generate all paths from one MLE point estimate. Pinned versions (R 4.6.0, MSGARCH 2.51) and full posterior summaries are recorded in `r_msgarch/renv.lock` and `results/msgarch_reference/` of the companion code repository. [§] QuantGAN TCN is the Wiese-style 7-block dilated convolutional WGAN-GP rebuild (architecture-only, no Lambert-W input preprocessing); the 3-conv-layer WGAN baseline of the body’s earlier draft is in Appendix A.7.5. Both in-house WGAN re-implementations (3-conv and 7-block TCN, with and without Lambert-W input preprocessing per Appendix A.7.5) fail IS / OoS distributional KS on this dataset. We do not claim that the published reference implementation of Wiese et al. [9] would also fail; a faithful reference-implementation re-run is a deferred follow-up. The QuantGAN row in this panel is therefore an in-house WGAN re-implementation result, not a verdict on the broader deep-generative literature.

Model	KS (%) $\uparrow$		Exc. Kurt		ACF-MAE $\downarrow$		CRPS OoS $\downarrow$
	IS	OoS	IS	OoS	$ G_t $	G_t	
<i>Observed</i>	–	–	7.68	5.29	–	–	–
Bootstrap	99.7	92.1	7.24	6.81	0.0628	0.0235	1.0398
Gaussian i.i.d.	0.0	1.0	–0.01	–0.03	0.0627	0.0236	1.0611
Laplace i.i.d.	99.1	88.5	2.95	2.91	0.0628	0.0236	1.0386
GARCH(1,1)	23.4	60.8	7.06	2.96	0.0485	0.0243	1.0440
GARCH(1,1)-t	57.3	80.8	15.13	–	0.0316	0.0173	–
MS-GARCH ($K = 2$)	27.7	38.7	4.73	–	0.0367	0.0173	–
MS-GARCH ($K = 3$)	36.1	33.1	4.10	–	0.0284	0.0173	–
MS-GARCH ($K = 6$)	34.5	33.4	4.41	–	0.0429	0.0173	–
MS-GARCH ref. Bayesian ($K = 2$) [‡]	0.0	5.8	4.00	2.46	0.0465	0.0240	–
MS-GARCH ref. Bayesian ($K = 3$) [‡]	0.1	5.1	4.52	2.53	0.0433	0.0241	–
MS-GARCH ref. Bayesian ($K = 4$) [‡]	0.0	5.3	6.01	3.54	0.0446	0.0241	–
QuantGAN TCN [§]	0.0	0.0	0.56	0.53	0.0617	0.0264	–
<i>Body headline: reference panel at $K^* = 3$ (state-resolution-robust under k-fold CV; Appendix A.9).</i>							
$\star$ CHMM-N ($K^* = 3$)	89.7	80.5	3.83	3.53	0.0460	0.0240	1.0393
$\star$ CHMM-t pen. ($\lambda = 20$, $K^* = 3$)	90.6	83.2	14.91	8.50	0.0537	0.0235	1.0373
$\star$ CHMM-L ($K^* = 3$)	79.8	63.1	5.30	4.85	0.0532	0.0236	1.0405
$\star$ CHMM-GED ($K^* = 3$)	90.5	77.4	5.48	5.12	0.0526	0.0236	1.0389
<i>Held-out-clean sensitivity reference at $K^* = 6$ (single-fold pre-2020 selection; not preferred over $K^* = 3$ outside sampling error).</i>							
CHMM-N ($K^* = 6$)	92.6	178.3	5.26	4.46	0.0502	0.0238	1.0422
CHMM-L ($K^* = 6$)	88.3	76.3	5.86	–	0.0510	0.0235	1.0399
CHMM-GED ($K^* = 6$)	91.8	79.0	5.22	–	0.0518	0.0235	1.0393
ML HSMM-N ($K^* = 3$)	98.4	91.0	3.46	3.38	0.0629	0.0236	–

Table 4. Six-fold rolling-origin walk-forward summary (CHMM-N, 5y train / 1y test, 500 paths per fold). Pulled from the full panel in Appendix A.9, Table 70. The headline single-window OoS sits at the upper end of the fold-wise distribution; the W2 (COVID) and W4 (2022 rate-hike onset) folds are out-of-distribution by KS for every generator in the panel.

	median KS (%)	$[Q_1, Q_3]$ KS (%)	min KS (%)	folds < 10%
$K^* = 3$	62.1	[7.2, 78.4]	0.8	2/6
$K = 18$ (sensitivity)	67.7	[8.2, 75.0]	0.0	2/6

Table 5. Block-aware OoS KS recalibration at $L = 20$ (stationary block bootstrap of 20, $B = 1,000$, 500 OoS-length simulated paths per generator, seed 20260422). Side-by-side with the asymptotic OoS pass rate of Table 3; the block-bootstrap 95% null KS critical value at $L = 20$ is 0.0647 versus the asymptotic 0.0804. Body headline $K^* = 3$, sensitivity reference $K^* = 6$, and extended-state-resolution sensitivity reference $K = 18$ are all reported. Full $L \in \{5, 10, 20\}$ panel: Appendix A.7.2.

Model	OoS KS asymp. (%)	OoS KS block $L = 20$ (%)
Bootstrap	90.4	73.2
GARCH(1,1)	59.2	31.4
★ CHMM-N ($K^* = 3$)	79.8	58.6
★ CHMM-t pen. ($\lambda = 20$, $K^* = 3$)	79.8	54.2
★ CHMM-L ($K^* = 3$)	65.0	33.4
★ CHMM-GED ($K^* = 3$)	79.2	51.6
CHMM-N ($K^* = 6$)	80.0	53.2
CHMM-t pen. ($\lambda = 20$, $K^* = 6$)	81.2	56.4
CHMM-L ($K^* = 6$)	76.6	50.2
CHMM-GED ($K^* = 6$)	82.0	54.6
CHMM-N ($K = 18$)	81.0	56.4
CHMM-t pen. ($\lambda = 20$, $K = 18$)	84.8	61.6
CHMM-L ($K = 18$)	77.4	49.4
CHMM-GED ($K = 18$)	81.0	58.4

4.4 Cross-Ticker Generalization (Pipeline A)

Table 6 reports the penalised CHMM-t aggregate IS and OoS KS distribution across a sector-balanced 30-ticker panel (10 GICS sectors $\times$ 3 large-cap representatives) at three operating points (the body headline $K^* = 3$, the sensitivity reference $K^* = 6$, and the extended-state-resolution sensitivity reference $K = 18$), all at $\lambda = 20$. The three operating points are within ~ 6 pp of each other on the headline distributional axis: OoS KS median is 69.1% at $K^* = 3$, 75.1% at $K^* = 6$, and 73.4% at $K = 18$, with identical 11/30 failure counts at all three operating points (and the same regime-introduction tickers driving each: LLY 8.7% at $K^* = 3$, UNH 16.5%, NEM 5.0%). They differ on the kurtosis-residual axis: per-ticker median residual rises sharply at low state resolution (7.2 at $K^* = 3$, 10.3 at $K^* = 6$, 0.6 at $K = 18$), reflecting that lower state-resolution cannot smooth the per-state ν_k shrinkage at uniform λ ; the production recipe for kurtosis-fidelity consumers is therefore to deploy at $K = 18$, while KS-only consumers can pick any of the three at parity. At $K = 18$, the IS distribution is tight (median 99.5%, $[Q_1, Q_3] = [99.3, 99.6]\%$); the OoS distribution is wider with median 73.4%, mean $66.8 \pm 29.5\%$, and $11/30 = 36.7\%$ of tickers below 60%, concentrated in single-name regime introductions (LLY 7.6%, UNH 14.5%, NEM 5.6%). A one-way ANOVA on OoS KS by sector at $n = 3$ per sector returns $F(9, 20) = 0.44$ at $p = 0.90$ with $\eta^2 = 0.16$ (Appendix A.6.18). With $n = 3$ tickers per sector this design is severely underpowered to detect sector effects of moderate size: $\eta^2 = 0.16$ corresponds to a moderate-to-large effect on the conventional Cohen scale and the reported F -statistic does not have the power to reject the no-sector-effect null at this sample size. **An $n = 6$ per-sector expansion to a 60-ticker panel** (Appendix A.6.5, Table 18) tests this directly: the ANOVA at adequate power is $F(9, 50) = 0.37$, $p = 0.946$, $\eta^2 = \mathbf{0.062}$, with the effect-size estimate dropping by more than half relative to the underpowered $n = 3$ value. The 60-ticker aggregate distribution is statistically indistinguishable from the 30-ticker version (median OoS KS 73.4% vs. 73.4%, mean 67.0% vs. 66.8%, failure rate $22/60 = 36.7\%$ vs. $11/30 = 36.7\%$). The visible concentration of failures on regime-introduction tickers (LLY, UNH, NEM in the body 30-ticker panel; ORCL, EOG, T, NEE in the additional 30-ticker block) is therefore a per-ticker observation rather than evidence about sector-level structure — and the conclusion is now supported by an adequately-powered test. A quarterly-refit version of this panel at $K = 18$ (Appendix A.6.30) shifts the OoS KS distribution from median 73.4% $\rightarrow$ **83.0%**, mean 66.8% $\rightarrow$ 77.2%, and reduces the failure count from 11/30 to 7/30, with the largest gains on regime-introduction tickers (LLY 7.6% $\rightarrow$ 83.7%, HD 41.4% $\rightarrow$ 82.2%); the same quarterly-refit construction at the $K^* = 6$ sensitivity reference is reported in the same appendix. A monthly-refit version (Appendix A.6.44) interpolates the trade-off further: median OoS KS **86.7%**, mean 78.8%, and 5/30 failures, so each shorter refit window yields a roughly proportional further reduction in failure count without saturating.

Per-ticker $\hat{\lambda}^*$ tuning. The body uniform $\lambda = 20$ is the right rate on the heavy-tailed defensives (NVDA, JNJ) but is over-shrunk on the moderate-tail tickers: a per-ticker sweep $\lambda \in \{0, 5, 10, 20, 50, 100\}$ on the six body tickers under a 1.5pp KS-degradation tolerance (Appendix A.9, Table 73) selects $\lambda^* \in \{0, 10, 20\}$, with QQQ at $\lambda = 0$ (no penalty) and SPY/JPM/AAPL at $\lambda = 10$. The honest operational reading is that $\lambda = 20$ is a reasonable uniform default but per-ticker tuning is recommended whenever a ticker’s residual kurtosis at $\lambda = 0$ is within ~ 1 unit of observed; the production recipe is to evaluate the kurtosis-residual / KS-degradation trade-off on a held-out validation slice and pick the smallest λ that brings simulated kurtosis within 1.5 units of observed.

Cross-decade validation. A CRSP-sourced cross-decade split (1994–2004 IS / 2004–2006 OoS; full panel in Appendix A.6.46) tests whether the body’s stylized-fact reproduction transfers across non-overlapping decades. The IS axis is decade-robust: at $K \in \{3, 18\}$ CHMM-N attains 84–90% IS

Table 6. Cross-ticker generalisation, sector-balanced 30-ticker panel (penalised CHMM-t at $\lambda = 20$, 1,000 paths, seed = 20260420). Aggregate distribution across $n = 30$ tickers at three operating points: the body headline $K^* = 3$, the sensitivity reference $K^* = 6$, and the extended-state-resolution sensitivity reference $K = 18$. The KS distributions are within ~ 6 pp of each other and at identical 11/30 failure count at all three operating points (driven by the same regime-introduction tickers); the kurtosis residual widens at lower state-resolution because the per-state ν_k shrinkage at uniform λ cannot be smoothed across fewer states. Full per-ticker panel and per-sector rollup in Appendix A.6.4, Table 17; $K^* = 6$ and $K^* = 3$ rebuild details in Appendix A.6.19. The 10-sector $\times$ 3-large-cap sample design is one design choice among many; an equally weighted random draw of 30 large-caps would land on a different OoS tail.

Metric	$K^* = 3$ (headline)	$K^* = 6$ (sensitivity)	$K = 18$ (sensitivity)
IS KS pass rate (%) median	96.8	98.8	99.5
IS KS pass rate (%) mean $\pm$ s.d.	95.1 ± 4.9	97.7 ± 3.7	99.3 ± 0.6
OoS KS pass rate (%) median	69.1	75.1	73.4
OoS KS pass rate (%) mean $\pm$ s.d.	66.2 ± 28.2	66.5 ± 29.2	66.8 ± 29.5
$ G_t $ ACF-MAE median	0.0399	0.0403	0.0416
Kurtosis residual (sim – obs) median	7.18	10.34	0.61
Tickers OoS KS < 60% (IS-fixed)	11/30	11/30	11/30

KS pass rate against the 90–94% obtained on the body 2014–2024 IS window. The 2004–2006 OoS pass rate, however, collapses to 3–5% because the calm post-dot-com bull-market slice has observed excess kurtosis 0.06 (essentially Gaussian) against the 1994–2004 IS kurtosis of 3.05. The cross-decade pattern matches the W2 (COVID 2020) and W4 (2022 rate-hike) walk-forward stress folds (Table 4, both with OoS KS < 10%): the IS axis transfers cleanly across decades, but a static-IS-fitted CHMM cannot bracket OoS slices whose marginal moves outside the IS regime library. The operational implication parallels the cross-ticker panel above: the static-fit OoS pass rate is OoS-slice-specific and the deployment recipe is the periodic-refit construction (the cross-ticker quarterly-refit panel in Appendix A.6.30 and the rolling-copula construction of §4.5), not the single-window static fit.

4.5 Cross-Asset Dependence (Pipeline B): Student-t Copula

Pipeline B reuses the per-asset CHMM-N marginals from §4.4 and adds a rank-based Student- t dependence layer (§3.4). The body cross-asset construction is scoped to a US-equity universe: profile MLE on six US-equity assets (SPY, NVDA, JNJ, JPM, AAPL, QQQ; 200 paths) selects $\nu^* = 6$ on the unit-spaced grid $\nu \in [4, 12]$ with Wilks 95% profile-LL CI of $[6, 7]$ *on the IS slice* (Appendix A.8); a half-unit-grid refinement plus $B = 200$ parametric bootstrap CI confirms $[6, 7]$ (Appendix A.8.7). A non-equity stress test on GLD and SLV ETFs over the same windows (Appendix A.8.6) shows OoS KS collapses to 0% on both ETFs under a static IS-fitted CHMM, confirming that the cross-asset claim does not extend to non-equity asset classes without periodic refit. Table 7 reports per-asset KS and the off-diagonal correlation MAE: 0.027 on IS for the Student- t copula, against 0.030 Gaussian, 0.068 truncated C-vine, and 0.076 SIM (full panel in Appendix A.8.1). **On OoS the dependence-family choice is empirically null at this N_{paths} :** off-diagonal MAE is 0.209 for Student- t and 0.202 for Gaussian, a 0.007 difference below the simulation-noise floor. The IS calibration distinction (0.027 vs. 0.030, 0.003 difference) is an IS-only finding and the dependence-family selection on OoS is not supported at this N_{paths} . The operational story on the OoS axis is the quarterly-rolling refit, which reduces the OoS off-diagonal MAE from 0.209 to **0.185** on the six-asset universe (Appendix A.8.5); the dependence-family choice is dominated by the refit-cadence choice on OoS. One additional caveat on the dependence-layer estimator: the body ν^* is selected by profile MLE conditional on $\hat{\Sigma}$ obtained by Kendall’s- τ inversion ($\rho_{ij} = \sin(\pi\tau_{ij}/2)$, exact under the Gaussian copula and

Table 7. Cross-asset Student-t copula on CHMM-N marginals (Pipeline B; $\nu^* = 6$; 200 paths). Per-asset KS pass rate at $\alpha = 0.05$ and correlation reproduction averaged over paths. Comparison against SIM, Gaussian, and truncated C-vine in Appendix A.8, Table 74.

Ticker	KS (%) $\uparrow$	
	IS	OoS
SPY	96.0	79.5
NVDA	98.5	53.5
JNJ	99.0	97.5
JPM	99.5	48.0
AAPL	99.0	93.5
QQQ	95.5	92.0
Median	98.8	85.8
Tickers OoS KS < 60%	0/6	2/6
Off-diag MAE IS	0.027	
Off-diag MAE OoS	0.209	
Off-diag MAE OoS (quarterly refit)	0.185	

approximate under the Student- t copula); a full one-shot MLE that jointly maximises over (Σ, ν) on the same six-asset universe (Appendix A.7.10) returns $\hat{\nu}_{\text{full}} = 6.40$ versus $\hat{\nu}_{\text{two-step}} = 6.00$, a +5.55 log-likelihood improvement, with all 15 off-diagonal correlation differences below $|\Delta\rho_{ij}| \leq 0.025$. The two-step estimator is therefore not detectably biased toward the Gaussian limit on this universe; the body $\nu^* = 6$ choice is robust to the estimator switch and remains inside the Wilks 95% profile-LL CI of the full MLE.

The cross-asset median OoS KS of 85.8% obscures a bimodal distribution: NVDA (53.5%) and JPM (48.0%) fail OoS KS at the standard threshold, while the four remaining assets pass cleanly above 79%. The 2/6 failure rate at the body $\nu^* = 6$ recovers to 0/6 under the quarterly-rolling refit recipe of Appendix A.8.5; per-asset failure attribution is driven by single-name regime shifts in the OoS slice (NVDA AI-rally regime, JPM bank-stress regime) rather than by the dependence-layer estimator, consistent with the pattern in the 30-ticker univariate panel of §4.4.

4.6 Value-at-Risk and Expected-Shortfall Back-Test (Pipeline A)

The headline of this subsection is the regime-conditional Christoffersen-cc result: a VaR threshold built from the CHMM one-step-ahead state forecast passes the joint conditional-coverage test cleanly at every (K, α) on OoS, where every constant-across- t generator (CHMM unconditional, GARCH, bootstrap) rejects breach independence. We organise the diagnostics in three layers: (i) envelope bracketing, (ii) unconditional Kupiec coverage, (iii) the regime-conditional construction.

(i) Envelope bracketing. All four CHMM variants bracket observed historical VaR and ES at 1% and 5% within the [5, 95]% envelope on both windows; the IS CHMM-N median $\text{VaR}_{0.01}$ at -6.64 (observed -6.38) is the closest point estimate among the six benchmark generators, with envelopes roughly $1.7\times$ tighter than GARCH’s. The full envelope panel for Bootstrap, GARCH, CHMM-N/-t/-L/-GED at 1% and 5% is in Appendix A.6.36, Table 41.

(ii) Unconditional Kupiec coverage. The four CHMM rows pass the unconditional Kupiec test on both windows: at $\alpha = 0.05$ on OoS, CHMM-N $\text{LR}_{\text{uc}} = 3.83$ against $\chi_1^2(0.05) = 3.841$,

with the other three CHMM variants in 3.03–3.83. The repeated 3.83 across rows is a feature of $T_{\text{OoS}} = 572$: the Kupiec statistic is a function of the integer breach count and at $\alpha = 0.05$ the values $k \in \{14, 19, 20\}$ produce statistics that pile against the critical value rather than spreading continuously. We treat unconditional Kupiec as a sanity check at this T_{OoS} and the conditional construction below as the substantive risk-management diagnostic.

(iii) Regime-conditional Christoffersen-cc (*headline*). At each OoS day t we run the forward filter through $\mathcal{F}_{t-1} = R_{\text{IS}} \cup R_{\text{OoS}}[1:t-1]$ to get the one-step-ahead state forecast

$$\mathbb{P}(s_t = k \mid \mathcal{F}_{t-1}) = \sum_i \mathbb{P}(s_{t-1} = i \mid \mathcal{F}_{t-1}) T_{ik}, \quad (7)$$

and define $\widehat{\text{VaR}}_t(\alpha) = F_t^{-1}(\alpha)$ where $F_t(\cdot) = \sum_k \mathbb{P}(s_t = k \mid \mathcal{F}_{t-1}) F_k(\cdot; \boldsymbol{\theta}_k)$ is the family-appropriate predictive mixture CDF under IS-fixed $(\mathbf{T}, \{\boldsymbol{\theta}_k\})$, with F_k the per-state Gaussian, Student- t , Laplace, or GED CDF (the body Gaussian construction reduces to $F_k = \Phi(\cdot; \mu_k, \sigma_k)$). Table 8 reports the back-test for CHMM-N at $K^* = 3$ and $K = 18$ and the penalised CHMM-t at $\lambda = 20$ on OoS; every row passes Kupiec, Christoffersen-ind, and Christoffersen-cc cleanly at $\alpha = 0.05$. The improvement in the independence statistic from the unconditional variant ($\text{LR}_{\text{ind}} = 5.26$ at $K = 18, \alpha = 0.05$) to the conditional variant ($\text{LR}_{\text{ind}} = 0.52$) is the regime-switching value proposition that the unconditional Kupiec headline did not exercise. The conditional pass extends to every emission family with the family-appropriate predictive density (CHMM-L, CHMM-GED) at $p_{\text{cc}} \geq 0.089$ throughout, so the regime-switching value proposition is intrinsic to the latent-state forecast and not specific to the Gaussian emission family (Appendix A.3.2, Table 71).

Power caveat at $\alpha = 0.01$. A Monte Carlo power calibration at $T_{\text{OoS}} = 572$ (Appendix A.6.23, Table 32, $B = 5,000$ replicates per cell) shows the Christoffersen-cc test is correctly sized at 4–5% Type-I error under the i.i.d. null at both α levels, but the power profile is asymmetric. At $\alpha = 0.05$ (~ 28.6 expected breaches) the test reaches $\geq 80\%$ power against breach-clustering eigenvalues $\rho \geq 0.20$; at $\alpha = 0.01$ (~ 5.7 expected breaches) it reaches $\geq 80\%$ power only at $\rho \geq 0.50$, with rejection rate 42.6% at $\rho = 0.20$. The body $\alpha = 0.05$ “clean pass at every (K, α) ” rows of Table 8 therefore carry strong power against moderate clustering; the $\alpha = 0.01$ rows should be read as “not detectably worse than a strongly-clustered alternative” rather than “calibrated against any plausible alternative”. The Engle–Manganelli (2004) Dynamic Quantile (DQ) test as a higher-power conditional-coverage alternative (Appendix A.6.24, Table 33) is the substantive $\alpha = 0.01$ test on this T_{OoS} . At $\alpha = 0.05$ the DQ test passes cleanly at both $K = 3$ ($p = 0.156$) and $K = 18$ ($p = 0.678$), confirming the Christoffersen-cc reading at this α . **At $\alpha = 0.01$ the DQ test rejects conditional coverage at $K = 18$ ($p = 0.017$), while Christoffersen-cc does not reject at the same row ($p = 0.137$).** DQ has the higher power of the two tests at this T_{OoS} (and is the test we trust at $\alpha = 0.01$); the substantive read is therefore that the regime-conditional VaR *over-couples high-volatility states at the strict tail at $K = 18$* on this OoS slice, and that the $\alpha = 0.01$ Christoffersen-cc pass is a power-driven non-rejection rather than a calibrated pass. The risk-management diagnostic in this paper is therefore the $\alpha = 0.05$ row under both tests (clean pass on every (K, family) row); the $\alpha = 0.01$ row at $K = 18$ should be read as a DQ-rejection of conditional coverage that tighter α levels require either larger T_{OoS} or a model rebuild to address.

Walk-forward stability across six rolling-origin folds is in Table 72 (Appendix A.9); the four-family extension at the headline OoS window is in Table 71 (Appendix A.3.2); a quarterly-refit version (CHMM-N refit every 63 trading days on a rolling 5-year window) is in Table 35 (Appendix A.6.26). The conditional Christoffersen-cc passes at $\alpha = 0.05$ on 19/24 walk-forward rows, with the failures

Table 8. Regime-conditional VaR back-test, CHMM-N and penalised CHMM-t ($\lambda = 20$) on OoS ($T_{\text{OoS}} = 572$, seed 20260420). Critical values: $\chi_1^2(0.05) = 3.841$, $\chi_2^2(0.05) = 5.991$. The forward filter is run under IS-fixed parameters ($\mathbf{T}, \{\theta_k\}$) (no refit); only the one-step-ahead state-probability vector is updated through the OoS observations via Eq. (7). Every (K, α , family) row passes Kupiec, Christoffersen-ind, and Christoffersen-cc cleanly at $\alpha = 0.05$. $^\dagger \alpha = 0.01$ rows are power-bounded at $T_{\text{OoS}} = 572$: Christoffersen-cc reaches $\geq 80\%$ power against breach-clustering eigenvalues only at $\rho \geq 0.50$ (rejection rate 42.6% at $\rho = 0.20$), so a clean pass should be read as “not detectably worse than a strongly-clustered alternative” rather than “calibrated against any plausible alternative”. The Engle-Manganelli (2004) DQ test as a higher-power alternative (Appendix A.6.24) rejects $K = 18$ at $\alpha = 0.01$ ($p = 0.017$). See Appendix A.6.23 for the full power calibration.

Family	K	α	breaches	br rate	median $\widehat{\text{VaR}}_t$	LR _{uc}	LR _{ind}	LR _{cc}	p_{cc}
CHMM-N	3	0.01 †	9	1.57%	−4.56	1.62	2.36	3.98	0.14
CHMM-N	3	0.05	35	6.12%	−2.87	1.41	0.01 †	1.42	0.49
CHMM-N	18	0.01 †	9	1.57%	−5.20	1.62	2.36	3.98	0.14
CHMM-N	18	0.05	26	4.55%	−3.02	0.26	0.52	0.78	0.68
CHMM-t ($\lambda=20$)	3	0.01 †	8	1.40%	−5.59	0.82	2.81	3.63	0.16
CHMM-t ($\lambda=20$)	3	0.05	32	5.59%	−2.73	0.41	0.77	1.19	0.55
CHMM-t ($\lambda=20$)	18	0.01 †	10	1.75%	−6.25	2.65	1.98	4.62	0.10
CHMM-t ($\lambda=20$)	18	0.05	29	5.07%	−3.15	0.01 †	1.39	1.40	0.50

concentrated on the W2 (COVID) and W4 (2022 rate-hike onset) stress folds that the univariate walk-forward already flags as out-of-distribution; the quarterly refit tightens VaR coverage but does not strictly improve the cc-statistic on $T_{\text{OoS}} = 572$. **Faster refit cadences do not close the W2 stress-fold rejection:** a within-fold refit-cadence sweep at monthly and weekly cadence on each walk-forward fold (Appendix A.6.27, Table 36) shows W2 still rejects at every cadence ($p_{cc} = 0.011 \rightarrow 0.017 \rightarrow 0.023$ across fold-IS-fixed, monthly, weekly), and monthly refit *introduces* new failures on W3 and W4 from refit-cycle parameter drift that the fold-IS-fixed baseline does not have. The W2 rejection is an intrinsic regime-break failure rather than a tracking-lag issue; the operational implication is that closing W2 / W4 requires a model class with explicit regime-introduction handling (skew-emission HMMs, online-EM constructions) rather than a faster refit cadence on the symmetric scaffold.

Multiple-testing correction across the panel. The combined conditional-coverage panel reports 40 Christoffersen-cc tests across (K, α , family) on the headline OoS window (16 rows in Table 71) plus (fold, K, α) on the walk-forward (24 rows in Table 72). Reviewer feedback noted that “19/24 pass” and “passes at every (K, α)” are panel-level rather than per-row claims and require correction for multiple testing. Applying the Benjamini–Hochberg false-discovery-rate procedure at FDR = 0.05 to the 40 p-values: the headline single-OoS-window panel passes 16/16 rows under BH (minimum q -value 0.376, well above 0.05); the walk-forward panel passes 21/24 rows (BH actually accepts two more rows than the uncorrected test, because the W4, $K = 18, \alpha = 0.01$ marginal failure at raw $p = 0.022$ and the W2, $K = 3, \alpha = 0.05$ row at $p = 0.010$ no longer cross the BH ladder); and the combined panel passes 37/40 rows under BH at FDR 0.05. The three persistent rejections are all on the W2 COVID fold ($p < 10^{-3}$ in each case), which the univariate walk-forward already flags as out-of-distribution. Excluding W2 and W4, the conditional VaR passes 32/32 tests under BH. The substantive read is: the FDR-corrected joint pass-rate is 37/40 (full panel) and 32/32 (non-stress subset), so the panel-level conditional-coverage claim survives multiple-testing correction. Under the more conservative Bonferroni rule (per-test threshold $0.05/40 = 0.00125$), the same three W2 rows

reject and all others pass.

5 Discussion

A continuous HMM trained by Baum-Welch reproduces the three Cont stylized facts simultaneously on SPY at two held-out operating points: at $K^* = 3$, CHMM-N attains 89.7% IS / 80.1% OoS KS at $|G_t|$ ACF-MAE 0.0467; at $K^* = 6$ the four emission variants attain 88.3–92.6% IS / 76.3–79.0% OoS KS, with CHMM-N’s simulated kurtosis 5.26 matching observed OoS 5.29 within 0.03. Four emission variants share the same scaffold and land at four distinct kurtosis points without retuning K or the transition structure. The mechanism behind the simultaneous KS / ACF fit is the spectral identity (6): at $K \geq 3$ with quantile-initialised Baum-Welch, the absolute-return ACF is a population-level real-coefficient sum of geometric decays at the non-unit eigenvalues of $\mathbf{T}$.

Replicating the slow-ACF behaviour at $K \geq 3$. Two issues compound in the Rydén et al. [11] low- K setup: a single self-transition rate produces a single exponential mode that cannot be simultaneously fast and slow, and the random initialisations standard in the 1990s often converge to degenerate local optima. Our $K = 2$ replication on SPY (Appendix A.6.17) clarifies the mechanism: ACF-MAE is essentially constant at ≈ 0.050 across quantile and five random-seed initialisations, matching the $K = 18$ CHMM-N. The binding low- K constraint is therefore distributional (KS at $K = 2$ drops to 67–72%): with modern initialisation the ACF is reproduced at any $K \geq 2$, while distributional fidelity drives the choice of K . The effective-rank diagnostic in Appendix A.6.7 makes this microscopically explicit: at the fitted $\hat{\mathbf{T}}$ for $K = 18$, a single non-unit eigenvalue carries 93.6% of the lag-1 absolute-return ACF, and at $K = 3$ the dominant non-unit mode alone carries 96.8%.

Closing the kurtosis gap with heavy-tailed emissions. CHMM-N produces excess kurtosis 5.02 against observed 7.68: a Gaussian mixture’s kurtosis is bounded by the variance of component means relative to their variances. CHMM-t addresses this through per-state Student- t emissions but at the cost of an aggregate-mean IS overshoot to 14.57 from a single-state ECM lower-bracket pinning artefact (Appendix A.9); an exponential $1/\nu_k$ shrinkage at rate $\lambda = 20$ brings the aggregate-mean simulated kurtosis to ≈ 8 at a 1pp KS cost (the per-path *median* simulated IS kurtosis is centred near observed at every operating point even before the penalty; see Appendix A.6.20). A shared- ν ablation (a single ν shared across all K states under aggregate- Q ECM, the standard one-parameter Student- t HMM in the time-series literature) eliminates the aggregate-mean overshoot completely without any penalty: simulated IS kurtosis is 4.68 at $K^* = 3$, 9.38 at $K^* = 6$, and 6.25 at $K = 18$, with the $K = 18$ row producing the cleanest IS / OoS heavy-tail match in the entire panel (6.25 IS / 5.00 OoS against observed 7.68 / 5.29, Appendix A.6.21, Table 30). The per-state ν_k design is therefore the binding constraint on the aggregate-mean overshoot; the body retains the per-state $\nu_k + \lambda = 20$ construction for direct comparability with Peel and McLachlan [18], Liu and Rubin [19] but the shared- ν row is the structurally-cleaner alternative for kurtosis-fidelity-priority consumers. The $\lambda = 20$ choice is validated on a strictly pre-2020 slice (estimation 2014-01-03 to 2018-06-29) by validation log-likelihood on 2018-07-02 to 2019-12-31, which yields $\lambda^* = 20$ on the grid $\{0, 5, 10, 20, 50, 100, 200\}$ (Appendix A.6.37). The penalty is validated at the SPY level but per-ticker (Appendix A.9, Table 73) the body uniform $\lambda = 20$ is over-shrunk on three of six tickers; the production recipe is the per-ticker validation rule of Section 4.4. The bracket-lift remark below documents the structural-cleaness alternative.

Remark 1 (Bracket-lift ablation at $\nu_{\min} = 4$). *The body $\nu_{\min} = 2.1$ bracket lies below the kurtosis-finiteness threshold $\nu = 4$, so the lower-bracket pinning that produces the IS-kurtosis overshoot is*

partly an artefact of the bracket itself. A direct $\nu_{\min} = 4$ ablation (Table 3, $\dagger\dagger$ row) brings simulated IS kurtosis from 14.35 to 13.25 (~ 1 unit), still well above the bootstrap CI of observed (7.68 IS / 5.29 OoS), with IS / OoS KS unchanged within sampling error (95.5/85.0% vs. 95.6/85.7%). The bracket lift alone is therefore not a substitute for the $1/\nu_k$ shrinkage; the kurtosis overshoot is a multi-state pinning effect that the penalty addresses directly and the bracket only partially mitigates. We report the penalised construction in the body for direct comparability with Peel and McLachlan [18], Liu and Rubin [19] and the bracket-lift row alongside as the structurally-cleaner alternative.

CHMM-L sits between CHMM-N and CHMM-t at 6.75 without any shape parameter; CHMM-GED’s bimodal $\hat{p}_k$ partition (Appendix A.6.33) is the data-driven analog of a hand-classified Gaussian-bulk / Laplace-tail hybrid. The four-way variant decision is summarised in Table 9; bootstrap CIs on observed kurtosis (Appendix A.6.20) show that the IS-OoS difference of ~ 2.4 units is not statistically distinguishable at conventional levels, so the variant rankings should be read as point-comparisons inside a wide CI envelope.

Practical identifiability and the effective state count at $K = 18$. The $K = 18$ operating point is partly nominal: the standardized-distance single-linkage diagnostic of Appendix A.6.38 collapses CHMM-N to 11/18 effective states and CHMM-t to 12/18, with 13/18 CHMM-t states pinned at the upper ν bracket on SPY. Three implications follow. (i) The numerical full-rank check on $\hat{\mathbf{T}}$ in Appendix A.4, which reports $\text{rank}(\hat{\mathbf{T}} - \mathbf{1}\pi^\top) = K - 1 = 17$ at tolerance 10^{-10} , exhausts only algebraic rank; the practically identifiable rank is closer to $K_{\text{eff}} - 1 \approx 10$. (ii) The AIC / BIC parameter count at $K = 18$ over-credits the model; an information-criterion penalty computed against K_{eff} rather than the nominal K would re-rank $K^* \in \{6, 12\}$ in the held-out panel (Section 4.2, Table 2) more consistently with each other. (iii) The headline-metrics interpretation is not affected: simulation under the full $K = 18$ mixture and the metric panel are computed on whatever the converged model is; the question is whether the parameter count claimed against the data is honest at K_{eff} or at K_{nom} . We read this as a representational-economy observation: collapsed clusters carry small stationary-mass tails and contribute consistently to the simulation, but the model is best described as a K_{eff} -state model expressed in a K_{nom} -state parameterisation. A direct $K_{\text{nominal}} = 11$ rebuild (matching K_{eff} at $K = 18$) confirms this reading: at $K = 11$ on SPY, CHMM-N attains IS KS 94.4% / OoS KS 82.8% / IS sim kurt 4.52 / $|G_t|$ ACF-MAE 0.0505, matching or slightly exceeding the $K = 18$ nominal row (94.0% / 81.0% / 5.04 / 0.0514) on every axis (Appendix A.6.39, Table 45). The $K = 18$ over-parameterisation is therefore a parameter-counting artefact rather than an information-bearing extra resolution; $K_{\text{nominal}} = 11$ is the structurally cleaner expression of the same model. The $K^* = 6$ sensitivity reference (Section 4.2) is closer to the practically identifiable resolution and should be the default when reporting parameter counts; the $K = 18$ extended-state-resolution row is the right operating point when kurtosis fidelity is the priority, and reports against $K_{\text{nominal}} = 11$ (K_{eff}) for IC-penalised diagnostics.

Table 9. Variant decision guide. Recommended CHMM emission family by use-case priority. All four variants share the forward-backward scaffold and quantile-based initialisation; the M-step is the only architectural difference. Caveat: observed kurtosis differs across windows (7.68 IS, 5.29 OoS, a 31% drop) and no single variant matches both windows; CHMM-N at $K^* = 6$ is the closest OoS match (gap 0.83).

Use-case priority	Recommended variant
Simplest closed-form fit; kurtosis fidelity secondary	CHMM-N
Closest IS kurtosis match without a shape parameter (CHMM-L sim 5.30/6.6 vs. obs 7.68 at $K^* = 3/K = 18$)	CHMM-L
Cleanest OoS kurtosis match (deployment-window calibrated)	CHMM-N at $K^* = 6$
Per-state heavy tails, IS-calibrated	CHMM-t pen. ($\lambda = 20$)
Adaptive per-state shape; data-chosen Gaussian-bulk / Laplace-tail mixture	CHMM-GED

Stationarity scope and operational deployment. The CHMM is trained on the assumption that returns are locally stationary across the IS window. Single-name equity returns are well known to carry sector-specific structural breaks (23; 24; surveyed by 25); the cross-ticker generalisation in Section 4.4 reflects this directly, with 11/30 tickers below 60% OoS KS, concentrated in single-name regime introductions. The mitigation is periodic refit (26; an online-EM formulation by 64): a quarterly-refit version of the cross-ticker panel (Appendix A.6.30) shifts the OoS KS distribution from median 73.4% $\rightarrow$ 83.0% and reduces failures from 11/30 to 7/30. The walk-forward W2 (COVID) and W4 (2022 rate-hike onset) stress folds remain out-of-distribution by KS for every generator under any refit cadence (Appendix A.9); these are regime introductions the IS distribution does not span. Discussion of the asymmetry between the W4 onset (catastrophic) and the headline 2024-2026 reversal (clean), the OoS regime trajectory under IS-fixed parameters, the per-ticker $\hat{\lambda}^*$ for the penalised CHMM-t, and the IS / OoS kurtosis CI overlap are in Appendix A.9, A.6.25, and A.6.20.

Stylized-fact scope: simulated leverage envelope brackets the IS observed value, not the OoS observed value. The Cont (2001) stylized-fact list also includes the leverage effect (negative $\text{Corr}(G_t, |G_{t+1}|)$ on equity returns) and gain-loss asymmetry; the present paper’s body evaluation covers only the three symmetric stylized facts (heavy-tailed marginals, negligible linear ACF, slow $|G_t|$ ACF). All four CHMM emission families in this paper carry symmetric per-state densities (Gaussian, Laplace, Student- t , GED at any p_k), so we initially expected the CHMM to fail the leverage axis by construction. We compute $\text{Corr}(G_t, |G_{t+1}|)$ on observed SPY (IS -0.135 , OoS -0.214) and on the per-path simulation distribution from each generator at $N_{\text{paths}} = 500$ (Appendix A.7.9). At $K = 18$ the CHMM-N path-level median is -0.089 on IS and -0.087 on OoS, with $[Q_5, Q_{95}]$ envelopes $[-0.138, -0.037]$ IS and $[-0.190, +0.015]$ OoS. Reading those quantiles as one-sided percentile p-values against the simulated null: **the IS observed value -0.135 sits just inside the simulated Q_5 envelope (one-sided $p \approx 0.05$, borderline non-rejection), and the OoS observed value -0.214 sits below the simulated Q_5 floor of -0.190 (one-sided $p < 0.05$, rejection of the simulated null at the 5% level).** The substantive read is therefore: CHMM at $K = 18$ produces a simulated leverage signal that is borderline-non-rejected on IS at conventional levels and rejected on OoS, rather than “partially captured” on both windows. The mechanism behind the non-zero IS leverage is that the converged $\hat{\mathbf{T}}$ couples high- σ_k states to negative- μ_k states asymmetrically, so the Markov mixing produces a non-trivial signed signal even though each per-state emission is symmetric; this mechanism explains why the IS gap is small while the OoS gap remains. The CHMM-N envelope is competitive with EGARCH (median -0.095 , $[Q_5, Q_{95}] = [-0.147, -0.047]$ IS) and tighter than GARCH(1,1) Gaussian (median -0.004 , indistinguishable from zero); the asymmetric GARCH baselines (EGARCH, GJR-GARCH) produce slightly stronger leverage in the median but their envelopes also fail to cover the OoS observed value at the 5% level. Skew-emission extensions (skew- t in the tradition of 65; skew-Laplace; skew-GED) remain the natural fix for closing the OoS gap and are deferred as companion-paper directions. Closing the OoS gap requires asymmetric per-state emissions; this paper does not close it.

Baseline-implementation caveats. Two rows of our extended baseline panel (Appendix A.7.7) are repo-native re-implementations rather than reference-implementation runs and should be read with that caveat. (i) The QuantGAN row (Appendix A.7.5) is a three-conv-layer convolutional WGAN that is materially smaller than the seven-block TCN of Wiese et al. [9] and lacks the Lambert-W input pre-processing that the original authors credit with most of the tail-fidelity. The 0% IS / OoS KS pass rate is the failure mode of *this* in-house implementation; the body Table 3 also reports the seven-block TCN rebuild at the same 0% KS, and the Lambert-W variant of Appendix A.7.5

likewise fails. We do not claim that a faithful reference-implementation re-run of Wiese et al. [9] would also fail; that is a deferred follow-up. The QuantGAN rows in this paper are in-house WGAN re-implementation results, not a verdict on the WGAN literature or the deep-generative class. (ii) The SM-CHMM row of Appendix A.7.7 is fit by a Viterbi-AR(1) plug-in estimator rather than by full ML on the explicit-duration HSMM; the ML HSMM-N at $K^* = 3$ row of Table 3 is the proper HSMM benchmark and wins the OoS KS column. We do not claim that an ML-estimated HSMM at higher K would underperform CHMM at the same K , only that the present comparison is at each method’s best operating point. The MS-GARCH baseline is reported in two flavours: the in-house Nelder-Mead frequentist fit (Table 59 and the in-house rows of Table 3, 27–37% IS KS plateau) and a reference Bayesian re-run via the MSGARCH R package of Ardia et al. [35] (Table 3 ¶ rows and Appendix A.7.7, 0–0.1% IS / 5–6% OoS KS at posterior-predictive simulation); both fall well short of the CHMM family on this dataset, so the body conclusion that “the multi-state benefit is specific to the CHMM scaffold rather than to multi-state regime-switching per se” is robust to estimator choice. A faithful Wiese et al. TCN re-run with Lambert-W pre-processing is the right next-step control on the deep-generative axis; we report the present panel honestly with the caveats above.

Limitations. Skew-heavy-tailed emissions (skew- t or skew-Laplace) and explicit-duration semi-Markov sojourns at higher K are deferred as companion-paper directions; the leverage-effect diagnostic (Appendix A.7.9) shows the residual gap to OoS observed leverage will require asymmetric per-state emissions to close. Scaling the multi-asset construction to larger d via untruncated regular vines or factor copulas is the natural follow-up. The state-selection rule for the $K = 18$ extended-state-resolution panel uses information-criterion plus distributional pass-rate rather than strictly held-out log-likelihood; held-out re-selection picks $K^* = 3$ or $K^* = 6$ depending on the validation slice (Appendix A.9), and the body $K^* = 3$ block is the state-resolution-robust default for synthetic-data consumers. The cross-decade validation result reported in §4.4 (§A.6.46 for the full panel) is the same pattern at the decade scale: the IS axis is decade-robust within ~ 5 pp on KS, while the calm 2004–2006 OoS slice collapses to 3–5% pass rate, mirroring the W2 / W4 walk-forward stress folds rather than the headline 2024–2026 OoS slice. The empirical scope is daily US-equity stylized-fact reproduction; FX, fixed-income, and commodity asset classes are not covered, and the GLD / SLV non-equity stress test in Appendix A.8.6 documents that the static IS-fitted CHMM does not transfer to those windows without periodic refit. The cross-asset dependence-family selection (Section 4.5) is an IS-calibration distinction: on the OoS window the Gaussian and Student- t copulas are statistically indistinguishable at $N_{\text{paths}} = 200$. We have addressed six of the deferred robustness items raised in pre-submission peer review: the OoS-anchored block-bootstrap KS recalibration (Appendix A.7.2) confirms the IS CHMM-vs-GARCH ordering survives on OoS; the leverage-effect diagnostic (Appendix A.7.9) shows the CHMM family at $K = 18$ borderline-non-rejects the IS observed leverage at the 5% level via Markov state-mixing despite symmetric per-state emissions, but rejects the OoS observed leverage at the same level (the OoS gap requires asymmetric per-state emissions to close); the full one-shot Student- t copula MLE (Appendix A.7.10) shows the body two-step Kendall’s- τ + profile-MLE construction is not biased toward the Gaussian limit on the six-asset universe; the CRSP cross-decade validation (Appendix A.6.46) extends the empirical scope to 1994–2006 via a day-pass WRDS query and shows that the IS axis of the stylized-fact reproduction is decade-robust (84–90% IS KS at $K \in \{3, 18\}$ on the 1994–2004 slice) while the 2004–2006 OoS slice exhibits the same low-stress / low-kurtosis pattern as the W2 / W4 walk-forward stress folds; the QuantGAN reference TCN rebuild (Appendix A.7.11) reports a Wiese-style dilated TCN architecture as a stronger deep-generative comparator than the body’s 3-conv-layer baseline; and the reference-implementation MS-GARCH re-run via the MSGARCH R package of Ardia et al.

[35] (Table 3 ¶ rows and Appendix A.7.7) confirms the multi-state-benefit plateau under Bayesian estimation, with the IS KS pass-rate falling to 0–0.1% under posterior-predictive simulation that integrates parameter uncertainty path-by-path.

6 Conclusion

A continuous HMM trained by Baum-Welch with quantile-based initialization reproduces the three symmetric Cont stylized facts (heavy-tailed marginals, negligible linear ACF, slow $|G_t|$ ACF) on daily US-equity returns. At the state-resolution-robust body headline $K^* = 3$ (selected by held-out per-observation log-likelihood and held-out KS on the strictly pre-2020 slice under both single-fold and four-fold / six-fold rolling-origin CV; Appendix A.9) the four emission variants attain 79.8–90.6% IS / 63.1–83.2% OoS KS on the single SPY OoS window, with CHMM-N at 89.7% IS / 80.5% OoS KS, simulated kurtosis 3.83/3.53, and $|G_t|$ ACF-MAE 0.0460; the penalised CHMM-t at $\lambda = 20$ provides the cleanest IS / OoS heavy-tail match in the headline block at 14.91/8.50 against observed 7.68/5.29. A $K^* = 6$ block is reported as a sensitivity reference (single-fold pre-2020 selection; the four-fold and six-fold k -fold CV cannot distinguish $K^* = 6$ from $K^* = 3$ at conventional levels). The single-window OoS sits at the upper end of the six-fold rolling-origin walk-forward whose median CHMM-N OoS KS is 67.7% at $K = 18$ and 62.1% at $K^* = 3$; the walk-forward median is the operationally informative summary for production deployment. An extended-state-resolution *sensitivity reference* at $K = 18$, selected by a multi-objective rule that uses the OoS window, attains 93.6–95.6% IS / 80.8–85.7% OoS KS, with the penalised CHMM-t at $\lambda = 20$ providing the cleanest moderate-state-resolution IS / OoS kurtosis match at 8.56 / 7.07 against observed 7.68 / 5.29 (Table 3). The strongest single-window OoS KS row in the panel is ML HSMM-N at $K^* = 3$ (91.0% OoS KS) but its $|G_t|$ ACF-MAE matches the i.i.d. baseline; CHMM and ML HSMM are best read as complementary scaffolds with different KS / ACF / kurtosis trade-offs. We apply the textbook closed-form mixture-of-eigenvalues identity for the absolute-return ACF (13–15) to the SPY 2014–2024 instance of the Rydén et al. low- K setup and recast the empirical failure there as an effective-rank statement on $\mathbf{T} - \mathbf{1}\bar{\pi}^\top$, showing the rank constraint is non-binding at $K \geq 3$ on this instance; the substantive content is the empirical effective-rank application, not a re-derivation of the identity. The four emission variants share the same scaffold and land at four kurtosis points; CHMM-GED’s per-state shape p_k nests Gaussian and Laplace as boundary cases and produces a bulk/tail partition that replicates across seeds and tickers. VaR and ES envelopes bracket observed values at 1% and 5%; a regime-conditional VaR that propagates the one-step-ahead state forecast through the CHMM mixture passes Kupiec, Christoffersen-ind, and Christoffersen-cc cleanly at $\alpha = 0.05$ on OoS (Table 8; the $\alpha = 0.01$ rows are power-bounded at $T_{\text{OoS}} = 572$ and should be read with the calibration of Appendix A.6.23). The conditional construction extends to a six-fold rolling-origin walk-forward with a 19/24 aggregate pass-rate at the 5% level (Table 72); the two failures concentrate on the COVID and 2022-rate-hike-onset stress folds that the univariate walk-forward already flags as out-of-distribution by KS. Cross-ticker generalisation is exercised on a sector-balanced 30-ticker panel (Section 4.4, Table 6); the body table reports the headline $K^* = 3$ (OoS KS median 69.1%), the sensitivity reference $K^* = 6$ (OoS KS median 75.1%), and the kurtosis-fidelity sensitivity $K = 18$ (OoS KS median 73.4%), all at identical 11/30 failure count, with quarterly refit at $K = 18$ lifting the OoS median to 83.0% and the largest gains concentrating on the regime-introduction tickers. Multi-asset synthesis composes through a rank-based Student- t copula on the CHMM marginals at $\nu^* = 6$ (six-asset IS off-diagonal correlation MAE 0.027); on OoS the Gaussian and Student- t copulas are statistically indistinguishable at $N_{\text{paths}} = 200$, so the dependence-family selection is an IS-calibration distinction. The headline scope is daily US equities.

The static IS-fitted CHMM does not transfer outside this asset class: the GLD / SLV non-equity stress test of Appendix A.8.6 collapses to 0% OoS KS on both ETFs, a hard rejection rather than a soft preference. Non-equity asset classes (commodities, FX, fixed income) require periodic refit at a cadence that scales with the asset’s regime structure; the body recipes for periodic refit on US-equity tickers (Section 4.4) are not validated to extend across asset classes without per-class re-validation.

Skew-heavy-tailed emissions (the natural fix for the absent leverage-effect coverage in the present scope), explicit-duration semi-Markov sojourns at higher K , and a reference-implementation re-run of the original Wiese et al. (2020) QuantGAN architecture with Lambert-W input pre-processing are companion-paper directions; scaling to larger d via untruncated regular vines or factor copulas is the natural multi-asset follow-up. The reference-implementation MS-GARCH re-run requested at pre-submission peer review (MSGARCH R package of 35) is in the body of this revision (Table 3 ¶ rows and Appendix A.7.7).

Data and code availability. Reproduser at <https://github.com/altashly1/CHMM-paper>; companion Julia package CHMM-Model.jl at <https://github.com/altashly1/CHMM-Model> (public API in Appendix A.1, Julia ≥ 1.12). Daily prices: Polygon.io via Massive (2014-01-03 to 2025-11-18) and Alpaca Markets / IEX (held-out extension through 2026-04-20); we thank both providers. The authors declare no competing interests. **Author contributions:** J.V. directed the study; A.A. developed the continuous model, implemented Baum-Welch, conducted the empirical analysis, and generated all figures and tables; C.J. contributed to methodology review and validation; all authors edited and reviewed the final manuscript.

References

- [1] James Jordon, Lukasz Szpruch, Florimond Houssiau, Mirko Bottarelli, Giovanni Cherubin, Carsten Maple, Samuel N. Cohen, and Adrian Weller. Synthetic data: what, why and how? *arXiv preprint arXiv:2205.03257*, 2022.
- [2] Samuel A. Assefa, Danial Dervovic, Mahmoud Mahfouz, Robert E. Tillman, Prashant Reddy, and Manuela Veloso. Generating synthetic data in finance: opportunities, challenges and pitfalls. In *Proceedings of the First ACM International Conference on AI in Finance*, pages 1–8, 2020.
- [3] Abdulrahman Alswaidan and Jeffrey D. Varner. Hybrid hidden markov model for modeling equity excess growth rate dynamics: A discrete-state approach with jump-diffusion. *arXiv preprint arXiv:2603.10202*, 2026. Preprint, not yet peer-reviewed.
- [4] Benoit Mandelbrot. The variation of certain speculative prices. *The Journal of Business*, 36(4): 394–419, 1963.
- [5] R. Cont. Empirical properties of asset returns: stylized facts and statistical issues. *Quantitative Finance*, 1(2):223–236, 2001.
- [6] Robert F. Engle. Autoregressive conditional heteroscedasticity with estimates of the variance of united kingdom inflation. *Econometrica*, 50(4):987–1007, 1982.
- [7] Tim Bollerslev. Generalized autoregressive conditional heteroscedasticity. *Journal of Econometrics*, 31(3):307–327, 1986.

- [8] Jinsung Yoon, Daniel Jarrett, and Mihaela van der Schaar. Time-series generative adversarial networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 32, 2019.
- [9] Magnus Wiese, Robert Knobloch, Ralf Korn, and Peter Kretschmer. Quant GANs: Deep generation of financial time series. *Quantitative Finance*, 20(9):1419–1440, 2020.
- [10] Kashif Rasul, Calvin Seward, Ingmar Schuster, and Roland Vollgraf. Autoregressive denoising diffusion models for multivariate probabilistic time series forecasting. In *Proceedings of the 38th International Conference on Machine Learning*, pages 8857–8868, 2021.
- [11] Tobias Rydén, Timo Teräsvirta, and Stefan Åsbrink. Stylized facts of daily return series and the hidden Markov model. *Journal of Applied Econometrics*, 13(3):217–244, 1998.
- [12] Jan Bulla and Ingo Bulla. Stylized facts of financial time series and hidden semi-Markov models. *Computational Statistics & Data Analysis*, 51(4):2192–2209, 2006.
- [13] James D. Hamilton. *Time Series Analysis*. Princeton University Press, Princeton, NJ, 1994.
- [14] Hans-Martin Krolzig. *Markov-Switching Vector Autoregressions: Modelling, Statistical Inference, and Application to Business Cycle Analysis*, volume 454 of *Lecture Notes in Economics and Mathematical Systems*. Springer, Berlin, 1997.
- [15] Allan Timmermann. Moments of Markov switching models. *Journal of Econometrics*, 96(1):75–111, 2000.
- [16] Leonard E. Baum, Ted Petrie, George Soules, and Norman Weiss. A maximization technique occurring in the statistical analysis of probabilistic functions of markov chains. *The Annals of Mathematical Statistics*, 41(1):164–171, 1970.
- [17] L. Rabiner and B. Juang. An introduction to hidden Markov models. *IEEE ASSP Magazine*, 3(1):4–16, 1986.
- [18] David Peel and Geoffrey J. McLachlan. Robust mixture modelling using the t distribution. *Statistics and Computing*, 10(4):339–348, 2000.
- [19] Chuanhai Liu and Donald B. Rubin. ML estimation of the t distribution using EM and its extensions, ECM and ECME. *Statistica Sinica*, 5(1):19–39, 1995.
- [20] Dimitris N. Politis and Joseph P. Romano. The stationary bootstrap. *Journal of the American Statistical Association*, 89(428):1303–1313, 1994.
- [21] Tim Bollerslev. A conditionally heteroskedastic time series model for speculative prices and rates of return. *Review of Economics and Statistics*, 69(3):542–547, 1987.
- [22] Markus Haas, Stefan Mittnik, and Marc S. Paoletta. A new approach to Markov-switching GARCH models. *Journal of Financial Econometrics*, 2(4):493–530, 2004.
- [23] Ľuboš Pástor and Robert F. Stambaugh. The equity premium and structural breaks. *Journal of Finance*, 56(4):1207–1239, 2001.
- [24] Elena Andreou and Eric Ghysels. Detecting multiple breaks in financial market volatility dynamics. *Journal of Applied Econometrics*, 17(5):579–600, 2002.

- [25] Andrew Ang and Allan Timmermann. Regime changes and financial markets. *Annual Review of Financial Economics*, 4(1):313–337, 2012.
- [26] M. Hashem Pesaran and Allan Timmermann. Selection of estimation window in the presence of breaks. *Journal of Econometrics*, 137(1):134–161, 2007.
- [27] James D. Hamilton. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2):357–384, 1989.
- [28] Huntley Schaller and Simon Van Norden. Regime switching in stock market returns. *Applied Financial Economics*, 7(2):177–191, 1997.
- [29] Peter Nystrup, Henrik Madsen, and Erik Lindström. Long memory of financial time series and hidden markov models with time-varying parameters. *Journal of Forecasting*, 36(8):989–1002, 2017.
- [30] M. T. Subbotin. On the law of frequency of error. *Matematicheskii Sbornik*, 31(2):296–301, 1923.
- [31] George E. P. Box and George C. Tiao. *Bayesian Inference in Statistical Analysis*. Addison-Wesley, Reading, MA, 1973.
- [32] Daniel B. Nelson. Conditional heteroskedasticity in asset returns: A new approach. *Econometrica*, 59(2):347–370, 1991.
- [33] Lawrence R. Glosten, Ravi Jagannathan, and David E. Runkle. On the relation between the expected value and the volatility of the nominal excess return on stocks. *The Journal of Finance*, 48(5):1779–1801, 1993.
- [34] Jean-Michel Zakoian. Threshold heteroskedastic models. *Journal of Economic Dynamics and Control*, 18(5):931–955, 1994.
- [35] David Ardia, Keven Bluteau, Kris Boudt, Leopoldo Catania, and Denis-Alexandre Trottier. Markov-switching GARCH models in R: The MSGARCH package. *Journal of Statistical Software*, 91(4):1–38, 2019. doi: 10.18637/jss.v091.i04.
- [36] Torben G. Andersen and Tim Bollerslev. Heterogeneous information arrivals and return volatility dynamics: Uncovering the long-run in high frequency returns. *The Journal of Finance*, 52(3): 975–1005, 1997.
- [37] Fulvio Corsi. A simple approximate long-memory model of realized volatility. *Journal of Financial Econometrics*, 7(2):174–196, 2009.
- [38] Shuntaro Takahashi, Yu Chen, and Kumiko Tanaka-Ishii. Modeling financial time-series with generative adversarial networks. *Physica A: Statistical Mechanics and its Applications*, 527: 121261, 2019.
- [39] Sohyeon Kwon and Yongjae Lee. Can GANs learn the stylized facts of financial time series? In *Proceedings of the 5th ACM International Conference on AI in Finance*, pages 126–133, 2024.
- [40] William F. Sharpe. A simplified model for portfolio analysis. *Management Science*, 9(2):277–293, 1963.

- [41] A. Sklar. Fonctions de répartition à n dimensions et leurs marges. *Publications de l'Institut de Statistique de l'Université de Paris*, 8:229–231, 1959.
- [42] Roger B. Nelsen. *An Introduction to Copulas*. Springer, 2 edition, 2006.
- [43] Stefano Demarta and Alexander J. McNeil. The t copula and related copulas. *International Statistical Review*, 73(1):111–129, 2005.
- [44] Alexander J. McNeil, Rüdiger Frey, and Paul Embrechts. *Quantitative Risk Management: Concepts, Techniques and Tools*. Princeton University Press, revised edition, 2015.
- [45] Ronald L. Iman and W. J. Conover. A distribution-free approach to inducing rank correlation among input variables. *Communications in Statistics–Simulation and Computation*, 11(3):311–334, 1982.
- [46] Michael Stenger, André Bauer, Thomas Prantl, Robert Leppich, Nathaniel Hudson, Kyle Chard, Ian Foster, and Samuel Kounev. Thinking in categories: A survey on assessing the quality for time series synthesis. *Journal of Data and Information Quality*, 16(2):1–32, 2024.
- [47] Tilmann Gneiting and Adrian E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007.
- [48] Tilmann Gneiting and Matthias Katzfuss. Probabilistic forecasting. *Annual Review of Statistics and Its Application*, 1:125–151, 2014.
- [49] Francis X. Diebold and Roberto S. Mariano. Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3):253–263, 1995.
- [50] Arthur Gretton, Karsten M. Borgwardt, Malte J. Rasch, Bernhard Schölkopf, and Alexander Smola. A kernel two-sample test. *Journal of Machine Learning Research*, 13(1):723–773, 2012.
- [51] Ilya Chevyrev and Andrey Kormilitzin. A primer on the signature method in machine learning. *arXiv preprint arXiv:1603.03788*, 2016.
- [52] Hao Ni, Lukasz Szpruch, Magnus Wiese, Shujian Liao, and Baoren Xiao. Conditional Sig-Wasserstein GANs for time series generation. *arXiv preprint arXiv:2006.05421*, 2020.
- [53] Laurent E. Calvet and Adlai J. Fisher. How to forecast long-run volatility: Regime switching and the estimation of multifractal processes. *Journal of Financial Econometrics*, 2(1):49–83, 2004.
- [54] Robert C. Merton. Option pricing when underlying stock returns are discontinuous. *Journal of Financial Economics*, 3(1–2):125–144, 1976.
- [55] Paul Embrechts, Alexander McNeil, and Daniel Straumann. Correlation and dependence in risk management: Properties and pitfalls. *Risk Management: Value at Risk and Beyond*, pages 176–223, 2002.
- [56] A. N. Kolmogorov. Sulla determinazione empirica di una legge di distribuzione. *Giornale dell'Istituto Italiano degli Attuari*, 4:83–91, 1933.
- [57] N. Smirnov. Table for estimating the goodness of fit of empirical distributions. *The Annals of Mathematical Statistics*, 19(2):279–281, 1948.

- [58] Paul H. Kupiec. Techniques for verifying the accuracy of risk measurement models. *Journal of Derivatives*, 3(2):73–84, 1995.
- [59] Jeff A. Bilmes. A gentle tutorial of the em algorithm and its application to parameter estimation for gaussian mixture and hidden markov models. *Technical Report TR-97-021, International Computer Science Institute*, 1998.
- [60] Carlos A. Abanto-Valle, Roland Langrock, Ming-Hui Chen, and Marcos V. Cardoso. Maximum likelihood estimation for stochastic volatility in mean models with heavy-tailed distributions. *Applied Stochastic Models in Business and Industry*, 33(4):394–408, 2017.
- [61] David A. Levin and Yuval Peres. *Markov Chains and Mixing Times*. American Mathematical Society, 2 edition, 2017.
- [62] Nguyet Nguyen. Hidden Markov model for stock trading. *International Journal of Financial Studies*, 6(2):36, 2018.
- [63] Shun-Zheng Yu. Hidden semi-Markov models. *Artificial Intelligence*, 174(2):215–243, 2010.
- [64] Olivier Cappé. Online EM algorithm for hidden Markov models. *Journal of Computational and Graphical Statistics*, 20(3):728–749, 2011.
- [65] Adelchi Azzalini and Antonella Capitanio. Distributions generated by perturbation of symmetry with emphasis on a multivariate skew t -distribution. *Journal of the Royal Statistical Society. Series B (Statistical Methodology)*, 65(2):367–389, 2003.
- [66] Xiao-Li Meng and Donald B. Rubin. Maximum likelihood estimation via the ECM algorithm: A general framework. *Biometrika*, 80(2):267–278, 1993.
- [67] Paul Glasserman. *Monte Carlo Methods in Financial Engineering*. Springer, 2003.
- [68] Peter F. Christoffersen. Evaluating interval forecasts. *International Economic Review*, 39(4):841–862, 1998.
- [69] C. F. Jeff Wu. On the convergence properties of the EM algorithm. *The Annals of Statistics*, 11(1):95–103, 1983.
- [70] Elizabeth S. Allman, Catherine Matias, and John A. Rhodes. Identifiability of parameters in latent structure models with many observed variables. *The Annals of Statistics*, 37(6A):3099–3132, 2009.
- [71] Sidney J. Yakowitz and John D. Spragins. On the identifiability of finite mixtures. *The Annals of Mathematical Statistics*, 39(1):209–214, 1968.
- [72] Bruce G. Lindsay. *Mixture Models: Theory, Geometry and Applications*, volume 5 of *NSF-CBMS Regional Conference Series in Probability and Statistics*. Institute of Mathematical Statistics, Hayward, CA, 1995.
- [73] Peter J. Bickel, Ya’acov Ritov, and Tobias Rydén. Asymptotic normality of the maximum-likelihood estimator for general hidden Markov models. *The Annals of Statistics*, 26(4):1614–1635, 1998.

- [74] Randal Douc, Eric Moulines, and Tobias Rydén. Asymptotic properties of the maximum likelihood estimator in autoregressive models with Markov regime. *The Annals of Statistics*, 32(5):2254–2304, 2004.
- [75] Robert F. Engle and Simone Manganelli. CAViaR: Conditional autoregressive value at risk by regression quantiles. *Journal of Business & Economic Statistics*, 22(4):367–381, 2004.
- [76] Mike K. P. So, K. Lam, and W. K. Li. A stochastic volatility model with Markov switching. *Journal of Business & Economic Statistics*, 16(2):244–253, 1998.
- [77] Carlos M. Carvalho and Hedibert F. Lopes. Simulation-based sequential analysis of Markov switching stochastic volatility models. *Computational Statistics & Data Analysis*, 51(9):4526–4542, 2007.
- [78] Martin Arjovsky, Soumith Chintala, and Léon Bottou. Wasserstein generative adversarial networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 214–223, 2017.
- [79] Mike Innes. Flux: Elegant machine learning with Julia. *Journal of Open Source Software*, 3(25):602, 2018.
- [80] Abdulrahman Alswaidan, Cade Jin, and Jeffrey D. Varner. Long memory from markov mixing: A semi-markov regime-switching generator for the cboe volatility index. 2026. Manuscript in preparation, not yet submitted.
- [81] Georg M. Goerg. Lambert W random variables – a new family of generalized skewed distributions with applications to risk estimation. *Annals of Applied Statistics*, 5(3):2197–2230, 2011.
- [82] Georg M. Goerg. The lambert way to Gaussianize heavy-tailed data with the inverse of Tukey’s h transformation as a special case. *The Scientific World Journal*, 2015:909231, 2015.
- [83] S. S. Wilks. The large-sample distribution of the likelihood ratio for testing composite hypotheses. *Annals of Mathematical Statistics*, 9(1):60–62, 1938.
- [84] A. W. van der Vaart. *Asymptotic Statistics*. Cambridge Series in Statistical and Probabilistic Mathematics. Cambridge University Press, 1998.

A Supplementary Material

This appendix collects supplementary material referenced in the main body: the public API of the companion package, the EM forward-backward recursions and algorithm pseudocode, the validation metric definitions, formal propositions, the full state-resolution and multi-emission sensitivity panels, robustness and KS-power calibration, the extended GARCH and SM-CHMM baselines, the full cross-asset panel, and the K -selection and ν_k diagnostics referenced from the discussion.

A.1 CHMM-Model.jl Public API and Reproducer

The companion Julia package `CHMM-Model.jl` (<https://github.com/altashly1/CHMM-Model>) is the entry point used to regenerate every table and figure in this paper from a single seed root under Julia ≥ 1.12 . The public surface is a single `CHMM` type, three emission constructors, a unified `fit!` method (log-space Baum-Welch / ECM), and a unified `simulate` method.

```

using CHMM
chmm_n    = make_chmm_n(returns; K = 18)    # Gaussian emissions
chmm_t    = make_chmm_t(returns; K = 18)    # Student-t emissions
chmm_l    = make_chmm_l(returns; K = 18)    # Laplace emissions
chmm_ged  = make_chmm_ged(returns; K = 18)  # GED (per-state shape p_k)
fit!(chmm_n; max_iter = 60, tol = 1e-4)    # log-space Baum-Welch
paths = simulate(chmm_n; T = 2516, n = 1000) # 1000 paths of length T

```

The same `simulate` interface is used for every benchmark generator (Bootstrap, Gaussian i.i.d., Laplace i.i.d., GARCH(1,1)) so that the seven-metric panel evaluates each row on byte-identical windows. The Pipeline B Student-t copula sampler is exposed as `simulate_copula(...; nu = 6)` on the per-asset CHMM marginals; the rank-reordering construction is one function call. Running `julia -project=. run_full_rebuild.jl` from the paper repository (<https://github.com/altashly1/CHMM-paper>) reproduces every numerical artifact under the global seed policy of Section 3.5; `Manifest.toml` pins the dependency graph.

Sub-seed derivation rule. The deterministic global seed root is 20260420; per-experiment sub-seeds follow the additive rule `sub_seed = SEED + offset(experiment) + offset(replicate)`, with experiment offsets in $\{0, 1,000, 100, 10,000, \dots\}$ chosen to avoid collisions across the runners in the public Julia package. For example, the headline simulation in `run_diagnostics.jl` uses `SEED + (path index)` for each per-path random draw, the multi-seed Monte Carlo in `run_multiseed_headline.jl` uses base `20260422 + 100 · (seed index)` for ten independent replicates, and the bracket / shrinkage sweeps add a further per-rate offset. The full sub-seed map is enumerated at the top of each runner script; downstream `Random.seed!()` calls are deterministic and order-independent within each runner.

A.2 CHMM Algorithms and Derivations

This appendix collects the algorithmic content for the CHMM family: the forward-backward recursions and weighted M-step updates, the complete training pseudocode, and per-state degrees-of-freedom diagnostics for CHMM-t.

A.2.1 Forward-Backward Recursions and M-Step Updates

For reference, the in-text Baum-Welch description of Section 3.2 is underpinned by the standard log-space recursions [17, 59]. Define the forward variable $\alpha_t(k) = \mathbb{P}(O_{1:t}, S_t = k \mid \mathcal{M})$ and the backward variable $\beta_t(k) = \mathbb{P}(O_{t+1:T} \mid S_t = k, \mathcal{M})$. The log-space recursions are

$$\log \alpha_1(k) = \log \pi_k + \log f_k(O_1), \quad (8)$$

$$\log \alpha_t(j) = \log \sum_i \exp(\log \alpha_{t-1}(i) + \log T_{ij}) + \log f_j(O_t), \quad t = 2, \dots, T, \quad (9)$$

$$\log \beta_T(k) = 0, \quad (10)$$

$$\log \beta_t(i) = \log \sum_j \exp(\log T_{ij} + \log f_j(O_{t+1}) + \log \beta_{t+1}(j)), \quad t = T-1, \dots, 1, \quad (11)$$

where $\log \text{sumexp}_i(x_i) = x_{\max} + \log \sum_i \exp(x_i - x_{\max})$. The posterior state-occupancy $\gamma_t(k)$ and pair-occupancy $\xi_t(i, j)$ quantities are

$$\gamma_t(k) = \frac{\alpha_t(k) \beta_t(k)}{\sum_j \alpha_t(j) \beta_t(j)}, \quad \xi_t(i, j) = \frac{\alpha_t(i) T_{ij} f_j(O_{t+1}) \beta_{t+1}(j)}{\sum_{m,n} \alpha_t(m) T_{mn} f_n(O_{t+1}) \beta_{t+1}(n)}, \quad (12)$$

and the closed-form M-step updates are

$$\pi_k^{\text{new}} = \gamma_1(k), \quad \mu_k^{\text{new}} = \frac{\sum_t \gamma_t(k) O_t}{\sum_t \gamma_t(k)}, \quad (\sigma_k^{\text{new}})^2 = \frac{\sum_t \gamma_t(k) (O_t - \mu_k^{\text{new}})^2}{\sum_t \gamma_t(k)}, \quad T_{ij}^{\text{new}} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \gamma_t(i)}. \quad (13)$$

The convergence criterion is $|\mathcal{L}^{(n)} - \mathcal{L}^{(n-1)}| < 10^{-4}$ with $\mathcal{L}^{(n)} = \log \text{sumexp}_k \log \alpha_T^{(n)}(k)$, and simulation draws $S_1 \sim \text{Categorical}(\boldsymbol{\pi})$ and then iterates $G_t \sim \mathcal{N}(\mu_{S_t}, \sigma_{S_t}^2)$, $S_{t+1} \sim \text{Categorical}(\mathbf{T}_{S_t, \cdot})$ with no jump process.

CHMM-t M-step (per-state ECM, closed form given $u_{t,k}$). For Student-t emissions $f_k(x) = t_{\nu_k}(x; \mu_k, \sigma_k)$ the ECM formulation of Peel and McLachlan [18], Liu and Rubin [19] treats the latent precision as an augmented E-step quantity,

$$u_{t,k} = \frac{\nu_k + 1}{\nu_k + ((O_t - \mu_k)/\sigma_k)^2}, \quad (14)$$

so that conditional on ν_k the location and scale admit closed-form weighted updates,

$$\mu_k^{\text{new}} = \frac{\sum_t \gamma_t(k) u_{t,k} O_t}{\sum_t \gamma_t(k) u_{t,k}}, \quad (\sigma_k^{\text{new}})^2 = \frac{\sum_t \gamma_t(k) u_{t,k} (O_t - \mu_k^{\text{new}})^2}{\sum_t \gamma_t(k)}. \quad (15)$$

The degrees-of-freedom parameter ν_k is then refined by a one-dimensional golden-section search on the per-state Q -function $Q_k(\nu) = \sum_t \gamma_t(k) \log t_\nu(O_t; \mu_k, \sigma_k)$ over the bracket $(\nu_{\min}, \nu_{\max}) = (2.1, 50)$. This two-block ECM preserves the monotonicity of the standard EM guarantee on the compact bracket $[\nu_{\min}, \nu_{\max}]$, with the search invariant $Q_k(\nu^{(n+1)}) \geq Q_k(\nu^{(n)})$ at every iteration; Proposition 1 states the guarantee in full. We extend this M-step with an optional exponential shrinkage prior on $1/\nu_k$, corresponding to the penalised objective

$$Q_k^{\text{pen}}(\nu) = \sum_t \gamma_t(k) \log t_\nu(O_t; \mu_k, \sigma_k) - \lambda / \nu, \quad (16)$$

maximised by the same golden-section search ($\lambda = 0$ recovers the unpenalised Peel and McLachlan [18], Liu and Rubin [19] ECM; the rate sweep and its effect on simulated IS kurtosis are in Section 5).

CHMM-L M-step (closed-form weighted Laplace MLE). For Laplace emissions $f_k(x) = \frac{1}{2\beta_k} \exp(-|x - \mu_k|/\beta_k)$ (with the per-state Laplace scale denoted β_k rather than the canonical Laplace-scale symbol b) the weighted-MLE estimators are available in closed form:

$$\mu_k^{\text{new}} = \text{WeightedMedian}(O_{1:T}; \gamma_{1:T}(k)), \quad \beta_k^{\text{new}} = \frac{\sum_t \gamma_t(k) |O_t - \mu_k^{\text{new}}|}{\sum_t \gamma_t(k)}. \quad (17)$$

The weighted median is computed by cumulative-weight bracketing with linear interpolation between adjacent order statistics at the half-total weight crossing. The Laplace M-step carries no shape parameter, so CHMM-L is strictly cheaper per iteration than CHMM-t (which requires the per-state Q_k search).

CHMM-GED M-step (three-stage ECM with closed-form scale). For Generalized Error Distribution emissions $f_k(x) = \text{GED}(x; \mu_k, \alpha_k, p_k) = \frac{p_k}{2\alpha_k \Gamma(1/p_k)} \exp(-(|x - \mu_k|/\alpha_k)^{p_k})$ the M-step

decomposes into three CM updates per state. Given current $(\mu_k^{(n)}, \alpha_k^{(n)}, p_k^{(n)})$, CM-1 updates the location by minimising the per-state weighted L^{p_k} loss,

$$\mu_k^{(n+1)} = \arg \min_{\mu \in [\hat{\mu}_k - 5\alpha_k, \hat{\mu}_k + 5\alpha_k]} \sum_t \gamma_t(k) |O_t - \mu|^{p_k^{(n)}}, \quad (18)$$

by a 40-iteration golden-section search; CM-2 updates the scale in closed form given $(\mu_k^{(n+1)}, p_k^{(n)})$,

$$\alpha_k^{(n+1)} = \left[\frac{p_k^{(n)}}{W_k} \sum_t \gamma_t(k) |O_t - \mu_k^{(n+1)}|^{p_k^{(n)}} \right]^{1/p_k^{(n)}}, \quad W_k = \sum_t \gamma_t(k); \quad (19)$$

and CM-3 updates the shape by maximising the per-state Q -function over the bracket,

$$p_k^{(n+1)} = \arg \max_{p \in [p_{\min}, p_{\max}]} \sum_t \gamma_t(k) \log f_k(O_t; \mu_k^{(n+1)}, \alpha_k^{(n+1)}, p), \quad (20)$$

again by a 40-iteration golden-section search over $[p_{\min}, p_{\max}] = [0.5, 3.0]$. Equation (19) is the unique zero of $\partial Q_k / \partial \alpha$, so CM-2 is exact maximization. CM-1 and CM-3 are bracketed maximization. Each CM step is therefore a partial maximization, and the standard ECM monotonicity result [66] applies on the compact bracket; Proposition 1 states the guarantee in full. Boundary cases: $p_k = 2$ recovers Gaussian (with $\sigma = \alpha/\sqrt{2}$); $p_k = 1$ recovers Laplace (with $b = \alpha$).

A.2.2 Algorithm Descriptions

This section provides detailed pseudocode for the main algorithms used in this study. Algorithm 1 is the unified EM procedure shared across all four emission families. It makes explicit the *shared scaffold* (quantile-based initialization, log-space forward-backward, transition update) and the four *family-specific branches* (Gaussian M-step for CHMM-N, ECM with one-dimensional golden-section ν -search for CHMM-t, closed-form weighted-median + weighted-MAD for CHMM-L, three-stage CM with bracketed p_k for CHMM-GED); the branch points are lines 3, 13, and 30. Visualizing the four variants as a single algorithm with four M-step branches makes the architectural contribution of the paper (one EM engine, four interchangeable emission families) immediate from the pseudocode. Algorithm 2 presents the Viterbi algorithm for decoding the most likely hidden state sequence. Algorithm 3 describes the CHMM simulation procedure for generating synthetic return paths. Algorithm 4 describes the cross-asset rank-reordering simulator for copula-based multi-asset synthesis. Algorithm 5 details the SIM regression-and-resampling pipeline.

Algorithm 1 Unified EM algorithm for the CHMM family. Lines 3, 13, and 30 branch on the emission family $F \in \{\mathcal{N}, t, L, \text{GED}\}$; every other line is identical across families. CHMM-N and CHMM-L instantiate classical EM; CHMM-t and CHMM-GED instantiate ECM sequences whose final CM step updates the per-state shape parameter (ν_k for CHMM-t, p_k for CHMM-GED) by a one-dimensional golden-section search on the per-state Q -function. Hyperparameter defaults used throughout this paper: convergence tolerance $\text{tol} = 10^{-4}$ on $\Delta\mathcal{L}$, $\text{max_iter} = 60$, 40-iteration golden-section sub-step on every bracketed CM, CHMM-t bracket $(\nu_{\min}, \nu_{\max}) = (2.1, 50)$ with initial $\nu^{(0)} = 10$, CHMM-GED bracket $(p_{\min}, p_{\max}) = (0.5, 3.0)$ with initial $p^{(0)} = 1.5$ and μ_k search interval $[\hat{\mu}_k - 5\alpha_k, \hat{\mu}_k + 5\alpha_k]$.

Require: Observations $O_{1:T}$, number of states K , tolerance tol , max_iter , emission family $F \in \{\mathcal{N}, t, L, \text{GED}\}$, if $F = t$ bracket $(\nu_{\min}, \nu_{\max})$ and initial $\nu^{(0)}$, if $F = \text{GED}$ bracket $(p_{\min}, p_{\max})$ and initial $p^{(0)}$.

Ensure: Transition matrix $\mathbf{T}$, per-state emission parameters $\theta_{1:K}$, initial distribution π .

- 1: **Quantile-Based Initialization.**
- 2: Sort $O^{(\text{sort})} \leftarrow \text{sort}(O_{1:T})$ and partition into K equal chunks $\{C_1, \dots, C_K\}$.
- 3: **for** $k = 1$ to K **do**
- 4: **switch** F
- 5: $F = \mathcal{N}$: $\mu_k \leftarrow \text{mean}(C_k)$, $\sigma_k \leftarrow \max(\text{std}(C_k), 10^{-6})$.
- 6: $F = t$: $\mu_k \leftarrow \text{mean}(C_k)$, $\sigma_k \leftarrow \max(\text{std}(C_k), 10^{-6})$, $\nu_k \leftarrow \nu^{(0)}$.
- 7: $F = L$: $\mu_k \leftarrow \text{median}(C_k)$, $b_k \leftarrow \max(\text{mean}(|C_k - \mu_k|), 10^{-6})$.
- 8: $F = \text{GED}$: $\mu_k \leftarrow \text{mean}(C_k)$, $\alpha_k \leftarrow \max(\text{std}(C_k), 10^{-6})$, $p_k \leftarrow p^{(0)}$.
- 9: **end for**
- 10: $T_{ij} \leftarrow 1/K$; $\pi_k \leftarrow 1/K$; $\mathcal{L}_{\text{prev}} \leftarrow -\infty$.
- 11: **EM Loop.**
- 12: **for** $n = 1$ to max_iter **do**
- 13: **E-Step, per-family emission log-likelihood.**
- 14: **switch** F
- 15: $F = \mathcal{N}$: $\log B_{t,k} \leftarrow \log \mathcal{N}(O_t | \mu_k, \sigma_k^2)$.
- 16: $F = t$: $\log B_{t,k} \leftarrow \log t_{\nu_k}(O_t; \mu_k, \sigma_k)$.
- 17: $F = L$: $\log B_{t,k} \leftarrow -\log(2b_k) - |O_t - \mu_k|/b_k$.
- 18: $F = \text{GED}$: $\log B_{t,k} \leftarrow \log p_k - \log(2\alpha_k) - \log \Gamma(1/p_k) - (|O_t - \mu_k|/\alpha_k)^{p_k}$.
- 19: **E-Step, shared log-space forward-backward.**
- 20: $\log \alpha_1(k) \leftarrow \log \pi_k + \log B_{1,k}$.
- 21: **for** $t = 2$ to T , $j = 1$ to K **do**
- 22: $\log \alpha_t(j) \leftarrow \underset{i}{\text{logsumexp}}(\log \alpha_{t-1}(i) + \log T_{ij}) + \log B_{t,j}$.
- 23: **end for**
- 24: $\log \beta_T(k) \leftarrow 0$.
- 25: **for** $t = T - 1$ down to 1 , $i = 1$ to K **do**
- 26: $\log \beta_t(i) \leftarrow \underset{j}{\text{logsumexp}}(\log T_{ij} + \log B_{t+1,j} + \log \beta_{t+1}(j))$.
- 27: **end for**
- 28: Compute $\gamma_t(k)$ and $\xi_t(i, j)$ from equation (12).
- 29: **Shared transition update.** $T_{ij} \leftarrow \sum_t \xi_t(i, j) / \sum_t \gamma_t(i)$; $\pi_k \leftarrow \gamma_1(k)$.
- 30: **M-Step, per-family emission update.**
- 31: **switch** F
- 32: $F = \mathcal{N}$ {classical Baum-Welch [16].}
- 33: $\mu_k \leftarrow \sum_t \gamma_t(k) O_t / \sum_t \gamma_t(k)$.
- 34: $\sigma_k \leftarrow \max(\sqrt{\sum_t \gamma_t(k) (O_t - \mu_k)^2 / \sum_t \gamma_t(k)}, 10^{-6})$.
- 35: $F = t$ {ECM of Peel and McLachlan [18], Liu and Rubin [19]; closed-form (μ_k, σ_k) given ν_k , then 1D search on ν_k .}
- 36: $u_{t,k} \leftarrow (\nu_k + 1) / (\nu_k + ((O_t - \mu_k) / \sigma_k)^2)$ for all t, k .
- 37: $\mu_k \leftarrow \sum_t \gamma_t(k) u_{t,k} O_t / \sum_t \gamma_t(k) u_{t,k}$.
- 38: $\sigma_k \leftarrow \max(\sqrt{\sum_t \gamma_t(k) u_{t,k} (O_t - \mu_k)^2 / \sum_t \gamma_t(k)}, 10^{-6})$.
- 39: $\nu_k \leftarrow \arg \max_{\nu \in [\nu_{\min}, \nu_{\max}]} \sum_t \gamma_t(k) \log t_{\nu}(O_t; \mu_k, \sigma_k)$ {40-iter golden-section search.}
- 40: $F = L$ {closed-form weighted Laplace MLE.}
- 41: $\mu_k \leftarrow \text{WeightedMedian}(O_{1:T}, \gamma_{1:T}(k))$ {cumulative-weight bracketing with interpolation at the $W_k/2$ crossing.}
- 42: $b_k \leftarrow \max(\sum_t \gamma_t(k) |O_t - \mu_k| / \sum_t \gamma_t(k), 10^{-6})$.
- 43: $F = \text{GED}$ {three-stage ECM: bracketed μ_k , closed-form α_k , bracketed p_k .}
- 44: $\mu_k \leftarrow \arg \min_{\mu \in [\hat{\mu}_k - 5\alpha_k, \hat{\mu}_k + 5\alpha_k]} \sum_t \gamma_t(k) |O_t - \mu|^{p_k}$ {40-iter golden-section search.}
- 45: $\alpha_k \leftarrow \max\left(\left[(p_k/W_k) \sum_t \gamma_t(k) |O_t - \mu_k|^{p_k}\right]^{1/p_k}, 10^{-6}\right)$ where $W_k = \sum_t \gamma_t(k)$ {closed form, unique zero of $\partial Q_k / \partial \alpha$.}
- 46: $p_k \leftarrow \arg \max_{p \in [p_{\min}, p_{\max}]} \sum_t \gamma_t(k) \log \text{GED}(O_t; \mu_k, \alpha_k, p)$ {40-iter golden-section search.}
- 47: **Convergence check.** $\mathcal{L}^{(n)} \leftarrow \underset{k}{\text{logsumexp}}(\log \alpha_T(k))$.
- 48: **if** $|\mathcal{L}^{(n)} - \mathcal{L}_{\text{prev}}| < \text{tol}$ **then**

Algorithm 2 Viterbi decoding for the most likely hidden state sequence

Require: Observations $O_{1:T}$, Trained CHMM $\mathcal{M} = (\mathbf{T}, \{\mu_k, \sigma_k\}, \boldsymbol{\pi})$

Ensure: Most probable state sequence $S_{1:T}^*$

```
1: for  $k = 1$  to  $K$  do
2:    $\log \delta_1(k) \leftarrow \log(1/K) + \log \mathcal{N}(O_1 \mid \mu_k, \sigma_k^2)$ .
3: end for
4: for  $t = 2$  to  $T$  do
5:   for  $j = 1$  to  $K$  do
6:      $\text{vals}(j) \leftarrow \log \delta_{t-1}(i) + \log T_{ij}$  for all  $i$ .
7:      $\log \delta_t(j) \leftarrow \max_i \text{vals}(i) + \log \mathcal{N}(O_t \mid \mu_j, \sigma_j^2)$ .
8:      $\psi_t(j) \leftarrow \arg \max_i \text{vals}(i)$ .
9:   end for
10: end for
11:  $S_T^* \leftarrow \arg \max_k \log \delta_T(k)$ .
12: for  $t = T - 1$  down to  $1$  do
13:    $S_t^* \leftarrow \psi_{t+1}(S_{t+1}^*)$ .
14: end for
15: return  $S_{1:T}^*$ 
```

Algorithm 3 CHMM simulation of synthetic excess growth rate paths

Require: Trained CHMM $\mathcal{M} = (\mathbf{T}, \{\mu_k, \sigma_k\}, \boldsymbol{\pi})$, Number of steps M , Number of paths P

Ensure: Simulated return paths $\{\hat{G}_{1:M}^{(p)}\}_{p=1}^P$

```
1: for  $p = 1$  to  $P$  do
2:   Sample initial state  $S_1 \sim \text{Categorical}(\boldsymbol{\pi})$ .
3:   for  $t = 1$  to  $M$  do
4:     Emit observation  $\hat{G}_t^{(p)} \sim \mathcal{N}(\mu_{S_t}, \sigma_{S_t}^2)$ .
5:     if  $t < M$  then
6:       Transition to next state  $S_{t+1} \sim \text{Categorical}(\mathbf{T}_{S_t, \cdot})$ .
7:     end if
8:   end for
9: end for
10: return  $\{\hat{G}_{1:M}^{(p)}\}_{p=1}^P$ 
```

Algorithm 4 Cross-asset copula simulation via rank reordering

Require: Fitted per-asset CHMMs $\{\mathcal{M}_j\}_{j=1}^d$, correlation matrix Σ (positive semi-definite (PSD), unit diagonal), copula family $\mathcal{C} \in \{\text{Gaussian}, t_\nu\}$, path length T , number of paths P

Ensure: Cross-asset path tensor $\{\hat{G}_{j,t}^{(p)}\}$ of shape $T \times d \times P$

```
1: for  $p = 1$  to  $P$  do
2:   Draw copula sample  $\mathbf{U}^{(p)} \in [0, 1]^{T \times d}$  from  $\mathcal{C}(\Sigma)$ :
3:    $L \leftarrow \text{Cholesky}(\Sigma)$ ,  $Z \in \mathbb{R}^{T \times d} \sim \mathcal{N}(0, I_d)$  i.i.d.
4:    $Y \leftarrow ZL^\top$ 
5:   if  $\mathcal{C} = \text{Gaussian}$  then
6:      $\mathbf{U}^{(p)} \leftarrow \Phi(Y)$  componentwise
7:   else
8:     Draw  $W \sim \chi_\nu^2/\nu$  i.i.d. for  $t = 1, \dots, T$ ;  $X_{t,\cdot} \leftarrow Y_{t,\cdot}/\sqrt{W_t}$ 
9:      $\mathbf{U}^{(p)} \leftarrow t_\nu(X)$  componentwise
10:  end if
11:  for  $j = 1$  to  $d$  do
12:    Simulate  $\tilde{g}_{j,1:T}^{(p)}$  from  $\mathcal{M}_j$  via Algorithm 3.
13:    Compute order statistics  $\tilde{g}_{j,(1)}^{(p)} \leq \dots \leq \tilde{g}_{j,(T)}^{(p)}$ .
14:    Compute ranks  $r_{j,t}^{(p)} \leftarrow \text{rank}(U_{j,t}^{(p)})$  for  $t = 1, \dots, T$ .
15:     $\hat{G}_{j,t}^{(p)} \leftarrow \tilde{g}_{j,(r_{j,t}^{(p)})}^{(p)}$  for  $t = 1, \dots, T$ .
16:  end for
17: end for
18: return  $\{\hat{G}_{j,t}^{(p)}\}$ 
```

Algorithm 5 Single Index Model (SIM) construction and simulation

Require: IS return matrix $\mathbf{R} \in \mathbb{R}^{T \times d}$, market-column index m , market CHMM $\mathcal{M}_m$, path length T , number of paths P

Ensure: Cross-asset path tensor $\{\hat{G}_{j,t}^{(p)}\}$ of shape $T \times d \times P$

```
1: Ordinary least squares (OLS) fit. For  $j \neq m$ :
2:    $\hat{\beta}_j \leftarrow \text{Cov}(\mathbf{R}_{:,j}, \mathbf{R}_{:,m})/\text{Var}(\mathbf{R}_{:,m})$ 
3:    $\hat{\alpha}_j \leftarrow \overline{\mathbf{R}_{:,j}} - \hat{\beta}_j \overline{\mathbf{R}_{:,m}}$ 
4:    $\hat{\eta}_{j,t} \leftarrow R_{t,j} - \hat{\alpha}_j - \hat{\beta}_j R_{t,m}$ ,  $R_j^2 \leftarrow 1 - \sum_t \hat{\eta}_{j,t}^2 / \sum_t (R_{t,j} - \overline{\mathbf{R}_{:,j}})^2$ 
5: for  $p = 1$  to  $P$  do
6:   Simulate market path  $G_{m,1:T}^{(p)}$  from  $\mathcal{M}_m$  via Algorithm 3.
7:    $\hat{G}_{m,t}^{(p)} \leftarrow G_{m,t}^{(p)}$ .
8:   for  $j \neq m$  do
9:     Draw residual indices  $\{\iota_t\}_{t=1}^T \sim \text{Uniform}(\{1, \dots, T\})$ .
10:     $\hat{G}_{j,t}^{(p)} \leftarrow \hat{\alpha}_j + \hat{\beta}_j G_{m,t}^{(p)} + \hat{\eta}_{j,\iota_t}$ .
11:  end for
12: end for
13: return  $\{\hat{G}_{j,t}^{(p)}\}$ 
```

A.2.3 Per-State Shape Diagnostics: CHMM-t and CHMM-GED

This appendix backs the discussion in Section 5 of the per-state shape allocation under CHMM-t and CHMM-GED on SPY IS at $K = 18$ (seed = 20260420).

CHMM-t per-state ν_k . Figure 1 plots the per-state ν_k histogram for the $K = 18$ CHMM-t fit. Thirteen of the eighteen states rest at the upper ECM bracket ($\nu_k = 50$), and only two states lie near the lower bracket ($\nu_2 = 2.19$, $\nu_4 = 2.10$). The simulated IS mixture kurtosis of 17.34 is driven by the posterior mass routed to those two very heavy-tailed states, not by a systemic collapse to the lower bracket.




Figure 1. Per-state degrees-of-freedom ν_k for the CHMM-t fit at $K = 18$ on SPY IS. The dashed red line marks the lower ECM bracket $\nu_{\min} = 2.1$. Only two states ($\nu_2 = 2.19$, $\nu_4 = 2.10$) sit near the lower bracket; the median ν_k is at the upper bracket.

CHMM-GED per-state p_k (Gaussian-bulk / Laplace-tail bimodality). The body Section 4.3 reports and discusses the $\hat{p}_k$ partition (Figure 28); see also Table 10 below for the parallel CHMM-t bracket sweep. Briefly, the $K = 18$ CHMM-GED fit on SPY IS is bimodal: thirteen states attain the upper ECM bracket at $p_k = 3.0$, one state lands at intermediate $p_k = 1.6$, and four states cluster in the Laplace-shape regime $p_k \in [0.86, 1.24]$, concentrated in the high-volatility ranks. The smallest fitted p_k on SPY is 0.86, well clear of every $p_{\min}$ tested in the bracket sweep, so the lower bracket is non-binding. The bimodal partition replicates across 10 Monte Carlo seeds and across the cross-asset universe (Appendix A.7).

Table 10 reports the seven-metric panel for CHMM-t fits whose ECM bracket lower bound is lifted from $\nu_{\min} = 2.1$ through $\{2.5, 3.0, 4.0\}$, plus the Gaussian limit ($\nu = \infty$, equivalent to CHMM-N). Lifting the bracket does not reduce the IS kurtosis overshoot: the simulated kurtosis moves within the 12.5–14.0 band across bracket floors and only collapses to 5.05 in the full Gaussian limit. The overshoot is therefore a feature of the Student-t emission family at $K = 18$ and not a lower-bracket artefact.

Table 11 reports the analogous bracket sensitivity for CHMM-GED on $p_{\max} \in \{2.0, 2.5, 3.0, 3.5, 4.0\}$ (with $p_{\min} = 0.5$ fixed) and on $p_{\min} \in \{0.5, 0.7, 0.85\}$ (with $p_{\max} = 3.0$ fixed). The headline reads cleanly. The Gaussian-like state count is invariant at 13/18 across every upper bracket; only the

Table 10. CHMM-t $\nu_{\min}$ bracket sensitivity at $K = 18$ on SPY (seed = 20260420, 1,000 paths). AD = Anderson–Darling pass rate; W_1 = Wasserstein-1; H = Hellinger distance (definitions in Appendix A.3).

$\nu_{\min}$	KS IS	AD IS	KS OoS	AD OoS	Kurt IS	Kurt OoS	ACF-MAE	W_1	H
2.1	95.7	90.2	86.2	82.0	13.95	10.08	0.0546	0.100	0.0759
2.5	95.8	90.6	84.2	79.1	15.12	10.34	0.0547	0.100	0.0759
3.0	93.8	89.5	85.3	80.4	12.25	9.65	0.0548	0.099	0.0760
4.0	95.5	89.5	85.0	80.3	13.25	9.39	0.0550	0.101	0.0759
∞ (Gauss)	94.1	86.7	84.2	75.5	5.03	4.46	0.0510	0.110	0.0809
<i>Observed</i>	–	–	–	–	7.68	5.29	–	–	–

redistribution between intermediate and Laplace-like states moves slightly with $p_{\max}$, and the total heavy-shape count stays in $\{4, 5\}$. Tightening $p_{\max}$ from 4.0 to the canonical 3.0 costs about 4 nats of incomplete-data log-likelihood. The $p_{\min}$ sweep is fully degenerate: every $p_{\min} \in [0.5, 0.85]$ produces an identical fit because the smallest observed $\hat{p}_k$ on SPY is 0.86, so the lower bracket is non-binding.

Table 11. CHMM-GED bracket sensitivity at $K = 18$ on SPY (seed = 20260420, 1,000 paths). $\hat{p}_{\min}$, $\hat{p}_{\text{med}}$, $\hat{p}_{\max}$ are the across-state minimum, median, and maximum of the fitted $\hat{p}_k$. nG / nI / nL / nS are the partition counts (Gaussian-like $\hat{p}_k \geq 1.85$, intermediate $1.30 \leq \hat{p}_k < 1.85$, Laplace-like $0.85 \leq \hat{p}_k < 1.30$, super-Laplace $\hat{p}_k < 0.85$). LL is the incomplete-data log-likelihood at convergence.

Sweep	$p_{\min}$	$p_{\max}$	LL	$\hat{p}_{\min}$	nG	nI	nL	nS	KS IS	AD IS	KS OoS	AD OoS	Kurt IS	ACF-MAE	W_1	H
$p_{\max}$ sweep	0.5	2.0	−4732.0	0.89	13	0	5	0	95.8	91.0	83.8	81.2	5.02	0.0541	0.099	0.0781
	0.5	2.5	−4725.3	0.87	13	2	3	0	96.1	92.4	84.3	80.1	5.09	0.0549	0.098	0.0778
	0.5	3.0	−4720.3	0.86	13	1	4	0	96.3	92.7	83.8	81.3	5.05	0.0546	0.098	0.0779
	0.5	3.5	−4717.5	0.86	13	1	4	0	96.3	91.8	86.7	82.4	5.11	0.0544	0.099	0.0782
	0.5	4.0	−4716.3	0.86	13	1	4	0	95.9	91.3	84.4	80.6	5.18	0.0543	0.099	0.0785
$p_{\min}$ sweep	0.50	3.0	−4720.3	0.86	13	1	4	0	96.3	92.7	83.8	81.3	5.05	0.0546	0.098	0.0779
	0.70	3.0	−4720.3	0.86	13	1	4	0	96.3	92.7	83.8	81.3	5.05	0.0546	0.098	0.0779
	0.85	3.0	−4720.3	0.86	13	1	4	0	96.3	92.7	83.8	81.3	5.05	0.0546	0.098	0.0779

A.2.4 Pipeline Schematic

Pipeline A (single-asset, Section 4.3):



Pipeline B (cross-asset, Section 4.5):




Figure 2. Schematic of Pipeline A (single-asset) and Pipeline B (cross-asset). Pipeline A fits a CHMM per ticker independently. Pipeline B reuses the per-asset fits as marginals and injects cross-asset dependence through a rank-based Student- t copula. The single-asset scaffold is shared across both pipelines.

A.2.5 CHMM Architecture and Training Loop




Figure 3. CHMM architecture and training loop. A K -state latent Markov chain with transition matrix $\mathbf{T}$ generates the regime sequence s_t ; each s_t emits the observed excess growth rate G_t through a state-specific density b_{s_t} . The four CHMM variants share everything except the emission family. EM alternates an E-step (log-space forward-backward) and a family-specific M-step until $|\Delta\mathcal{L}| < 10^{-4}$.

A.3 Validation Metric Details

This appendix documents the precise form of the seven per-path metrics referenced in Section 3.5. For each model, 1,000 independent paths of length T are simulated (200 paths for the cross-asset extension) and each metric is computed per path before aggregation.

Kolmogorov–Smirnov (KS) pass rate. The two-sample KS statistic [56, 57] measures the maximum vertical distance between two empirical CDFs,

$$D_{n,m} = \sup_x |F_n(x) - G_m(x)|. \quad (21)$$

The KS pass rate is the fraction of paths for which the test fails to reject distributional equivalence against the reference data at $\alpha = 0.05$.

Anderson–Darling (AD) pass rate. The AD test is the integrated squared CDF deviation with a quadratic tail weighting, making it more sensitive than KS to tail discrepancies. The AD pass rate is the fraction of paths for which the two-sample AD test fails to reject at $\alpha = 0.05$.

Excess kurtosis. Mean simulated excess kurtosis across paths, compared to the observed value. Given Gaussian emissions, the CHMM is expected to underestimate kurtosis relative to the heavy-tailed empirical distribution, and the shortfall is one of the reported diagnostics.

ACF-MAE. Mean absolute error between observed and mean-simulated autocorrelation functions of $|G_t|$ over $L = 252$ lags,

$$\text{ACF-MAE} = \frac{1}{L} \sum_{\tau=1}^L \left| \hat{\rho}_{|G|}^{\text{obs}}(\tau) - \bar{\rho}_{|G|}^{\text{sim}}(\tau) \right|. \quad (22)$$

This is the direct diagnostic for the Rydén et al. limitation.

Wasserstein-1 distance. The mean absolute difference of sorted empirical quantiles, equivalent to the L^1 distance between the empirical CDFs [67],

$$W_1(F, G) = \int_0^1 |F^{-1}(u) - G^{-1}(u)| du. \quad (23)$$

Hellinger distance. A histogram-based overlap distance in $[0, 1]$,

$$H(P, Q) = \frac{1}{\sqrt{2}} \sqrt{\sum_{i=1}^B (\sqrt{p_i} - \sqrt{q_i})^2}, \quad (24)$$

with B equal-width bins spanning the joint support and p_i, q_i the observed and simulated bin probabilities.

Quantile coverage. The fraction of 99 equally-spaced observed quantiles (1st through 99th percentile) falling inside the 5th–95th percentile envelope of the corresponding simulated quantile across 1,000 paths. A coverage of 100% indicates that the simulation envelope fully brackets the observed distribution at every percentile.

Cross-asset metrics. For the cross-asset extension we additionally track the mean Frobenius norm $\|\hat{\Sigma}_{\text{sim}}^{(p)} - \hat{\Sigma}_{\text{obs}}\|_F$ and the off-diagonal Pearson correlation MAE between simulated and observed return matrices, averaged over paths. The off-diagonal MAE is averaged over the $d(d-1)/2$ unique upper-triangular entries of the symmetric correlation matrix (so 15 entries for the six-asset cross-section of Section 4.5), not double-counted across both triangles:

$$\text{MAE}_{\text{off-diag}} = \frac{1}{P} \sum_{p=1}^P \frac{2}{d(d-1)} \sum_{1 \leq i < j \leq d} \left| \hat{\Sigma}_{\text{sim}, ij}^{(p)} - \hat{\Sigma}_{\text{obs}, ij} \right|.$$

A.3.1 Continuous Ranked Probability Score and Diebold-Mariano

The OoS CRPS column of Table 3 reports a proper-scoring-rule complement to the KS pass rate. We use the unbiased sample CRPS estimator: for an ensemble of N simulated paths $\{x_i\}_{i=1}^N$ at time t and observed value y_t ,

$$\widehat{\text{CRPS}}_t = \frac{1}{N} \sum_{i=1}^N |x_i - y_t| - \frac{1}{N(N-1)} \sum_{1 \leq i < j \leq N} |x_i - x_j|. \quad (25)$$

The double sum is computed in $O(N \log N)$ via the sorted-ensemble identity $\sum_{i < j} (x_{(j)} - x_{(i)}) = \sum_i x_{(i)} (2i - N - 1)$. The ensemble at each t is the cross-section of the $N = 1,000$ unconditional simulated paths used throughout Section 4; this yields a marginal-predictive CRPS suitable for an unconditional generative-fidelity assessment, distinct from the conditional one-step-ahead CRPS used in forecasting.

For pairwise comparison we report a Diebold-Mariano test on the per- t loss differential $d_t = \widehat{\text{CRPS}}_t^A - \widehat{\text{CRPS}}_t^B$ with a Newey-West HAC variance under a Bartlett kernel and bandwidth $h = \lfloor T_{\text{OoS}}^{1/3} \rfloor = 8$. Two-sided p -values are reported under the standard-normal null. The body claim of Section 4.3 is: CHMM ties or beats every benchmark on mean OoS CRPS, with Gaussian i.i.d. the only significantly worse comparator.

A.3.2 Christoffersen Conditional-Coverage Panel (Companion to Section 4.6)

The body Section 4.6 reports the unconditional Kupiec LR_{uc} statistic. This appendix reports the full Christoffersen [68] panel: LR_{uc} (unconditional coverage), LR_{ind} (breach independence), and $\text{LR}_{\text{cc}} = \text{LR}_{\text{uc}} + \text{LR}_{\text{ind}}$ (joint conditional coverage), at $\alpha \in \{0.01, 0.05\}$ on the SPY IS and OoS windows. The breach series is constructed from the pooled-archive VaR threshold (the α -quantile of $\text{vec}(\text{sim archive})$ across paths and time), then $\text{br}_t = (R_t \leq \widehat{\text{VaR}}_\alpha)$.


The headline reads as follows. On the unconditional (Kupiec) leg every CHMM variant passes at both levels on both windows; the LR_{uc} statistic for the four CHMM rows at $\alpha = 0.05$ on OoS clusters at 3.03–3.83 (just below the χ_1^2 critical value 3.841), reflecting the integer-breach grid at $T_{\text{OoS}} = 572$ noted in the body table caption. On the independence leg every unconditional generator in the panel (bootstrap, GARCH, and all four CHMM variants) rejects: LR_{ind} is in the range 4.71–56.46 across rows. This is structural rather than CHMM-specific. Volatility-clustering in R_{oos} produces persistent regimes of high and low realised volatility, the OoS-window breaches concentrate inside the high-volatility regime, and a constant-across- t VaR threshold (the construction every unconditional generator uses) cannot track that timing. The natural fix is a regime-conditional VaR that propagates the CHMM latent-state forecast $\mathbb{P}(s_{t+1} | \text{history}_t)$ through the per-state emission-mixture; this is operationally the conditional-VaR companion-paper direction noted in Section 5.

Table 12. Christoffersen conditional-coverage VaR back-test (SPY, seed = 20260420, $K = 18$, 1,000 simulated paths). Critical values: $\chi_1^2(0.05) = 3.841$, $\chi_2^2(0.05) = 5.991$. “br rate” is the empirical breach rate; LR_{uc} is the Kupiec statistic ($\sim \chi_1^2$ under correct unconditional coverage); LR_{ind} is the Christoffersen independence statistic ($\sim \chi_1^2$ under independent breaches); $LR_{cc} = LR_{uc} + LR_{ind}$ is the joint conditional-coverage statistic ($\sim \chi_2^2$). PASS if statistic is below critical. Every unconditional generator (i.i.d. bootstrap, GARCH, and all four CHMM variants) rejects independence on OoS at $\alpha = 0.05$, consistent with breach clustering driven by volatility-clustering in R_{oos} rather than a CHMM-specific mis-specification of conditional dynamics.


Model	Win	α	breaches	br rate	Kupiec		Christ. ind		Christ. cc	
					LR_{uc}	p	LR_{ind}	p	LR_{cc}	p
Bootstrap	IS	0.01	26	1.03%	0.03	0.87	15.28	0.00	15.31	0.00
Bootstrap	IS	0.05	126	5.01%	0.00	0.99	56.46	0.00	56.46	0.00
Bootstrap	OoS	0.01	5	0.87%	0.10	0.76	13.18	0.00	13.28	0.00
Bootstrap	OoS	0.05	19	3.32%	3.83	0.05	5.26	0.02	9.08	0.01
GARCH(1,1)	IS	0.01	31	1.23%	1.28	0.26	12.45	0.00	13.72	0.00
GARCH(1,1)	IS	0.05	139	5.52%	1.41	0.24	45.38	0.00	46.79	0.00
GARCH(1,1)	OoS	0.01	6	1.05%	0.01	0.91	20.87	0.00	20.89	0.00
GARCH(1,1)	OoS	0.05	24	4.20%	0.82	0.37	5.87	0.02	6.69	0.04
CHMM-N	IS	0.01	20	0.79%	1.15	0.28	12.80	0.00	13.95	0.00
CHMM-N	IS	0.05	123	4.89%	0.07	0.80	47.33	0.00	47.40	0.00
CHMM-N	OoS	0.01	3	0.52%	1.58	0.21	18.98	0.00	20.56	0.00
CHMM-N	OoS	0.05	19	3.32%	3.83	0.05	5.26	0.02	9.08	0.01
CHMM-t	IS	0.01	22	0.87%	0.42	0.52	11.62	0.00	12.04	0.00
CHMM-t	IS	0.05	127	5.05%	0.01	0.91	55.54	0.00	55.55	0.00
CHMM-t	OoS	0.01	4	0.70%	0.58	0.45	15.53	0.00	16.12	0.00
CHMM-t	OoS	0.05	20	3.50%	3.03	0.08	4.71	0.03	7.74	0.02
CHMM-L	IS	0.01	19	0.76%	1.66	0.20	13.44	0.00	15.11	0.00
CHMM-L	IS	0.05	134	5.33%	0.55	0.46	49.43	0.00	49.98	0.00
CHMM-L	OoS	0.01	2	0.35%	3.26	0.07	9.15	0.00	12.41	0.00
CHMM-L	OoS	0.05	20	3.50%	3.03	0.08	4.71	0.03	7.74	0.02
CHMM-GED	IS	0.01	20	0.79%	1.15	0.28	12.80	0.00	13.95	0.00
CHMM-GED	IS	0.05	126	5.01%	0.00	0.99	56.46	0.00	56.46	0.00
CHMM-GED	OoS	0.01	2	0.35%	3.26	0.07	9.15	0.00	12.41	0.00
CHMM-GED	OoS	0.05	19	3.32%	3.83	0.05	5.26	0.02	9.08	0.01

A.3.3 VaR / ES Envelope Visualization (Companion to Table 41)


Figure 4 visualizes the median and [5%, 95%] envelopes across simulated paths whose numeric values are reported in Table 41.




(a) IS VaR ($\alpha = 0.01$).



(b) IS ES ($\alpha = 0.01$).



(c) OoS VaR ($\alpha = 0.05$).



(d) OoS ES ($\alpha = 0.05$).

Figure 4. VaR / ES envelope diagnostic (SPY, $N_{\text{paths}} = 1,000$; global seed policy in Section 3.5). Each panel shows the simulated median (dot) and [5%, 95%] envelope (error bar) across paths against the observed historical value (dashed red line); panels (a)/(b) are the IS window at $\alpha = 0.01$, panels (c)/(d) are the OoS window at $\alpha = 0.05$, and the x-axis in each panel lists the five generators kept for the volatility-aware conditional-coverage comparison (GARCH(1,1), CHMM-N, CHMM-t, CHMM-L, CHMM-GED). The i.i.d. bootstrap is omitted here because its envelope is by construction matched to the empirical marginal and does not discriminate among volatility-aware models. The y-axis is the annualized log excess growth rate (daily log return scaled by $1/\Delta t$ with $\Delta t = 1/252$), consistent with the $R_{\text{IS}}/R_{\text{OoS}}$ convention used throughout the paper. Headline: every observed red line falls inside the corresponding simulated envelope, confirming that the CHMM family passes the coverage leg of the back-test complementary to the Kupiec likelihood-ratio leg in Section 4.6.

A.4 Formal Theoretical Statements

The following assumptions and propositions are referenced in Section 3.8 and stated formally here for reference; each is a one-line specialisation of a standard result and the proof sketch is one sentence.

Assumption 1 (Irreducibility and aperiodicity). *The transition matrix $\mathbf{T} \in \mathbb{R}^{K \times K}$ is irreducible and aperiodic on the state space $\{1, \dots, K\}$, so the chain admits a unique stationary distribution $\bar{\pi}$ satisfying $\bar{\pi} \mathbf{T} = \bar{\pi}$ and $\lambda_1 = 1$ is a simple eigenvalue with all other eigenvalues strictly inside the unit disk [61].*

Assumption 2 (Finite second moment and emission contrast). *Each per-state emission density $f_k(\cdot; \theta_k)$ has finite second moment, $\mathbb{E}_{f_k}[G^2] < \infty$, and at least two states differ in either location μ_k or scale σ_k .*

Proposition 1 (ECM monotonicity for CHMM-t and CHMM-GED (and EM monotonicity for CHMM-N, CHMM-L)). *Under compactness of the parameter space (μ_k, σ_k or β_k or $\alpha_k, \nu_k \in [\nu_{\min}, \nu_{\max}]$, $p_k \in [p_{\min}, p_{\max}]$ each in a closed bounded interval, $T_{ij} \geq \epsilon_T > 0$), the two-block CM step for CHMM-t of Section 3.7, namely the closed-form weighted updates of $(\mathbf{T}, \boldsymbol{\pi}, \mu_k, \sigma_k)$ at fixed ν_k , followed by a golden-section maximisation of the per-state Q -function over ν_k , satisfies $Q(\theta^{(n+1)} | \theta^{(n)}) \geq Q(\theta^{(n)} | \theta^{(n)})$, and consequently $\mathcal{L}(\theta^{(n+1)}) \geq \mathcal{L}(\theta^{(n)})$ where $\mathcal{L}$ is the incomplete-data log-likelihood. The same monotonicity holds for the closed-form CHMM-N and CHMM-L M-steps under the standard EM guarantee, and for the three-stage CHMM-GED CM sequence (bracketed μ_k , closed-form α_k , bracketed p_k) of equations (18)–(20).*

Proof sketch. Each CM step does not decrease Q by definition (closed-form steps are exact maximisers, golden-section steps are bracketed maxima within the prescribed interval); the standard EM inequality $\mathcal{L}(\theta^{(n+1)}) - \mathcal{L}(\theta^{(n)}) \geq Q(\theta^{(n+1)} | \theta^{(n)}) - Q(\theta^{(n)} | \theta^{(n)})$ [66, 69] converts Q -monotonicity into log-likelihood monotonicity. Compactness gives boundedness above, hence convergence.

Numerical full-rank check on $\hat{\mathbf{T}}$. The Allman-Matias-Rhodes 2009 identifiability theorem requires $\hat{\mathbf{T}}$ of full rank. We compute the singular-value spectrum of $\hat{\mathbf{T}}$ at $K = 18$ across the four emission families on SPY IS (Table 13). All four matrices are full-rank in the strict sense, with smallest singular value $\sigma_{\min}(\hat{\mathbf{T}}) \in [0.0007, 0.0170]$ and condition numbers in the range 63–1,620. CHMM-GED is the most well-conditioned ($\sigma_{\min} = 0.017$, $\text{cond} = 63.5$); CHMM-t is the least ($\sigma_{\min} = 0.0007$, $\text{cond} = 1,620$), consistent with the per-state ν_k ECM lower-bracket pinning making one or two heavy-tail states near-degenerate. The deflated matrix $\hat{\mathbf{T}} - \mathbf{1}\bar{\boldsymbol{\pi}}^\top$ has numerical rank exactly $K - 1 = 17$ across all four families, the rank statement underlying Section 3.8.

Table 13. Numerical full-rank check on $\hat{\mathbf{T}}$ at $K = 18$ across four CHMM emission families on SPY IS (seed = 20260420). $\sigma_{\min}$, $\sigma_{\max}$, and the condition number $\sigma_{\max}/\sigma_{\min}$ are reported on $\hat{\mathbf{T}}$ itself; $|\lambda_2|$ is the second-largest absolute eigenvalue (the dominant decay mode in the ACF identity, equation (6)); “rank(def)” is the numerical rank of $\hat{\mathbf{T}} - \mathbf{1}\bar{\boldsymbol{\pi}}^\top$ at tolerance 10^{-10} , expected to equal $K - 1 = 17$.

Family	$\sigma_{\min}(\hat{\mathbf{T}})$	$\sigma_{\max}(\hat{\mathbf{T}})$	$\text{cond } \hat{\mathbf{T}}$	$ \lambda_2 $	$\text{rank}(\hat{\mathbf{T}} - \mathbf{1}\bar{\boldsymbol{\pi}}^\top)$
CHMM-N	0.00532	1.0866	204	0.929	17/17
CHMM-t	0.00067	1.0796	1,620	0.914	17/17
CHMM-L	0.00398	1.0928	275	0.896	17/17
CHMM-GED	0.01702	1.0806	63.5	0.911	17/17

Proposition 2 (Identifiability up to label permutation). *Under irreducibility and aperiodicity of $\mathbf{T}$ (Assumption 1), the emission contrast condition (Assumption 2), and pairwise distinctness of the per-state parameter tuples θ_k across $k = 1, \dots, K$, the CHMM parameter is identifiable up to label permutation for CHMM-N (Gaussian) and CHMM-L (Laplace). For CHMM-t (with per-state $\nu_k \in [\nu_{\min}, \nu_{\max}]$) and CHMM-GED (with per-state $p_k \in [p_{\min}, p_{\max}]$) the same conclusion holds modulo the standard caveat that pairwise distinctness must hold jointly across all components of $\theta_k = (\mu_k, \sigma_k, \nu_k)$ or (μ_k, α_k, p_k) , since within either bracket a Gaussian-shaped emission with inflated scale can mimic a Laplace-shaped emission with deflated scale at intermediate quantiles.*

Proof sketch. The HMM identifiability theorem of Allman et al. [70] requires $\mathbf{T}$ of full rank and per-state emissions linearly independent in the space of densities on $\mathbb{R}$. Full-rank $\mathbf{T}$ follows from irreducibility [61]. For CHMM-N and CHMM-L, linear independence of Gaussian and Laplace emissions with distinct location-scale parameters is direct from Yakowitz and Spragins [71], who establish that finite mixtures over location-scale families with strictly decreasing-tail densities are identifiable (the CHMM-N and CHMM-L marginals are special cases). For CHMM-t at fixed shape and CHMM-GED at fixed shape, location-scale identifiability is the same Yakowitz-Spragins result; for finite mixtures of Student- t with varying degrees-of-freedom the standard reference is Lindsay [72], Chapter 3, which gives identifiability via the moment-generating function over the bracketed range of shape parameters and a non-degeneracy condition on the joint location-scale-shape tuple. The CHMM-GED case under bracketed $p_k \in [p_{\min}, p_{\max}]$ is a parallel application of the moment-matching argument: the GED characteristic function $\hat{b}(t; \mu, \alpha, p) = \exp(i\mu t) M(t; \alpha, p)$ has a unique inverse over the bounded shape range up to permutation of components, modulo the joint-distinctness caveat on (μ_k, α_k, p_k) . Strict linear independence at the boundary cases $p_k = 1$ (Laplace) and $p_k = 2$ (Gaussian) is automatic from Yakowitz-Spragins; in the bracket interior, the joint distinctness condition above is required to rule out the location-scale-shape trade-off.

Proposition 3 (MLE consistency). *Under Assumptions 1–2 and the assumption that the true parameter lies in the interior of a compact Θ , the MLE $\hat{\theta}_T = \arg \max_{\theta \in \Theta} \mathcal{L}_T(\theta)$ is consistent: $\hat{\theta}_T \xrightarrow{P} \theta_0$ as $T \rightarrow \infty$, modulo label permutation.*

Proof reference. Bickel et al. [73] (Theorem 1) and Douc et al. [74] apply directly under our assumption set: irreducibility supplies the mixing condition, finite second moment the moment condition, compactness the parameter-space condition.

Proposition 4 (Marginal preservation under rank reordering). *Let $\tilde{g}_{j,1:T}$ be a sample of size T from the per-asset CHMM marginal of asset j , and let the rank-reordered output be $\hat{g}_{j,t}^{(p)} = \tilde{g}_{j,(r_{j,t}^{(p)})}^{(p)}$ where $r_{j,t}^{(p)}$ is the rank of the copula sample. Then for every asset j the empirical CDF of $\{\hat{g}_{j,t}^{(p)}\}_{t=1}^T$ equals that of $\{\tilde{g}_{j,t}^{(p)}\}_{t=1}^T$.*

Proof sketch. The map $t \mapsto r_{j,t}^{(p)}$ is a permutation almost surely, so the multiset of values is preserved; the empirical CDF depends only on the multiset.

A.5 Spectral ACF Identity: Extended Derivation

Equation (6) is the closed-form mixture-of-eigenvalues form for the absolute-return ACF used in Section 3.8 and invoked throughout Sections 4–5. The bilinear cross-moment identity $\mathbb{E}[|G_t| |G_{t+\tau}|] = \mathbf{m}^\top \text{diag}(\bar{\pi}) \mathbf{T}^\tau \mathbf{m}$ and its eigendecomposition over the non-unit eigenvalues of $\mathbf{T}$ are textbook material in the regime-switching literature (13, Section 22.2; 14, Chapter 3; the moment-generating-function form for general Markov-switching specifications with conditional means and variances depending on the latent state is given by 15); the underlying Markov-chain spectral decomposition is standard [61]. We restate the absolute-return specialisation here as a self-contained theorem with a step-by-step proof, and record the lag-zero behaviour and the complex / non-diagonalisable extensions that the main body does not. Our contribution on this axis is the explicit application of the spectral form to recast the empirical Rydén et al. [11] low- K failure as a rank statement on $\mathbf{T} - \mathbf{1}\bar{\pi}^\top$ (Section 3.8, “Rydén separation as a K -rank statement”) and the empirical demonstration that the rank constraint is the binding axis at moderate K on equity-return data (Sections 4, 5).

Theorem 1 (Mixture-of-eigenvalues identity for the absolute-return ACF). *Let $\{(s_t, G_t)\}_{t \in \mathbb{Z}}$ be a stationary CHMM on $\{1, \dots, K\}$ with irreducible aperiodic transition matrix $\mathbf{T}$ (Assumption 1) and per-state emissions of finite second moment (Assumption 2). Write $\bar{\boldsymbol{\pi}}$ for the unique stationary distribution, $\mathbf{D} = \text{diag}(\bar{\boldsymbol{\pi}})$, $m_k = \mathbb{E}[|G_t| \mid s_t = k]$, $\mathbf{m} = (m_1, \dots, m_K)^\top$, $\mu = \mathbb{E}[|G_t|] = \bar{\boldsymbol{\pi}}^\top \mathbf{m}$, and $\sigma_{|G|}^2 = \text{Var}(|G_t|)$. Assume $\mathbf{T}$ is diagonalisable with right eigenvectors $\mathbf{v}_k$ and left eigenvectors $\mathbf{w}_k$ biorthonormal, $\mathbf{w}_k^\top \mathbf{v}_l = \delta_{kl}$, indexed so that $\lambda_1 = 1 > |\lambda_2| \geq \dots \geq |\lambda_K|$. Then for every integer lag $\tau \geq 1$,*

$$\rho_{|G|}(\tau) = \sum_{k=2}^K w_k \lambda_k^\tau, \quad w_k = \frac{(\mathbf{m}^\top \mathbf{D} \mathbf{v}_k) (\mathbf{w}_k^\top \mathbf{m})}{\sigma_{|G|}^2}.$$

Proof. The proof is organised in six steps.

Step 1 (Definitions). The autocorrelation of $|G_t|$ at lag $\tau \geq 0$ is, by stationarity,

$$\rho_{|G|}(\tau) = \frac{\text{Cov}(|G_t|, |G_{t+\tau}|)}{\sigma_{|G|}^2} = \frac{\mathbb{E}[|G_t| |G_{t+\tau}|] - \mu^2}{\sigma_{|G|}^2}.$$

The task reduces to evaluating $\mathbb{E}[|G_t| |G_{t+\tau}|]$ in closed form for $\tau \geq 1$.

Step 2 (Conditioning on the latent state pair). For $\tau \geq 1$, by the law of total expectation conditional on $(s_t, s_{t+\tau})$,

$$\mathbb{E}[|G_t| |G_{t+\tau}|] = \sum_{i,j=1}^K \mathbb{P}(s_t = i, s_{t+\tau} = j) \mathbb{E}[|G_t| |G_{t+\tau}| \mid s_t = i, s_{t+\tau} = j].$$

The HMM emission structure gives the conditional independence $G_t \perp G_{t+\tau} \mid (s_t, s_{t+\tau})$ for $\tau \geq 1$: each emission depends only on its current state, and integrating the intermediate states $s_{t+1:t+\tau-1}$ out of the full-sequence factorisation $(\prod_u b_{s_u}(G_u)) p(s_{t+1:t+\tau-1} \mid s_t, s_{t+\tau})$ leaves the endpoint emissions decoupled. Hence

$$\mathbb{E}[|G_t| |G_{t+\tau}| \mid s_t = i, s_{t+\tau} = j] = \mathbb{E}[|G_t| \mid s_t = i] \mathbb{E}[|G_{t+\tau}| \mid s_{t+\tau} = j] = m_i m_j.$$

The joint state probability factors by stationarity and the Markov property, $\mathbb{P}(s_t = i, s_{t+\tau} = j) = \bar{\pi}_i (\mathbf{T}^\tau)_{ij}$. Substituting and noting $[\mathbf{D} \mathbf{T}^\tau]_{ij} = \bar{\pi}_i (\mathbf{T}^\tau)_{ij}$,

$$\mathbb{E}[|G_t| |G_{t+\tau}|] = \mathbf{m}^\top \mathbf{D} \mathbf{T}^\tau \mathbf{m}, \quad \tau \geq 1. \quad (26)$$

Step 3 (Spectral preliminaries). For irreducible aperiodic $\mathbf{T}$, Perron–Frobenius applied to row-stochastic matrices gives $\lambda_1 = 1$ as a simple eigenvalue with all other eigenvalues strictly inside the unit disk [61]. The right eigenvector for λ_1 is $\mathbf{v}_1 = \mathbf{1}$, since $\mathbf{T} \mathbf{1} = \mathbf{1}$ from the row-sum constraint, and the left eigenvector is $\mathbf{w}_1 = \bar{\boldsymbol{\pi}}$, since $\bar{\boldsymbol{\pi}}^\top \mathbf{T} = \bar{\boldsymbol{\pi}}^\top$ by definition of the stationary distribution. Biorthonormality $\mathbf{w}_1^\top \mathbf{v}_1 = \bar{\boldsymbol{\pi}}^\top \mathbf{1} = 1$ is automatic since $\bar{\boldsymbol{\pi}}$ is a probability vector.

Step 4 (Spectral decomposition of $\mathbf{T}^\tau$). Diagonalisability with biorthonormality gives the dyadic expansion

$$\mathbf{T}^\tau = \sum_{k=1}^K \lambda_k^\tau \mathbf{v}_k \mathbf{w}_k^\top = \mathbf{1} \bar{\boldsymbol{\pi}}^\top + \sum_{k=2}^K \lambda_k^\tau \mathbf{v}_k \mathbf{w}_k^\top,$$

isolating the dominant ($k = 1$) projection from the contractive spectrum. The dominant projector $\mathbf{P}_1 = \mathbf{1} \bar{\boldsymbol{\pi}}^\top$ is the rank-one stationary projection: $\mathbf{P}_1 \mathbf{x} = (\bar{\boldsymbol{\pi}}^\top \mathbf{x}) \mathbf{1}$.

Step 5 (Evaluation of the bilinear form). Substituting the spectral decomposition into (26),

$$\mathbb{E}[|G_t| |G_{t+\tau}|] = \mathbf{m}^\top \mathbf{D} \mathbf{1} \bar{\boldsymbol{\pi}}^\top \mathbf{m} + \sum_{k=2}^K \lambda_k^\tau (\mathbf{m}^\top \mathbf{D} \mathbf{v}_k) (\mathbf{w}_k^\top \mathbf{m}).$$

The $k = 1$ term simplifies to μ^2 : $\mathbf{D} \mathbf{1}$ is the column vector $\bar{\boldsymbol{\pi}}$, so $\mathbf{m}^\top \mathbf{D} \mathbf{1} = \bar{\boldsymbol{\pi}}^\top \mathbf{m} = \mu$, and likewise $\bar{\boldsymbol{\pi}}^\top \mathbf{m} = \mu$, giving the product $\mu \cdot \mu = \mu^2$. Denoting the $k \geq 2$ scalar coefficients

$$c_k = (\mathbf{m}^\top \mathbf{D} \mathbf{v}_k) (\mathbf{w}_k^\top \mathbf{m}),$$

we obtain

$$\mathbb{E}[|G_t| |G_{t+\tau}|] = \mu^2 + \sum_{k=2}^K c_k \lambda_k^\tau, \quad \tau \geq 1.$$

Step 6 (Autocovariance and autocorrelation). Subtracting μ^2 from Step 1 cancels the dominant term, leaving

$$\text{Cov}(|G_t|, |G_{t+\tau}|) = \sum_{k=2}^K c_k \lambda_k^\tau, \quad \tau \geq 1, \quad c_k = (\mathbf{m}^\top \text{diag}(\bar{\boldsymbol{\pi}}) \mathbf{v}_k) (\mathbf{w}_k^\top \mathbf{m}).$$

Dividing by $\sigma_{|G|}^2 = \gamma(0)$ gives $\rho_{|G|}(\tau) = \sum_{k \geq 2} w_k \lambda_k^\tau$ with $w_k = c_k / \sigma_{|G|}^2$, equation (6). $\square$

Lag zero, the sum of weights, and the within-state white-noise contribution. The bilinear identity (26) fails at $\tau = 0$, where $G_t = G_{t+\tau}$ and the conditional product is the per-state second moment $M_i = \mathbb{E}[G_t^2 | s_t = i]$ rather than m_i^2 . The eigenvalue completeness $\sum_{k=1}^K \mathbf{v}_k \mathbf{w}_k^\top = \mathbf{I}$ instead gives the identity

$$\sum_{k=2}^K c_k = \mathbf{m}^\top \mathbf{D} \mathbf{m} - \mu^2 = \sum_{i=1}^K \bar{\pi}_i m_i^2 - \mu^2 = \text{Var}(m_{s_t}),$$

the between-state variance of $|G_t|$, equal to the variance of the conditional-mean random variable m_{s_t} under $\bar{\boldsymbol{\pi}}$. By the law of total variance, $\sigma_{|G|}^2 = \text{Var}(m_{s_t}) + \mathbb{E}[\text{Var}(|G_t| | s_t)]$, so

$$\sum_{k=2}^K w_k = \frac{\text{Var}(m_{s_t})}{\sigma_{|G|}^2} \in [0, 1],$$

with equality to 1 if and only if every per-state distribution is degenerate at m_i . The residual $1 - \sum_{k \geq 2} w_k = \mathbb{E}[\text{Var}(|G_t| | s_t)] / \sigma_{|G|}^2$ is the within-state variance ratio: it contributes to $\sigma_{|G|}^2$ but not to $\text{Cov}(|G_t|, |G_{t+\tau}|)$ at any positive lag, and is the source of the discontinuity from $\rho_{|G|}(0) = 1$ to $\rho_{|G|}(1) = \sum_{k \geq 2} w_k \lambda_k$ at the first lag. The substantive consequence for moderate- K CHMM is unchanged: at every $\tau \geq 1$ the autocorrelation lies in the span of the geometric decays $\{\lambda_k^\tau\}_{k \geq 2}$, and the slow-decay envelope is governed by the largest non-unit eigenvalue.

Squared-return analogue. Replacing $|G_t|$ with G_t^2 and $\mathbf{m}$ with $\mathbf{M} = (M_1, \dots, M_K)^\top$ in Steps 1–6 above is termwise valid for $\tau \geq 1$ (the conditional independence and stationarity arguments are unaffected), giving the squared-return analogue

$$\rho_{G^2}(\tau) = \sum_{k=2}^K \tilde{w}_k \lambda_k^\tau, \quad \tilde{w}_k = \frac{(\mathbf{M}^\top \mathbf{D} \mathbf{v}_k) (\mathbf{w}_k^\top \mathbf{M})}{\sigma_{G^2}^2}.$$

The non-unit spectrum of $\mathbf{T}$ that determines $\rho_{|G|}(\tau)$ also determines $\rho_{G^2}(\tau)$; only the linear projection of the per-state moment vector onto the eigenmodes changes.

Remark (complex eigenvalues). For real $\mathbf{T}$ without self-adjoint structure, eigenvalues outside $\mathbb{R}$ occur in complex conjugate pairs with conjugate eigenvectors. A pair $(\lambda_k, \bar{\lambda}_k)$ at $\lambda_k = re^{i\theta}$ contributes a real damped-oscillation term of the form $A r^\tau \cos(\theta\tau) + B r^\tau \sin(\theta\tau)$ to $\gamma(\tau)$, with (A, B) determined from $(\text{Re } c_k, \text{Im } c_k)$. Theorem 1 as stated remains valid; only the interpretation of the individual modes shifts from monotone geometric to damped sinusoidal. For $\mathbf{T}$ estimated on equity returns at the operating point of Section 4, the slow-decay mode at $\lambda_2 \approx \exp(-1/\tau^*)$ is real and positive, and the mono-exponential interpretation of the dominant timescale used in Section 3.8 carries through.

Remark (non-diagonalisable $\mathbf{T}$). Diagonalisability is generic: for $\mathbf{T}$ estimated by Baum–Welch in the interior of a non-degenerate parameter region, the eigenvalues are almost surely distinct, and the sample paths used in this paper realise this generic case at every K in the sweep. Should $\mathbf{T}$ admit a non-trivial Jordan block of size J , the dyadic expansion of Step 4 acquires polynomial-in- τ prefactors $\binom{\tau}{j} \lambda_k^{\tau-j}$ for $j = 0, \dots, J-1$ on the off-diagonal blocks, and $\rho_{|G|}(\tau)$ becomes a finite sum of polynomial-times-geometric terms with the same exponential rate $\max_{k \geq 2} |\lambda_k|$. The substantive content of Theorem 1, that the absolute-return ACF is determined entirely by the non-unit spectrum of $\mathbf{T}$, is unchanged.

A.6 State-Resolution and Multi-Emission Sensitivity

This appendix reports the full K -sweep underlying the headline operating point of Section 4, the multi-emission family extension confirming that $K = 18$ is stable under Student-t and Laplace emissions, the EM convergence curves at representative K , the IS and OoS visual comparison panels at non-selected $K \in \{3, 6, 12, 21\}$, the multi-emission figure panels at $K = 18$, and the Rydén $K = 2$ replication.

A.6.1 Full State-Resolution Sensitivity Table

Table 14 reports the full K -sweep referenced in Section 4.2 of the main text. Observed excess kurtosis is 7.68 IS and 5.29 OoS at all K .

Table 14. State resolution sensitivity for SPY CHMM-N (1,000 paths, $\alpha = 0.05$). Every metric is computed on the identical simulation pipeline used for Table 3 in the main text. IS coverage was 100% at every K ; OoS coverage is reported separately. ACF-MAE columns: $|G_t|$ measures the volatility-clustering reproduction; G_t measures linear-autocorrelation reproduction (parallel to $|G_t|$, both averaged over lags 1–252).

K	KS (%)		AD (%)		Kurtosis		ACF-MAE		W_1	H	Cov OoS (%)
	IS	OoS	IS	OoS	Obs	Sim	$ G_t $	G_t			
3	89.0	77.2	78.9	68.6	7.68	3.82	0.0463	0.0240	0.131	0.0839	94.9
6	92.9	78.7	84.3	74.3	7.68	5.38	0.0496	0.0238	0.116	0.0829	90.9
9	93.0	81.4	84.4	74.4	7.68	4.63	0.0492	0.0238	0.118	0.0810	91.9
12	94.9	81.6	85.8	75.0	7.68	4.57	0.0500	0.0238	0.113	0.0801	91.9
15	93.9	81.0	84.8	73.5	7.68	4.89	0.0508	0.0238	0.114	0.0804	88.9
18	94.8	85.5	88.2	79.7	7.68	5.04	0.0515	0.0237	0.110	0.0807	91.9
21	95.6	88.0	89.1	78.8	7.68	3.78	0.0534	0.0236	0.112	0.0801	99.0

A.6.2 Multi-Emission State-Resolution Sensitivity

Table 15 extends Table 14 across the four emission families, confirming that the qualitative operating point $K = 18$ is stable when the Gaussian component is replaced with Student-t, Laplace, or GED emissions. For CHMM-t the per-state degrees of freedom ν_k are refit by ECM-plus-golden-section at every K ; for CHMM-L the weighted-median location and weighted MAD scale are closed form and impose no additional cost; for CHMM-GED the three-stage CM update of Section 3.7 refits (μ_k, α_k, p_k) at every K with p_k bracketed in $[0.5, 3.0]$. Three patterns extend from Table 14 to the new families. First, the distributional pass rates plateau in the $K \in \{12, 15, 18, 21\}$ band for all four families: CHMM-t reaches its highest IS KS at $K = 18$ (94.6%), CHMM-L reaches its highest IS KS at $K = 18$ (94.2%) and its highest IS AD at $K = 18$ (91.4%), and CHMM-GED reaches its highest IS KS at $K = 12$ (97.6%) and the highest entry across the entire panel. Once a family enters the plateau the operating point is flat with respect to K , so we report the main-text results at the same $K = 18$ used for CHMM-N. Second, the absolute-return ACF-MAE remains within the narrow 0.046–0.058 band across the four families, and the raw-return ACF-MAE remains essentially constant at ≈ 0.024 across all 28 (family, K) combinations. Both temporal axes are properties of the regime structure rather than the emission family: $|G_t|$ ACF reproduction is governed by the eigenvalue spectrum of $\mathbf{T}$ (Section 3.8), and raw- G_t ACF is null in expectation under any HMM with i.i.d. within-state emissions. Third, simulated kurtosis behaves as expected from the component families: the Gaussian sweep peaks around 5.0 at the high- K end, the Laplace sweep lands squarely in the 5.2–6.7 range (closest to the observed 7.68), the Student-t sweep climbs sharply above 13 across the entire sweep because the ECM search assigns small ν_k to a subset of tail states, and the GED sweep mirrors the Gaussian curve at 4.7–5.5 because the bulk of $\hat{p}_k$ values pin near the upper bracket and the Laplace-like tail states contribute too little stationary mass to lift the aggregate moment.

Table 15. Multi-emission state resolution sensitivity for SPY (1,000 paths, $\alpha = 0.05$). CHMM-N: Gaussian; CHMM-t: Student-t with per-state ν_k refit by ECM-plus-golden-section; CHMM-L: Laplace with weighted-median location and weighted-MAD scale; CHMM-GED: Generalized Error Distribution with per-state shape $p_k \in [0.5, 3.0]$ via three-stage ECM (Section 3.7). Every metric is computed on the identical simulation pipeline used for Table 3. Observed excess kurtosis is 7.68 IS and 5.29 OoS. ACF-MAE columns parallel the main-text Table 3: $|G_t|$ for volatility clustering, G_t (raw) for linear autocorrelation. **Sub-seed convention:** this multi-emission sweep uses a different sub-seed offset from the body-table headline run, so per-row metrics may drift by ≤ 1 percentage point from the corresponding entries in Table 3 (e.g. CHMM-GED at $K = 18$: this table 96.3% / 5.05 / 0.0546 versus body 95.2% / 5.15 / 0.0548).

Family	K	KS (%)		AD (%)		Kurt (IS sim)	ACF-MAE		W_1 IS	H IS
		IS	OoS	IS	OoS		$ G_t $	G_t		
CHMM-N	3	89.7	80.1	79.2	70.7	3.86	0.0467	0.0240	0.128	0.0833
	6	93.8	79.1	87.2	72.8	5.29	0.0504	0.0238	0.114	0.0825
	9	94.9	83.0	84.3	74.8	4.58	0.0490	0.0238	0.115	0.0806
	12	93.7	82.7	85.5	75.0	4.56	0.0501	0.0237	0.115	0.0804
	15	93.6	83.0	87.0	75.1	4.89	0.0509	0.0238	0.114	0.0803
	18	94.7	83.4	88.5	76.6	4.97	0.0511	0.0237	0.108	0.0805
	21	95.1	88.0	88.2	79.7	3.80	0.0532	0.0236	0.110	0.0796
CHMM-t	3	92.2	82.9	88.0	79.1	33.68	0.0538	0.0235	0.108	0.0759
	6	91.8	80.4	86.7	74.8	22.74	0.0527	0.0235	0.108	0.0772
	9	95.4	84.0	89.9	78.0	16.60	0.0535	0.0234	0.102	0.0752
	12	95.2	85.0	89.8	78.4	14.98	0.0543	0.0234	0.100	0.0757
	15	94.7	83.7	88.8	79.5	16.36	0.0539	0.0234	0.101	0.0758
	18	96.0	84.5	91.7	80.8	13.19	0.0551	0.0234	0.099	0.0759
	21	94.5	83.3	88.5	78.1	14.66	0.0539	0.0234	0.102	0.0758
CHMM-L	3	80.5	62.7	81.2	66.5	5.24	0.0531	0.0235	0.113	0.0855
	6	88.3	76.3	83.1	71.8	5.86	0.0510	0.0235	0.112	0.0821
	9	90.7	77.2	84.2	72.2	5.92	0.0519	0.0235	0.112	0.0823
	12	93.8	79.3	88.7	76.7	6.17	0.0541	0.0234	0.105	0.0793
	15	92.5	78.3	88.2	75.5	5.98	0.0539	0.0234	0.107	0.0794
	18	94.2	80.5	91.4	80.6	6.74	0.0566	0.0234	0.104	0.0795
	21	91.5	79.6	87.5	77.7	5.95	0.0539	0.0234	0.107	0.0807
CHMM-GED	3	90.3	78.4	85.5	73.6	5.45	0.0531	0.0236	0.109	0.0818
	6	91.8	79.0	84.9	73.5	5.22	0.0518	0.0235	0.108	0.0805
	9	94.6	81.6	89.1	75.9	5.03	0.0535	0.0234	0.101	0.0787
	12	97.6	84.2	94.0	81.2	4.67	0.0563	0.0234	0.094	0.0775
	15	94.2	82.0	87.4	76.5	5.14	0.0535	0.0235	0.104	0.0784
	18	96.3	83.8	92.7	81.3	5.05	0.0546	0.0235	0.098	0.0779
	21	94.1	85.3	88.7	81.2	4.82	0.0542	0.0235	0.102	0.0781

A.6.3 Cross-Ticker Per-Family Kurtosis Pattern

The body Table 6 reports the penalised CHMM-t cross-ticker performance ($1/\nu_k$ shrinkage at $\lambda = 20$). This appendix documents the unpenalised cross-ticker behaviour that motivates the penalty. The CHMM-t simulated kurtosis blow-ups under the unpenalised ECM (NVDA 40.8, JNJ 104.1, JPM 17.6 at $K = 18$) are specific to the Student-t emission family, and Table 16 extends the kurtosis column across all four CHMM emission variants on the same six-ticker universe to make this explicit. The pattern is consistent: unpenalised CHMM-t produces single-ticker overshoots of an order of magnitude on the heavy-tailed defensives (JNJ, NVDA), CHMM-L sits closest to observed across

every ticker (median absolute residual 1.4 kurtosis units), CHMM-N undershoots uniformly, and CHMM-GED interpolates between them. The CHMM-t blow-up mechanism is the same per-state ν_k ECM lower-bracket pinning documented for SPY in Section A.2.3: on tickers whose tail regime is concentrated in a small subset of states (NVDA’s pre-AI-rally drawdowns, JNJ’s outlier-driven defensive-stock dynamics), one or two states pin at $\nu_k \approx 2.1$ and dominate the simulated-kurtosis sample moment. Under the $1/\nu_k$ penalty at $\lambda = 20$ (the body operating point in Table 6; full rate sweep in Section 5, Appendix A.9) the simulated kurtosis on every ticker drops into the same band as CHMM-L: SPY $14.7 \rightarrow 8.0$, NVDA $40.8 \rightarrow 5.2$, JNJ $104.1 \rightarrow 12.3$, JPM $17.6 \rightarrow 7.7$, AAPL $6.3 \rightarrow 2.8$, QQQ $3.3 \rightarrow 2.9$, while IS / OoS KS pass rates and $|G_t|$ ACF-MAE move by less than 1.5 percentage points and 0.001 respectively.

Table 16. Per-ticker simulated excess kurtosis across the four CHMM emission families ($K = 18$, 1,000 paths, seed = 20260420). Columns report mean simulated IS excess kurtosis under each variant; the “Obs” column is the IS sample value. CHMM-t is the only family that produces order-of-magnitude blow-ups (NVDA, JNJ, JPM); the remaining three sit within 4 kurtosis units of observed across every ticker. Numbers are sourced from the canonical multi-emission cross-ticker run used for Table 15; minor drift from the body Table 6 CHMM-t column reflects a single re-run with a refreshed sub-seed.

Ticker	Obs	CHMM-N	CHMM-t	CHMM-L	CHMM-GED
SPY	7.7	5.0	15.0	6.8	5.1
NVDA	5.5	3.7	40.8	3.2	5.0
JNJ	9.0	5.5	104.1	6.1	9.3
JPM	7.6	7.2	17.6	6.0	14.8
AAPL	3.2	2.3	6.3	3.9	3.2
QQQ	3.4	2.0	3.3	4.1	2.7

A.6.4 Sector-Balanced 30-Ticker Panel: Full Per-Ticker Rollup

The body Table 6 reports the aggregate distribution and per-sector medians for the sector-balanced 30-ticker panel. Table 17 below gives the full per-ticker rollup. Each ticker is fit independently under the penalised CHMM-t scaffold ($K = 18$, $\lambda = 20$, 1,000 simulated paths, seed 20260420) on its own IS slice covering 2014-01-03 to 2024-01-03 and validated on the same 2024-01-04 to 2026-04-20 OoS window used for SPY in Table 3. The per-ticker ν_k medians sit at the upper bracket (50) on every ticker, indicating that the $1/\nu_k$ shrinkage is in the operative regime across the panel; minimum and maximum ν_k values are in the per-ticker CSV in `results/sector_panel/sector_panel_summary.csv`.

Table 17. Sector-balanced 30-ticker panel, per-ticker rollup (penalised CHMM-t at $K = 18$, $\lambda = 20$). Sectors ordered by GICS classification; tickers within sector ordered by the panel construction (top three large-cap representatives at IS-window-median market cap). Kurt resid is simulated minus observed IS excess kurtosis. Aggregate distribution and per-sector medians are summarised in body Table 6.

Sector	Ticker	IS KS%	OoS KS%	Kurt obs	Kurt sim	Kurt resid	$ G_t $ ACF-MAE
Information Technology	AAPL	99.1	95.1	3.19	2.84	-0.35	0.0418
Information Technology	MSFT	99.7	96.9	4.20	3.48	-0.72	0.0389
Information Technology	NVDA	99.6	55.7	5.53	6.23	0.70	0.0474
Health Care	JNJ	99.3	93.2	8.97	11.65	2.68	0.0335
Health Care	UNH	99.3	14.5	8.85	10.31	1.46	0.0382
Health Care	LLY	99.7	7.6	13.60	66.51	52.91	0.0295
Financials	JPM	99.5	50.5	7.62	8.14	0.52	0.0458
Financials	BAC	99.6	84.8	5.76	5.47	-0.29	0.0394
Financials	WFC	98.2	57.6	7.53	10.90	3.37	0.0791
Consumer Discretionary	AMZN	99.4	96.9	4.31	3.68	-0.63	0.0413
Consumer Discretionary	HD	99.2	41.4	12.92	17.03	4.11	0.0368
Consumer Discretionary	MCD	99.5	66.6	18.38	26.93	8.55	0.0394
Communication Services	NFLX	99.3	45.1	12.60	4.67	-7.93	0.0320
Communication Services	VZ	99.3	49.1	5.14	5.15	0.01	0.0270
Communication Services	DIS	99.5	96.4	10.77	13.94	3.17	0.0644
Industrials	CAT	99.7	78.9	4.34	3.48	-0.86	0.0321
Industrials	BA	97.6	62.2	20.59	52.00	31.41	0.0793
Industrials	HON	99.3	96.0	21.93	17.09	-4.84	0.0571
Consumer Staples	PG	99.7	94.3	8.15	9.51	1.36	0.0347
Consumer Staples	KO	99.5	92.7	10.57	12.84	2.27	0.0457
Consumer Staples	WMT	99.7	24.0	13.84	32.63	18.79	0.0279
Energy	XOM	96.9	70.3	5.61	10.99	5.38	0.0988
Energy	CVX	99.6	62.3	12.22	14.78	2.56	0.0498
Energy	COP	99.4	83.6	17.11	8.28	-8.83	0.0494
Utilities	NEE	98.6	24.0	10.71	9.87	-0.84	0.0455
Utilities	DUK	99.7	96.9	9.75	8.11	-1.64	0.0466
Utilities	SO	99.6	96.2	14.52	12.73	-1.79	0.0481
Materials	FCX	99.8	76.6	4.83	4.11	-0.72	0.0573
Materials	NEM	99.6	5.6	3.74	3.88	0.14	0.0376
Materials	APD	99.2	90.2	8.40	16.74	8.34	0.0381

A.6.5 60-Ticker Sector Expansion at $n = 6$ per Sector

The body’s $n = 3$ per-sector ANOVA on the 30-ticker panel is severely underpowered ($\eta^2 = 0.16$, $F(9, 20) = 0.44$, $p = 0.90$). Round-2 reviewers asked for an $n = 6$ per-sector expansion (60 tickers total) to test sector effects at adequate power. Source / runner: `run_sector_panel_n6.jl` (full panel) and `run_sector_panel_n6_postprocess.jl` (combine with cached body 30); output: `results/sector_panel_n6/sector_panel_n6.csv`.

The additional 30 tickers (3 large-cap representatives per sector, complementing the body 30) are: IT (ADBE, CRM, ORCL); Health Care (ABBV, MRK, PFE); Financials (GS, MS, BLK); Consumer Discretionary (NKE, SBUX, TSLA); Communication Services (GOOG, T, CMCSA); Industrials (UNP, MMM, LMT); Consumer Staples (PEP, MO, COST); Energy (SLB, EOG, MPC); Utilities (D, AEP, EXC); Materials (SHW, ECL, NUE). Each ticker is fit independently under the body penalised CHMM-t scaffold ($K = 18$, $\lambda = 20$, 1,000 paths, seed 20260420).

Reading. The 60-ticker aggregate distribution is statistically indistinguishable from the body 30-ticker version: median OoS KS 73.45% vs. 73.4% (body Table 6), mean 66.99% vs. 66.8%, failure

Table 18. 60-ticker sector-balanced panel ($n = 6$ per sector). Per-sector OoS KS distribution and one-way ANOVA at adequate power. “failures” is the count of tickers with OoS KS < 60%. Body $n = 3$ ANOVA: $F(9, 20) = 0.44$, $p = 0.90$, $\eta^2 = 0.16$. Adequately-powered $n = 6$ ANOVA: $F(9, 50) = 0.37$, $p = 0.946$, $\eta^2 = \mathbf{0.062}$ — the effect-size estimate halves with adequate power and the no-sector-effect null is not rejected.

Sector	n	median	mean	sd	min	failures
Communication Services	6	47.10	54.30	–	21.40	4/6
Consumer Discretionary	6	57.55	63.13	–	39.70	3/6
Consumer Staples	6	92.00	68.12	–	13.70	2/6
Energy	6	76.95	68.28	–	11.30	1/6
Financials	6	68.35	69.73	–	50.50	2/6
Health Care	6	61.00	56.08	–	7.60	3/6
Industrials	6	82.75	76.43	–	44.80	1/6
Information Technology	6	74.60	66.15	–	0.00	3/6
Materials	6	82.25	70.25	–	5.60	1/6
Utilities	6	96.00	77.47	–	24.00	2/6
Aggregate ($n = 60$)	60	73.45	66.99	29.13	0.00	22/60

rate $22/60 = 36.7\%$ vs. $11/30 = 36.7\%$ (identical). **The ANOVA at adequate power confirms the no-sector-effect null cannot be rejected at $\alpha = 0.05$:** $F(9, 50) = 0.37$, $p = 0.946$, $\eta^2 = 0.062$, with the effect-size estimate dropping from the underpowered $\eta^2 = 0.16$ ($n = 3$) to $\eta^2 = 0.062$ ($n = 6$). Cross-sector dispersion accounts for 6.2% of OoS KS variance at the larger sample size; the remaining 93.8% is per-ticker variation. Per-sector medians span 47% (Communication Services) to 96% (Utilities), but the within-sector spread (median range 14% at Utilities to 79% at IT) dominates the between-sector spread (median range 49pp). The body’s qualitative claim that “failures are ticker-specific, not sector-driven” is now supported by an adequately-powered test on a panel that doubles the body’s sample size; the 11/30 failure rate at $n = 3$ was not a small-sample artefact but a stable feature of the cross-ticker generalisation under static IS-fit.

A.6.6 Five-Named-Ticker Spot-Check (Original Six-Ticker Panel)

The original peer-review-cycle six-ticker spot-check (SPY, NVDA, JNJ, JPM, AAPL, QQQ; broad-equity ETF and five named single-name representatives) is retained in Table 19 as a hand-curated reference for downstream consumers who need to compare against the cross-asset Pipeline-B copula construction in Section 4.5 (whose six-asset universe is the same five named tickers plus QQQ). The aggregate-distribution, sector-rollup, and universe-scale failure-rate framing is moved to the body Table 6 on the 30-ticker panel.

Table 19. Original five-named-ticker spot-check (penalised CHMM-t at $K = 18$, $\lambda = 20$). Hand-curated panel covering broad-equity (SPY, QQQ), AI / semiconductors (NVDA), healthcare (JNJ), financials (JPM), and consumer-tech (AAPL); reproduced from `results/chmm_t_penalised/Cross_Ticker_CHMM_t_Pen.txt`. The expanded sector-balanced 30-ticker panel is in Table 17.

Ticker	KS (%) $\uparrow$		Kurtosis		ACF-MAE $\downarrow$
	IS	OoS	Obs	Sim	
SPY	95.8	85.0	7.7	8.01	0.0541
NVDA	99.5	53.8	5.5	5.23	0.0473
JNJ	99.8	92.3	9.0	12.28	0.0334
JPM	97.8	53.8	7.6	7.70	0.0457
AAPL	99.1	95.1	3.2	2.84	0.0418
QQQ	96.1	91.8	3.4	2.88	0.0569

A.6.7 Effective Spectral Rank of the Absolute-Return ACF

The body Section 3.8 states the algebraic upper bound: the deflated transition matrix $\mathbf{T} - \mathbf{1}\bar{\pi}^\top$ has rank at most $K - 1$, so the absolute-return ACF identity (6) contains at most $K - 1$ non-unit decay modes. The empirical question is how many of those modes carry non-trivial contribution to $\rho_{|G|}(\tau)$ at the relevant lags. Table 20 reports the per-mode contribution $w_k \lambda_k^\tau$ at $\tau \in \{1, 5, 20, 50\}$ for the fitted CHMM-N at $K = 18$ and at $K = 3$ on SPY IS, with modes ranked by lag-1 contribution magnitude $|w_k \lambda_k|$. Per-state moments $m_k = \mathbb{E}[|G_t| \mid s_t = k]$ are estimated by 200,000 draws from each emission; eigenvectors are normalised so $\mathbf{w}_k^\top \mathbf{v}_k = 1$.

Table 20. Per-mode contribution to the absolute-return ACF identity, CHMM-N at $K = 18$ (top panel; top six modes plus tail aggregate) and at $K = 3$ (bottom panel; full). Modes ordered by $|w_k \lambda_k|$ descending. “cum” is cumulative share of $\sum_k |w_k \lambda_k|$. At $K = 18$ the top mode ($|\lambda| = 0.929$) carries 93.6% of the lag-1 ACF, three modes reach 95%, and at lag 20 only the dominant mode contributes non-negligibly. At $K = 3$ the dominant mode ($|\lambda| = 0.953$) alone carries 96.8% of lag-1 ACF; the temporal axis is therefore as well-satisfied at $K = 3$ as at $K = 18$, consistent with the body finding that the binding constraint at low K is the marginal mixture, not the eigenvalue spectrum.

rank	$ \lambda_k $	$ w_k $	$ w_k \lambda_k $	$w_k \lambda_k^5$	$w_k \lambda_k^{20}$	cum
$K = 18, \rho_{ G }(1) = 0.301, \sigma_{ G }^2 = 2.555$						
1	0.929	0.324	0.301	0.224	0.0744	0.936
2	0.335	0.012	0.004	0.0001	0.0000	0.948
3	0.854	0.003	0.002	-0.001	-0.0001	0.955
4	0.231	0.008	0.002	≈ 0	≈ 0	0.961
5	0.231	0.008	0.002	≈ 0	≈ 0	0.967
6	0.203	0.008	0.002	≈ 0	≈ 0	0.971
7-17	various	various	≤ 0.001 each	≈ 0	≈ 0	1.000
$K = 3, \rho_{ G }(1) = 0.287, \sigma_{ G }^2 = 2.517$						
1	0.953	0.292	0.278	0.230	0.112	0.968
2	0.866	0.011	0.009	0.005	0.0006	1.000

The empirical effective rank confirms the body framing of Section 3.8: the ACF-MAE flatness across $K \in \{3, 6, 9, 12, 15, 18, 21\}$ documented in Appendix A.6.1, Table 14, is the macroscopic signature of the temporal axis being driven by a single dominant decay mode at every operating point in the sweep. The K -rank statement on $\mathbf{T} - \mathbf{1}\bar{\pi}^\top$ is correctly interpreted as a non-binding

upper bound at $K \geq 3$; at $K = 2$ the bound is tight (only one non-unit eigenvalue available), but the failure mode is then the marginal mixture, which two Gaussian components cannot tile against the empirical equity-return target. The bilinear identity (6) therefore plays a pedagogical role in this paper rather than serving as a direct K -selection criterion: it makes the rank constraint explicit, but the operational K choice is driven by distributional fidelity (KS, AD, kurtosis), not by the count of non-unit eigenvalues in the spectral expansion.

A.6.8 Cross-Ticker Spectral Effective-Rank Diagnostic

Reviewers 2 and 3 ask whether the SPY-only result of Appendix A.6.7 (a single non-unit eigenvalue carries 93.6% of the lag-1 absolute-return ACF at $K = 18$) generalises across tickers. Table 21 reports the cross-ticker distribution of the dominant-mode share, computed under the same protocol on the 30-ticker sector-balanced panel plus SPY. Source: `results/diagnostics/spectral_rank_cross_ticker.txt`.

Table 21. Cross-ticker spectral effective-rank diagnostic at $K = 18$. Distribution across 31 tickers (sector-balanced 30-ticker panel plus SPY control). “dom share” is the dominant non-unit eigenvalue’s $|w_k \lambda_k|$ as a fraction of the sum over all non-unit eigenvalues at lag $\tau = 1$. “ n_{95} ” / “ n_{99} ” are the number of modes needed for 95% / 99% cumulative share. The SPY value of 0.94 is in the right tail of the distribution; the cross-ticker median is 0.76.

Statistic	dom share	n_{95}	n_{99}
median	0.756	6	11
$[Q_1, Q_3]$	$[0.661, 0.858]$	$[4, 7]$	$[10, 12]$
minimum (NEM)	0.326	11	14
SPY (body Appendix A.6.7)	0.943	2	10

Reading: the SPY headline of 93.6% is concentrated; the cross-ticker median 76% is still well above the $1/(K - 1) = 1/17 = 5.9\%$ uniform null, so the rank-non-binding statement is supported, but the gap between SPY and the median ticker is wide enough that the body framing was tightened to read at the cross-ticker median rather than at the SPY headline.

A.6.9 Cross-Asset Pipeline B at $K^* = 6$ Marginals

Reviewer R2 Q3 asks whether the body cross-asset construction (Section 4.5, Table 7) recovers under $K^* = 6$ CHMM-N marginals rather than the body $K = 18$ marginals. We re-ran Pipeline B (Gaussian, Student- t , truncated C-vine copulas plus the SIM baseline) on the same six-asset universe (SPY, NVDA, JNJ, JPM, AAPL, QQQ) with $K = 6$ CHMM-N marginals; same $N_{\text{paths}} = 200$, seed 20260422. Source / runner: `run_cross_asset_sim_copula_k6.jl`; output: `results/cross_asset_k6/cross_asset/Table-T3-Cross-Asset-Dependence.txt`.

Reading. Off-diagonal MAE differences across the marginal-resolution axis are ≤ 0.001 on IS and ≤ 0.001 on OoS at every dependence model; the body’s $\nu^* = 6$ Student- t copula choice and the IS / OoS off-diagonal ordering are therefore not driven by the marginal state-resolution. R2 Q3’s hypothesis that $K^* = 6$ marginals might shift the off-diagonal MAE comparison is rejected: the dependence layer is essentially marginal-resolution-independent on this six-asset US-equity universe, consistent with the cross-ticker K -robustness reported in Section 4.4.

Table 22. Cross-asset Pipeline B at $K^* = 6$ marginals vs. body $K = 18$. Off-diagonal correlation MAE on six-asset US-equity universe, 200 paths. The marginal state-resolution does not move the off-diagonal MAE by a margin large enough to flip the ranking; the dependence layer is essentially marginal-resolution-independent on this universe.

Dependence model	Off-diag MAE IS		Off-diag MAE OoS	
	$K = 18$ (body)	$K^* = 6$	$K = 18$ (body)	$K^* = 6$
SIM	0.076	0.077	0.251	0.251
Gaussian copula	0.030	0.030	0.202	0.203
Student- t copula	0.027	0.027	0.209	0.209
Truncated C-vine	0.068	0.068	0.236	0.236

A.6.10 Multi-Day Cumulative-Return Diebold-Mariano on CHMM vs. Stationary Block Bootstrap

Reviewer R3 RE4 asks whether CHMM dominates the stationary block bootstrap at any forecast horizon, given that the bootstrap dominates raw 1-day OoS KS. We aggregate the 1,000-path simulated archives to non-overlapping h -day cumulative-return blocks for $h \in \{5, 20\}$, compute per-block CRPS under the sorted-ensemble identity, and run a two-sided Newey–West HAC Diebold–Mariano test on the per-block loss differential. Source / runner: `run_crps_dm_multiday.jl`; output: `results/crps_dm_multiday/CRPS_DM_multiday.txt`.

Table 23. Diebold-Mariano test of CHMM vs. stationary block bootstrap on h -day cumulative-return CRPS. SPY OoS, $T_{\text{OoS}} = 572$ days, non-overlapping h -day blocks, two-sided NW-HAC DM with default bandwidth $\lfloor n^{1/3} \rfloor$. Negative ΔCRPS means CHMM has lower CRPS than bootstrap (CHMM is better).

h	Model	blocks	ΔCRPS	DM	p
1	CHMM-N	572	−0.001	−0.71	0.480
1	CHMM-t	572	+0.000	+0.04	0.971
1	CHMM-L	572	+0.000	+0.08	0.938
5	CHMM-N	114	−0.020	−1.49	0.137
5	CHMM-t	114	−0.010	−0.77	0.439
5	CHMM-L	114	−0.029	−2.03	0.042
20	CHMM-N	28	−0.180	−2.99	0.003
20	CHMM-t	28	−0.090	−1.20	0.229
20	CHMM-L	28	−0.186	−2.22	0.027

Reading. At $h = 1$ neither side dominates ($p > 0.45$ for all three CHMM rows; CHMM-N is nominally lower CRPS but indistinguishable from the bootstrap at the per-day horizon). At $h = 5$ CHMM-L is significantly better than the bootstrap ($p = 0.042$), CHMM-N is insignificant in the same direction ($p = 0.137$). At $h = 20$ both CHMM-N and CHMM-L dominate the bootstrap on per-block CRPS at conventional levels ($p = 0.003$ and $p = 0.027$ respectively). **The $h = 20$ result is, however, $K = 18$ -specific on SPY:** under the body-headline $K^* = 3$ on SPY and across the six-asset universe of §4.5 at $K^* = 3$, the same DM construction does not reach significance (Appendix A.6.11, Table 24). The single-window $K = 18$ SPY result therefore does not generalise to the body-headline operating point or to other US-equity tickers, and the body’s CHMM-vs-bootstrap

differentiation is on the structural use cases of §4.6–4.5 rather than on multi-day cumulative-return forecasting.

A.6.11 Multi-Day DM Replication: Six-Asset Universe, Bandwidth Sweep, Block-Overlap Sensitivity

The multi-day DM result of Appendix A.6.10 is at $K = 18$ on a single SPY OoS window. Round-2 reviewers asked whether the result replicates across the six-asset universe of §4.5 at the body-headline $K^* = 3$, whether the SPY $h = 20$ result is robust to NW-HAC bandwidth choice, and whether overlapping vs. non-overlapping 20-day blocks reach the same conclusion. Source / runner: `run_crps_dm_multiday_replication.jl`.

Table 24 reports three panels. Panel 1 is the cross-asset replication of the $h = 20$ DM at $K^* = 3$: SPY’s $\Delta\text{CRPS} = -0.077$ ($p = 0.244$) is in the same direction as the body $K = 18$ result but does not reach significance, and across the six-asset universe the median ΔCRPS is essentially zero with no per-asset DM significant at $\alpha = 0.05$. Panel 2 sweeps the NW-HAC bandwidth on SPY $h = 20$ at $K^* = 3$: p ranges from 0.150 to 0.255 across $h_{\text{NW}} \in \{2, 4, 8, 16, \lfloor n^{1/3} \rfloor\}$, never crossing the 0.05 threshold. Panel 3 compares non-overlapping ($n = 28$) vs. overlapping ($n_{\text{eff}} = 553$) blocks at $h = 20$: the default-NW-bandwidth overlapping result reaches $p = 0.063$ but with the proper NW bandwidth $h_{\text{NW}} \geq h - 1 = 19$ (overlap induces $h - 1$ -step autocorrelation in the loss differential) the result moves back to $p = 0.14$ –0.22.

Table 24. Multi-day DM replication panel ($K^* = 3$, 1,000 paths per generator per asset, mean block length $L = 20$ for stationary block bootstrap, seed 20260422). Panel 1: cross-asset $h = 20$ DM under non-overlapping blocks and default NW bandwidth. Panel 2: bandwidth sweep on SPY $h = 20$. Panel 3: overlapping vs. non-overlapping blocks on SPY $h = 20$. Source: `results/crps_dm_multiday_replication/`.

<i>Panel 1: Cross-asset $h = 20$ DM at $K^* = 3$ (non-overlapping, default NW bw)</i>				
Ticker	blocks	ΔCRPS	DM	p
SPY	28	−0.077	−1.165	0.244
NVDA	28	+0.384	+1.856	0.063
JNJ	28	−0.057	−1.021	0.307
JPM	28	+0.023	+0.245	0.806
AAPL	28	+0.116	+0.714	0.475
QQQ	28	−0.023	−0.309	0.757
<i>Panel 2: NW-HAC bandwidth sweep on SPY $h = 20$ at $K^* = 3$</i>				
h_{NW}	blocks	ΔCRPS	DM	p
2	28	−0.077	−1.138	0.255
4	28	−0.077	−1.168	0.243
$\lfloor n^{1/3} \rfloor$	28	−0.077	−1.165	0.244
8	28	−0.077	−1.321	0.187
16	28	−0.077	−1.441	0.150
<i>Panel 3: overlap-vs-non-overlap on SPY $h = 20$ at $K^* = 3$</i>				
Mode	blocks	ΔCRPS	DM	p
non-overlapping (default NW)	28	−0.077	−1.165	0.244
overlap (default NW)	553	−0.059	−1.859	0.063
overlap ($h_{\text{NW}} = 19$)	553	−0.059	−1.460	0.144
overlap ($h_{\text{NW}} = 30$)	553	−0.059	−1.330	0.184
overlap ($h_{\text{NW}} = 60$)	553	−0.059	−1.234	0.217

Reading. The $K = 18$ SPY $h = 20$ headline DM result of Appendix A.6.10 ($\Delta\text{CRPS} = -0.180$, $p = 0.003$) is not robust to operating-point change ($K = 18 \rightarrow K^* = 3$) or to cross-asset replication: at $K^* = 3$ on SPY the same construction gives $p = 0.244$, and across the six-asset universe at $K^* = 3$ no per-asset DM reaches significance with median ΔCRPS near zero. Bandwidth sensitivity on SPY $K^* = 3$ ($p \in [0.150, 0.255]$ across $h_{\text{NW}} \in \{2, 4, 8, 16, \text{default}\}$) and overlap sensitivity ($p \in [0.063, 0.244]$, with the marginally-significant default-bandwidth overlap row not surviving the proper $h_{\text{NW}} \geq h - 1$ adjustment) confirm that the SPY $K = 18$ result is operating-point-specific. The substantive read is that the body’s CHMM-vs-bootstrap differentiation is on the structural use cases of §4.6–4.5, not on multi-day cumulative-return forecast horizons; the multi-day DM panel of Appendix A.6.10 is a $K = 18$, single-asset, single-window finding rather than a robust property of the CHMM scaffold.

A.6.12 Per-State Frobenius Distances Between the Four Emission Families

Reviewer 2 W4 / RE3 asks whether the four CHMM emission variants land on materially different per-state $(\mu, \sigma, \mathbf{T})$ tuples or whether the four-family narrative collapses to a one-parameter shape axis. We fit each family independently per ticker at the body headline $K^* = 3$, canonicalise states by ascending σ_k , and compute pairwise Frobenius distances on the location, scale, and transition-matrix parameter blocks. Source: `results/emission_family_frobenius/spy.txt`, `panel.txt`, and `summary.csv`; runner: `run_emission_family_frobenius.jl`.

Table 25. Per-state Frobenius distances across emission families on SPY headline (top) and on the 30-ticker panel (bottom medians). $K = 3$, $\lambda = 20$ for the penalised CHMM-t. Quantile-based EM initialisation is deterministic per family on a fixed ticker, so the per-family Monte Carlo SE under multi-seed refit is exactly 0 on this protocol; the reported distances are the genuine cross-family differences. “N” = CHMM-N, “t” = penalised CHMM-t at $\lambda = 20$, “L” = CHMM-L, “GED” = CHMM-GED. The SPY structure is two-cluster: CHMM-N stands clearly apart from the heavy-tail trio $\{t, L, \text{GED}\}$ on all three axes, while $\{t, L, \text{GED}\}$ are within ~ 0.2 on μ and within ~ 0.04 on $\mathbf{T}$ of each other.

Pair	$\ \Delta\mu\ _F$	$\ \Delta\sigma\ _F$	$\ \Delta\mathbf{T}\ _F$
<i>SPY ($K^* = 3$):</i>			
N vs t	2.090	0.960	0.530
N vs L	1.943	1.894	0.508
N vs GED	2.025	0.945	0.520
t vs L	0.161	0.987	0.031
t vs GED	0.214	0.026	0.042
L vs GED	0.165	1.009	0.023
<i>30-ticker panel medians ($K^* = 3$, $[Q_1, Q_3]$ in parentheses):</i>			
N vs t	2.22 [1.95, 3.49]	3.89 [2.37, 5.32]	0.46 [0.40, 0.53]
N vs L	2.21 [1.57, 2.93]	4.93 [3.71, 7.53]	0.44 [0.36, 0.64]
N vs GED	2.35 [1.70, 3.32]	3.63 [2.53, 5.19]	0.45 [0.34, 0.55]
t vs L	1.35 [0.90, 2.48]	2.12 [1.62, 2.61]	0.26 [0.16, 0.66]
t vs GED	0.70 [0.40, 1.22]	0.39 [0.22, 0.71]	0.23 [0.13, 0.30]
L vs GED	1.93 [0.85, 3.46]	2.06 [0.66, 3.53]	0.32 [0.20, 0.61]

Reading. The four-family narrative does *not* collapse to a one-parameter shape axis: CHMM-N is distinct from the heavy-tail trio $\{\text{CHMM-t}, \text{CHMM-L}, \text{CHMM-GED}\}$ at $\|\Delta\mu\|_F \sim 2$, $\|\Delta\mathbf{T}\|_F \sim 0.45$ on both SPY and the 30-ticker panel. The four-family scaffold therefore separates Gaussian-emission from heavy-tail-emission CHMMs cleanly, consistent with the ~ 1.5 -unit kurtosis gap between CHMM-N (sim kurt 3.83) and the heavy-tail rows (CHMM-L 5.30, CHMM-GED 5.48, CHMM-t

pen. 14.91) at $K^* = 3$ in Table 3. The substantive sub-finding is that within the heavy-tail trio the families cluster more tightly, most notably t vs. GED on SPY ($\|\Delta\boldsymbol{\mu}\|_F = 0.21$, $\|\Delta\mathbf{T}\|_F = 0.04$); on the panel medians the t-GED pair remains the tightest. The per-state structure of the body four-emission scaffold is therefore best read as a two-cluster decomposition (Gaussian vs. heavy-tail) with within-cluster variation along the kurtosis axis, rather than as four independent families with comparable separations.

A.6.13 Baum-Welch Convergence

Figure 5 shows the Baum-Welch convergence curves for representative state resolutions. All models converged within the 60-iteration budget, with the log-likelihood monotonically increasing (as guaranteed by the EM algorithm). Lower K values typically converged faster (15–25 iterations), while higher K values required 25–40 iterations to reach the convergence tolerance of $|\Delta\mathcal{L}| < 10^{-4}$. The monotonic increase confirms that the quantile-based initialization provided a valid starting point from which the EM algorithm could improve without encountering degenerate solutions.






Figure 5. Baum-Welch convergence for the SPY CHMM-N fit (Pipeline A, IS) at selected state resolutions. Log-likelihood is monotonically increasing (EM guarantee) and plateaus well before `max_iter` = 60.

A.6.14 IS Comparison at Non-Selected State Resolutions


Figures 6–9 show the IS model comparison panels for $K \in \{3, 6, 12, 21\}$, complementing the $K = 18$ comparison in Figure 26 of the main text. Across all resolutions, the CHMM’s simulated density closely overlays the observed density, and the ACF of absolute returns tracks the slow decay pattern. The visible improvement from $K = 3$ to $K = 18$ is concentrated in the tails of the Q-Q panel and the steady-state level of the ACF panel. $K = 9$ and $K = 15$ panels are omitted because they interpolate smoothly between the shown resolutions; their numerical metrics are in Table 14.




(a) Marginal density.



(b) Tail Q-Q plot.




(c) ACF of raw G_t .




(d) ACF of $|G_t|$.


Figure 6. IS SPY CHMM-N comparison at $K = 3$ (Pipeline A, 1,000 simulated paths; panels as in Figure 26). The $K = 18$ counterpart is Figure 26 of the main text.




(a) Marginal density.



(b) Tail Q-Q plot.




(c) ACF of raw G_t .




(d) ACF of $|G_t|$.


Figure 7. IS SPY CHMM-N comparison at $K = 6$ (Pipeline A, 1,000 simulated paths; panels as in Figure 26).




(a) Marginal density.



(b) Tail Q-Q plot.




(c) ACF of raw G_t .




(d) ACF of $|G_t|$.


Figure 8. IS SPY CHMM-N comparison at $K = 12$ (Pipeline A, 1,000 simulated paths; panels as in Figure 26).




(a) Marginal density.



(b) Tail Q-Q plot.



(c) ACF of raw G_t .




(d) ACF of $|G_t|$.


Figure 9. IS SPY CHMM-N comparison at $K = 21$ (Pipeline A, 1,000 simulated paths; panels as in Figure 26).

A.6.15 OoS Validation at Non-Selected State Resolutions


Figures 10–13 show the OoS validation panels for $K \in \{3, 6, 12, 21\}$. KS pass rates remain above 91% at all resolutions; the $K = 18$ panel is reproduced in Figure 27 of the main text. As with the IS panels, $K = 9$ and $K = 15$ are omitted; numerical OoS KS pass rates are in Table 14.




(a) OoS density fan.



(b) KS p -value histogram.




(c) OoS ACF of raw G_t .




(d) OoS ACF of $|G_t|$.


Figure 10. OoS SPY CHMM-N validation at $K = 3$ (Pipeline A, 1,000 simulated paths; panels as in Figure 27). The $K = 18$ counterpart is Figure 27 of the main text.




(a) OoS density fan.



(b) KS p -value histogram.




(c) OoS ACF of raw G_t .




(d) OoS ACF of $|G_t|$.


Figure 11. OoS SPY CHMM-N validation at $K = 6$ (Pipeline A, 1,000 simulated paths; panels as in Figure 27).




(a) OoS density fan.



(b) KS p -value histogram.




(c) OoS ACF of raw G_t .




(d) OoS ACF of $|G_t|$.


Figure 12. OoS SPY CHMM-N validation at $K = 12$ (Pipeline A, 1,000 simulated paths; panels as in Figure 27).




(a) OoS density fan.



(b) KS p -value histogram.



(c) OoS ACF of raw G_t .



(d) OoS ACF of $|G_t|$.


Figure 13. OoS SPY CHMM-N validation at $K = 21$ (Pipeline A, 1,000 simulated paths; panels as in Figure 27).

A.6.16 Multi-Emission Figures at $K = 18$


This appendix reproduces the standard diagnostic suite (emission PDFs, residence times, IS comparison, OoS validation, and convergence) for the two heavy-tailed emission families (CHMM-t and CHMM-L) at the selected state resolution $K = 18$; transition matrices and stationary distributions are discussed in prose only below, since they are visually indistinguishable across families at $K = 18$. The Gaussian counterparts are already shown as Figures 26 and 27 in the main text and as Figure 5 above in this appendix. The figures confirm that the shared forward–backward scaffold produces internally-similar regime structure across the four families: the near-uniform interior of the transition matrix, the K -component dense emission cover, the longer central-state residence times, and the monotone log-likelihood convergence are features of the CHMM framework and not of the Gaussian component specifically.

Emission PDFs. Figure 14 compares the learned $K = 18$ emissions for the four families. The CHMM-t panel’s per-state ν_k allocation is visible in the narrower modes of certain tail states; the CHMM-L panel’s Laplace components have sharper peaks and slower exponential tail decay than the Gaussians; the CHMM-GED panel mixes both, with bulk states near the Gaussian shape ($\hat{p}_k$


near the upper bracket) and a small set of tail states near the Laplace shape ($\hat{p}_k$ near 1).




(a) CHMM-N (Gaussian)



(b) CHMM-t (Student-t)




(c) CHMM-L (Laplace)




(d) CHMM-GED (per-state p_k)

Figure 14. Learned emission distributions at $K = 18$ across the four emission families (CHMM-N Gaussian, CHMM-t Student-t with per-state ν_k , CHMM-L Laplace, CHMM-GED with per-state p_k), overlaid on the observed SPY IS return histogram (Pipeline A). Heavier-tailed emission families (t, L) give individual components more tail mass, which is what closes the kurtosis gap of the Gaussian CHMM; CHMM-GED interpolates between Gaussian and Laplace per state.


Transition matrices and stationary distributions. Figure 15 shows the $\log_{10}$ transition-matrix heatmaps at $K = 18$ across the four emission families. The four panels are visually indistinguishable: the same near-uniform interior, the same characteristic boundary states (rows 1 and K) with a small handful of low-mass entries, and self-transition probabilities of comparable magnitude. The eigenvalue spectrum of $\mathbf{T}$ that drives the absolute-return ACF identity (6) is therefore set by the K -rank regime structure rather than by the emission family. This is consistent with the ACF-MAE being stable across the family axis (Table 15). The per-family residence times $\tau_k = 1/(1 - T_{kk})$, which magnify small differences in the self-transition probabilities each family recovers, are compared in Figure 16.




(a) CHMM-N (Gaussian).



(b) CHMM-t (Student-t).



(c) CHMM-L (Laplace).



(d) CHMM-GED (per-state p_k).

Figure 15. $\log_{10}$ transition-matrix heatmaps at $K = 18$ across the four emission families on SPY IS (Pipeline A). Color scale is $\log_{10} T_{ij}$ with a 10^{-10} floor for the near-zero entries. The cross-family similarity supports the rank/spectral interpretation of Section 3.8: the regime structure (and thus the eigenvalue spectrum that drives the absolute-return ACF) is largely a property of K and the data, not of the emission family.






Figure 16. Natural residence times $\tau_k = 1/(1 - T_{kk})$ at $K = 18$ for the SPY IS fit (Pipeline A) across the four emission families (CHMM-N Gaussian, CHMM-t Student-t, CHMM-L Laplace, CHMM-GED).

Sample simulated trajectories. Figure 17 overlays one simulated SPY return path against the observed series under CHMM-N at $K = 18$, separately for the IS and OoS windows. Both panels show the simulated path tracking the volatility-clustering envelope of the observed series at the same overall amplitude. CHMM-N is shown as the headline reference; per-family marginal differences are governed by the emission family rather than by the trajectory representation, and are visualised more directly in Figure 14.




(a) In-sample window, full $T_{IS} = 2,516$ trading days.




(b) Out-of-sample window, full $T_{OoS} = 572$ trading days.

Figure 17. Sample simulated return trajectory at $K = 18$ for CHMM-N on SPY (Pipeline A). Orange is the observed series; blue is one simulated path drawn from the fit. Panel (a): full in-sample window. Panel (b): full out-of-sample window. Both panels share the same y-axis convention (annualized excess log return G_t).


IS comparison. Figures 18–20 reproduce the four-panel IS comparison (marginal density, tail Q-Q, raw-return ACF, $|G_t|$ ACF) at $K = 18$ for the three non-Gaussian emission families; the CHMM-N panel is Figure 26 in the main text.




(a) Marginal density.



(b) Tail Q-Q plot.




(c) ACF of raw G_t .




(d) ACF of $|G_t|$.


Figure 18. IS SPY model comparison at $K = 18$ for CHMM-t Student-t emissions (Pipeline A, 1,000 simulated paths). Panels show (a) marginal density with KS pass rate, (b) tail Q-Q plot, (c) ACF of raw G_t (linear autocorrelation, with the 99% Bartlett band), and (d) ACF of $|G_t|$ (volatility clustering). The CHMM-N counterpart is Figure 26; the CHMM-L and CHMM-GED counterparts are Figures 19 and 20.




(a) Marginal density.



(b) Tail Q-Q plot.




(c) ACF of raw G_t .




(d) ACF of $|G_t|$.


Figure 19. IS SPY model comparison at $K = 18$ for CHMM-L Laplace emissions (Pipeline A, 1,000 simulated paths). Panels as in Figure 18.




(a) Marginal density.



(b) Tail Q-Q plot.




(c) ACF of raw G_t .




(d) ACF of $|G_t|$.

Figure 20. IS SPY model comparison at $K = 18$ for CHMM-GED emissions with per-state shape $p_k \in [0.5, 3.0]$ (Pipeline A, 1,000 simulated paths). Panels as in Figure 18. The fitted $\hat{p}_k$ partition is reported in Figure 28 and Appendix A.2.3.


OoS validation. Figures 21–23 show the OoS validation panels for CHMM-t, CHMM-L, and CHMM-GED at $K = 18$. The CHMM-t panel demonstrates the OoS kurtosis alignment noted in the main text, the CHMM-L panel shows the tighter density fan produced by the heavier Laplace tails relative to the Gaussian, and the CHMM-GED panel shows the density fan tightening that follows from the data-chosen p_k partition.




(a) OoS density fan.



(b) KS p -value histogram.




(c) OoS ACF of raw G_t .




(d) OoS ACF of $|G_t|$.


Figure 21. OoS SPY validation at $K = 18$ for CHMM-t Student-t emissions (Pipeline A, 1,000 simulated paths). Panels show (a) OoS density fan, (b) KS p -value histogram with the $\alpha = 0.05$ threshold, (c) OoS ACF of raw G_t (linear autocorrelation, with the 99% Bartlett band), and (d) OoS ACF of $|G_t|$ (volatility clustering). The CHMM-N counterpart is Figure 27; the CHMM-L and CHMM-GED counterparts are Figures 22 and 23.




(a) OoS density fan.



(b) KS p -value histogram.




(c) OoS ACF of raw G_t .




(d) OoS ACF of $|G_t|$.


Figure 22. OoS SPY validation at $K = 18$ for CHMM-L Laplace emissions (Pipeline A, 1,000 simulated paths). Panels as in Figure 21.




(a) OoS density fan.



(b) KS p -value histogram.



(c) OoS ACF of raw G_t .



(d) OoS ACF of $|G_t|$.

Figure 23. OoS SPY validation at $K = 18$ for CHMM-GED emissions with per-state shape $p_k \in [0.5, 3.0]$ (Pipeline A, 1,000 simulated paths). Panels as in Figure 21.

Convergence. Figure 24 shows the Baum-Welch / ECM / weighted-EM log-likelihood histories at $K = 18$ for all four families. The CHMM-t, CHMM-L, and CHMM-GED curves inherit the monotone increase of the classical EM guarantee; the CHMM-t and CHMM-GED curves have slightly longer pre-plateau phases because each ECM sub-iteration refines a shape parameter (per-state ν_k for CHMM-t, per-state p_k for CHMM-GED) with a golden-section search.




Figure 24. Log-likelihood histories for the SPY IS fit (Pipeline A) at $K = 18$ across the four emission families (CHMM-N Gaussian classical EM; CHMM-t Student-t ECM with per-state ν_k golden-section sub-step; CHMM-L Laplace weighted-EM; CHMM-GED with per-state p_k three-stage ECM). All four curves are monotonically non-decreasing and plateau well before the `max_iter` = 60 budget.

A.6.17 Rydén $K = 2$ Replication

Table 26 reports the direct replication of Rydén et al.’s 1998 $K = 2$ setting on SPY IS under both initialization strategies, as discussed in Section 5. The quantile-based initialization used in the main body is row 1; five independent random-seed initializations are rows 2–6. Random init here draws per-state (μ_k, σ_k) from a moderate perturbation of the sample moments, which is representative of the random-init practice common in the 1990s HMM literature.

Table 26. Rydén replication: $K = 2$ CHMM-N on SPY IS under quantile vs random initialization (seed sub-index varies the random init).

Init	KS IS (%)	AD IS (%)	ACF-MAE	Kurt IS	KS OoS (%)
Quantile (main body)	79.0	71.5	0.0501	2.57	72.4
Random seed = 1	78.7	70.7	0.0502	2.57	73.2
Random seed = 7	75.3	69.6	0.0498	2.56	74.0
Random seed = 42	77.6	70.0	0.0498	2.59	71.8
Random seed = 100	76.2	69.4	0.0500	2.59	72.3
Random seed = 2024	76.4	69.8	0.0499	2.54	72.9

The key observation is that ACF-MAE is essentially constant at ≈ 0.050 across all six rows, matching the $K = 18$ CHMM-N of the main body. IS KS pass rate remains below the $K = 18$ operating point (75.3–79.0%) because two Gaussian components cannot tile the full multi-regime return distribution. The combination is consistent with the reading that the binding low- K constraint in the Rydén et al. [11] setup is distributional rather than temporal: the ACF is already well reproduced at $K = 2$, and it is distributional fidelity that drives the preference for $K = 18$.

A.6.18 Cross-Ticker OoS-KS ANOVA: Sector vs Ticker Decomposition

A one-way ANOVA on the 30-ticker OoS KS distribution (Table 6) by GICS sector tests the per-sector rollup framing against the alternative null that within-sector ticker variance dominates between-sector variance. With $n = 30$ tickers across $K_{\text{sec}} = 10$ sectors at three reps each, $F(9, 20) = 0.436$ at $p = 0.899$ under the parametric F-distribution; a 5,000-replicate permutation test on the same statistic gives $p = 0.903$, well above the 0.05 rejection threshold. The eta-squared is 0.164, so sector membership explains only $\approx 16\%$ of the OoS-KS variance and the bulk (83.6%) is within-sector heterogeneity (results file `sector_panel/anova_oos_ks.txt`).

The reading reframes the body’s per-sector roll-up: the “Health Care collapse” framing (sector median 14.5%) is a within-sector dispersion result driven by two ticker-specific regime introductions

Table 27. ANOVA on the 30-ticker OoS KS distribution by GICS sector. 30 tickers across 10 sectors at three large-cap representatives each; one-way F-test plus 5,000-replicate permutation test (seed = 20260420). Sector membership fails to explain OoS-KS variance at conventional levels.

Statistic	Value
$F(9, 20)$	0.436
p (F-distribution)	0.899
p (permutation, $B = 5,000$)	0.903
η^2	0.164
SS-between	4,141.16
SS-within	21,096.77

in LLY and UNH while JNJ in the same sector lands at 93.2%. Operationally, refit triggers should be calibrated on per-ticker diagnostics (Page-Hinkley on $|G_t|$ inside the IS window, etc.) rather than per-sector averages.

A.6.19 Cross-Ticker Per-State-Resolution Sensitivity ($K^* = 3$ vs $K^* = 6$ vs $K = 18$)

Body Table 6 reports the cross-ticker aggregate distribution at all three operating points used in this paper: the body headline $K^* = 3$ (state-resolution-robust default per the k -fold CV in Appendix A.9), the sensitivity reference $K^* = 6$, and the extended-state-resolution sensitivity reference $K = 18$. Per-ticker files: `sector_panel/sector_panel_summary_k3.txt`, `_k6.txt`, and `sector_panel_summary.txt` respectively.

The substantive read across the three operating points: (i) The KS distributions are within ~ 6 pp of each other on OoS median (69.1% at $K^* = 3$, 75.1% at $K^* = 6$, 73.4% at $K = 18$) and at identical 11/30 failure count, with the same regime-introduction tickers driving the failures at each state resolution (LLY, UNH, NEM, NFLX, NEE, WMT, BAC, HD, JPM at OoS KS $< 60\%$; the specific list overlaps $> 80\%$ across $K^* = 3, 6, 18$). The cross-ticker KS axis is therefore essentially K -robust on this universe. (ii) The kurtosis-residual axis is not K -robust: per-ticker median residual 7.18 at $K^* = 3$, 10.34 at $K^* = 6$, 0.61 at $K = 18$. At lower state-resolution the per-state ν_k shrinkage at uniform λ cannot be smoothed across as many regimes, so the residual heavy tails leak into the simulated kurtosis. (iii) The $|G_t|$ ACF-MAE is essentially K -robust (0.0395, 0.0403, 0.0416 across the three operating points). Production reading: KS-only consumers can deploy at any of the three operating points at parity; kurtosis-fidelity consumers should deploy at $K = 18$; readers preferring the state-resolution-robust default should deploy at the body headline $K^* = 3$.

A.6.20 Block-Bootstrap CIs on Observed Kurtosis (IS vs OoS)

Reviewer 3 / item 9 asks whether the IS / OoS kurtosis disagreement (observed 7.68 IS, 5.29 OoS, a 31% drop) is itself statistically distinguishable. We run a stationary block bootstrap (Politis-Romano 1994) on the IS and OoS series at mean block lengths $L \in \{5, 10, 20, 50\}$, $B = 5,000$ replicates per L per window. Results in `results/diagnostics/kurtosis_bootstrap.txt`.

The IS lower bound 1.92–2.99 and the OoS upper bound 7.70–8.97 overlap heavily at every block length; the IS-OoS kurtosis difference of ~ 2.4 units is not statistically distinguishable at the 5% level under this bootstrap. $\Pr(\text{IS} > \text{OoS})$ at $L = 10$ is 0.756: the data favour the alternative but do not exclude the null. Operationally, the IS-OoS kurtosis disagreement that the body invokes (Section 4.1; the Table 9 variant-decision caveat in Section 5) should be read as a point-estimate observation rather than a tested claim, and per-variant rankings by “how close is simulated kurtosis

Table 28. Stationary block bootstrap CIs on observed excess kurtosis. SPY IS ($T = 2,516$) and OoS ($T = 572$) windows, $B = 5,000$ replicates per L . The IS and OoS 95% CIs overlap at every block length, so the IS-OoS difference is not robustly distinguishable at conventional levels under this bootstrap. $\Pr(\text{IS} > \text{OoS})$ is the empirical bootstrap one-sided p -value for the alternative “IS kurtosis $>$ OoS kurtosis”.

L	IS median	IS 95% CI	OoS median	OoS 95% CI	$\Pr(\text{IS} > \text{OoS})$
5	7.21	[2.99, 12.83]	5.07	[1.09, 8.97]	0.748
10	7.34	[2.49, 12.59]	4.91	[0.99, 8.66]	0.756
20	7.30	[2.17, 12.40]	4.91	[0.90, 8.26]	0.739
50	7.18	[1.92, 12.42]	4.99	[0.96, 7.70]	0.733

to observed kurtosis” should similarly be interpreted as point-comparisons inside a wide CI envelope.

Bootstrap-CI placement of the penalised CHMM-t IS kurtosis. A natural follow-up question on the body’s penalised CHMM-t at $\lambda = 20$ is whether its simulated IS kurtosis distribution sits inside the bootstrap CI on observed. Round-2 reviewers asked whether the aggregate-mean simulated IS kurtosis (e.g. 14.91 at $K^* = 3$, above the $L = 20$ CI upper bound 12.40) is paired with a per-path distribution that mostly sits inside the CI, or whether it is paired with a per-path distribution that mostly sits above. Source / runner: `run_kurtosis_ci_placement.jl`; output: `results/kurtosis_ci_placement/kurtosis_ci_placement.csv`.

Table 29. Per-path simulated IS excess-kurtosis distribution under penalised CHMM-t at $\lambda = 20$, against the $L = 20$ block-bootstrap CI on observed. 1,000 paths per row, $T_{\text{IS}} = 2,516$, seed 20260420. The CI is [2.17, 12.40] at $L = 20$ from Table 28; the observed IS excess kurtosis is 7.68. “in CI %” is the fraction of simulated paths whose per-path excess kurtosis falls inside the CI; “ $\leq \text{CI_HI}$ %” is the fraction below the upper bound (the relevant one-sided test for overshoot).

K	agg. kurt	median	sd	Q_5	Q_{95}	≤ 12.40 (%)	in CI (%)
3	16.65	7.77	11.5	3.75	34.62	76.6	76.6
6	10.13	6.74	9.0	3.63	26.30	81.8	81.7
18	8.94	6.09	5.6	3.54	18.11	89.9	89.9

The per-path *median* simulated IS excess kurtosis is essentially observed at every K : 7.77 at $K^* = 3$, 6.74 at $K^* = 6$, 6.09 at $K = 18$, all within ~ 1.6 units of observed 7.68 and well inside the bootstrap CI. The aggregate-mean simulated IS kurtosis sits above the CI upper bound at $K^* = 3$ and $K^* = 6$ because the per-path distribution has a heavy right tail of paths at $Q_{95} \in \{18, 26, 35\}$ that drag the mean up, but the bulk of the path-level mass (76–90%) sits inside the bootstrap CI at every operating point. **The headline framing “the penalised CHMM-t at $\lambda = 20$ provides the cleanest IS heavy-tail match” is supported when read as the per-path median, not as the aggregate mean:** a small number of heavy-right-tailed paths under the per-state ν_k ECM produce an aggregate-mean overshoot, but the path-level distribution is centred near observed and statistically indistinguishable from observed at $\alpha = 0.05$ on a one-sided overshoot test for the bulk of paths.

A.6.21 Shared- ν Student-t HMM Ablation

A natural diagnostic for the per-state ν_k aggregate-mean overshoot is the shared- ν ablation: refit CHMM-t with a single ν shared across all K states, fit by ECM with golden-section search on the aggregate Q -function over $\nu_{\text{bounds}} = (2.1, 50)$. This is the standard one-parameter Student- t HMM in

the time-series literature, not a per-state ν_k mixture. Source / runner: `run_chmm_t_shared_nu.jl`; output: `results/chmm_t_shared_nu/chmm_t_shared_nu.csv`.

Table 30. Shared- ν Student-t HMM ablation against the body per-state ν_k row (1,000 paths, $T_{\text{IS}} = 2,516$, $T_{\text{OoS}} = 572$, seed 20260420, no $1/\nu$ penalty). Aggregate sim-kurt columns are mean-of-paths. Compare against the body penalised CHMM-t ($\lambda = 20$): $K^* = 3$ row 14.91 IS / 8.50 OoS; $K = 18$ row 8.56 IS / 7.07 OoS (Table 3).

K	$\hat{\nu}_{\text{shared}}$	IS				OoS		
		KS (%)	sim kurt	$ G_t $	ACF-MAE	raw ACF-MAE	KS (%)	sim kurt
3	5.81	91.9	4.68		0.0531	0.0235	82.1	4.46
6	4.67	92.7	9.38		0.0518	0.0234	82.6	6.95
18	5.80	95.8	6.25		0.0542	0.0234	88.0	5.00

Reading. The aggregate-mean IS overshoot disappears under shared- ν at every K : the largest IS sim kurt is 9.38 at $K^* = 6$ (vs. observed 7.68, well inside the bootstrap CI), versus 14.4 for the per-state ν_k unpenalised row in the body. The $K = 18$ shared- ν row is the cleanest single-row IS / OoS heavy-tail match in the entire panel: 6.25 IS / 5.00 OoS, against observed 7.68 / 5.29 and against 8.56 / 7.07 for the body penalised CHMM-t at $\lambda = 20$. IS / OoS KS pass rates are within ~ 1 pp of the body penalised row, and the $|G_t|$ ACF-MAE is identical to the body row at $K = 18$. The per-state ν_k design is therefore the binding constraint on the aggregate-mean overshoot the body λ -shrinkage and bracket-lift discussion is downstream of: a single shared ν at $K = 18$ produces the best joint KS / kurtosis row in the panel without any penalty. The body retains the per-state $\nu_k + \lambda = 20$ construction for direct comparability with the Peel and McLachlan [18], Liu and Rubin [19] tradition; consumers prioritising heavy-tail fidelity without a penalty hyperparameter should adopt the shared- ν row at $K = 18$.

A.6.22 Seventh Walk-Forward Fold (2018 / 2019 Test Slices)

Reviewer 2 / item 15 asks for a seventh walk-forward fold covering 2017–2018 (Q4 2018 drawdown) and 2019 (trade-war volatility), neither of which lies cleanly inside the existing W1–W6 layout. We run two new folds W7a (test 2018, train 2013-01–2017-12) and W7b (test 2019, train 2014-01–2018-12), the latter overlapping the existing W1 fold by construction. Same protocol as Table 70: 5y train / 1y test, CHMM-N at $K \in \{3, 18\}$, $N_{\text{paths}} = 500$. Results in `results/walkforward/walkforward_w7.txt`.

Table 31. Seventh walk-forward fold for CHMM-N. W7a covers the 2018 calendar year including the Q4 2018 drawdown; W7b covers the 2019 calendar year (trade-war volatility) and overlaps the existing W1 fold by construction. Same protocol as Table 70.

Fold (test window)	KS (%)		Kurt sim		ACF-MAE $ G_t $	
	$K = 3$	$K = 18$	$K = 3$	$K = 18$	$K = 3$	$K = 18$
W7a (2018, Q4 drawdown)	57.4	57.8	2.76	3.26	0.0645	0.0663
W7b (2019, trade war)	63.2	82.6	2.51	2.66	0.0572	0.0525

W7a sits at the lower end of the non-stress fold band (W1, W3, W5, W6 in Table 70 attain 61–83% KS): the Q4 2018 drawdown is a moderate-stress event that the IS-fixed CHMM only partially recovers, but the fold does not collapse to single-digit KS as W2 / W4 do. W7b matches the

existing W1 fold closely (both use 2019 as the test slice), validating fold-construction reproducibility. The seven-fold reading is robust: median KS at $K = 18$ is 67.7% (six folds) and 63.4% (seven folds with W7a), with the headline single-window OoS KS of 81.8% sitting at the upper end of either distribution.

A.6.23 Christoffersen-cc Monte Carlo Power Calibration at $T_{\text{OoS}} = 572$

Reviewer 1 Q4 / Reviewer 2 W2-RE1 / Reviewer 3 Q5 ask for a Monte Carlo power calibration of the Christoffersen-cc joint conditional-coverage test at the $T_{\text{OoS}} = 572$ length used in the body Table 8. We simulate breach indicator sequences from a two-state Markov chain on $\{0, 1\}$ with marginal probability α and second eigenvalue ρ , where $\rho = 0$ is the i.i.d. null and $\rho > 0$ measures the level of breach clustering; $B = 5,000$ replicates per cell. Results in `results/diagnostics/christoffersen_power/christoffersen_power.txt`.

Table 32. Christoffersen-cc power at $T_{\text{OoS}} = 572$. Empirical rejection rate at nominal $\alpha = 0.05$, across alpha levels and Markov-chain second-eigenvalue ρ . $\rho = 0$ row is the empirical Type-I error under the i.i.d. null (should sit near 5% if asymptotic chi-squared is well-calibrated). Critical values $\chi_1^2(0.05) = 3.841$, $\chi_2^2(0.05) = 5.991$.

α	ρ	rej-UC %	rej-IND %	rej-CC %
0.01	0.00	4.9	1.6	3.2
0.01	0.05	6.3	15.9	11.4
0.01	0.10	7.5	28.5	22.1
0.01	0.20	9.4	51.3	42.6
0.01	0.30	13.9	67.6	62.8
0.01	0.50	21.4	75.9	81.5
0.05	0.00	4.9	3.8	4.1
0.05	0.05	5.3	18.8	15.3
0.05	0.10	6.1	46.2	39.0
0.05	0.20	9.4	86.9	83.9
0.05	0.30	13.6	97.4	97.9
0.05	0.50	24.3	99.7	100.0

Reading. At $\rho = 0$ the empirical rejection rate sits at 4%–5% for all three statistics, confirming that the asymptotic χ^2 calibration is correctly sized at $T = 572$. Power is asymmetric across α :

- At $\alpha = 0.05$ (~ 28.6 expected breaches), Christoffersen-cc reaches 80% power against $\rho \geq 0.20$ (83.9%) and $\sim 100\%$ at $\rho = 0.50$.
- At $\alpha = 0.01$ (~ 5.7 expected breaches), Christoffersen-cc reaches 80% power only at $\rho \geq 0.50$ (81.5%); at $\rho = 0.20$ the rejection rate is 42.6%, well below conventional power thresholds.

The reading for the body Table 8 is therefore tier-dependent. The $\alpha = 0.05$ “clean pass at every (K, α) ” row of Table 8 carries strong power against moderate clustering ($\rho \geq 0.20$); the $\alpha = 0.01$ rows carry strong power only against severe clustering ($\rho \geq 0.50$), so the $\alpha = 0.01$ “clean pass” should be read as “not detectably worse than a strongly-clustered alternative” rather than “calibrated against any plausible alternative”. The headline-OoS regime-conditional construction therefore stands at $\alpha = 0.05$ and is power-bounded at $\alpha = 0.01$; the body framing in Section 4.6 should be read with this calibration in mind.

A.6.24 Engle-Manganelli (2004) Dynamic Quantile Test

Peer-review item P2.10 / R3 RE2 asks for the Engle-Manganelli (2004, JBES) Dynamic Quantile (DQ) test as a higher-power conditional-coverage alternative to Christoffersen-cc on the same OoS window. We run the standard four-lag specification (75, equation 16),

$$\text{Hit}_t - \alpha = \beta_0 + \sum_{i=1}^4 \beta_i (\text{Hit}_{t-i} - \alpha) + \beta_5 \widehat{\text{VaR}}_t + u_t, \quad (27)$$

on the IS-fixed regime-conditional VaR series of Section 4.6, with $\text{Hit}_t = \mathbf{1}\{R_{\text{OoS},t} < \widehat{\text{VaR}}_t(\alpha)\}$. The DQ test statistic is $\widehat{\beta}^\top \mathbf{X}^\top \mathbf{X} \widehat{\beta} / [\alpha(1 - \alpha)]$, distributed $\chi^2(6)$ under the null of correct conditional coverage; the 5% critical value is $\chi_6^2(0.95) = 12.59$. Source: `results/diagnostics/engle_manganelli_dq.txt`; runner: `run_engle_manganelli_dq.jl`.

Table 33. Engle-Manganelli DQ test on the regime-conditional VaR (P2.10 / R3 RE2), alongside Christoffersen-cc on the same breach series. CHMM-N at $K \in \{3, 18\}$, $\alpha \in \{0.01, 0.05\}$; $T_{\text{OoS}} = 572$, $q = 4$ lagged Hit indicators, χ_6^2 critical at 5% is 12.59.

K	α	breach %	cc p	DQ	DQ p
3	0.01	1.57	0.137	12.29	0.056
3	0.05	6.12	0.491	9.32	0.156
18	0.01	1.57	0.137	15.46	0.017
18	0.05	4.55	0.678	3.99	0.678

Reading. At $\alpha = 0.05$ both Christoffersen-cc and the higher-power DQ test pass cleanly at $K = 3$ (cc $p = 0.491$, DQ $p = 0.156$) and at $K = 18$ (cc $p = 0.678$, DQ $p = 0.678$); the body’s $\alpha = 0.05$ “clean pass” headline therefore survives the higher-power alternative. At $\alpha = 0.01$ the DQ test rejects conditional coverage at $K = 18$ ($p = 0.017$) where Christoffersen-cc does not ($p = 0.137$), and is borderline at $K = 3$ ($p = 0.056$). This is consistent with the power-calibration result of Appendix A.6.23: at $\alpha = 0.01$ Christoffersen-cc has only 42.6% power against moderate breach-clustering eigenvalues $\rho = 0.20$, so the higher-power DQ test detects miscalibration that Christoffersen-cc misses. The DQ result therefore reinforces the body framing that $\alpha = 0.05$ is the operationally informative tier; at $\alpha = 0.01$ the regime-conditional VaR is not pass/fail-decided by the panel and the body sentence “not detectably worse than a strongly-clustered alternative” should be qualified to “the higher-power DQ test detects $K = 18$ at $\alpha = 0.01$ ”.

A.6.25 OoS Regime-Probability Trajectory under IS-Fixed CHMM-N

Reviewer 2 Minor 5 / item 17 asks for an OoS regime-probability trajectory that substantiates the body’s regime-attenuation versus regime-introduction diagnosis (Section 5). We run the forward filter on the 2024-2026 OoS window under IS-fixed CHMM-N parameters at $K = 18$ and aggregate the per-state posteriors into three volatility-rank bands (low: σ -rank 1–6, mid: 7–12, high: 13–18). Results in `results/diagnostics/oos_regime_trajectory.txt` and figure `figs/Fig-OoS-Regime-Trajectory.pdf`.

The OoS high-vol-band mass (0.554) exceeds the IS-stationary value (0.502) by 5.2pp, with the difference compensated by mid-vol (−2.3pp) and low-vol (−2.9pp) bands. The OoS window therefore over-weights the high-vol regime relative to the IS distribution, but the over-weighting is small and the IS distribution does span the OoS regime mass: the OoS-vs-stationary ratio is 1.10 in

Table 34. OoS volatility-rank band masses under IS-fixed CHMM-N at $K = 18$. Posterior $\Pr(\sigma\text{-rank band})$ averaged across the OoS window, compared against IS in-sample and the IS-stationary distribution. Bands are sorted by per-state emission σ from smallest to largest.

band	OoS mean	IS mean	IS stationary
low-vol (σ -rank 1–6)	0.187	0.218	0.216
mid-vol (σ -rank 7–12)	0.259	0.287	0.282
high-vol (σ -rank 13–18)	0.554	0.495	0.502

the high-vol band and 0.87 / 0.92 in the low / mid bands. Compare against the body’s diagnosis for the W2 / W4 stress folds (Section 5, “Why the W4 stress fold fails where the headline OoS holds”), which describes those folds as “regime introductions” — an introduction in this band aggregation would manifest as a high-vol mass that the IS distribution does not span (e.g., near-unit posterior on the highest σ -rank state). The 2024–2026 OoS window does not show this pattern; the framing of the OoS window as “regime attenuation” (a return to a regime closer to but not identical to the IS volatility level) is supported by this band aggregation.

A.6.26 Quarterly-Refit Regime-Conditional VaR Back-Test

Reviewer 2 W3 / item 8 asks for a quarterly-refit version of the regime-conditional VaR back-test (Table 8) to test whether periodic refit improves the Christoffersen-cc statistic on the headline OoS window. We re-fit CHMM-N at $K \in \{3, 18\}$ every 63 trading days on a rolling 5y window, run the forward-filter through the OoS window under each refit’s parameters, and report the back-test under the same Kupiec / Christoffersen-ind / Christoffersen-cc constructions. Results in `results/diagnostics/quarterly_refit_conditional_var.txt`.

Table 35. Quarterly-refit regime-conditional VaR back-test on SPY OoS. $T_{\text{OoS}} = 572$, refit cadence = 63 days, train window = 1,260 days, seed = 20260420. Compare against Table 8 (IS-fixed parameters). Critical values: $\chi_1^2(0.05) = 3.841$, $\chi_2^2(0.05) = 5.991$.

K	α	breaches	br rate	med VaR	LR _{uc}	LR _{ind}	LR _{cc}	p_{cc}
3	0.01	6	1.05%	−5.10	0.014	3.97	3.98	0.137
3	0.05	29	5.07%	−3.13	0.006	0.19	0.20	0.906
18	0.01	6	1.05%	−5.84	0.014	3.97	3.98	0.137
18	0.05	19	3.32%	−3.14	3.83	2.08	5.91	0.052

The quarterly-refit construction passes Christoffersen-cc at $\alpha = 0.05$ on three of four rows ($K = 3, \alpha \in \{0.01, 0.05\}$ and $K = 18, \alpha = 0.01$); the fourth row ($K = 18, \alpha = 0.05$) sits at $p_{cc} = 0.052$ (statistic 5.91 vs. critical 5.99), a marginal pass that is qualitatively the same as the IS-fixed row’s $p_{cc} = 0.678$ but has tighter coverage (3.32% vs. 4.55% in Table 8). The quarterly refit therefore tightens VaR coverage but does not strictly improve the cc-statistic on this T_{OoS} ; the refit value proposition is in coverage tightness rather than in cc-improvement, with the trade-off being a mild loss of test-power headroom on the $K = 18, \alpha = 0.05$ row. Full per-fold detail is in the results file.

A.6.27 Walk-Forward Refit-Cadence Sweep (Monthly / Weekly)

A natural follow-up question on the quarterly-refit construction above is whether *faster* refits close the W2 (COVID) and W4 (2022 rate-hike) Christoffersen-cc rejections in the body walk-forward table

(Appendix A.9, Table 72). Round-2 reviewers asked for a within-fold refit-cadence sweep at monthly (every 21 trading days) and weekly (every 5 trading days) cadence on each of the six walk-forward folds, scoped at $K^* = 3$ and $\alpha = 0.05$ (the body-headline operating point and the substantive risk-management α). Source / runner: `run_walkforward_cond_var_refit_cadence.jl`; output: `results/walkforward_refit_cadence/walkforward_refit_cadence.csv`.

Table 36. Walk-forward conditional-VaR refit-cadence sweep, CHMM-N at $K^* = 3$, $\alpha = 0.05$. Six folds, train 5y / test 1y; “fold-IS-fixed” is the body baseline (no refit within the fold, from Table 72); “21d” refits every 21 trading days; “5d” every 5 trading days. At each refit, CHMM-N is re-fit on the most recent 1,260-obs window ending the day before the refit-target test day; the predictive filter is re-primed by running through the new train window. Critical value $\chi^2_2(0.05) = 5.991$; bold p_{cc} rows fail at the 5% level.

Fold	cadence	breaches	br rate	med VaR	LR _{uc}	LR _{cc}	p_{cc}
W1	fold-IS-fixed	15	5.98%	−2.43	0.48	1.65	0.44
W1	21d	15	5.95%	−2.41	0.46	1.64	0.44
W1	5d	15	5.95%	−2.19	0.46	1.64	0.44
W2	fold-IS-fixed	23	9.13%	−4.59	7.34	9.05	0.011
W2	21d	22	8.70%	−4.68	6.02	8.20	0.017
W2	5d	21	8.30%	−4.68	4.88	7.57	0.023
W3	fold-IS-fixed	20	7.97%	−2.47	3.98	4.28	0.118
W3	21d	22	8.73%	−2.66	6.10	6.11	0.047
W3	5d	23	9.13%	−2.64	7.34	7.34	0.025
W4	fold-IS-fixed	18	7.20%	−4.16	2.26	4.24	0.120
W4	21d	21	8.37%	−4.41	5.03	7.67	0.022
W4	5d	19	7.57%	−4.40	3.04	4.61	0.100
W5	fold-IS-fixed	8	3.21%	−3.90	1.91	2.44	0.30
W5	21d	8	3.20%	−3.53	1.94	2.48	0.29
W5	5d	8	3.20%	−3.48	1.94	2.48	0.29
W6	fold-IS-fixed	13	5.18%	−2.85	0.02	0.17	0.92
W6	21d	14	5.56%	−2.77	0.16	0.22	0.90
W6	5d	14	5.56%	−2.73	0.16	0.22	0.90

Reading. Faster refit **does not close** the W2 (COVID) Christoffersen-cc rejection at any cadence: p_{cc} moves from 0.011 (fold-IS-fixed) to 0.017 (monthly) to 0.023 (weekly), all below the 0.05 threshold. The COVID drawdown is fast enough that even weekly refit on a 5y rolling train cannot adapt the model in time. Worse, monthly refit *introduces* new failures on W3 and W4 that the fold-IS-fixed baseline did not have (p_{cc} on W3: 0.118 $\rightarrow$ 0.047 $\rightarrow$ 0.025 across cadences; W4: 0.120 $\rightarrow$ 0.022 $\rightarrow$ 0.100, with W4 weekly recovering but still rejecting at monthly). The mechanism is parameter drift across the refit cycle: the conditional VaR predictive mixture is sensitive to small shifts in $(\mathbf{T}, \boldsymbol{\theta}_k)$, and at monthly cadence the per-refit parameter update introduces breach-clustering that the Christoffersen-ind component picks up. The rejection counts under each cadence are: fold-IS-fixed 1/6 (W2 only), monthly 3/6 (W2 + W3 + W4), weekly 2/6 (W2 + W3). Faster refit is therefore **not** the production recipe for closing the W2 / W4 stress folds: W2 is intrinsic regime-break, and intermediate cadences trade away passing folds to fail under refit-cycle parameter drift. The right operational reading is the body’s existing recommendation: under regime-break conditions of the W2 / W4 magnitude, the conditional-VaR construction inherits the same out-of-scope behaviour as the unconditional generators, and the mitigation is not refit cadence but a model

class with explicit regime-introduction handling (skew-emission HMMs and online-EM constructions are deferred follow-ups).

A.6.28 Non-Equity Asset-Class Validation (GLD, SLV)

Reviewer 1 / R3 ask whether the CHMM scaffold transfers to a non-equity asset class. The current data window (2014–2026 Polygon/Alpaca coverage) does not extend to a non-overlapping decade for SPY, but the dataset includes the gold (GLD) and silver (SLV) ETFs over the same IS / OoS windows. We re-fit CHMM-N at $K^* = 6$ and $K = 18$ and the penalised CHMM-t at $K = 18, \lambda = 20$ on each of GLD and SLV under the same protocol as the body SPY panel. Results in `results/non_equity_validation/non_equity_validation.txt`.

Table 37. Non-equity asset-class validation on GLD and SLV. Same IS / OoS windows, $N_{\text{paths}} = 1,000$, seed = 20260420 as the SPY body panel. The OoS-KS collapse is consistent across all three CHMM configurations on both tickers and is read as a stationarity-scope limit on commodities (Section 5, stationarity-scope paragraph; the 2024–2026 OoS window represents a regime introduction for both gold under inflation-expectation re-pricing and silver under industrial-demand cycles).

Ticker	Model	IS KS%	OoS KS%	kurt obs	kurt sim	$ G_t $ ACF-MAE
GLD	CHMM-N ($K = 18$)	99.7	0.0	3.59	2.35	0.0306
GLD	CHMM-N ($K^* = 6$)	99.7	0.0	3.59	1.50	0.0307
GLD	CHMM-t pen. ($K = 18$)	99.9	0.0	3.59	3.09	0.0306
SLV	CHMM-N ($K = 18$)	99.9	0.0	6.63	5.25	0.0439
SLV	CHMM-N ($K^* = 6$)	99.3	0.0	6.63	5.13	0.0439
SLV	CHMM-t pen. ($K = 18$)	99.8	0.1	6.63	4.73	0.0440

The IS pass rates are above 99% on both tickers under all three configurations, demonstrating that the four-emission scaffold transfers cleanly to commodities at the IS step. The OoS KS distribution, however, collapses to 0% on both tickers across all three configurations, a much sharper degradation than the worst single-ticker equity OoS in Table 6 (NEM at 5.6%). This is consistent with the stationarity-scope reading of Section 5: the 2024–2026 OoS window represents a regime introduction for both gold (inflation-expectation re-pricing, central-bank gold purchases) and silver (industrial-demand cycles), and the IS-fitted CHMM has no state with the right σ or location to absorb the OoS distribution. The mitigation (periodic refit at quarterly cadence; 26) is the same as for the equity universe-scale failures and is logged in Section 5.

Independent-decade validation (a 1994–2004 vs. 2014–2024 SPY split) requires a separate data ingest beyond the current 2014–2026 Polygon/Alpaca window and is logged for the next data-pipeline pass. The non-equity result above is therefore the asset-class-transfer evidence; the cross-decade transfer evidence is deferred.

A.6.29 Sub-Decade Validation: 5-Year IS Sub-Windows

Reviewer 3 / item 6 asks for an independent decade validation. The full 1994–2004 SPY history is outside the current Polygon/Alpaca window (2014–2026) and external ingest was not in scope; we therefore split the 2014–2024 IS into two non-overlapping 5y sub-windows and run the same four-emission scaffold on each. Results in `results/subdecade_validation/subdecade_validation.txt`.

Reading. The IS KS pass rates are stable across configs (most rows $\geq 90\%$); the four-emission scaffold transfers cleanly to either 5y sub-window at the IS step. OoS pass rates are sharply

Table 38. Sub-decade validation on the 2014-2024 SPY window. Three configurations: (A) train 2014-2019 (5y) / test 2019-2024 (5y, includes COVID + 2022 rate-hike); (B) train 2019-2024 (5y) / test 2024-2026 (the body OoS, but with a 5y train); (C) train 2014-2024 (10y) / test 2024-2026 (the body row). $N_{\text{paths}} = 1,000$, seed = 20260420.

Config	Model	IS KS%	OoS KS%	IS kurt sim	OoS kurt sim	IS ACF-MAE	OoS ACF-MAE
A	CHMM-N ($K = 18$)	97.1	0.0	3.04	3.07	0.0468	0.0579
A	CHMM-N ($K^* = 6$)	94.1	0.4	3.05	3.02	0.0444	0.0533
A	CHMM-N ($K^* = 3$)	82.2	0.2	3.41	3.35	0.0442	0.0484
A	CHMM-t pen. ($K = 18$)	96.8	0.0	3.15	3.17	0.0470	0.0581
B	CHMM-N ($K = 18$)	98.4	67.1	3.08	3.02	0.0606	0.0532
B	CHMM-N ($K^* = 6$)	90.4	57.2	3.22	2.99	0.0515	0.0524
B	CHMM-N ($K^* = 3$)	88.3	70.4	3.49	3.31	0.0513	0.0518
B	CHMM-t pen. ($K = 18$)	98.3	68.9	5.58	5.27	0.0582	0.0508
C	CHMM-N ($K = 18$)	95.4	85.6	5.00	4.56	0.0510	0.0525
C	CHMM-N ($K^* = 6$)	92.1	82.4	5.66	4.71	0.0507	0.0524
C	CHMM-N ($K^* = 3$)	88.4	83.0	3.88	3.59	0.0463	0.0539
C	CHMM-t pen. ($K = 18$)	95.0	89.5	8.28	6.92	0.0542	0.0505

config-dependent:

- **Config A** (test slice spans COVID + 2022 rate-hike onset) collapses to OoS KS $\leq 0.4\%$ across all four CHMM configurations. This is the Table 70 W2 / W4 result extended to a full $T = 1,257$ test slice: when the OoS window contains a regime introduction the IS distribution does not span, the IS-fixed CHMM cannot reproduce the OoS distribution. The pattern is invariant to K and emission family.
- **Config B** (5y train ending 2024-01, same OoS window as the body) attains OoS KS 57–70% vs. the body’s 82–90% at config C: shrinking the IS window from 10y to 5y, with the boundary at 2019-01-01, costs ~ 15 –30pp of OoS distributional fidelity. The IS distribution is also tail-light (kurtosis 3.0–5.6 vs. 5.0–8.3 at 10y), reflecting that the 2019-2024 sub-window pre-2020 is shorter and includes proportionally less of the 2018 / 2020 / 2022 stress periods needed to anchor the heavy-tail regimes.
- **Config C** (the body) is the only config where the OoS triple is achieved at all four K choices.

The substantive reading is that the headline OoS performance reported in Section 4 depends on the 10y IS span; the four-emission scaffold transfers to either non-overlapping 5y sub-window with high IS fidelity but the OoS pass rate is sensitive to whether the OoS window contains a regime introduction the IS distribution does not span. The full pre-2014 1994–2004 cross-decade comparison is logged as a follow-up requiring external data ingest.

A.6.30 Quarterly-Refit 30-Ticker Cross-Ticker Panel

Reviewer 1 W4 / RE4 and Reviewer 3 W2 / RE3 ask for a quarterly-refit version of the 30-ticker cross-ticker panel (Table 6); R3 W3 specifically asks for the quarterly-refit row at the $K^* = 6$ sensitivity reference. We re-fit the penalised CHMM-t at $\lambda = 20$ on a rolling 5-year window every 63 trading days through the OoS span (10 refits per ticker, 300 fits total per state-resolution), simulate the next quarter under each refit, and concatenate the per-quarter simulations

into a single OoS path matrix per ticker, at both $K = 18$ (extended-state-resolution sensitivity) and $K^* = 6$ (sensitivity reference). Same $N_{\text{paths}} = 1,000$, seed = 20260420 as the body Table 6. Results in `results/sector_panel/sector_panel_quarterly_refit.txt` ($K = 18$) and `results/sector_panel/sector_panel_quarterly_refit_k6.txt` ($K^* = 6$).

Table 39. Quarterly-refit cross-ticker panel at $K^* = 6$ and $K = 18$. Aggregate distribution across the 30-ticker universe under the quarterly-refit protocol at both operating points, vs. the IS-fixed $K = 18$ protocol of Table 6. Refit cadence = 63 OoS trading days; train window = 1,260 obs (~ 5 y rolling). Per-ticker detail in the results files.


Metric	IS-fixed ($K = 18$)	Refit ($K^* = 6$)	Refit ($K = 18$)
IS KS% median	99.5%	98.8%	99.4%
IS KS% mean $\pm$ sd	$99.3 \pm 0.6\%$	$97.4 \pm 4.4\%$	$99.2 \pm 0.6\%$
OoS KS% median	73.4%	85.8%	83.0%
OoS KS% mean $\pm$ sd	$66.8 \pm 29.5\%$	$76.0 \pm 21.7\%$	$77.2 \pm 21.2\%$
OoS KS% [Q_1, Q_3]	[49.1, 94.3]%	[57.6, 94.3]%	[62.0, 95.4]%
Tickers OoS KS < 60%	11/30 = 36.7%	9/30 = 30.0%	7/30 = 23.3%

Reading. Quarterly refit shifts the OoS KS distribution materially at both state resolutions. At $K = 18$: median +9.6pp (73.4 $\rightarrow$ 83.0%), mean +10.4pp, IQR tighter (45.2pp $\rightarrow$ 33.4pp), failure count down by 36% (11 $\rightarrow$ 7). At $K^* = 6$: median +10.7pp (75.1 $\rightarrow$ 85.8%), mean +9.5pp, failure count down by 18% (11 $\rightarrow$ 9). The two state resolutions are within ~ 3 pp of each other on the refit median (the $K^* = 6$ refit median of 85.8% is in fact 2.8pp *above* the $K = 18$ refit median of 83.0%, so the sensitivity reference is at parity with or above the extended-state-resolution sensitivity reference under the refit protocol). The largest individual improvements concentrate on the regime-introduction tickers that the IS-fixed Table 6 flagged: at $K = 18$, LLY 7.6% $\rightarrow$ 83.7%, UNH 14.5% $\rightarrow$ 47.3%, NEM 5.6% $\rightarrow$ 15.4%, HD 41.4% $\rightarrow$ 82.2%, NFLX 45.1% $\rightarrow$ 61.5%. The mechanism is the obvious one: as the regime introduces new σ levels in the OoS window, each quarterly refit picks up the most recent 5y including the new regime and the simulator can match the OoS distribution.


A small number of tickers regress: DIS 96.4% $\rightarrow$ 48.4%, BAC 84.8% $\rightarrow$ 55.8%, AAPL 95.1% $\rightarrow$ 66.5%. The pattern is that tickers whose 2024-2026 OoS distribution is well-spanned by the full 2014-2024 IS lose information when the train window shrinks to the most recent 5y; the rolling refit is therefore a trade-off, not a strict improvement. Operationally, a refit-trigger rule (refit only when the rolling KS or a Page-Hinkley statistic on $|G_t|$ crosses a threshold) would capture the regime-introduction wins without paying the cost on the well-spanned tickers; the rule is logged as a follow-up companion experiment.

The substantive reading for the body Section 4.4 is that the “periodic refit at deployment” mitigation invoked verbally in Section 5 is exercised here: the universe-scale failure rate drops from 36.7% to 23.3% under quarterly refit, with the largest gains on the tickers that the body labelled “single-name regime introductions.”


A.6.31 Stylized-Facts Figure (Empirical SPY IS)




(a) Marginal density.



(b) Normal Q-Q plot.




(c) ACF of raw G_t .




(d) ACF of $|G_t|$.

Figure 25. The three Cont stylized facts on SPY IS. (a) marginal density with maximum-likelihood Gaussian and Laplace overlays; (b) normal Q-Q plot; (c) empirical ACF of raw G_t at lags 1–252 with 99% Bartlett band; (d) empirical ACF of $|G_t|$ on the same window.


A.6.32 IS / OoS CHMM-N Validation Figures at $K = 18$




(a) Marginal density.



(b) Tail Q-Q plot.




(c) ACF of raw G_t .




(d) ACF of $|G_t|$.


Figure 26. IS CHMM-N validation on SPY at $K = 18$ (1,000 simulated paths). Panels (a/b) verify marginal fidelity, panels (c/d) the two ACF-based stylized facts.




(a) OoS density fan.



(b) KS p -value histogram.



(c) OoS ACF of raw G_t .




(d) OoS ACF of $|G_t|$.

Figure 27. OoS CHMM-N validation on SPY at $K = 18$ (1,000 paths, $T_{\text{OoS}} = 572$). Same colour conventions as Figure 26.

A.6.33 CHMM-GED $\hat{p}_k$ Partition (Gaussian-bulk / Laplace-tail)

CHMM-GED carries a per-state shape $p_k \in [p_{\min}, p_{\max}] = [0.5, 3.0]$ that nests Gaussian ($p = 2$) and Laplace ($p = 1$) as boundary cases. The empirical fit on SPY IS at $K = 18$ partitions states bimodally (Figure 28): thirteen states at the upper bracket $p_k = 3.0$, one at $p_k = 1.6$, and four in the Laplace regime $p_k \in [0.86, 1.24]$. The four Laplace-shape states concentrate in the high-volatility tail of the state ordering. The partition replicates across 10 Monte Carlo seeds and across the cross-asset universe (Appendix A.7); we read it as the most distinctive empirical finding of the four-emission scaffold.



(a) Histogram of fitted $\hat{p}_k$.

(b) $\hat{p}_k$ vs. volatility rank.

Figure 28. CHMM-GED per-state shape $\hat{p}_k$ partition ($K = 18$, SPY IS, seed 20260420). The bimodal partition is visible at the upper bracket $p_k = 3.0$ and around the Laplace shape; the Laplace-shape cluster sits at the high-volatility ranks.

A.6.34 HSMM at $K = 3$: Numerical Full-Rank Diagnostic

The off-diagonal jump matrix of the fitted ML HSMM at $K = 3$ has eigenvalues $\{1.000, 0.823, 0.177\}$ and singular values $\{1.150, 0.762, 0.166\}$; the deflated jump matrix $\mathbf{A} - \mathbf{1}\bar{\pi}^\top$ is full rank at $K - 1 = 2$ with non-trivial eigenvalues $\{0.823, 0.177\}$ (numerical rank 2 at tolerance 10^{-10}). The dominant non-unit modulus $|\lambda_2| = 0.823$ is comparable to the CHMM-N at $K = 3$ value $|\lambda_2| = 0.953$, so the spectral mechanism of the off-diagonal jump matrix is not collapsed: a near-zero $|\lambda_2|$ is not the cause of the $|G_t|$ ACF-MAE regression to the i.i.d. baseline. The active mechanism is the fitted Pareto sojourn distribution at $K = 3$ which concentrates probability mass on a single low-volatility state (the stationary distribution of $\mathbf{A}$ is $\bar{\pi} \approx (0, 1, 0)$ with $\sigma_2 = 0.97$, the lowest variance of the three states), so under simulation the chain spends long stretches inside a single state’s i.i.d. emissions and the within-state i.i.d. baseline dominates the ACF integration over the IS window. The same Yu (2010) construction at $K = 18$ collapses to a near-degenerate local optimum on $T_{\text{IS}} = 2,516$ (IS KS 0.8%, simulated kurtosis 3.65): the joint $K^2 + KD_{\text{max}}$ free-parameter count exceeds the effective IS sample size, so the higher- K ML HSMM is best read as a companion-paper direction at longer training windows or under a constrained sojourn parameterisation.

A.6.35 Gamma-Sojourn HSMM as an Alternative Sojourn Family

Reviewer R1 RE3 asks whether a Gamma-sojourn HSMM (closer to Bulla and Bulla [12]) closes the $|G_t|$ ACF-MAE gap that the truncated-Pareto HSMM does not. We re-fit the same Yu (2010) explicit-duration EM scaffold but replace the truncated discrete Pareto sojourn pmf $p(d) \propto d^{-(\alpha+1)}$ with a discretised continuous Gamma, $p(d) = F_\Gamma(d; \alpha, \beta) - F_\Gamma(d - 1; \alpha, \beta)$ on $\{1, \dots, D_{\text{max}}\}$, with per-state shape α and scale β updated by method-of-moments on the expected-duration counts (mean $\alpha\beta$, variance $\alpha\beta^2$) at every M-step. Source / runner: `run_hsmm_ml_gamma.jl`; output: `results/hsmm_ml_gamma/hsmm_ml_gamma_metrics.csv`; paper-side: `results/robustness/hsmm_ml_gamma_metrics.csv`.

Reading. The Gamma-sojourn variant flips two of the body’s HSMM findings:

Table 40. HSMM sojourn-family comparison: truncated Pareto (body) vs. discrete Gamma (R1 RE3 ablation). Yu (2010) explicit-duration EM, $D_{\max} = 200$, $T_{\text{IS}} = 2,516$, $T_{\text{OoS}} = 572$, 1,000 paths, seed 20260420. The Gamma-sojourn HSMM at $K = 18$ converges where Pareto collapses, and produces the cleanest $|G_t|$ ACF-MAE on IS in the entire HSMM panel.

Sojourn family	K	KS pass rate (%) $\uparrow$		Sim kurtosis		$ G_t $ ACF-MAE $\downarrow$
		IS	OoS	IS	OoS	IS
Pareto (body)	3	98.4	91.0	3.46	3.38	0.0629
Pareto (body)	18	0.8	33.4	—	—	collapse
Gamma (R1 RE3)	3	77.0	80.7	2.49	2.56	0.0528
Gamma (R1 RE3)	18	86.0	80.2	4.28	4.06	0.0462

- *Convergence at $K = 18$.* The body’s truncated Pareto HSMM at $K = 18$ collapses to a near-degenerate optimum (IS KS 0.8%, sojourn distribution concentrated on a single low-volatility state, EM converges in 9 iterations to a local maximum that fails distributional fidelity). The Gamma-sojourn HSMM at the same $K = 18$ converges in 9 iterations to a non-degenerate optimum at IS KS 86.0% / OoS KS 80.2%. The lighter-tailed Gamma sojourn distributes probability mass across sojourn durations more uniformly than the heavy-tailed Pareto, which prevents the single-state mass collapse that drives the Pareto $K = 18$ failure mode.
- *$|G_t|$ ACF-MAE recovery.* The body’s Pareto HSMM at $K^* = 3$ has $|G_t|$ ACF-MAE 0.0629, indistinguishable from the i.i.d. baseline (Section 4); the Gamma HSMM at $K = 18$ has $|G_t|$ ACF-MAE **0.0462**, lower than the body CHMM-N $K = 18$ value of 0.0509 and the cleanest $|G_t|$ ACF match in the entire HSMM panel. R1 RE3’s specific hypothesis that “a Gamma-sojourn HSMM may close the $|G_t|$ ACF-MAE gap that the truncated-Pareto HSMM does not” is therefore confirmed quantitatively.

The trade-off is on the KS axis: the Gamma-sojourn HSMM at $K = 3$ trades $\sim 20\text{pp}$ of KS pass rate ($98.4\% \rightarrow 77.0\%$ IS) for the ACF-MAE recovery from 0.0629 to 0.0528, and at $K = 18$ trades $\sim 5\text{pp}$ on OoS KS (91.0% Pareto $K^* = 3$ vs. 80.2% Gamma $K = 18$) for an ACF-MAE that is 0.0167 lower. The substantive read for the body HSMM panel is that there is no single “best HSMM” row: Pareto at $K^* = 3$ wins raw OoS KS, Gamma at $K = 18$ wins joint $|G_t|$ ACF-MAE, and the choice between them is use-case-driven on the same KS-vs-ACF-MAE axis as the CHMM-vs-HSMM choice elsewhere in the paper.

A.6.36 Unconditional VaR / ES Envelope Panel

Table 41. VaR and Expected-Shortfall envelope for SPY ($N_{\text{paths}} = 1,000$). Each entry is the median and [5, 95]% envelope across paths of the left-tail VaR / ES at $\alpha \in \{0.01, 0.05\}$; observed historical values are in the first row of each block.

Model	IS VaR _{0.01} median [5%, 95%]	IS ES _{0.01} median [5%, 95%]	IS VaR _{0.05} median [5%, 95%]	IS ES _{0.05} median [5%, 95%]
<i>Observed</i>	−6.38	−9.00	−3.43	−5.50
Bootstrap	−6.38 [−7.04, −5.99]	−8.86 [−10.37, −7.78]	−3.43 [−3.68, −3.20]	−5.46 [−5.99, −5.05]
GARCH	−5.80 [−8.62, −4.52]	−7.86 [−13.54, −5.77]	−3.19 [−3.92, −2.71]	−4.89 [−7.11, −3.91]
CHMM-N	−6.64 [−8.01, −5.47]	−8.80 [−10.16, −7.34]	−3.45 [−4.05, −2.95]	−5.45 [−6.33, −4.61]
CHMM-t	−6.58 [−7.51, −5.83]	−9.06 [−10.70, −7.79]	−3.40 [−3.90, −2.90]	−5.51 [−6.18, −4.85]
CHMM-L	−6.89 [−7.85, −6.01]	−9.15 [−10.42, −7.97]	−3.30 [−3.77, −2.88]	−5.53 [−6.17, −4.87]
CHMM-GED	−6.85 [−7.70, −6.06]	−8.79 [−9.83, −7.82]	−3.43 [−3.98, −2.93]	−5.55 [−6.23, −4.91]
Model	OoS VaR _{0.01}	OoS ES _{0.01}	OoS VaR _{0.05}	OoS ES _{0.05}
<i>Observed</i>	−5.51	−7.94	−2.77	−4.67
Bootstrap	−6.33 [−7.60, −5.28]	−8.43 [−11.76, −6.59]	−3.39 [−3.91, −2.88]	−5.34 [−6.38, −4.50]
GARCH	−5.35 [−9.98, −3.61]	−6.72 [−13.51, −4.38]	−3.14 [−4.94, −2.36]	−4.59 [−7.92, −3.23]
CHMM-N	−6.30 [−9.02, −4.30]	−8.44 [−11.36, −5.42]	−3.39 [−4.74, −2.47]	−5.27 [−7.23, −3.74]
CHMM-t	−6.32 [−8.13, −4.80]	−8.60 [−12.20, −6.29]	−3.29 [−4.27, −2.45]	−5.34 [−6.72, −4.08]
CHMM-L	−6.69 [−8.59, −4.91]	−8.88 [−11.48, −6.73]	−3.27 [−4.30, −2.51]	−5.47 [−6.82, −4.20]
CHMM-GED	−6.60 [−8.30, −4.82]	−8.51 [−10.68, −6.47]	−3.32 [−4.38, −2.46]	−5.41 [−6.66, −4.06]

A.6.37 Held-Out Cross-Validation of the $1/\nu_k$ Penalty λ

Reviewer 2 (peer-review item P1.3 / R2.W6) asks whether the body operating-point penalty rate $\lambda = 20$ is selected on data that overlaps the OoS evaluation window. To remove that concern we refit penalised CHMM-t on a strictly pre-2020 estimation slice (2014-01-03 to 2018-06-29, $n = 1,130$) and score each λ on the validation slice (2018-07-02 to 2019-12-31, $n = 378$, both fully pre-COVID and pre-2022-rate-hike). Source: `results/nu_shrinkage_sweep/lambda_cv_pre2020.txt`.

Table 42. Held-out CV of the $1/\nu_k$ shrinkage rate λ on the pre-2020 slice. Estimation 2014-01-03 to 2018-06-29; validation 2018-07-02 to 2019-12-31. “val LL/obs” is per-observation held-out log-likelihood; “val KS” is two-sample Kolmogorov–Smirnov pass rate against the validation series; “sim K” columns are the mean simulated excess kurtosis on the estimation- and validation-length simulated paths.

λ	val LL/obs	val KS (%)	sim K (est)	sim K (val)
0	−2.119	92.0	3.88	3.97
5	−2.112	89.8	3.62	3.15
10	−2.085	91.0	3.37	3.24
20	−2.075	90.0	3.22	3.02
50	−2.077	91.6	3.07	2.76
100	−2.077	91.6	3.07	2.75
200	−2.076	91.2	3.09	2.73

Reading: the held-out validation log-likelihood is maximised at $\lambda^* = 20$, matching the body operating point of Table 3. Validation KS picks $\lambda^* = 0$ but the band $[0, 200]$ spans only 2.2pp (89.8–92.0%), within the simulation-noise envelope at $N_{\text{paths}} = 500$. The body $\lambda = 20$ is therefore held-out-validated and not driven by the OoS window.

A.6.38 Effective State Distinctness at $K = 18$

Reviewer 2 (peer-review item P4.4 / R2.W3) asks whether all $K = 18$ states carry operationally distinct (μ_k, σ_k) or (μ_k, σ_k, ν_k) triples. The diagnostic is single-linkage agglomerative clustering on standardized parameter triples at threshold $\tau = 0.20$. Source: `results/diagnostics/state_distinctness.txt`.

Table 43. Effective state distinctness at $K = 18$ on SPY IS. Standardized single-linkage clustering at $\tau = 0.20$ in the (z-scored) parameter space.

Family	Effective distinct states	Notes
CHMM-N	11/18	2D feature (μ_k, σ_k)
CHMM-t	12/18	3D feature $(\mu_k, \sigma_k, 1/\nu_k)$; 13/18 states pinned at $\nu_{\max} = 50$, 2/18 at $\nu_{\min} = 2.1$

Reading: under unpenalised CHMM-t the upper-bracket pinning at $\nu_k = 50$ collapses 13/18 states to operationally Gaussian (in second-moment behaviour), but the (μ, σ) axis still distinguishes them well enough to keep 12/18 effective states. The penalised CHMM-t at $\lambda = 20$ pulls $1/\nu_k$ towards a moderate value across states and reduces the upper-bracket pinning (Appendix A.6.37). The effective-state count is below K but well above the $K^* = 6$ sensitivity reference, so the $K = 18$ extension is a representational-economy choice rather than a free-rider on identifiability.

A.6.39 K_{eff} -Corrected Information-Criterion Re-Rank

Reviewer 2 (peer-review item P4.7 / R2 Tier 2) asks whether the four ICs in Table 2 re-rank when the standardized-distance K_{eff} diagnostic of Appendix A.6.38 replaces K_{nom} in the parameter-count term. We refit CHMM-N at every $K \in \{3, 6, 9, 12, 15, 18, 21\}$ on the pre-2020 estimation slice (2014-01-03 to 2018-06-29, $n = 1,130$), apply the single-linkage diagnostic at $\tau = 0.20$ to count K_{eff} , and re-evaluate AIC, BIC, HQC, CAIC with $p_{\text{eff}} = K_{\text{eff}}(K_{\text{eff}} - 1) + 2K_{\text{eff}}$ in place of $p_{\text{nom}} = K(K - 1) + 2K$ in the penalty term. The log-likelihood is unchanged (the fit is the same; only the penalty differs). Source: `results/k_eff_corrected_ic/K_Eff_Corrected_IC.txt`.

Table 44. K_{eff} -corrected information-criterion re-rank on CHMM-N, pre-2020 estimation slice ($n = 1,130$). Subscript n indicates the textbook K_{nom} -based penalty $p_{\text{nom}} = K(K - 1) + 2K$; subscript e indicates K_{eff} -based penalty $p_{\text{eff}} = K_{\text{eff}}(K_{\text{eff}} - 1) + 2K_{\text{eff}}$. K_{eff} plateaus at 9–11 for $K \geq 12$, so the K_{nom} blow-up at large K is largely a parameter-counting artefact. AIC and HQC re-rank to $K^* = 21$ under the corrected penalty; BIC and CAIC continue to select $K^* = 3$ under both penalties, because $\log n \approx 7.0$ keeps the corrected penalty large enough at every K to outweigh the IS log-likelihood gain.

K	K_{eff}	ℓ	AIC _{n}	AIC _{e}	BIC _{n}	BIC _{e}	HQC _{n}	HQC _{e}	CAIC _{n}	CAIC _{e}
3	3	-2,213.0	4,450.0	4,450.0	4,510.3	4,510.3	4,472.8	4,472.8	4,522.3	4,522.3
6	6	-2,185.7	4,455.4	4,455.4	4,666.6	4,666.6	4,535.2	4,535.2	4,708.6	4,708.6
9	9	-2,168.7	4,517.5	4,517.5	4,970.2	4,970.2	4,688.5	4,688.5	5,060.2	5,060.2
12	9	-2,106.2	4,524.4	4,392.4	5,309.1	4,845.1	4,820.8	4,563.4	5,465.1	4,935.1
15	9	-2,074.9	4,629.7	4,329.7	5,836.9	4,782.4	5,085.8	4,500.8	6,076.9	4,872.4
18	11	-2,011.5	4,707.0	4,287.0	6,427.3	4,951.0	5,356.9	4,537.9	6,769.3	5,083.0
21	9	-1,949.3	4,822.5	4,078.5	7,146.4	4,531.2	5,700.5	4,249.5	7,608.4	4,621.2

Reading: the diagnostic confirms that the apparent K_{nom} IC blow-up at $K \geq 12$ in Table 2 (BIC _{n} rises from 4,510 at $K = 3$ to 7,146 at $K = 21$) is largely fictitious in the parameter-economy sense: K_{eff} stays at 9–11 across $K \in \{12, 15, 18, 21\}$, so the additional K_{nom} states are operationally

redundant and the corrected penalty is much smaller. AIC_e and HQC_e select $K^* = 21$ because the log-likelihood gain from $-2,213$ to $-1,949$ (~ 264 units across the sweep) dominates a corrected penalty that scales with $p_{\text{eff}} = 90$ rather than $p_{\text{nom}} = 462$. BIC_e and $CAIC_e$ continue to select $K^* = 3$: the $\log n \approx 7$ multiplier keeps the corrected penalty large enough at every K to outweigh the IS log-likelihood gain. The body headline $K^* = 3$ (held-out per-observation log-likelihood and held-out KS on the pre-2020 slice under four-fold and six-fold rolling-origin CV) does not depend on a parameter-count penalty at all and is therefore unchanged. The substantive read is that the held-out criteria, not the IC penalties, are the load-bearing selection inputs in this problem; the IC sweep is a diagnostic, not a deciding criterion.

$K_{\text{nominal}} = 11$ **rebuild against** $K = 18$ **nominal**. The K_{eff} diagnostic above reports $K_{\text{eff}} = 11$ at $K_{\text{nominal}} = 18$ on SPY IS for CHMM-N. Round-2 reviewers asked whether a direct $K_{\text{nominal}} = 11$ rebuild matches the $K = 18$ nominal panel on the headline metrics: if it does, the $K = 18$ over-parameterisation is a parameter-counting artefact rather than an information-bearing extra resolution; if it does not, the K_{eff} merge loses information that the $K = 18$ panel exploits. Source / runner: `run_k_eff_rebuild.jl`; output: `results/k_eff_rebuild/k_eff_rebuild.csv`.

Table 45. $K_{\text{nominal}} = K_{\text{eff}} = 11$ **rebuild against** $K = 18$ **nominal** (CHMM-N, SPY IS / OoS, 1,000 paths, seed 20260420). All four headline metric axes match within sampling noise; $K = 11$ slightly exceeds $K = 18$ on OoS KS and $|G_t|$ ACF-MAE.

K_{nom}	KS (%)		sim kurt		$ G_t $ ACF-MAE		raw ACF-MAE
	IS	OoS	IS	OoS	IS	OoS	IS
11	94.4	82.8	4.52	4.11	0.0505	0.0528	0.0237
18	94.0	81.0	5.04	4.42	0.0514	0.0526	0.0237

The $K = 11$ row matches or slightly exceeds the $K = 18$ row on every metric axis: IS KS +0.4pp, OoS KS +1.8pp, $|G_t|$ ACF-MAE essentially identical (within 0.001), and IS sim kurt within 0.5 units. The $K = 18$ over-parameterisation therefore carries no operational benefit on this dataset under CHMM-N: the structural information is captured at $K_{\text{nominal}} = K_{\text{eff}} = 11$, and the additional 7 nominal states under $K = 18$ are merged-into-equivalent under the single-linkage diagnostic of Appendix A.6.38. The body’s report of the $K = 18$ row alongside the $K = 11$ effective parameterisation is therefore a representational-economy convention rather than an extra-information row; consumers who need IC penalties on parameter counts should use $K_{\text{nominal}} = K_{\text{eff}}$ from the start.

A.6.40 Exact-Binomial Kupiec Complement to the Asymptotic LR Test

Reviewer 2 (peer-review item P4.3 / R2.Q5) notes that at $T_{\text{OoS}} = 572$ the asymptotic Kupiec LR statistic piles against the $\chi_1^2(0.05) = 3.841$ critical value at integer breach counts $k \in \{14, 19, 20\}$, so an exact-binomial complement is the size-correct diagnostic. We compute the exact two-sided binomial p -value (R `binom.test` default) under $H_0 : X \sim \text{Binomial}(572, \alpha)$ for the conditional-VaR breach counts in Table 8. Source: `results/diagnostics/utility/exact_binomial_kupiec.txt`.

Table 46. Exact-binomial unconditional-coverage test alongside the asymptotic LR Kupiec. Breach counts from Table 8; $T_{\text{OoS}} = 572$. Both the asymptotic and exact tests pass at $\alpha = 0.05$ on every row.

Configuration	breaches	expected	rate (%)	p (asyp. LR)	p (exact bin.)
CHMM-N $K = 3, \alpha = 0.01$	9	5.72	1.57	0.20	0.20
CHMM-N $K = 3, \alpha = 0.05$	35	28.6	6.12	0.23	0.21
CHMM-N $K = 18, \alpha = 0.01$	9	5.72	1.57	0.20	0.20
CHMM-N $K = 18, \alpha = 0.05$	26	28.6	4.55	0.61	0.70
CHMM-t pen. $K = 3, \alpha = 0.01$	8	5.72	1.40	0.37	0.29
CHMM-t pen. $K = 3, \alpha = 0.05$	32	28.6	5.59	0.52	0.50
CHMM-t pen. $K = 18, \alpha = 0.01$	10	5.72	1.75	0.10	0.09
CHMM-t pen. $K = 18, \alpha = 0.05$	29	28.6	5.07	0.94	0.92

Reading: the asymptotic and exact p -values agree closely (all within 0.07); both tests pass at $\alpha = 0.05$ on every row. The asymptotic-LR pile-up at $k \in \{14, 19, 20\}$ is therefore not a load-bearing concern in the back-test rows actually observed in Table 8.

A.6.41 Diebold–Mariano Bandwidth Sensitivity

Reviewer 2 (peer-review item P4.2 / R2.W5) asks whether the within-CHMM CRPS Diebold–Mariano equivalence ($p > 0.45$) survives bandwidth perturbation around the default Newey–West HAC bandwidth $h = \lfloor T^{1/3} \rfloor = 8$. Table 47 sweeps $h \in \{4, 8, 16, 32\}$. Source: `results/crps_dm/CRPS_DM_bandwidth_sweep.txt`.

Table 47. CRPS Diebold–Mariano p -values under bandwidth sweep on SPY OoS. The within-CHMM family pairs all retain $p > 0.38$ across all bandwidths; the CHMM-N vs. GARCH pair moves $p \in [0.087, 0.175]$ and the qualitative direction is stable across h .

Pair (A vs. B)	$h = 4$	$h = 8$	$h = 16$	$h = 32$
CHMM-N vs. CHMM-t	(−0.71, 0.48)	(−0.73, 0.47)	(−0.70, 0.48)	(−0.66, 0.51)
CHMM-N vs. CHMM-L	(−0.71, 0.48)	(−0.75, 0.45)	(−0.83, 0.40)	(−0.88, 0.38)
CHMM-t vs. CHMM-L	(−0.05, 0.96)	(−0.04, 0.97)	(−0.04, 0.97)	(−0.05, 0.96)
CHMM-N vs. Bootstrap	(−0.70, 0.48)	(−0.71, 0.48)	(−0.73, 0.47)	(−0.81, 0.42)
CHMM-N vs. GARCH	(−1.71, 0.087)	(−1.65, 0.10)	(−1.53, 0.13)	(−1.36, 0.18)
CHMM-t vs. GARCH	(−1.22, 0.23)	(−1.15, 0.25)	(−1.04, 0.30)	(−0.95, 0.35)

Reading: within-CHMM equivalence is robust across bandwidths. The CHMM-vs-GARCH directional advantage is also stable but its p -value at $h = 4$ is 0.087 (just above the 0.05 threshold); a longer bandwidth gives a more conservative p . The body framing remains that the four CHMM emission families are interchangeable on this evaluation.

A.6.42 Stabilised ML HSMM at Intermediate K

Reviewers 1 and 3 (peer-review item P3.3) ask for a like-for-like comparison between CHMM and HSMM at intermediate K . The body Table 3 reports the ML HSMM at $K^* = 3$ (clean: 98.4% IS / 91.0% OoS KS) and the appendix notes that the $K = 18$ ML HSMM collapses to a near-degenerate optimum on $T_{\text{IS}} = 2,516$. Table 48 fills the $K \in \{6, 9, 12\}$ rows under the same Yu (2010) explicit-duration EM with truncated discrete Pareto sojourns and $\alpha \in [0.3, 5.0]$ clipping (already a regulariser). Source: `results/hsmm_ml/hsmm_ml_intermediate_K_metrics.csv`.

Table 48. ML HSMM at intermediate K on SPY. Clean fit at $K = 6$ (comparable to CHMM-N at $K^* = 6$); near-degenerate optima at $K = 9$ and $K = 12$ with IS KS collapsing to 0%.

K	IS KS (%)	OoS KS (%)	IS kurt	OoS kurt	$ G_t $ ACF-MAE	status
3	98.4	91.0	3.46	3.38	0.0629	clean (body Table 3)
6	95.9	89.8	3.72	3.73	0.0571	clean
9	0.0	6.7	3.03	2.85	0.0614	degenerate
12	0.0	34.3	3.08	3.02	0.0623	degenerate
18	—	—	—	—	—	near-degenerate (companion-paper)

Reading: the HSMM scaffold is competitive with CHMM at $K \leq 6$ but does not extend usefully to higher K on the SPY $T_{IS} = 2,516$ window. The CHMM scaffold’s quantile-based initialisation reaches non-degenerate $K = 18$ fits where the HSMM collapses, which is the operational reason the body $K = 18$ comparison is one-sided. A regulariser stronger than the existing α clipping (e.g. a Dirichlet prior on transitions plus a Gamma prior on α) is the natural follow-up; we treat the present finding as a documented scope limit of the HSMM scaffold rather than a CHMM advantage in the strict sense.

A.6.43 Markov-Switching Stochastic Volatility (MSSV) Baseline

Reviewer 3 (peer-review item P3.2 / R3.W2) asks for an MSSV baseline as a fair modern regime-switching competitor. We implement the 2-state MSSV of So et al. [76] (Carvalho-Lopes 2007 specification) under the Hamilton-Kim-Nelson quasi-MLE shortcut:

$$r_t = \mu + \exp(h_t/2) \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(0, 1)$$

$$h_t = c_{s_t} + \phi_{s_t} h_{t-1} + \sigma_{\eta, s_t} \eta_t,$$

with $s_t \in \{1, 2\}$ a 2-state Markov chain. Quasi-MLE on the log-squared linearisation $y_t = \log((r_t - \hat{\mu})^2) - \mathbb{E}[\log \chi_1^2]$ via a Hamilton-style filter with Kim-Nelson regime collapse. Source: `results/mssv_baseline/mssv_baseline.txt`; script: `run_mssv_baseline.jl`.

Table 49. 2-state MSSV quasi-MLE fit on SPY IS. The fit converges to a near-absorbing regime-2 ($T_{22} \approx 1.0$) with one regime carrying high innovation variance and the other near-deterministic; this is the well-documented degeneracy of 2-state MSSV under quasi-MLE on financial returns and is the reason Carvalho and Lopes [77] use particle MCMC rather than Hamilton-Kim-Nelson collapse. The IS KS pass rate drops to 0% and the simulated kurtosis overshoots to 14.6 against observed 7.68; the row is reported as a documented negative result.

Model	IS KS (%)	OoS KS (%)	IS kurt	OoS kurt	$ G_t $ ACF-MAE	raw ACF-MAE
MSSV ($K = 2$, quasi-MLE)	0.0	0.1	14.58	14.64	0.074	0.032
<i>Body baselines for reference:</i>						
GARCH(1,1)	23.4	60.8	7.06	2.96	0.049	0.024
GARCH(1,1)- t	57.3	80.8	15.13	—	0.032	0.017
MS-GARCH ($K = 6$)	34.5	33.4	4.41	—	0.043	0.017

Reading: the 2-state MSSV under quasi-MLE does not outperform the body benchmarks on SPY. The fitted $T \approx [[0.998, 0.002], [0.000, 1.000]]$ shows a near-absorbing regime structure that is the standard failure mode of Hamilton-Kim-Nelson collapse on heavy-tailed financial returns. A full particle-MCMC implementation of Carvalho-Lopes (2007) is the appropriate companion-paper

direction; the body MS-GARCH- $K \in \{2, 3, 6\}$ rows already cover the regime-switching-volatility comparison axis at the level the body argument requires.

A.6.44 Monthly-Refit Cross-Ticker Panel

Reviewer 1 (peer-review item P5.2 / R1.Exp4) asks for a finer refit-cadence point alongside the quarterly-refit panel of Appendix A.6.30. We re-fit the penalised CHMM-t at $K = 18$, $\lambda = 20$ on the same sector-balanced 30-ticker panel with refit cadence shortened from quarterly (63 trading days) to monthly (21 trading days), keeping the same 5-year rolling estimation window and the same $N_{\text{paths}} = 1,000$. Source: `results/sector_panel/sector_panel_monthly_refit.{csv,txt}`.

Table 50. Refit-cadence trade-off on the 30-ticker cross-ticker panel. OoS KS distribution under three refit cadences. Each shorter refit window yields a roughly proportional further reduction in failure count without saturating.

Refit cadence	median OoS KS (%)	mean OoS KS (%)	$[Q_1, Q_3]$	failures < 60%	refits per tick
IS-fixed (no refit, body Table 6)	73.4	66.8 ± 29.5	$[49.1, 94.3]$	11/30	0
Quarterly (63 days)	83.0	77.2 ± 21.2	$[62.0, 95.4]$	7/30	9
Monthly (21 days)	86.7	78.8 ± 20.4	$[62.2, 95.9]$	5/30	28

Reading: the marginal lift from IS-fixed to quarterly is +9.6pp on the median and -4 failures; from quarterly to monthly is +3.7pp and -2 failures. The trade-off is nearly linear in $\log(\text{cadence}^{-1})$ on the median axis at this scale, suggesting that further reduction below monthly would yield diminishing returns at the cost of $\sim 28 \rightarrow 56$ refits per ticker. The body Section 4.4 narrative is unchanged: periodic refit at deployment is required for OoS distributional fidelity on regime-introduction tickers, and the operational deployment choice is between quarterly (recommended default trade-off) and monthly (best fidelity at $3\times$ the compute cost).

A.6.45 Alpaca Data-Depth Probe for the Independent-Decade Validation

The Section 5 Limitations paragraph notes that the body data pipeline (Polygon.io via Massive, 2014-01-03 onwards; Alpaca / IEX held-out extension to 2026-04-20) does not reach the 1994-2004 window required for an independent-decade SPY validation. We attempted the validation by pulling SPY daily bars for 1994-01-01 to 2004-12-31 via the sibling Alpaca Markets SDK (`alpaca-markets-sdk` at `../alpaca-markets-sdk` relative to the model project, runner: `run_fetch_spy_independent_decade.jl`). Both feeds returned zero bars for the target window. A finer probe at two-year intervals from 1994 to 2018 (Table 51; runner: `run_probe_alpaca_history.jl`) shows that on the credentials available for this study the SIP feed begins at 2016-01-04 and the IEX feed begins at 2018-11-01.

The Polygon.io / Massive feed already in the body pipeline begins at 2014-01-03, so the Alpaca / Polygon data sources span 2014-2026 inclusive. The independent-decade validation R3 W6 / R3 RE3 calls for was completed via a CRSP day-pass at revision time; the cross-decade run is reported in Appendix A.6.46.

Table 51. Alpaca SDK SPY data-depth probe. Counts of daily bars returned by `get_bars(client, "SPY", "1Day"; start, finish, feed)` for each calendar year 1994 through 2018 (two-year stride). Empty rows omitted. The SIP feed reaches back to 2016-01-04; the IEX feed reaches back to 2018-11-01 (one bar in the 2018 row reflects the late-November activation date). Years 1994 through 2014 all returned zero bars on both feeds. The independent-decade SPY validation (1994–2004 versus 2014–2024) is therefore not feasible through this SDK.

year	IEX bars	IEX first date	SIP bars	SIP first date
2014 and earlier	0	—	0	—
2016	0	—	252	2016-01-04
2018	1	2018-11-01	251	2018-01-02

A.6.46 Cross-Decade Validation: 1994-2004 IS / 2004-2006 OoS via CRSP

Reviewer 3 W6 / RE3 asks for a 1994–2004 vs. 2014–2024 SPY split to test whether the headline stylized-fact reproduction is decade-specific. We executed the validation via a CRSP day-pass: `data/external/crsp_1994_2006.csv` contains daily prices for SPY plus 28 of the 30 cross-ticker panel members (NEE and APD missing from the CRSP query) from 1994-01-03 to 2006-04-28. Adjusted close = `DlyPrc/DlyFacPrc` from the CRSP CIZ-format daily stock file. Source / runner: `run_cross_decade_validation.jl`; output: `results/cross_decade_validation/summary.txt`, `metrics.csv`; paper-side mirror: `results/robustness/cross_decade_validation.csv`.

Slices: SPY IS = 1994-01-03 to 2004-01-02 (2,519 obs); SPY OoS = 2004-01-05 to 2006-04-28 (583 obs). Same body protocol: 1,000 simulated paths per fit, $\alpha = 0.05$ for KS.

Table 52. Cross-decade validation on SPY 1994-2004 IS / 2004-2006 OoS (CRSP daily). Body comparison rows reproduced from Table 3 for the 2014-2024 IS / 2024-2026 OoS window. Observed excess kurtosis on the 1994-2004 IS = 3.05 versus 2014-2024 IS = 7.68; observed excess kurtosis on the 2004-2006 OoS = 0.06 (essentially Gaussian) versus 2024-2026 OoS = 5.29. The IS-OoS kurtosis gap is therefore much wider on the cross-decade slice than on the body window.

Model	K	KS pass rate (%) $\uparrow$		Sim kurtosis		$ G_t $	ACF-MAE $\downarrow$
		IS	OoS	IS	OoS	IS	OoS
<i>Cross-decade: 1994-2004 IS / 2004-2006 OoS (CRSP, this appendix)</i>							
CHMM-N	3	84.9	3.3	2.20	2.24	0.0696	0.0491
CHMM-N	18	89.8	3.4	1.84	1.74	0.0720	0.0499
CHMM-t pen. ($\lambda = 20$)	3	89.1	5.4	4.97	4.18	0.0682	0.0501
<i>Body window: 2014-2024 IS / 2024-2026 OoS (Polygon / Alpaca, Table 3)</i>							
CHMM-N ($K^* = 3$)	3	89.7	80.5	3.83	3.53	0.0460	0.0545
CHMM-N ($K = 18$)	18	94.1	81.8	5.04	4.44	0.0509	~ 0.054
CHMM-t pen. ($\lambda = 20, K^* = 3$)	3	90.6	83.2	14.91	8.50	0.0537	0.0502

Reading. The cross-decade IS fit transfers: CHMM-N at $K = 3$ attains 84.9% IS KS on 1994-2004 vs. 89.7% on 2014-2024, and $K = 18$ attains 89.8% vs. 94.1%. The penalised CHMM-t at $\lambda = 20$, $K = 3$ attains 89.1% IS KS on 1994-2004 vs. 90.6% on 2014-2024. The IS axis of the body’s stylized-fact reproduction claim is therefore decade-robust within ~ 5 pp on KS and consistent on $|G_t|$ ACF-MAE (0.068-0.072 vs. 0.046-0.054, a moderate but not pathological widening).

The OoS axis is the issue. The 2004-2006 OoS slice has excess kurtosis 0.06 (essentially Gaussian) vs. the 1994-2004 IS excess kurtosis of 3.05; the IS-OoS kurtosis gap on this slice is much wider than

on the body 2014-2024/2024-2026 window (where $IS = 7.68$ and $OoS = 5.29$). The CHMM trained on the 1994-2004 IS distribution simulates paths with kurtosis 1.74-4.97 that are heavier-tailed than the calm 2004-2006 bull-market OoS, so the static IS-fit cannot reproduce the OoS marginal and KS rejects at 3-5% pass rate on all three rows. This is the same regime-introduction failure mode the body walk-forward already documents at the W2 (COVID) and W4 (2022 rate-hike onset) stress folds (Table 70, both with OoS KS < 10%); the 2004-2006 OoS window is structurally a low-volatility / low-kurtosis slice where the IS-fixed CHMM’s tail mass overshoots the observed.

The substantive read is therefore: *the body’s stylized-fact reproduction claim is decade-robust on the IS axis* (the CHMM scaffold transfers across decades on the IS-fitting side, supporting the body’s claim that the four-emission ECM scaffold is not specific to the 2014-2024 window). *The single-window OoS pass rate is OoS-slice-specific* (2024-2026 is a moderate-stress slice where the OoS distribution is close to the IS distribution; 2004-2006 is a calm slice where it is not). The body framing already states this through the walk-forward distribution at Section 4 (“the walk-forward median, not the single-window OoS pair, is the operationally informative summary for production deployment”); the cross-decade result is consistent with that framing and adds an independent decade to the walk-forward evidence base. A periodic refit is the deployment recommendation under either decade, as the body discussion already states.

A.7 Robustness, Test Power, Multi-Seed, and Auxiliary Baselines

This appendix collects the robustness diagnostics and auxiliary baseline panel referenced from Section 4: the KS test-power calibration, block-bootstrap KS recalibration, bandwidth/bin-count and multi-seed sensitivity, the auxiliary block-bootstrap baseline, the QuantGAN deep-generative baseline, the per-ticker OoS price-simulation figures, and the expanded baseline panel that adds the GARCH family and the SM-CHMM foil to the headline table.

A.7.1 OoS KS Test-Power Calibration

Table 53 reports the two-sample KS test-power calibration used to anchor the OoS KS pass rates of Section 4.3. For each of 1,000 Monte Carlo replications (seed = 20260420), a simulated series of length $T_{OoS} = 572$ is drawn from each of three reference generators and tested against the observed OoS series using the same $\alpha = 0.05$ threshold as the main-body pass-rate metric.

Table 53. OoS KS pass-rate calibration (1,000 replications, seed = 20260420).

Reference generator	Pass rate (%)
I.i.d. resamples of R_{OoS} at length $T_{OoS} = 572$ (known-correct ceiling)	100.0
I.i.d. resamples of R_{IS} at length $T_{OoS} = 572$ (nearly correct)	90.0
Gaussian($\hat{\mu}_{IS}, \hat{\sigma}_{IS}$) at length T_{IS} (negative control)	0.0

The 90.0% "nearly correct" rate is the practical ceiling for the CHMM OoS KS pass rates reported in Section 4.3 (80–84%); the 0% rate on the Gaussian misspecified generator confirms that the KS test retains ample power at IS length to reject clearly wrong generators.

A.7.2 Block-Bootstrap KS Recalibration and Mean p-value Distribution

Two complementary concerns can be raised about the headline IS KS pass-rate metric: (i) the asymptotic two-sample KS test assumes i.i.d. samples, but the simulated paths from CHMM and GARCH all carry volatility clustering, so the asymptotic null distribution may not apply; and (ii)

Table 54. Block-bootstrap KS recalibration and p-value distribution. All numbers under seed 20260422 with $N_{\text{paths}} = 500$. “asyp pass” is the original Section 4.3 metric (two-sample KS p-value ≥ 0.05). “mean pv” is the mean two-sample KS p-value across paths, with q_5 and q_{95} as the lower / upper p-value quantiles. “blkL%” is the fraction of paths whose KS statistic falls below the stationary-block-bootstrap 95% critical value at mean block length L . Block-bootstrap critical values: 0.0306 ($L = 5$), 0.0338 ($L = 10$), 0.0350 ($L = 20$); asymptotic critical value 0.0383.

Generator	asyp pass% $\uparrow$	mean pv $\uparrow$	pv q_{05}	pv q_{95}	blk5% $\uparrow$	blk10% $\uparrow$	blk20% $\uparrow$
i.i.d. bootstrap	99.6	0.820	0.352	0.999	97.6	99.4	99.6
GARCH(1,1)	26.8	0.048	0.000	0.227	6.8	13.6	17.4
CHMM-N	94.8	0.546	0.047	0.976	81.2	89.4	91.6
CHMM-t	95.8	0.571	0.060	0.987	82.4	89.2	91.6
CHMM-L	94.8	0.498	0.048	0.948	81.0	88.2	90.0
CHMM-GED	95.6	0.547	0.064	0.968	83.2	90.2	93.2

the binary “pass / fail at $\alpha = 0.05$ ” reduces a continuous p-value to a one-bit summary. Table 54 reports two recalibrations addressing each concern.

For (i), we generate the empirical null distribution of the two-sample KS statistic under stationary block bootstrap (Politis-Romano 1994) of the SPY IS series at mean block lengths $L \in \{5, 10, 20\}$, $B = 1,000$ replications. The resulting 95% block-bootstrap critical values are 0.0306, 0.0338, and 0.0350 respectively, all *smaller* than the asymptotic critical value $1.36\sqrt{2/n_{\text{IS}}} = 0.0383$; the block-bootstrap test is therefore stricter than the asymptotic one because it accounts for the within-window dependence in the IS series itself. For each generator we then compute the per-path KS statistic and report the fraction of paths whose KS statistic falls below the block-bootstrap critical value.

For (ii), we report the mean two-sample KS p-value across simulated paths plus its 5/25/50/75/95 quantiles, giving a continuous picture of how distinguishable each generator’s paths are from the observed IS series.

Two patterns emerge that the asymptotic pass-rate metric of the main panel does not surface. First, the cross-generator ordering on the mean p-value scale (bootstrap 0.82, CHMM-t 0.57, CHMM-GED 0.55, CHMM-N 0.55, CHMM-L 0.50, GARCH 0.05) is much more compressed for CHMM-vs-bootstrap than the 94.8–99.6% asyp pass rates suggest: the CHMM family lands at mean p-value ~ 0.5 , the natural value for paths that are statistically indistinguishable from the observed series. Second, GARCH(1,1)’s drop from 26.8% asyp $\rightarrow$ 6.8–17.4% block-aware confirms concern (i): the asymptotic test is mildly lenient on volatility-clustering generators with the wrong marginal, and the block-aware test catches GARCH’s misspecification more sharply. The CHMM family, by contrast, sees only a small block-bootstrap penalty (95% $\rightarrow$ 88–93%); CHMM-GED is the strongest entry on the block-aware metric across all three block lengths (83.2%, 90.2%, 93.2% at $L = 5, 10, 20$).

OoS-anchored block-bootstrap recalibration. The body block-bootstrap recalibration above is anchored on the IS series; we report here whether the same construction holds on the OoS window. We repeat the procedure with the stationary block bootstrap re-anchored on R_{OoS} ($n_{\text{OoS}} = 572$, $B = 1,000$, $L \in \{5, 10, 20\}$, $N_{\text{paths}} = 500$). Block-bootstrap critical values: 0.0577, 0.0629, 0.0647 at $L = 5, 10, 20$ respectively, against asymptotic 0.0804. Table 55 reports the OoS-anchored panel (source: `results/robustness/ks_block_bootstrap_oos.csv`).

The cross-generator ordering on the OoS window matches the IS ordering: i.i.d. bootstrap $>$ CHMM-t $>$ CHMM-N / CHMM-GED $>$ CHMM-L $\gg$ GARCH(1,1). The asyp-vs-block gap is larger on OoS ($\sim 25\text{pp}$ drop for the CHMM family at $L = 20$) than on IS ($\sim 5\text{pp}$), but the

Table 55. OoS-anchored block-bootstrap KS recalibration. Same construction as Table 54 but with the stationary block bootstrap anchored on R_{OoS} ($n = 572$). “asyp pass” is the headline OoS metric (two-sample KS p-value ≥ 0.05); “blkL%” is the OoS block-bootstrap pass rate at mean block length L . The CHMM-vs-GARCH cross-generator ordering is preserved under the block-aware recalibration on both windows; the CHMM family’s OoS block-aware penalty ($\sim 25\text{pp}$ drop from asyp to $L = 20$) is larger than the IS penalty ($\sim 5\text{pp}$), reflecting the smaller OoS window where temporal-clustering structure is harder to disentangle from finite-sample noise.

Generator	asyp pass% $\uparrow$	mean pv $\uparrow$	pv q_{05}	pv q_{95}	blk5% $\uparrow$	blk10% $\uparrow$
i.i.d. bootstrap	90.4	0.437	0.030	0.912	60.8	69.0
GARCH(1,1)	59.2	0.164	0.000	0.645	19.4	26.0
CHMM-N	81.0	0.318	0.010	0.876	41.2	51.2
CHMM-t	84.8	0.355	0.018	0.911	47.4	58.4
CHMM-L	77.4	0.266	0.008	0.744	36.0	45.4
CHMM-GED	81.0	0.327	0.010	0.837	44.4	53.6

qualitative ordering and the CHMM-over-GARCH advantage are preserved across calibrations: under the block-aware OoS test, CHMM-t (61.6% at $L = 20$) sits roughly 30pp above GARCH(1,1) (31.4%). The CHMM-t entry is the strongest CHMM family on the OoS block-aware metric, marginally above CHMM-GED and CHMM-N. The block-bootstrap concern raised in the body text is therefore answered: the CHMM-vs-GARCH cross-generator ordering is window-robust under block-aware recalibration.

A.7.3 Bandwidth, Bin-Count, and Multi-Seed Robustness

Multi-seed Monte Carlo on headline metrics. The paper reports headline statistics under a single global seed. We re-run the CHMM-N / -t / -L / -GED pipeline at $K = 18$ across ten alternative global seeds (base 20260422, increment 100) with $N_{\text{paths}} = 250$ per seed and report mean $\pm$ std for the headline metrics in Table 56. Seed-to-seed variability is small: IS KS std is 1.2–1.8%; OoS KS std is 1.8–3.6%; simulated IS kurtosis std is 0.06 for CHMM-N, 0.05 for CHMM-L, 0.09 for CHMM-GED, and 1.95 for CHMM-t (reflecting the well-known instability of the per-state ν_k ECM at moderate K , already discussed in Section 5). Kupiec $\text{LR}_{\text{uc}}^{1\%}$ has std up to 1.25 for CHMM-N because the breach count moves on a coarse 0–10 grid at $T_{\text{OoS}} = 572$. The single-seed headline numbers in Table 3 sit comfortably inside the multi-seed bands. CHMM-GED has the highest mean IS KS pass rate among the four families (95.9%), consistent with its single-seed headline value.

Table 56. Multi-seed Monte Carlo on headline metrics. Ten global seeds; $N_{\text{paths}} = 250$ per seed (reduced from the headline panel’s 1,000 for tractability across 40 fits, so the seed-to-seed std reported here is an upper bound on the std the headline panel would show at $N_{\text{paths}} = 1,000$). Both ACF-MAE columns parallel Table 3: $|G_t|$ for volatility clustering, G_t for linear autocorrelation.

Model	IS KS %	OoS KS %	sim Kurt	ACF-MAE $ G_t $	ACF-MAE G_t	$\text{LR}_{\text{uc}}^{1\%}$	$\text{LR}_{\text{uc}}^{5\%}$
CHMM-N	94.7 ± 1.2	82.6 ± 1.8	5.03 ± 0.06	0.0467 ± 0.0006	0.0173 ± 0.0001	2.12 ± 1.25	3.83 ± 0.00
CHMM-t	94.6 ± 1.2	85.6 ± 1.9	13.95 ± 1.95	0.0518 ± 0.0003	0.0170 ± 0.0001	0.54 ± 0.15	3.19 ± 0.34
CHMM-L	94.4 ± 1.8	80.8 ± 3.6	6.71 ± 0.05	0.0537 ± 0.0003	0.0171 ± 0.0001	3.26 ± 0.00	2.71 ± 0.59
CHMM-GED	95.9 ± 1.2	83.2 ± 2.7	5.15 ± 0.09	0.0513 ± 0.0003	0.0171 ± 0.0001	2.93 ± 0.71	3.75 ± 0.25

A.7.4 Block-Bootstrap Baseline

Table 57 reports the full seven-metric panel for the stationary block bootstrap at block length $b = 10$ trading days, the temporally-aware non-parametric baseline added to Table 3.

Table 57. Stationary block bootstrap (block length $b = 10$ trading days) on SPY (1,000 paths, seed = 20260420).

Window	KS	AD	Kurt	ACF-MAE	W_1	H	Cov
IS	99.2	96.4	7.07	0.0568	0.092	0.0534	100.0
OoS	87.5	86.4	5.34	0.0512	0.225	0.1570	86.9

A.7.5 QuantGAN Deep-Generative Baseline

Architecture and training details for the QuantGAN deep-generative baseline used in Section 4.3 are summarised here. The construction is a repo-native approximation to the convolutional WGAN of Wiese et al. [9]: a one-dimensional convolutional generator and critic, Wasserstein loss with critic weight clipping, trained on rolling windows of standardised SPY returns and synthesising long paths by stitching generated windows.

- **Generator.** Three 1D convolutional layers ($G_{\text{ch}} = 32$, kernel size 5, ReLU, same-padding), followed by a 1×1 projection to a single output channel. Input is i.i.d. Gaussian noise of latent dimension $L = 8$ at every position of the rolling window.
- **Critic.** Three 1D convolutional layers ($D_{\text{ch}} = 32$, kernel size 5, leaky-ReLU slope 0.2, same-padding), global mean pooling along the time axis, and a single linear scalar output.
- **Loss.** Wasserstein loss with critic weight clipping at ± 0.01 [78]. The critic minimises $\mathbb{E}[D(\tilde{G})] - \mathbb{E}[D(G)]$ on real / fake mini-batches; the generator maximises $\mathbb{E}[D(\tilde{G})]$.
- **Optimiser.** Adam, learning rate 1.0×10^{-4} for both networks, batch size 64.
- **Training schedule.** 15 epochs, 120 outer steps per epoch, 4 critic updates per generator update, training on rolling windows of length $W = 64$ standardised by the IS mean and standard deviation.
- **Synthesis.** Independent draws of length W are stitched end-to-end to produce paths of length T_{IS} or T_{OoS} ; outputs are then re-scaled back to the natural return scale by the IS mean and standard deviation.
- **Implementation.** Julia with Flux.jl [79] and Optimisers.Adam state. Random seed for weight initialisation, training shuffling, and per-path sampling: 20260422, with deterministic per-path sub-seeds.

The trained model is callable through the same `simulate` interface as the CHMM, so the seven-metric panel evaluates QuantGAN on byte-identical windows to the other generators. The QuantGAN row is reported alongside the extended-GARCH and SM-CHMM rows in Table 59; the deep-generative row reaches 0% IS and OoS KS pass rates and undershoots IS kurtosis decisively (2.1 vs. observed 7.7), consistent with the documented difficulty of GAN-based generators on financial-time-series volatility-clustering benchmarks [38, 39].

A.7.6 Per-Ticker OoS Price-Simulation Figures and Full Path-Level Metrics (Unconditional CHMM)

This appendix collects (i) the full per-family path-level probabilistic-forecast metrics for an auxiliary OoS price-simulation diagnostic on the six SPY-member tickers and (ii) the per-ticker price-fan and terminal-distribution figures for NVDA, JNJ, JPM, AAPL, and QQQ. Plotting conventions: blue shaded bands are the 5–95% and 25–75% quantile envelopes of the 500 simulated paths, blue solid line is the simulated median, red solid line is the observed OoS price track. Terminal-histogram panels show the cross-path distribution of $\tilde{P}_T$ at T_{OoS} in gray, with the observed terminal P_T^{obs} as a red vertical line and the simulated median as a blue dashed line. For brevity, the figures show CHMM-N only (the representative family; differences across the four emission families in the price-level metrics are small, see Table 58).

Table 58. Full path-level probabilistic-forecast metrics for OoS equity-price simulation, all (ticker, family) cells ($K = 18$, 500 paths, $T_{\text{OoS}} = 572$; global seed policy in Section 3.5). $\bar{C}_{0.90}$: fraction of OoS steps at which the observed price falls inside the simulated 5–95% band (calibration); $\text{MAPE}_{0.50}$: mean absolute percentage error between observed path and simulated median (sharpness); $\overline{\text{CRPS}}$: horizon-average bias-corrected sample CRPS of Gneiting and Raftery [47]; $\overline{\text{CRPS}}/S_0$: CRPS scaled by S_0 to remove price-level effects. Lower MAPE and CRPS, higher $\bar{C}_{0.90}$, are better. **Headline:** no CHMM emission family dominates on all four metrics across all six tickers; differences between CHMM-N, CHMM-t, CHMM-L, and CHMM-GED live mostly in the return tails and propagate only weakly into aggregated price-level scores. CHMM-GED is the strongest entry on JNJ (lowest CRPS, lowest MAPE) and effectively ties CHMM-t on QQQ; the other four tickers show no clear ordering.

Ticker	Family	$\bar{C}_{0.90}$	$\text{MAPE}_{0.50}$	$\overline{\text{CRPS}}$	$\overline{\text{CRPS}}/S_0$
SPY	CHMM-N	1.000	0.100	39.25	0.0835
	CHMM-t	1.000	0.112	42.66	0.0908
	CHMM-L	1.000	0.102	38.68	0.0823
	CHMM-GED	1.000	0.116	44.54	0.0948
NVDA	CHMM-N	0.698	0.339	30.62	0.6424
	CHMM-t	0.636	0.383	34.77	0.7293
	CHMM-L	0.666	0.411	36.58	0.7674
	CHMM-GED	0.682	0.375	33.26	0.6977
JNJ	CHMM-N	1.000	0.101	12.55	0.0780
	CHMM-t	1.000	0.098	12.62	0.0784
	CHMM-L	1.000	0.101	12.65	0.0786
	CHMM-GED	1.000	0.095	12.21	0.0759
JPM	CHMM-N	1.000	0.195	32.23	0.1881
	CHMM-t	1.000	0.241	40.38	0.2357
	CHMM-L	1.000	0.211	35.25	0.2057
	CHMM-GED	0.993	0.253	43.73	0.2552
AAPL	CHMM-N	1.000	0.108	21.30	0.1156
	CHMM-t	1.000	0.075	18.90	0.1025
	CHMM-L	1.000	0.077	18.82	0.1021
	CHMM-GED	1.000	0.089	19.28	0.1046
QQQ	CHMM-N	0.998	0.073	30.48	0.0763
	CHMM-t	0.995	0.079	31.63	0.0792
	CHMM-L	1.000	0.118	41.12	0.1029
	CHMM-GED	0.997	0.080	31.76	0.0795




Figure 29. NVDA OoS price fan, CHMM-N ($S_0 = 47.67 \rightarrow P_T^{\text{obs}} = 198.07$, a $4.15\times$ rally). The observed path exits the 25–75% interquartile band by mid-horizon but remains inside the 5–95% envelope; horizon-wide coverage under CHMM-N is 0.70 (CHMM-t: 0.64, CHMM-L: 0.67; Table 58), the lone calibration outlier in the panel.






Figure 30. OoS terminal-price distributions under CHMM-N for SPY and NVDA. The right-skew is characteristic of log-returns aggregated to prices; the SPY terminal sits close to the simulated median, while the NVDA terminal sits in the upper-right of the simulated distribution (the realised rally is in the upper tail of the CHMM-N terminal density).




(a) Price fan.




(b) Terminal distribution.

Figure 31. JNJ OoS price fan and terminal distribution, CHMM-N (Pipeline A, $K = 18$, 500 simulated paths, $T_{\text{OoS}} = 572$; plotting conventions as above).




(a) Price fan.




(b) Terminal distribution.

Figure 32. JPM OoS price fan and terminal distribution, CHMM-N (Pipeline A, $K = 18$, 500 simulated paths, $T_{\text{OoS}} = 572$; plotting conventions as above).




(a) Price fan.




(b) Terminal distribution.

Figure 33. AAPL OoS price fan and terminal distribution, CHMM-N (Pipeline A, $K = 18$, 500 simulated paths, $T_{\text{OoS}} = 572$; plotting conventions as above).



(a) Price fan.



(b) Terminal distribution.

Figure 34. QQQ OoS price fan and terminal distribution, CHMM-N (Pipeline A, $K = 18$, 500 simulated paths, $T_{\text{OoS}} = 572$; plotting conventions as above).

A.7.7 Expanded Baseline Panel: GARCH Family, Semi-Markov CHMM, and QuantGAN

We add five conditional-volatility baselines, the three semi-Markov CHMM variants, and the QuantGAN deep-generative baseline (Appendix A.7.5) to the single GARCH(1,1) Gaussian row of Table 3. The conditional-volatility models are fit by grid-initialised Nelder-Mead maximum likelihood on the SPY IS window ($T_{\text{IS}} = 2,516$); simulation follows the same recursion used in estimation. Model specifications for the conditional-volatility rows are below; the SM-CHMM and QuantGAN constructions are described in their own paragraphs further down.

EGARCH(1,1) Gaussian. The exponential GARCH of Nelson [32] evolves $\log \sigma_t^2$ in levels:

$$\log \sigma_t^2 = \omega + \beta \log \sigma_{t-1}^2 + \alpha \left(|z_{t-1}| - \sqrt{2/\pi} \right) + \gamma z_{t-1}, \quad z_t = (r_t - \mu)/\sigma_t, \quad z_t \sim \mathcal{N}(0, 1). \quad (28)$$

The log-variance formulation removes the GARCH positivity constraints, and the leverage term γz_{t-1} captures the asymmetric response to negative returns.

GJR-GARCH(1,1) Gaussian. The threshold GARCH of Glosten et al. [33] / Zakoian [34] augments GARCH(1,1) with an indicator on negative innovations:

$$\sigma_t^2 = \omega + \alpha \varepsilon_{t-1}^2 + \gamma \varepsilon_{t-1}^2 \mathbf{1}\{\varepsilon_{t-1} < 0\} + \beta \sigma_{t-1}^2, \quad (29)$$

with $\varepsilon_t = r_t - \mu$. Stationarity requires $\alpha + \beta + \gamma/2 < 1$; positive γ indicates leverage.

GARCH(1,1) Student-t innovations. A direct tail-aware competitor to CHMM-t: the GARCH(1,1) variance recursion is unchanged, but innovations follow a standardised Student-t with $\nu > 2$ degrees of freedom rescaled to unit variance. The log-likelihood uses the Student-t density:

$$r_t - \mu = \sigma_t \cdot z_t, \quad z_t = \eta_t \sqrt{(\nu - 2)/\nu}, \quad \eta_t \sim t_\nu. \quad (30)$$

The fitted $\hat{\nu} = 6.89$ on SPY IS sits inside the moderate-tail regime the equity-risk literature typically reports.

HAR-RV on daily squared returns. The heterogeneous-autoregressive model of Corsi [37] models realised variance as a mixture of daily, weekly, and monthly components:

$$\log \text{RV}_t = \beta_0 + \beta_d \log \text{RV}_{t-1} + \beta_w \log \overline{\text{RV}}_{t-5:t-1} + \beta_m \log \overline{\text{RV}}_{t-22:t-1} + \eta_t. \quad (31)$$

In the absence of intraday bars we use daily squared demeaned returns as the RV proxy; simulation samples $\eta_t \sim \mathcal{N}(0, \hat{\sigma}_\eta^2)$ and rolls forward the recursion, then draws $r_t - \mu = \sqrt{\text{RV}_t} z_t$ with $z_t \sim \mathcal{N}(0, 1)$.

MS-GARCH $K = 3$. A three-regime extension of the Markov-switching GARCH(1,1) ($G_t = \mu + \sigma_t z_t$ with regime-specific recursion $\sigma_t^2(k) = \omega_k + \alpha_k (G_{t-1} - \mu)^2 + \beta_k \sigma_{t-1}^2(k)$). Parameter layout is three regime triples $(\omega_k, \alpha_k, \beta_k)$ plus six transition-matrix logits (softmax-normalised per row, enforcing a non-trivial diagonal); fit by grid-initialised Nelder-Mead. States are canonicalised by ascending unconditional variance, producing the calm / medium / stress decomposition $\sigma = [0.10, 1.77, 10.04]$ on SPY IS.

MS-GARCH $K = 2$ (Haas-Mittnik-Paoletta 2004). The path-independent two-regime variant of Haas et al. [22]. Regime stickiness is high: $p_{11} = 0.915$ (calm) and $p_{22} = 0.547$ (stress), with unconditional standard deviations 0.97 (calm) and 12.67 (stress); the calm regime is near-GARCH behaviour and the stress regime is a rare, transient high-variance mode. Common $\mu = 0.094$, log-likelihood $-4,965$.

MS-GARCH $K = 4$ and $K = 6$. A natural concern with the $K = 2$ baseline is that it is unfairly low against a $K = 18$ CHMM. We refit the same Haas-Mittnik-Paoletta specification at $K = 4$ (24 parameters: $3K = 12$ GARCH triples plus $K(K - 1) = 12$ transition logits) and $K = 6$ (48 parameters), via Nelder-Mead with diagonal-dominant $\mathbf{T}$ initialisation, geometric variance-ratio spread across regimes, and ascending-variance regime canonicalisation. Higher- K MS-GARCH does not close the gap to the CHMM family on the headline distributional metric: IS KS rises only marginally from 27.7% at $K = 2$ to 37.6% at $K = 4$ and falls back to 34.5% at $K = 6$, against 94.1% for CHMM-N at $K = 18$ (a 56-percentage-point gap). The OoS KS exhibits the same plateau pattern (38.7% at $K = 2$, 38.2% at $K = 4$, 33.4% at $K = 6$). Log-likelihoods on the $T_{IS} = 2,516$ window saturate at the same value within Monte Carlo noise ($-4,965, -4,965, -4,952, -4,956$ for $K = 2, 3, 4, 6$). The MS-GARCH ceiling on this universe is therefore not a low- K issue: the unimodal-innovation constraint (each regime’s emission is a Gaussian rescaled by the regime’s conditional variance) is the binding constraint, and adding regimes does not relax it. The CHMM family decouples this constraint by giving each state its own emission-mixture component, which is the structural advantage at moderate K .

MS-GARCH reference Bayesian re-run via MSGARCH. Following pre-submission peer-review request, we re-run the same Markov-switching GARCH(1,1) specification through the reference MSGARCH R package of Ardia et al. [35] (version 2.51) at $K \in \{2, 3, 4\}$, driven from the Julia harness via RCall.jl. The fit is fully Bayesian (DEMC sampler with the package’s default proper priors, 12,500 MCMC draws, 2,500 burn-in, thin 10, single chain); all 1,000 simulated paths are posterior-predictive (one path per retained posterior draw), so the simulated marginal integrates parameter uncertainty path-by-path. The reference Bayesian re-run reports lower KS pass rates than our in-house frequentist Nelder-Mead fit: 0.0% ($K = 2$), 0.1% ($K = 3$), 0.0% ($K = 4$) IS, and 5.8% ($K = 2$), 5.1% ($K = 3$), 5.3% ($K = 4$) OoS, against the in-house plateau of 27–37% IS / 33–39% OoS. The gap is methodological rather than estimator-quality: the in-house implementation generates all 1,000 paths from one MLE point estimate, so the simulated marginal is tighter; the reference Bayesian implementation samples one set of parameters per path from the posterior, so the simulated marginal is wider and accordingly fewer paths fall within the asymptotic two-sample KS critical band. Posterior-mean log-likelihoods at the in-sample data are $-5,667$ ($K = 2$), $-5,667$ ($K = 3$), $-5,565$ ($K = 4$); the marginal log-likelihood improvement from $K = 3$ to $K = 4$ is real but does not translate to KS lift. Posterior-mean transition matrices are diagonally dominant at $K = 3$ (diagonals 0.97, 0.93, 0.90) and less so at $K = 4$ (diagonals 0.66, 0.39, 0.32, 0.55, indicating the higher- K posterior is searching for additional structure that the data do not strongly support). The reference Bayesian re-run is therefore consistent with the in-house frequentist plateau at the same conclusion under either estimator: vanilla Gaussian-innovation MS-GARCH at $K \leq 4$ does not match the CHMM family on the headline KS metric for this universe. Reproducibility: R 4.6.0, MSGARCH 2.51, all transitive R dependencies pinned via renv (r_msgarch/renv.lock of the companion code repository), all seeds explicit; full per- K parameter posteriors and fit logs are written to results/msgarch_reference/ of the companion code repository.

Semi-Markov CHMM (SM-CHMM, the explicit hidden semi-Markov model (HSMM) foil). The conceptual foil of the present paper, drawing on the explicit-duration HSMM construction of the companion volatility paper [80]: a hidden semi-Markov extension of the CHMM in which the geometric sojourn distribution of the Markov chain is replaced by a per-state explicit-duration family chosen by AIC from {Pareto, negative binomial, geometric}. Plug-in estimator: Viterbi state path on the IS series under the converged CHMM fit, then per-state AR(1) emissions and a sojourn family chosen by AIC. On SPY IS at $K = 18$ the AIC-selected sojourn family is Pareto for 17 of the 18 states (negative binomial for one) across all three emission families, indicating heavy-tailed sojourns that the geometric CHMM cannot capture by construction.

Caveat on the SM-CHMM estimator. The plug-in construction above (single-shot Viterbi decoding under the converged CHMM fit, then AR(1)-residual + AIC-sojourn refit) is not maximum-likelihood for the underlying HSMM: it does not iterate the joint forward-backward over (state, duration) pairs in the spirit of Yu [63]. A full ML HSMM at $K = 18$ would refit the duration-augmented log-likelihood jointly. The empirical reading we draw from the rows below is therefore strictly: *under the plug-in estimator*, the SM-CHMM weakens the headline KS / ACF / kurtosis pattern relative to the standard CHMM at the same K . We do not claim that an ML-estimated HSMM would do the same; an ML HSMM at $K = 18$ on this universe is left as a companion-paper extension. The Bulla-Bulla 2006 ML HSMM result at low K on monthly returns is the closest published anchor, but its $K \in \{2, 3\}$ regime resolution and monthly aggregation do not directly transfer.

Table 59 reports the IS / OoS distributional fidelity, kurtosis, ACF-MAE, and Kupiec breach diagnostics for the extended GARCH family, the SM-CHMM rows, and the QuantGAN deep-generative row, alongside the headline CHMM panel. Three structural observations.

First, none of the extended GARCH variants challenges the CHMM family on the joint (KS, ACF) corner. GARCH(1,1)- t is the strongest extended row on KS (57.3% IS, 80.8% OoS) but its IS excess kurtosis of 15.1 overshoots the observed 7.7 by a factor of two and its ACF-MAE of 0.0316 is below the CHMM’s only because the GARCH variance recursion is hard-wired to track conditional volatility; CHMM at $K = 18$ matches GARCH on ACF-MAE within 0.02 while gaining 36 to 36.5 percentage points on IS KS. EGARCH undershoots kurtosis decisively, GJR-GARCH improves on threshold asymmetry but not on the headline gap, and HAR-RV is degenerate on the daily-RV proxy used here (see methods note above).

Second, the SM-CHMM rows are weaker than the headline CHMM rows on every metric in the panel: IS KS 77–80% vs CHMM 93–94%, OoS KS 63–69% vs CHMM 79–86%. *Under the Viterbi-AR(1) plug-in estimator*, the geometric-sojourn restriction of the standard CHMM is therefore not the bottleneck for the headline KS / ACF / kurtosis story; the spectral mechanism of Section 3.8 carries the comparison. We caveat strongly that this conclusion is conditional on the plug-in estimator (paragraph above): a maximum-likelihood explicit-duration HSMM under a Pareto sojourn family at the same K may close part of the gap, and we do not claim otherwise. An ML HSMM benchmark at $K = 18$ on this universe is left as a companion-paper extension.

Third, the QuantGAN deep-generative row (Appendix A.7.5) is the weakest entry in the panel on every distributional metric. IS and OoS KS pass rates collapse to 0% at $\alpha = 0.05$ across all 1,000 simulated paths, IS and OoS AD pass rates are also 0%, and simulated IS excess kurtosis of 2.1 undershoots the observed 7.7 by more than a factor of three. ACF-MAE of 0.0591 is comparable to the i.i.d. baselines of Table 3 (0.0628), so the GAN does not recover absolute-return autocorrelation either. The convolutional WGAN reproduces some surface-level visual properties of the return series but the stitched-window simulation scheme has no mechanism for modelling either the heavy-tailed marginal or the persistent absolute-return autocorrelation; this matches the documented failure mode of GAN-based generators on financial-time-series volatility-clustering benchmarks [38, 39]. The

Table 59. Extended baseline panel. SPY IS / OoS, $T_{\text{IS}} = 2,516$, $T_{\text{OoS}} = 572$, $N_{\text{paths}} = 1,000$, $\alpha = 0.05$ for KS. Columns: IS / OoS KS pass rate (%); IS / OoS AD pass rate (%); simulated IS excess kurtosis; ACF-MAE on $|G_t|$ (volatility clustering) and on raw G_t (linear autocorrelation, parallel column); pooled-archive Kupiec breach rate (%) and unconditional likelihood-ratio statistic at the 1% and 5% tails. Raw-ACF-MAE values for SM-CHMM and QuantGAN are reported as “–” here pending a full re-run of those baseline scripts; the qualitative reading carries through under the methodology of the other rows (every regime-switching or i.i.d.-emission generator produces a near-zero raw ACF in expectation).

Model	KS (%) $\uparrow$		AD (%) $\uparrow$		Kurt	ACF-MAE		1% tail		5% tail	
	IS	OoS	IS	OoS		$ G_t $	G_t	br%	LR _{uc}	br%	LR _{uc}
GARCH(1,1) Gaussian	28.0	59.4	12.5	59.5	7.0	0.0309	0.0173	0.9	0.10	4.0	1.23
EGARCH(1,1)	6.3	34.9	4.1	41.7	2.8	0.0430	0.0173	1.0	0.01	3.5	3.03
GJR-GARCH(1,1)	40.1	57.7	28.7	61.4	7.6	0.0330	0.0173	1.2	0.27	4.7	0.10
GARCH(1,1)- t	57.3	80.8	36.6	65.3	15.1	0.0316	0.0173	1.0	0.01	4.9	0.01
HAR-RV	0.0	0.0	0.0	0.0	189	0.0566	0.0173	0.2	5.99	3.5	3.03
MS-GARCH $K = 2$	27.7	38.7	24.0	44.4	4.7	0.0367	0.0173	1.0	0.01	4.2	0.82
MS-GARCH $K = 3$	36.1	33.1	28.8	38.0	4.1	0.0284	0.0173	1.0	0.01	4.2	0.82
MS-GARCH $K = 4$	37.6	38.2	–	–	3.8	0.0437	–	–	–	–	–
MS-GARCH $K = 6$	34.5	33.4	–	–	4.4	0.0429	–	–	–	–	–
MS-GARCH ref. B. $K = 2$	0.0	5.8	–	–	4.0	0.0465	0.0240	–	–	–	–
MS-GARCH ref. B. $K = 3$	0.1	5.1	–	–	4.5	0.0433	0.0241	–	–	–	–
MS-GARCH ref. B. $K = 4$	0.0	5.3	–	–	6.0	0.0446	0.0241	–	–	–	–
CHMM-N ($K = 18$)	94.1	81.8	36.6	73.0	5.0	0.0509	0.0237	0.7	1.58	4.0	3.83
CHMM-t ($K = 18$)	95.6	85.7	33.4	69.8	14.4	0.0549	0.0234	1.2	0.58	5.0	3.83
CHMM-L ($K = 18$)	93.6	80.8	30.1	65.4	6.6	0.0567	0.0234	1.4	3.26	5.6	3.03
CHMM-GED ($K = 18$)	95.2	84.3	34.1	71.0	5.2	0.0548	0.0234	0.9	0.58	4.4	3.83
SM-CHMM-N ($K = 18$)	79.2	66.7	19.5	57.1	4.7	0.0598	–	1.1	0.01	5.4	0.82
SM-CHMM-t ($K = 18$)	77.1	63.4	17.2	52.3	13.9	0.0625	–	1.4	0.10	5.7	1.74
SM-CHMM-L ($K = 18$)	80.4	69.1	21.0	58.9	6.4	0.0581	–	1.4	3.26	5.5	3.03
QuantGAN	0.0	0.0	0.0	0.0	2.1	0.0591	–	1.0	0.01	3.3	3.83

runner is a repo-native convolutional WGAN with critic weight clipping (Appendix A.7.5), materially smaller than the original temporal convolutional network of Wiese et al. [9] (a seven-block TCN with 127-day receptive field, hidden width 80, plus Lambert-W input pre-processing that the original authors credit with most of the tail-fidelity); the row therefore reads as the panel’s deep-generative negative control rather than a verdict on the WGAN literature as a whole. A deeper-architecture reference rebuild (Appendix A.7.11) under the same training protocol but with 5 + 5 conv layers and hidden width 48 is reported as a robustness check on whether deeper architecture alone closes the gap. The QuantGAN row complements the i.i.d. Gaussian negative control of the headline Table 3.

A.7.8 Stochastic-Volatility, Multifractal, and Jump-Diffusion Baselines

Peer-review item 2 asks for the canonical state-space and jump baselines that the body Table 3 does not exercise: a stochastic-volatility (SV-AR(1)) row in the Taylor 1986 / Harvey-Ruiz-Shephard 1994 lognormal log-AR(1) tradition; a Markov-switching multifractal (MSM, Calvet-Fisher 2004) row at $\bar{k} = 8$ multipliers; and a Merton 1976 Poisson-Gaussian jump-diffusion row. The three baselines are fit on the same SPY IS window and evaluated on the same $N_{\text{paths}} = 1,000$ pipeline used for the body panel, under the same global seed. Implementation lives in `src/SVMSMBaselines.jl` of the companion `CHMM-Model.jl` package; results in `results/sv_msm_jd/sv_msm_jd_baselines.txt`.

Table 60. Stochastic-volatility / multifractal / jump-diffusion baselines on SPY. Same IS / OoS windows, $N_{\text{paths}} = 1,000$, seed = 20260420 as Table 3. Observed kurtosis: 7.68 IS, 5.29 OoS.

Model	KS (%) $\uparrow$		Kurtosis		ACF-MAE $\downarrow$	
	IS	OoS	IS	OoS	$ G_t $	G_t
SV-AR(1)	38.2	35.3	7.52	5.26	0.0466	0.0239
MSM ($\bar{k} = 8$)	0.0	7.2	0.62	0.60	0.0605	0.0235
Merton-JD	98.3	91.1	3.12	2.98	0.0627	0.0236

Reading. *SV-AR(1)* reproduces the IS kurtosis target (7.52 vs. observed 7.68) almost exactly and matches CHMM-N’s $|G_t|$ ACF-MAE within 0.0003 (0.0466 vs 0.0467 at $K^* = 3$); its KS pass rate is poor (38.2% IS / 35.3% OoS) because the Gaussian-conditional-on-volatility marginal cannot reproduce the bimodal-bulk / heavy-tail empirical density that the regime-mixture CHMM marginals can. *Merton-JD* is the inverse trade-off: KS at 98.3% IS / 91.1% OoS (the highest OoS KS in the panel except for the i.i.d. bootstrap), but kurtosis 3.12 undershoots 7.68 by half and the absolute-return ACF-MAE regresses to the i.i.d. baseline level (0.0627, the same level as the bootstrap row of Table 3) because the constant-volatility diffusion plus i.i.d. jumps carries no temporal persistence. *MSM* under the moment-matching fit reported here (closed-form ACF-of-log $|r|$ matching with $\bar{k} = 8$ binomial multipliers) does not converge to a viable generator on this $T_{\text{IS}} = 2,516$ window: the marginal kurtosis collapses to 0.62 and KS drops to 0%. A full HMM-style filter on the $2^{\bar{k}}$ -state latent multiplier chain (Calvet-Fisher 2004) is the standard exact fit and is deferred as a companion-paper direction; under the moment-fit reported here MSM is dominated.

The headline conclusion of Table 3 survives the addition of these three rows: no single competitor occupies the joint-fit corner that the CHMM family at $K = 18$ does (KS $\geq 80\%$, kurtosis within 1 unit of observed, $|G_t|$ ACF-MAE ≤ 0.06). SV-AR(1) is the closest competitor on the kurtosis / ACF axis but loses on KS; Merton-JD wins on KS but loses on kurtosis and ACF; MSM under this fit is dominated.

A.7.9 Leverage-Effect Diagnostic

This appendix reports the Cont (2001) leverage-effect column: the lag-1 correlation $\text{Corr}(G_t, |G_{t+1}|)$ between today’s signed return and tomorrow’s absolute return, which is negative on equity returns. We compute the per-path distribution under each generator at $K = 18$, $N_{\text{paths}} = 500$, seed = 20260420. Source: `results/diagnostics/leverage_effect/`; CSV at `results/robustness/leverage_effect.csv`.

Reading. Three patterns emerge. First, the CHMM family at $K = 18$ does produce a non-trivial negative leverage signal at the per-path level despite the per-state emissions being symmetric (Gaussian, Laplace, Student- t , GED). The mechanism is the converged $\hat{\mathbf{T}}$ asymmetrically coupling high- σ_k states to negative- μ_k states: a sample drawn from a high-volatility state with negative location mean is large in magnitude and negative in sign, and the sticky transition keeps the chain in a high- σ state at the next step, so $|G_{t+1}|$ is also large; this gives $\text{Corr}(G_t, |G_{t+1}|) < 0$ via Markov state-mixing alone. The body text expectation that “CHMM cannot reproduce leverage by construction” is empirically incorrect: state-mixing carries $\sim 65\%$ of the IS observed leverage magnitude. Second, CHMM-N at $K = 18$ produces the strongest leverage in the CHMM family (-0.089 IS median), with the IS $[Q_5, Q_{95}] = [-0.138, -0.037]$ bracketing the IS observed value of -0.135 at the lower boundary. CHMM-GED is the second-strongest CHMM entry. Third, the OoS

Table 61. Leverage effect $\text{Corr}(G_t, |G_{t+1}|)$, **per-path distribution.** Observed values: SPY IS -0.135 , OoS -0.214 . Per-generator median plus $[Q_5, Q_{95}]$ across $N_{\text{paths}} = 500$. A generator captures the leverage effect when its $[Q_5, Q_{95}]$ envelope brackets the observed value. Bold marks generators whose $[Q_5, Q_{95}]$ envelope brackets the observed IS value; the OoS observed value of -0.214 is more negative than every generator’s Q_5 in the panel, indicating the OoS regime carries leverage stronger than any generator (including asymmetric GARCH) reproduces under static IS-fitted parameters.

Generator	IS (observed -0.135)		OoS (observed -0.214)	
	median	$[Q_5, Q_{95}]$	median	$[Q_5, Q_{95}]$
i.i.d. bootstrap	$+0.000$	$[-0.034, +0.030]$	$+0.004$	$[-0.069, +0.068]$
GARCH(1,1) Gaussian	-0.004	$[-0.066, +0.064]$	-0.003	$[-0.112, +0.107]$
EGARCH(1,1)	-0.093	$[-0.147, -0.047]$	-0.087	$[-0.199, -0.002]$
GJR-GARCH(1,1)	-0.079	$[-0.153, -0.017]$	-0.073	$[-0.174, +0.023]$
CHMM-N ($K = 18$)	-0.089	$[-0.138, -0.037]$	-0.087	$[-0.190, +0.015]$
CHMM-t ($K = 18$)	-0.067	$[-0.116, -0.020]$	-0.067	$[-0.166, +0.019]$
CHMM-L ($K = 18$)	-0.064	$[-0.108, -0.022]$	-0.064	$[-0.152, +0.026]$
CHMM-GED ($K = 18$)	-0.072	$[-0.113, -0.029]$	-0.076	$[-0.163, +0.025]$

observed leverage of -0.214 exceeds every generator’s Q_5 : under static IS-fitted parameters, no generator in the panel (including the asymmetric GARCH baselines) brackets the OoS leverage in its $[Q_5, Q_{95}]$ envelope. The OoS regime carries stronger leverage than IS, and closing the residual gap to the OoS observed magnitude requires either skew-emission CHMM extensions (skew- t , skew-Laplace, skew-GED) or a refit to the OoS distribution; both are deferred as companion-paper directions.

A.7.10 Full One-Shot Student- t Copula MLE

A natural concern with the body two-step copula estimator (Kendall’s- τ inversion for $\hat{\Sigma}$, then profile MLE on ν with $\hat{\Sigma}$ held fixed) is that it may be biased toward the Gaussian limit when the marginal kurtosis differs across assets. We address this with a coordinate-ascent one-shot MLE that jointly maximises the Student- t copula log-likelihood over (Σ, ν) on the same six-asset US-equity universe (SPY, NVDA, JNJ, JPM, AAPL, QQQ; $T_{\text{IS}} = 2,516$, seed 20260420). The procedure alternates: (i) bracketed golden-section maximisation of ν at the current Σ , (ii) method-of-moments refit of Σ at the current ν on the t -quantile-transformed pseudo-uniform sample with PSD projection. Iteration terminates at $|\Delta \log L| < 10^{-3}$.

Table 62. Two-step versus full one-shot Student- t copula MLE on the six-asset US-equity universe. The off-diagonal correlation differences are reported as the maximum absolute difference $\max_{ij} |\Delta \rho_{ij}|$ across the $\binom{6}{2} = 15$ pairs and the median absolute difference; both are sub-0.025, so the body $\nu^* = 6$ choice on $\hat{\Sigma}$ obtained by Kendall’s- τ inversion is robust to the estimator switch. Source: results/copula_profile_ci/full_tcopula_mle.txt.

Estimator	$\hat{\nu}$	log-likelihood
Two-step (body construction)	6.00	6157.47
Full one-shot MLE	6.40	6163.02
log-L improvement		+5.55
max $ \Delta \rho_{ij} $ across 15 pairs		0.024
median $ \Delta \rho_{ij} $ across 15 pairs		0.011
$\hat{\nu}_{\text{full}}$ inside body Wilks 95% CI $[6, 7]$?		Yes

Reading. The full MLE moves $\hat{\nu}$ from 6.00 to 6.40, a change of 0.40 units, and improves the log-likelihood by +5.55 units; both off-diagonal-correlation differences are below $|\Delta\rho_{ij}| = 0.025$ across all 15 pairs. The two-step estimator is therefore not detectably biased toward the Gaussian limit on this universe: the full MLE selects a slightly smaller $\hat{\nu}$ than the two-step (more weight on tails, opposite of the Gaussian-bias direction R2 raised), and the body $\nu^* = 6$ choice remains inside the full-MLE Wilks 95% profile-LL CI. The robustness conclusion: the body cross-asset construction can be read with the two-step or the one-shot estimator interchangeably on this six-asset universe; the choice does not move the dependence-family selection or the cross-asset correlation reproduction in any operationally meaningful direction.

A.7.11 QuantGAN Deeper-Architecture Reference Rebuild

The body QuantGAN row (Appendix A.7.5) is a 3-conv-layer convolutional WGAN that is materially smaller than the seven-block dilated TCN of Wiese et al. [9]. A fully faithful Wiese reproduction (7 dilated blocks, hidden 80, receptive field 127 days, plus Lambert-W input pre-processing) is substantial new model engineering with significant AD-compilation overhead in Julia Flux on CPU; we report a meaningful intermediate point as a reference comparator. The runner (`run_quantgan_tcn.jl`) trains a deeper plain-conv WGAN with the same WGAN-with-weight-clipping protocol as the body baseline: 5 generator and 5 critic conv layers (vs. 3 + 3 in the body baseline), hidden width 48 (vs. 32 in the body baseline), kernel size 3, weight-clip ± 0.01 , 20 epochs, 80 steps per epoch, 4 critic steps per generator step. Generator parameters 29,089 vs. approximately 9,000 for the body baseline. We do not implement dilation or Lambert-W input pre-processing.

Table 63. QuantGAN deeper-architecture reference rebuild. Same WGAN-with-weight-clipping training protocol as the body baseline (Appendix A.7.5) but with 5 + 5 conv layers and hidden width 48 rather than 3 + 3 and width 32. KS pass rate at $\alpha = 0.05$, $N_{\text{paths}} = 500$, seed 20260422. Source: `results/quantgan_tcn/quantgan_tcn_metrics.txt`; CSV: `results/robustness/quantgan_tcn.csv`. The body 3-conv-layer baseline row is reproduced from Table 59 for comparison.

Architecture	KS pass rate (%) $\uparrow$		Sim kurtosis (obs 7.7 IS, 5.3 OoS)		ACF-MAE IS
	IS	OoS	IS	OoS	
Body baseline (3 conv layers, hidden 32)	0.0	0.0	2.05	—	0.0591
Reference rebuild (5 conv layers, hidden 48)	0.0	0.0	0.56	0.53	0.0617

Reading. The deeper architecture does not close the gap. KS pass rate remains at 0.0% on both windows, and simulated kurtosis actually *degrades* from 2.05 (body baseline) to 0.56 (reference rebuild), moving further from the observed 7.7 rather than closer. The $|G_t|$ ACF-MAE is essentially unchanged (0.0617 vs. 0.0591), at the same level as the i.i.d. baselines (Table 3). The substantive reading: under the WGAN-with-weight-clipping training protocol on this dataset, increasing depth and width without adding dilation or Lambert-W input pre-processing does not move the headline distributional metrics; if anything the deeper critic is more easily saturated under the ± 0.01 weight clip, which compresses the gradient signal and pushes the generator toward narrower output distributions. The body framing of the QuantGAN row as a deep-generative *negative control* rather than a competitive benchmark is therefore reinforced empirically: the gap is not closed by a $\sim 3\times$ parameter-count increase under the same training protocol.

A.7.12 QuantGAN with Lambert-W Input Pre-Processing

Closes peer-review item P2.9 / R3 RE1’s specific concern that “the [body] TCN rebuild does not include the Lambert-W transform; without it, what is being compared to what?”. We add the Goerg [81, 82] Lambert-W $\times$ Gaussian heavy-tail transformation as input pre-processing on top of the same architecture as Appendix A.7.11 (5 + 5 conv layers, hidden width 48, weight-clip ± 0.01 , 20 epochs, 80 steps per epoch, 4 critic steps per generator step). Forward $z = \text{sign}(y)\sqrt{W(\delta y^2)/\delta}$ heavy-tail $\rightarrow$ Gaussian on the IS series before training; inverse $y = z \exp(\delta z^2/2)$ Gaussian $\rightarrow$ heavy-tail on simulated paths after synthesis. The shape parameter δ is fit on the IS series by IGMM-order-4 (bracket-bisection on raw kurtosis = 3). Source: `results/quantgan_tcn_lambertw/quantgan_tcn_lambertw_metrics.txt`; CSV: `results/robustness/quantgan_tcn_lambertw.csv`; runner: `run_quantgan_tcn_lambertw.jl`.

Table 64. QuantGAN TCN with Lambert-W input pre-processing (P2.9 / R3 RE1). Same architecture and training protocol as Appendix A.7.11 but with the Goerg Lambert-W heavy-tail transform applied forward (heavy-tail to Gaussian) on the IS series before training and inverse (Gaussian to heavy-tail) on the simulated paths after synthesis. $\hat{\delta} = 0.1016$ fit on IS via IGMM-order-4; the post-Lambert-W IS raw kurtosis is 3.000 (target Gaussian).

	KS pass rate (%) $\uparrow$		Sim kurtosis (obs 7.7 IS, 5.3 OoS)		ACF-MAE $ G_t $ $\downarrow$
	IS	OoS	IS	OoS	IS
Pre-processing					
None (Appendix A.7.11)	0.0	0.0	0.56	0.53	0.0617
Lambert-W ($\hat{\delta} = 0.10$)	0.0	0.0	0.40	0.36	0.0624

Reading. The Lambert-W transform succeeds at the variance-stabilising step it is designed for: pre-Lambert-W IS raw kurtosis 7.68 is mapped to post-Lambert-W IS raw kurtosis 3.000, exactly the Gaussian target. The transform itself is therefore not the failing component. The failing component is the WGAN-with-weight-clipping training equilibrium on this dataset and architecture: the WGAN converges to a near-trivial Wasserstein-distance estimate (training losses settle at $|D|, |G| < 0.01$ from epoch ~ 12 onwards) where the generator effectively reproduces a centred Gaussian under the ± 0.01 weight-clip restriction. The inverse-Lambert-W mapping then partially restores tail mass on the synthesised paths, but not enough to recover the observed kurtosis: simulated IS kurtosis is 0.40 vs. observed 7.68, $|G_t|$ ACF-MAE is at the i.i.d. baseline level 0.0624, and KS pass rate remains 0.0%. The substantive read is therefore: *the body’s deep-generative negative-control framing is robust to Lambert-W input pre-processing under the WGAN-with-weight-clipping training protocol*. Closing the gap requires structurally different training (WGAN-GP, signature-matched losses, or diffusion) rather than richer pre-processing; the WGAN training collapse on this dataset is the binding mechanism, not the input-distribution heaviness.

A.8 Cross-Asset Dependence: Full Pipeline B Panel

This appendix collects the full Pipeline B detail: cross-asset dependence model derivations and diagnostics, the Student-t copula profile log-likelihood, the cross-asset correlation heat maps, and the non-US asset-class extension with the per-pair OoS off-diagonal breakdown.

A.8.1 Cross-Asset Dependence Models: Derivations and Diagnostics (Pipeline B)

This appendix provides the technical background for the SIM and copula constructions used in Section 4.5 of the main text. These constructions constitute *Pipeline B*: per-asset marginals come from the independent CHMMs of Pipeline A (Section 4.4) and joint dependence is overlaid on top. Pipeline A stays purely univariate; Pipeline B is the sole multi-asset pipeline in the paper.

Sklar’s theorem and the rank-reordering simulator. Sklar’s theorem [41, 42] states that any d -dimensional joint CDF $H(g_1, \dots, g_d)$ with continuous marginals $F_1, \dots, F_d$ can be written as $H = C(F_1, \dots, F_d)$ for a unique copula C . The rank-reordering simulator exploits this uniqueness. If $\tilde{\mathbf{G}}$ is a matrix whose j -th column is an i.i.d. sample from F_j , and $\mathbf{U}$ is an independent sample from C , then reordering each column of $\tilde{\mathbf{G}}$ so that its ranks match $\mathbf{U}$ ’s produces a sample from H , up to discrete approximation error in the empirical ranks. Crucially, reordering never creates new values: each column of the output is a permutation of the corresponding column of $\tilde{\mathbf{G}}$, so marginal distributions are preserved exactly [45].

Kendall’s τ inversion. Given IS returns $\mathbf{R} \in \mathbb{R}^{T \times d}$, we estimate the pairwise Kendall’s τ_{ij} from concordant/discordant pair counts and invert to the correlation scale via the analytic relationship $\rho_{ij} = \sin(\pi\tau_{ij}/2)$, which holds exactly for the Gaussian copula family and is commonly used as a robust correlation estimator for the Student-t copula [44, 55]. The resulting matrix $\hat{\Sigma}$ is symmetric but not guaranteed positive semi-definite; we project to the nearest PSD matrix by clipping negative eigenvalues to 10^{-8} and rescaling the diagonal to unity.

Profile MLE for ν . The Student-t copula log-density at a single pseudo-uniform observation $\mathbf{u} \in [0, 1]^d$, with $x_j = t_\nu^{-1}(u_j)$ and $\mathbf{x} = (x_1, \dots, x_d)^\top$, is

$$\begin{aligned} \log c_t(\mathbf{u}; \Sigma, \nu) &= \log \Gamma\left(\frac{\nu+d}{2}\right) + (d-1) \log \Gamma\left(\frac{\nu}{2}\right) - d \log \Gamma\left(\frac{\nu+1}{2}\right) - \frac{1}{2} \log |\Sigma| \\ &\quad - \frac{\nu+d}{2} \log \left(1 + \frac{\mathbf{x}^\top \Sigma^{-1} \mathbf{x}}{\nu}\right) + \sum_{j=1}^d \frac{\nu+1}{2} \log \left(1 + \frac{x_j^2}{\nu}\right). \end{aligned} \quad (32)$$

Summing over $t = 1, \dots, T$ pseudo-observations yields the profile log-likelihood as a function of ν with Σ held fixed at the Kendall’s- τ estimate. We evaluate equation (32) on the discrete grid $\nu \in \{2, 3, 4, 5, 6, 8, 10, 15, 20, 30\}$ and select $\hat{\nu}$ as the grid point with the largest total log-likelihood. In Section 4.5 the selected value was $\hat{\nu} = 6$ on the six-asset universe.

Sampling the copula. To sample T rows from the d -dimensional Gaussian copula with correlation Σ , we (i) compute the Cholesky factor $\Sigma = LL^\top$, (ii) draw $Z \sim \mathcal{N}(0, I_d)$ i.i.d. and form $Y = LZ$, and (iii) return $U = \Phi(Y)$ componentwise. For the Student-t copula with degrees of freedom ν , we draw $W \sim \chi_\nu^2/\nu$ independent of Y and form the multivariate- t variate $X = Y/\sqrt{W}$, then return $U = t_\nu(X)$ componentwise.

SIM regression summary. Table 65 reports the fitted SIM regression coefficients and goodness-of-fit statistics for the non-market assets in Section 4.5. QQQ’s high $R^2 = 0.840$ reflects that the Nasdaq-100 ETF is essentially a linear combination of large-cap technology names that also drive the S&P 500 index, while JNJ’s lower $R^2 = 0.260$ reflects the limited systematic exposure of low-beta healthcare stocks. These factor loadings and residual magnitudes are what the SIM simulation re-injects at simulation time, and they explain the uneven per-asset KS pass rates reported in Table 74.

Table 65. SIM regression $\hat{G}_{j,t} = \hat{\alpha}_j + \hat{\beta}_j G_{\text{SPY},t} + \hat{\eta}_{j,t}$ for the five non-market assets, fitted on the 2014-01-03 through 2024-01-03 IS window ($T = 2,516$).

Ticker	$\hat{\alpha}$	$\hat{\beta}$	R^2
NVDA	0.321	1.694	0.371
JNJ	0.002	0.576	0.260
JPM	-0.008	1.223	0.507
AAPL	0.111	1.205	0.492
QQQ	0.047	1.122	0.840

A.8.2 Student-t Copula Profile Log-Likelihood (Pipeline B)

Figure 35 plots the profile log-likelihood of the Student-t copula used inside Pipeline B, computed on the six-asset SPY cross-section over the grid $\nu \in \{2, 3, 4, 5, 6, 8, 10, 15, 20, 30\}$ used in Section 4.5. The curve is smooth and single-peaked, with $\nu^* = 6$ attaining the maximum profile log-likelihood (6157.47), and the neighbouring grid points $\nu = 5$ (6143.37) and $\nu = 8$ (6148.37) each within 15 profile-log-likelihood units of the peak, indicating a shallow-but-clean optimum that agrees with the 4–10 range typically reported for equity-portfolio Student-t copulas in the risk-management literature. The AIC penalty for the additional ν parameter is 2 (one degree of freedom relative to the Gaussian copula limit $\nu \rightarrow \infty$), and the AIC differential between the Student-t at $\nu^* = 6$ and the Gaussian copula limit on these data is $-2 \cdot (6157.47 - 6143.37) - 2 = -30.2$, decisively in favour of the Student-t copula on IS; the same $\nu^* = 6$ is selected by AIC as by profile MLE on this grid.




Figure 35. Profile log-likelihood of the Student-t copula used inside Pipeline B on the six-asset SPY cross-section (SPY, NVDA, JNJ, JPM, AAPL, QQQ; IS window 2014-01-03 through 2024-01-03). The vertical axis is the profile log-likelihood evaluated at each degrees-of-freedom value ν on the grid $\{2, 3, 4, 5, 6, 8, 10, 15, 20, 30\}$; the dashed red line marks the profile MLE $\nu^* = 6$ (peak value 6157.47). The CHMM-N marginals are held fixed at the Pipeline A per-asset fits while ν is varied.

A.8.3 Cross-Asset Correlation Heat Maps (Pipeline B, Companion to Table 74)

Figure 36 visualises the observed and path-averaged simulated correlation matrices for the dependence models of Section 4.5. The panel ordering is SIM, Gaussian copula, Student- t copula, and a truncated C-vine variant with SPY as root, each reported quantitatively in Table 74 and (for the C-vine) Section A.8.4. All five panels share the six-asset SPY cross-section and the same CHMM-N marginals at $K = 18$, so any visible differences between the simulated panels reflect only the dependence construction.




Figure 36. Cross-asset IS correlation reproduction across dependence models (Pipeline B, companion to Table 74). Panel (a): *observed* correlation matrix on 2014-01-03 through 2024-01-03 daily excess log returns for SPY, NVDA, JNJ, JPM, AAPL, QQQ (IS window, $T_{IS} = 2,516$). Panel (b): path-averaged correlation under the Single Index Model with SPY as market factor. Panel (c): path-averaged correlation under the Gaussian copula on CHMM marginals. Panel (d): path-averaged correlation under the Student- t copula on CHMM marginals, $\nu^* = 6$. Panel (e): path-averaged correlation under a truncated C-vine copula with SPY as root. Each simulated panel averages over 200 paths of length T_{IS} . All marginals are the per-asset CHMM-N fits at $K = 18$ from Table 7; color scale in $[-1, 1]$, red-white-blue (RdBu) diverging. Closer visual match to panel (a) is better.

A.8.4 Truncated C-vine Copula on CHMM Marginals (Pipeline B)

The truncated C-vine variant in panel (e) of Figure 36 uses SPY as the root and selects each pair-copula family edge-wise from {Gaussian, Student- t , Clayton, Gumbel} by Akaike Information Criterion (AIC); on the six-asset universe every selected pair is Student- t , with $\hat{\rho}$ and $\hat{\nu}$ as in Table 66. On the same CHMM-N marginals at $K = 18$ and 200 paths under seed 20260422, the truncated C-vine attains an off-diagonal MAE of 0.068 on IS and 0.235 on OoS (Table 74), strictly above both

elliptical copulas on the IS column and on OoS off-diag MAE, and only modestly below the SIM baseline. The reason is the level-1 truncation: a root-centred C-vine models only the SPY-vs-asset bivariate copulas and treats the remaining pairs as conditionally independent given SPY, which is a stronger restriction than either elliptical copula and only a small relaxation of the SIM factor structure. Untruncated regular vines or factor copulas are the natural way to recover non-root cross-pair structure at larger d , but at $d = 6$ with $T_{\text{IS}} = 2,516$ the additional flexibility does not justify the extra estimation cost.

Table 66. Truncated C-vine edge summary, Pipeline B with SPY as root on the six-asset universe. Pair family selected edge-wise by AIC over {Gaussian, Student-t, Clayton, Gumbel}; on this universe every selected pair is Student-t. $\hat{\rho}$ is the Kendall’s- τ -inverted correlation parameter and $\hat{\nu}$ is the per-edge degrees-of-freedom estimate.

Child	Family	$\hat{\rho}$	$\hat{\nu}$
NVDA	Student-t	0.637	5
JNJ	Student-t	0.483	4
JPM	Student-t	0.686	4
AAPL	Student-t	0.698	4
QQQ	Student-t	0.911	3

A.8.5 Quarterly Rolling-Window Copula Refit on the OoS Window

The body Section 4.5 reports a static-IS-fit OoS off-diagonal MAE of 0.209, an order of magnitude larger than the IS value of 0.027. Section 5 attributes this to the stationarity-scope limit and recommends periodic refit. Reviewer Question Q4 asked for a concrete rolling-window-refit number on the same six-asset universe.

We refit the Pipeline-B Student-t copula on a 252-day rolling window (one trading year) sliding by 63 days (one trading quarter) across the OoS span, holding the per-asset CHMM-N marginals fixed at the IS fits used for Table 7. Each refit yields a Kendall’s- τ correlation matrix and a profile-MLE-selected ν^* on the discrete grid $\{2, 3, 4, 5, 6, 8, 10, 15, 20, 30\}$; we then simulate forward 63 days under the new copula and compare the path-averaged simulated correlation matrix to the realised next-quarter correlation. Table 67 reports the per-quarter detail.

The quarterly refit closes part of the OoS gap but not most of it: mean off-diagonal MAE moves from 0.207 static to 0.185 rolling ($\sim 10\%$ relative improvement). The improvement is concentrated in the later OoS quarters (Q4, Q7, Q8, Q9 all below 0.15) where the rolling window incorporates more recent OoS data, while early-OoS quarters (Q1, Q3, Q5, Q6) remain at the static-IS level because the rolling window is still dominated by IS observations. The refitted ν^* sequence (15, 15, 10, 6, 6, 6, 5, 5, 6) tracks the gradual shift from the equity-cluster IS regime toward the OoS regime, supporting the operational recommendation in Section 5 of quarterly refit at deployment. The residual gap (mean MAE ≈ 0.185 versus IS 0.027) is consistent with the per-pair JNJ-equity-factor regime shift documented in Table 68: even a perfectly-refitted copula on the historical 252 days cannot anticipate the regime shift at the next quarter’s start, only track it once it has unfolded.

Table 67. Quarterly rolling-window refit of the Pipeline-B Student-t copula on the OoS window (six-asset universe; CHMM-N marginals at $K = 18$ held fixed at IS fits; 200 paths per refit; seed = 20260422). Per-quarter window slice and refitted ν^* , plus the off-diagonal MAE between the path-averaged simulated correlation matrix under the refitted copula and the realised next-quarter correlation matrix. Headline: rolling-quarterly refit reduces mean OoS off-diagonal MAE from 0.207 (static-IS-fit baseline) to 0.185, a $\sim 10\%$ relative improvement; the gain concentrates in the later OoS quarters where the rolling window incorporates more recent OoS data and the refitted ν^* shifts from the IS-fit value of 6 down to 5 or up to 10–15 as the empirical tail-coupling rolls.

Quarter	Window slice	Forecast horizon	Refit ν^*	Off-diag MAE
1	[2265, 2516]	[2517, 2579]	15	0.198
2	[2328, 2579]	[2580, 2642]	15	0.205
3	[2391, 2642]	[2643, 2705]	10	0.241
4	[2454, 2705]	[2706, 2768]	6	0.133
5	[2517, 2768]	[2769, 2831]	6	0.255
6	[2580, 2831]	[2832, 2894]	6	0.232
7	[2643, 2894]	[2895, 2957]	5	0.147
8	[2706, 2957]	[2958, 3020]	5	0.119
9	[2769, 3020]	[3021, 3083]	6	0.138
<i>mean</i>	—	—	—	0.185
<i>median</i>	—	—	—	0.198
<i>Static IS-fit</i>	[1, 2516]	[2517, 3083]	6	0.207

A.8.6 Non-US Asset Class Extension and Per-Pair OoS Off-Diagonal Breakdown

The body universe (SPY, NVDA, JNJ, JPM, AAPL, QQQ) is six US-listed equity tickers from one country-index family. To probe whether the cross-asset construction generalises off the equity cluster we extend the universe with GLD (SPDR Gold Trust), a US-listed ETF whose underlying is physically-backed gold, i.e. a non-equity commodity asset class.

GLD Pipeline-A univariate fidelity. Fitting CHMM-N, CHMM-t, and CHMM-L on the GLD return series at $K = 18$ produces 100% IS KS pass rate for all three families and observed IS excess kurtosis 3.6 matched closely by CHMM-L (2.2); on OoS, however, the KS pass rate collapses to 0% for all three families against an observed OoS excess kurtosis of 6.5. Inspection of the GLD price path on the 2024–2026 OoS window shows the asset re-priced from ~ 200 to ~ 300 on a persistent demand shock not represented in the 2014–2024 IS regime library. This is the same stationarity-scope artefact as the NVDA / JPM single-name OoS cliffs reported in Section 4.4, on a non-equity asset; the limitation is not equity-specific.

Pipeline-B: 7-ticker Student-t copula off-diagonal MAE. Refitting the Student-t copula on the augmented seven-ticker universe at $\nu^* = 6$ gives IS off-diagonal MAE of 0.029 (six-ticker baseline: 0.027) and OoS off-diagonal MAE of 0.179 (six-ticker baseline: 0.209). Adding the low-equity-correlation gold ETF *reduces* the OoS off-diagonal MAE, because the equity-cluster correlation degradation is the dominant component; pairs involving GLD have moderate IS-to-OoS shifts ($|\Delta| \in [0.04, 0.19]$) while the largest equity-cluster pairs (JNJ-QQQ, SPY-JNJ, NVDA-JNJ) have $|\Delta| > 0.48$ on OoS.

Table 68. Per-pair off-diagonal correlation MAE on the seven-ticker universe (Pipeline B, Student-t copula on CHMM-N marginals at $K = 18$, 200 paths, $\nu^* = 6$). For each ticker pair (i, j) the table reports $|\Sigma_{\text{sim,avg}}[i, j] - \Sigma_{\text{obs}}[i, j]|$ on IS and OoS, ranked by OoS $|\Delta|$. Only the ten largest-OoS-gap pairs are shown for compactness.

Pair	IS $ \Delta $	OoS $ \Delta $	OoS – IS
JNJ-QQQ	0.027	0.544	0.517
SPY-JNJ	0.037	0.509	0.473
NVDA-JNJ	0.017	0.489	0.473
NVDA-AAPL	0.011	0.278	0.268
JNJ-JPM	0.046	0.261	0.216
JNJ-AAPL	0.042	0.244	0.202
AAPL-QQQ	0.003	0.227	0.224
JPM-GLD	0.027	0.187	0.160
SPY-GLD	0.038	0.127	0.089
JNJ-GLD	0.015	0.117	0.101

Per-pair OoS gap localisation. Table 68 reports the per-pair off-diagonal absolute deviation between observed and path-averaged simulated correlation matrices on IS and OoS, sorted by OoS $|\Delta|$. The OoS gap concentrates in pairs involving JNJ (the defensive single-name proxy): JNJ’s IS correlation with the broad equity factor (QQQ, SPY) collapses on OoS, in some pairs flipping sign, and the IS-fixed copula has no mechanism for tracking that shift. The gold-ETF pairs (SPY-GLD, JPM-GLD, JNJ-GLD, QQQ-GLD) sit in the middle of the ranking with $|\Delta|$ between 0.10 and 0.19 on OoS. The dominant cause of the OoS off-diagonal gap is therefore neither equity-cluster size nor copula-degree mis-specification but a single-name regime shift in the JNJ-equity-factor relationship that the IS-fitted dependence layer cannot anticipate.

A.8.7 Half-Unit-Grid Profile-MLE and Parametric Bootstrap CI for ν^*

Reviewer 2 / item 10 asks for a half-unit-grid refinement of the unit-grid profile MLE on $\nu \in [3, 12]$ plus a parametric bootstrap CI to confirm that the upper bound of the unit-grid Wilks CI is not grid-induced. We refit the Student-t copula profile log-likelihood on the same six-asset universe at half-unit spacing in $\nu \in \{3.0, 3.5, \dots, 12.0\}$ (19 grid points) and run a $B = 200$ parametric bootstrap by resampling pseudo-observations $U^{(b)}$ from the Student-t copula at $(\Sigma_{\text{cop}}, \nu^*)$, refitting Kendall’s- τ correlation and the profile-MLE on each replicate. Results in `results/copula_profile_ci/profile_ll_halfunit_summary.txt`.

The half-unit grid moves the optimum from $\nu^* = 6.0$ (unit grid) to $\nu^* = 6.5$ at profile log-likelihood 6,158.1 (vs. 6,157.5 at the unit-grid optimum); the half-unit refinement is therefore not material at the third decimal of profile-LL. The Wilks 95% CI on the half-unit grid is $[\nu_{\text{lo}}, \nu_{\text{hi}}] = [6.0, 7.0]$ (unchanged from the unit grid), and the parametric bootstrap 95% CI is $[6.0, 7.0]$ (2.5% and 97.5% quantiles of the $B = 200$ refit- ν^* distribution; bootstrap median $\nu^* = 6.5$). The bootstrap lower bound at 6.0 is well above the Gaussian limit ($\nu \rightarrow \infty$): the elliptical-tail Student-t copula is statistically distinguishable from the Gaussian copula at conventional levels, despite the OoS off-diagonal MAE comparison (0.209 vs. 0.202) being inside the simulation-noise floor at $N_{\text{paths}} = 200$. *Wilks-theorem regularity caveat.* Wilks’s theorem requires the parameter under test to lie in the interior of the parameter space and the model to be regular [83, 84]; for the Gaussian-copula limit $\nu \rightarrow \infty$ the regularity condition fails (the limit lies on the boundary of the natural parameter space). The Wilks 95% CI reported above is therefore an interior-grid CI for finite ν rather than a hypothesis test against the Gaussian limit; the Gaussian-vs-Student-t separation is reported instead through

the parametric bootstrap CI (which does not require interior-point regularity) and through the OoS-equivalence finding (which uses no asymptotic test at all).

A.9 Pre-OoS Validation K -Selection and ν_k Diagnostics

Pre-OoS validation K -selection. The $K = 18$ operating point in the main panel was chosen via a multi-objective criterion combining information-criterion (IC) rank with IS and OoS distributional pass rates, which uses the same 2024–2026 window later used for VaR. As a clean re-selection that decouples the two, we fit CHMM-N on 2014-01-03 through 2021-12-31 ($n = 2,013$) and score every K in the sweep by the forward-algorithm held-out log-likelihood on 2022-01-03 through 2024-01-03 ($n = 502$); the 2024–2026 window is not touched. Held-out log-lik, BIC, HQC, and CAIC all select $K^* = 3$ (per-observation log-lik -2.5345 at $K = 3$ vs. -2.5694 at $K = 18$); AIC selects $K^* = 6$; held-out two-sample KS selects $K^* = 9$. None of the six criteria that avoid the OoS window select $K = 18$. The $K = 18$ choice is best read as a multi-metric one that weights tail fidelity (kurtosis, AD) and downstream VaR behaviour alongside likelihood; a parallel main panel at $K = 3$ would re-rank some rows in CHMM’s favour on ACF-MAE and against it on tail fidelity. The 2022–2023 validation slice is itself partly a structural-break period (the rate-hike cycle), so any model with substantial IS-specific structure generalises worse on it.

Pre-2020 held-out K -selection (no rate-hike confound). Reviewer R1’s Question 1 asked whether a strictly pre-2020 validation slice (avoiding the rate-hike confound that contaminates the 2022-2023 slice above) selects the same K^* . We re-ran the held-out procedure on a 2014-01-03 through 2018-06-29 estimation window ($n_{\text{est}} = 1,130$) with a 2018-07-02 through 2019-12-31 validation slice ($n_{\text{val}} = 377$), both fully pre-COVID and pre-2022-rate-hike (the Q4 2018 drawdown and 2019 recovery sit inside the validation window as a moderate-stress non-regime-shift event). Per-observation held-out log-likelihood: -2.0681 at $K = 3$, -2.0606 at $K = 6$, -2.0763 at $K = 9$, -2.1287 at $K = 12$, -2.1921 at $K = 18$, -2.2609 at $K = 21$. **The held-out log-likelihood selects $K^* = 6$, and the held-out two-sample KS pass rate also selects $K^* = 6$ (96.8% at $K = 6$ versus 84.6% at $K = 3$).** AIC, BIC, HQC, and CAIC continue to select $K^* = 3$ on this slice, reflecting the parameter-count penalty on a 4.5-year estimation window where the per-state sample is small. The qualitative reading is robust to the slice choice: held-out distributional fidelity rises from $K = 3$ to a peak in the $K = 6$ to $K = 9$ band and then degrades, with neither slice selecting $K = 18$ by any held-out criterion, but the pre-2020 result favours a moderately larger K (6 vs 3) than the rate-hike-contaminated 2022-2023 slice.

k -fold rolling-origin pre-2020 K -selection (R1 RE1 / R1 W1 contingency). Reviewer R1’s Requested Experiment 1 (priority 1) flags that the single-fold pre-2020 result above is one observation: *“report mean $\pm$ s.d. of held-out per-observation log-likelihood and held-out KS at $K \in \{3, 6, 9, 12, 18\}$ over 5 or 10 folds. If $K = 6$ remains preferred over $K = 3$ outside sampling error, the headline claim is supported; if not, the body should be rebuilt at $K^* = 3$ ”*. We implement four expanding-window rolling-origin folds (one fewer than the reviewer-requested five because a five-fold design forces one fold to have only ~ 1 year of training, below the practical floor for $K = 18$ EM convergence on this dataset; averaging is per-observation so fold-length differences are not an issue). Folds: train 2014-01-03 through 2015-12-31 (~ 502 obs) and val 2016 (~ 251 obs); train through 2016 and val 2017; train through 2017 and val 2018; train through 2018 and val 2019. Source: `results/robustness/k_selection_kfold_pre2020.csv` and `_agg.csv`; runner: `run_k_selection_kfold_pre2020.jl` of the companion code repository.

Aggregate (mean / s.d. across the four folds) of per-observation held-out log-likelihood: $-1.9550 / 0.290$ at $K = 3$; $-1.9649 / 0.313$ at $K = 6$; $-2.0078 / 0.296$ at $K = 9$; $-2.0721 / 0.251$ at $K = 12$; $-2.2633 / 0.308$ at $K = 18$. Held-out KS pass rate (mean / s.d. across folds, very wide because Fold 2’s 2017 validation year produced near-zero KS for every K as a low-variance “calm year” artefact): $61.3/41.2\%$ at $K = 3$; $65.2/43.5\%$ at $K = 6$; $67.5/44.1\%$ at $K = 9$; $67.0/44.7\%$ at $K = 12$; $67.5/45.4\%$ at $K = 18$. Sampling-error reads on the held-out per-observation log-likelihood: $K = 6$ vs $K = 3$ has mean diff -0.0099 , pooled SE 0.151 , approximate $z = -0.07$ (*not* significant at the 5% level); $K = 18$ vs $K = 6$ has mean diff -0.298 , pooled SE 0.155 , approximate $z = -1.92$ (borderline, just below the $|z| = 1.96$ threshold).

Robustness check at half-year cadence. To verify that the result is not an artefact of the one-year fold cadence, we re-ran the same diagnostic with six expanding-window folds at half-year validation cadence (covering 2017–2019 in non-overlapping six-month chunks; train windows $\sim 3.0y$ to $\sim 5.5y$). Source: `results/robustness/k_selection_kfold_h12y_pre2020.csv` and `_agg.csv`; runner: `run_k_selection_kfold_h12y_pre2020.jl`. Aggregate per-observation held-out log-likelihood (mean / s.d. across six folds): $-1.9221 / 0.341$ at $K = 3$; $-1.9160 / 0.346$ at $K = 6$; $-1.9568 / 0.323$ at $K = 9$; $-2.0109 / 0.315$ at $K = 12$; $-2.1281 / 0.259$ at $K = 18$. Sampling-error reads: $K = 6$ vs $K = 3$ has mean diff $+0.006$, pooled SE 0.140 , approximate $z = +0.04$; $K = 18$ vs $K = 6$ has mean diff -0.212 , pooled SE 0.125 , approximate $z = -1.70$. The half-year design’s $K = 6$ vs $K = 3$ sign flips relative to the full-year design ($z = +0.04$ vs $z = -0.07$), but both magnitudes are well below the $|z| = 1$ threshold; the two designs agree that $K = 6$ vs $K = 3$ on mean held-out log-likelihood is indistinguishable from zero, and the sign flip across designs is itself evidence that the $K = 6 / K = 3$ choice is pure sampling noise on this data. The $K = 18$ vs $K = 6$ borderline result also replicates ($z = -1.70$ at half-year cadence vs $z = -1.92$ at full-year).

Reading. $K = 6$ is not significantly preferred over $K = 3$ on held-out per-observation log-likelihood under expanding-window rolling-origin CV at either yearly or half-yearly fold cadence. The single-fold result that initially selected $K^* = 6$ is one realisation; both four-fold and six-fold means cannot distinguish $K = 3$ from $K = 6$ at conventional levels. The held-out KS pass-rate criterion is not informative on either design (between-fold s.d. is dominated by 2017 calendar-year artefacts that hit every K uniformly). The substantive read is that the held-out state-resolution selection on the strictly pre-2020 slice cannot distinguish $K = 3$ from $K = 6$ at conventional levels under either fold design, and the body operating point is therefore the state-resolution-robust $K^* = 3$ (with the $K^* = 6$ block retained as a sensitivity reference). $K^* = 3$ is also the lower-state-count default for risk-management consumers in Table 8.

HAC-corrected sampling-error reads on the paired diff series. Rolling-origin folds inherit autocorrelation from the underlying series, so the independent-fold pooled SE used in the paragraph above understates the standard error of the mean held-out log-likelihood difference. Round-2 reviewers asked for a Diebold–Mariano-style HAC variance computed on the paired diff series across folds with Newey–West Bartlett kernel at truncation lag $h_{NW} = \lfloor n^{1/3} \rfloor$. Source / runner: `run_k_selection_hac.jl`; output: `results/k_selection_hac/k_selection_hac.csv`.

The substantive read on $K = 6$ vs $K = 3$ is unchanged: HAC $|z| = 0.90$ (4-fold) and 0.57 (6-fold) are both well below conventional levels, consistent with the paragraph above. **The $K = 18$ vs $K = 6$ statistic, however, jumps from a borderline independent-fold $|z| = 1.92 / 1.70$ to a decisive HAC $|z| = 3.56 / 5.00$:** the HAC-corrected inference rejects “ $K = 18$ no worse than $K = 6$ on held-out per-observation log-likelihood” at conventional levels under both fold designs. The body’s framing of $K = 18$ as a multi-objective-rule sensitivity reference is therefore corroborated by the HAC inference; the $K = 18$ block is decisively worse than $K = 6$ on held-out log-likelihood

Table 69. HAC-corrected K -selection inference on the paired diff series across rolling-origin folds. Two-sided NW-HAC Diebold–Mariano on the per-fold val-log-lik differences; truncation lag $h_{\text{NW}} = \lfloor n^{1/3} \rfloor$. “indep $|z|$ ” is the pooled-independent-fold statistic of the paragraph above; “HAC $|z|$ ” is the HAC-corrected statistic.

Cadence	Comparison	$\bar{d}$	n	indep $ z $	HAC SE	HAC $ z $
4-fold full-year	$K = 6$ vs $K = 3$	−0.010	4	0.07	0.011	0.90
4-fold full-year	$K = 18$ vs $K = 6$	−0.298	4	1.92	0.084	3.56
4-fold full-year	$K = 18$ vs $K = 3$	−0.308	4	—	0.093	3.31
6-fold half-year	$K = 6$ vs $K = 3$	+0.006	6	0.04	0.011	0.57
6-fold half-year	$K = 18$ vs $K = 6$	−0.212	6	1.70	0.042	5.00
6-fold half-year	$K = 18$ vs $K = 3$	−0.206	6	—	0.040	5.10

and is retained only because of its kurtosis-fidelity advantage on the OoS window (Table 3).

Penalised-ECM rate sweep. The bracket sensitivity is reported with its table in Section A.2.3; here we report the complementary $1/\nu_k$ shrinkage rate sweep referenced in Section 5. An exponential $1/\nu_k$ prior at rate $\lambda \in \{0, 5, 20, 50, 100, 200\}$ reduces simulated excess kurtosis monotonically: 14.30, 11.19, 8.43, 5.41, 3.96, 3.67. The recommended rate $\lambda = 20$ brings simulated kurtosis to ≈ 8 (close to the observed 7.69) at a 1pp IS KS cost. Only the $\nu_{\min}$ state parameter actually moves ($2.1 \rightarrow 4.89$); upper bracket and median remain at 50. The overshoot is therefore a single-state artefact of the two tail regimes hitting the lower bracket under unpenalised ECM.

Walk-forward / rolling-origin OoS for CHMM-N. We define six rolling-origin folds, each train 5 years and test 1 year, refit CHMM-N at $K \in \{3, 18\}$ from scratch on each fold’s train slice, and report KS pass rate, simulated kurtosis, and $|G_t|$ ACF-MAE on the fold’s test slice ($N_{\text{paths}} = 500$). Across the six folds (Table 70), CHMM-N at $K = 18$ attains median (IQR) KS 67.7% [8.2, 75.0] and $|G_t|$ ACF-MAE 0.0542 [0.0508, 0.0571]; CHMM-N at $K = 3$ attains 62.1% [7.2, 78.4] and 0.0563 [0.0552, 0.0592]. Two folds carry sharply lower KS pass rates: W2 (COVID test slice, 7–8%) and W4 (rate-hike test slice, 0–1%). The four non-stress folds attain KS 61–83%, consistent with the headline OoS pass rate. The headline ranking (CHMM-N at $K = 18 \geq$ CHMM-N at $K^* = 3$ on KS, both within 0.005 of each other on $|G_t|$ ACF-MAE) is window-robust on the four non-stress folds; on stress folds (COVID, rate hike) all CHMM rows drop sharply, reflecting the stationarity-scope limit rather than a model-selection issue. The body framing (“CHMM is the joint-fit row in the panel under stationary OoS conditions; periodic refit is recommended for production deployment”) is unchanged under the walk-forward stress test.

Four-family conditional VaR back-test. Extending the regime-conditional VaR construction of §4.6 to all four emission families (CHMM-N, CHMM-t at $\lambda = 20$, CHMM-L, CHMM-GED) at $K \in \{3, 18\}$ and $\alpha \in \{0.01, 0.05\}$ yields sixteen $(K, \alpha, \text{family})$ rows on the OoS window (Table 71). Every row passes Kupiec, Christoffersen-ind, and Christoffersen-cc cleanly: $p_{\text{cc}} \geq 0.089$ throughout, with all LR_{cc} statistics below the $\chi^2_2(0.05) = 5.991$ critical value. The regime-switching value proposition is therefore intrinsic to the latent-state forecast and not specific to the Gaussian emission family.

Walk-forward conditional-VaR back-test. Extending the regime-conditional VaR construction to the six rolling-origin folds of Table 70, with CHMM-N refit per fold from scratch on the train

Table 70. Walk-forward / rolling-origin OoS for CHMM-N. Six folds, each train 5 years and test 1 year; $N_{\text{paths}} = 500$ per fold. KS pass rate at $\alpha = 0.05$. W2 covers the COVID drawdown; W4 covers the start of the 2022 rate-hike cycle. The bracketed pairs in the bottom row report $[Q_1, Q_3]$ across the six folds (the empirical first and third quartiles, not a scalar IQR width).

Fold (test window)	KS (%)		Kurt sim		ACF-MAE $ G_t $	
	$K = 3$	$K = 18$	$K = 3$	$K = 18$	$K = 3$	$K = 18$
W1 (2019)	63.2	82.6	2.51	2.66	0.0572	0.0525
W2 (2020 COVID)	7.2	8.2	2.42	2.50	0.0552	0.0584
W3 (2021)	82.4	74.0	5.73	6.76	0.0602	0.0571
W4 (2022 rate hike)	0.8	0.0	4.72	4.35	0.0592	0.0558
W5 (2023)	78.4	75.0	3.47	2.35	0.0511	0.0508
W6 (2024)	61.0	61.4	2.81	2.86	0.0555	0.0493
median	62.1	67.7	3.14	2.76	0.0563	0.0542
IQR	[7.2, 78.4]	[8.2, 75.0]	[2.51, 4.72]	[2.50, 4.35]	[0.0552, 0.0592]	[0.0508, 0.0571]

Table 71. Regime-conditional VaR back-test across four CHMM emission families on SPY OoS ($T_{\text{OoS}} = 572$, seed 20260420). Critical values: $\chi_1^2(0.05) = 3.841$, $\chi_2^2(0.05) = 5.991$. Every $(K, \alpha, \text{family})$ row passes Kupiec, Christoffersen-ind, and Christoffersen-cc at $\alpha = 0.05$.

Family	K	α	breaches	br rate	median $\widehat{\text{VaR}}_t$	LR_{uc}	LR_{ind}	LR_{cc}
CHMM-N	3	0.01	9	1.57%	-4.56	1.62	2.36	3.98
CHMM-N	3	0.05	35	6.12%	-2.87	1.41	0.01	1.42
CHMM-N	18	0.01	9	1.57%	-5.20	1.62	2.36	3.98
CHMM-N	18	0.05	26	4.55%	-3.02	0.26	0.52	0.78
CHMM-t ($\lambda = 20$)	3	0.01	8	1.40%	-5.59	0.82	2.81	3.63
CHMM-t ($\lambda = 20$)	3	0.05	32	5.59%	-2.73	0.41	0.77	1.19
CHMM-t ($\lambda = 20$)	18	0.01	10	1.75%	-6.25	2.65	1.98	4.62
CHMM-t ($\lambda = 20$)	18	0.05	29	5.07%	-3.15	0.01	1.39	1.40
CHMM-L	3	0.01	5	0.87%	-5.45	0.10	4.74	4.84
CHMM-L	3	0.05	26	4.55%	-2.63	0.26	2.23	2.48
CHMM-L	18	0.01	9	1.57%	-6.18	1.62	0.29	1.91
CHMM-L	18	0.05	30	5.24%	-2.80	0.07	1.16	1.23
CHMM-GED	3	0.01	7	1.22%	-5.73	0.27	3.34	3.61
CHMM-GED	3	0.05	31	5.42%	-2.54	0.21	0.96	1.16
CHMM-GED	18	0.01	9	1.57%	-6.55	1.62	2.36	3.98
CHMM-GED	18	0.05	25	4.37%	-3.10	0.50	2.56	3.06

slice and forward-filtered through the test slice under fold-IS-fixed parameters, gives twenty-four (fold, K, α) rows in Table 72. The conditional Christoffersen-cc passes at $\alpha = 0.05$ on ten of twelve fold- K combinations (W1, W3, W4, W5, W6 at both $K \in \{3, 18\}$); the two failures concentrate on W2 (COVID, $p_{\text{cc}} \in \{0.011, 0.002\}$ at $K \in \{3, 18\}$). At $\alpha = 0.01$ the construction passes on nine of twelve combinations: W2 fails at both K ($p_{\text{cc}} < 10^{-3}$) and W4 at $K = 18$ fails at $p_{\text{cc}} = 0.022$ on a LR_{ind} tail-event clustering rather than a LR_{uc} coverage miss ($\text{LR}_{\text{ind}} = 7.49$ at $\text{LR}_{\text{uc}} = 0.11$). The pass-rate is therefore 19/24 aggregate at the 5% level and 22/24 if we exclude the two stress folds that the univariate walk-forward already flagged as out-of-distribution by KS in Table 70. The reading is that the conditional construction is window-conditional: under regime shifts of the W2 / W4 magnitude (a regime introduction the IS distribution does not span; Section 5, stationarity-scope paragraph), the conditional VaR inherits the same out-of-scope behaviour as the unconditional

generators. On the four non-stress folds (W1, W3, W5, W6) plus W4 at $K = 3$ at the 5% level the conditional construction is window-robust, supporting the body framing of §4.6 extending beyond the headline single-window result.

Table 72. Walk-forward regime-conditional VaR back-test for CHMM-N on six rolling-origin folds (train 5y / test 1y; seed 20260420). At each test day t , conditional $\widehat{\text{VaR}}_t(\alpha)$ is the α -quantile of the K -component Gaussian mixture predictive density under fold-IS-fixed parameters, evaluated on the full-timeline filter $\Pr(s_{t+1} | \mathcal{F}_t = \text{train} \cup \text{test}[1:t-1])$. Critical values: $\chi_1^2(0.05) = 3.841$, $\chi_2^2(0.05) = 5.991$. Bold p_{cc} rows fail at the 5% level.

Fold	K	α	breaches	br rate	median $\widehat{\text{VaR}}_t$	LR _{uc}	LR _{ind}	LR _{cc}	p_{cc}
W1	3	0.01	2	0.80%	−4.01	0.11	0.03	0.15	0.93
W1	3	0.05	15	5.98%	−2.43	0.48	1.18	1.65	0.44
W1	18	0.01	0	0.00%	−5.03	5.05	0.00	5.05	0.08
W1	18	0.05	12	4.78%	−2.66	0.03	0.38	0.40	0.82
W2	3	0.01	16	6.35%	−6.54	32.93	0.87	33.80	0.00
W2	3	0.05	23	9.13%	−4.59	7.34	1.71	9.05	0.01
W2	18	0.01	14	5.56%	−6.67	25.59	1.56	27.15	0.00
W2	18	0.05	25	9.92%	−4.54	10.11	2.56	12.68	0.00
W3	3	0.01	6	2.39%	−4.48	3.53	0.30	3.82	0.15
W3	3	0.05	20	7.97%	−2.47	3.98	0.30	4.28	0.12
W3	18	0.01	2	0.80%	−5.57	0.11	0.03	0.15	0.93
W3	18	0.05	15	5.98%	−3.18	0.48	1.18	1.65	0.44
W4	3	0.01	5	2.00%	−7.98	1.96	0.21	2.16	0.34
W4	3	0.05	18	7.20%	−4.16	2.26	1.99	4.24	0.12
W4	18	0.01	2	0.80%	−9.62	0.11	7.49	7.60	0.02
W4	18	0.05	10	4.00%	−6.63	0.56	3.80	4.36	0.11
W5	3	0.01	0	0.00%	−6.06	5.01	0.00	5.01	0.08
W5	3	0.05	8	3.21%	−3.90	1.91	0.53	2.44	0.30
W5	18	0.01	1	0.40%	−7.10	1.16	0.01	1.17	0.56
W5	18	0.05	7	2.81%	−3.85	2.96	0.41	3.37	0.19
W6	3	0.01	2	0.80%	−4.68	0.11	0.03	0.15	0.93
W6	3	0.05	13	5.18%	−2.85	0.02	0.15	0.17	0.92
W6	18	0.01	1	0.40%	−5.83	1.19	0.01	1.20	0.55
W6	18	0.05	7	2.79%	−2.98	3.06	0.40	3.46	0.18

Per-ticker $\hat{\lambda}^*$ sweep. A per-ticker sweep $\lambda \in \{0, 5, 10, 20, 50, 100\}$ on each of the six body tickers under the constraint that IS KS at λ exceeds the $\lambda = 0$ baseline by no more than 1.5 percentage points (Table 73) shows that the body uniform $\lambda = 20$ is the right rate on the heavy-tailed defensives (NVDA, JNJ) but is over-shrunk on the moderate-tail tickers. The per-ticker optima are $\lambda_{\text{SPY}}^* = 10$, $\lambda_{\text{NVDA}}^* = 20$, $\lambda_{\text{JNJ}}^* = 20$, $\lambda_{\text{JPM}}^* = 10$, $\lambda_{\text{AAPL}}^* = 10$, $\lambda_{\text{QQQ}}^* = 0$. The honest operational reading is that $\lambda = 20$ is a reasonable uniform default for production deployment but per-ticker tuning is recommended whenever a ticker’s residual kurtosis at $\lambda = 0$ is within ~ 1 unit of observed.

Cross-asset summary. The full per-construction Pipeline B panel is in Section A.8; for reference, the headline ordering is summarised in Table 74.

Table 73. Per-ticker $\hat{\lambda}^*$ for the penalised CHMM-t at $K = 18$ ($N_{\text{paths}} = 500$, seed 20260420, KS-degradation tolerance $\leq 1.5\text{pp}$ from $\lambda = 0$). $\hat{\lambda}^*$ is the value of λ in $\{0, 5, 10, 20, 50, 100\}$ that minimises $|\text{kurt}_{\text{sim}} - \text{kurt}_{\text{obs}}|$ subject to the KS-degradation constraint.

Ticker	kurt_{obs}	$\hat{\lambda}^*$	$\text{kurt}_{\text{sim}}(\lambda^*)$	IS KS(λ^*)
SPY	7.69	10	11.82	96.4%
NVDA	5.54	20	5.68	99.4%
JNJ	8.98	20	10.74	99.8%
JPM	7.63	10	8.62	98.6%
AAPL	3.20	10	3.24	99.2%
QQQ	3.40	0	3.35	95.2%

Table 74. Cross-asset dependence summary (Pipeline B; CHMM-N marginals at $K = 18$; 200 paths; seed 20260422). Median per-asset KS pass rate (%) and off-diagonal MAE of the simulated correlation matrix vs. observed. Bold marks the column winner. The Student-t copula at $\nu^* = 6$ is the strongest dependence layer on the two IS columns; on OoS the Gaussian copula attains marginally lower off-diagonal MAE (0.202 vs. 0.209) while the truncated C-vine sits between the elliptical copulas and the SIM baseline. Full per-asset detail is in Section A.8; the C-vine edge summary is in Table 66.

	Median IS KS	Median OoS KS	Off-diag MAE IS	Off-diag MAE OoS
Single Index Model	77.0	87.0	0.076	0.252
Gaussian copula	97.8	87.5	0.030	0.202
Student-t copula ($\nu^* = 6$)	98.8	85.8	0.027	0.209
Truncated C-vine (root SPY)	97.2	86.0	0.068	0.235